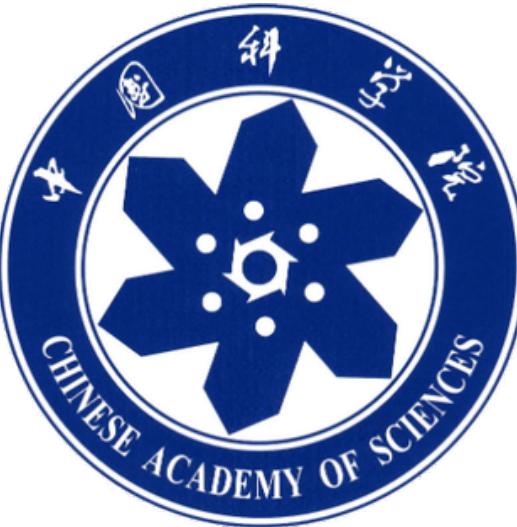
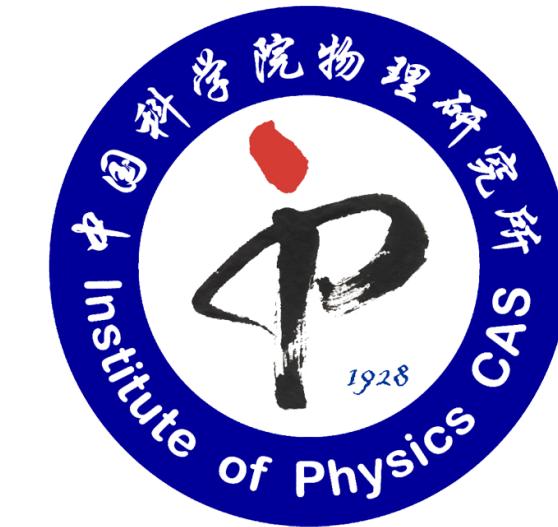


Unlocking the power of the variational free-energy principle with deep generative models

Lei Wang (王磊)

Institute of Physics, CAS

<https://wangleiphy.github.io>



From the 2020 edition
of this school

generative model
for sampling
quantum state
tomography

force
field

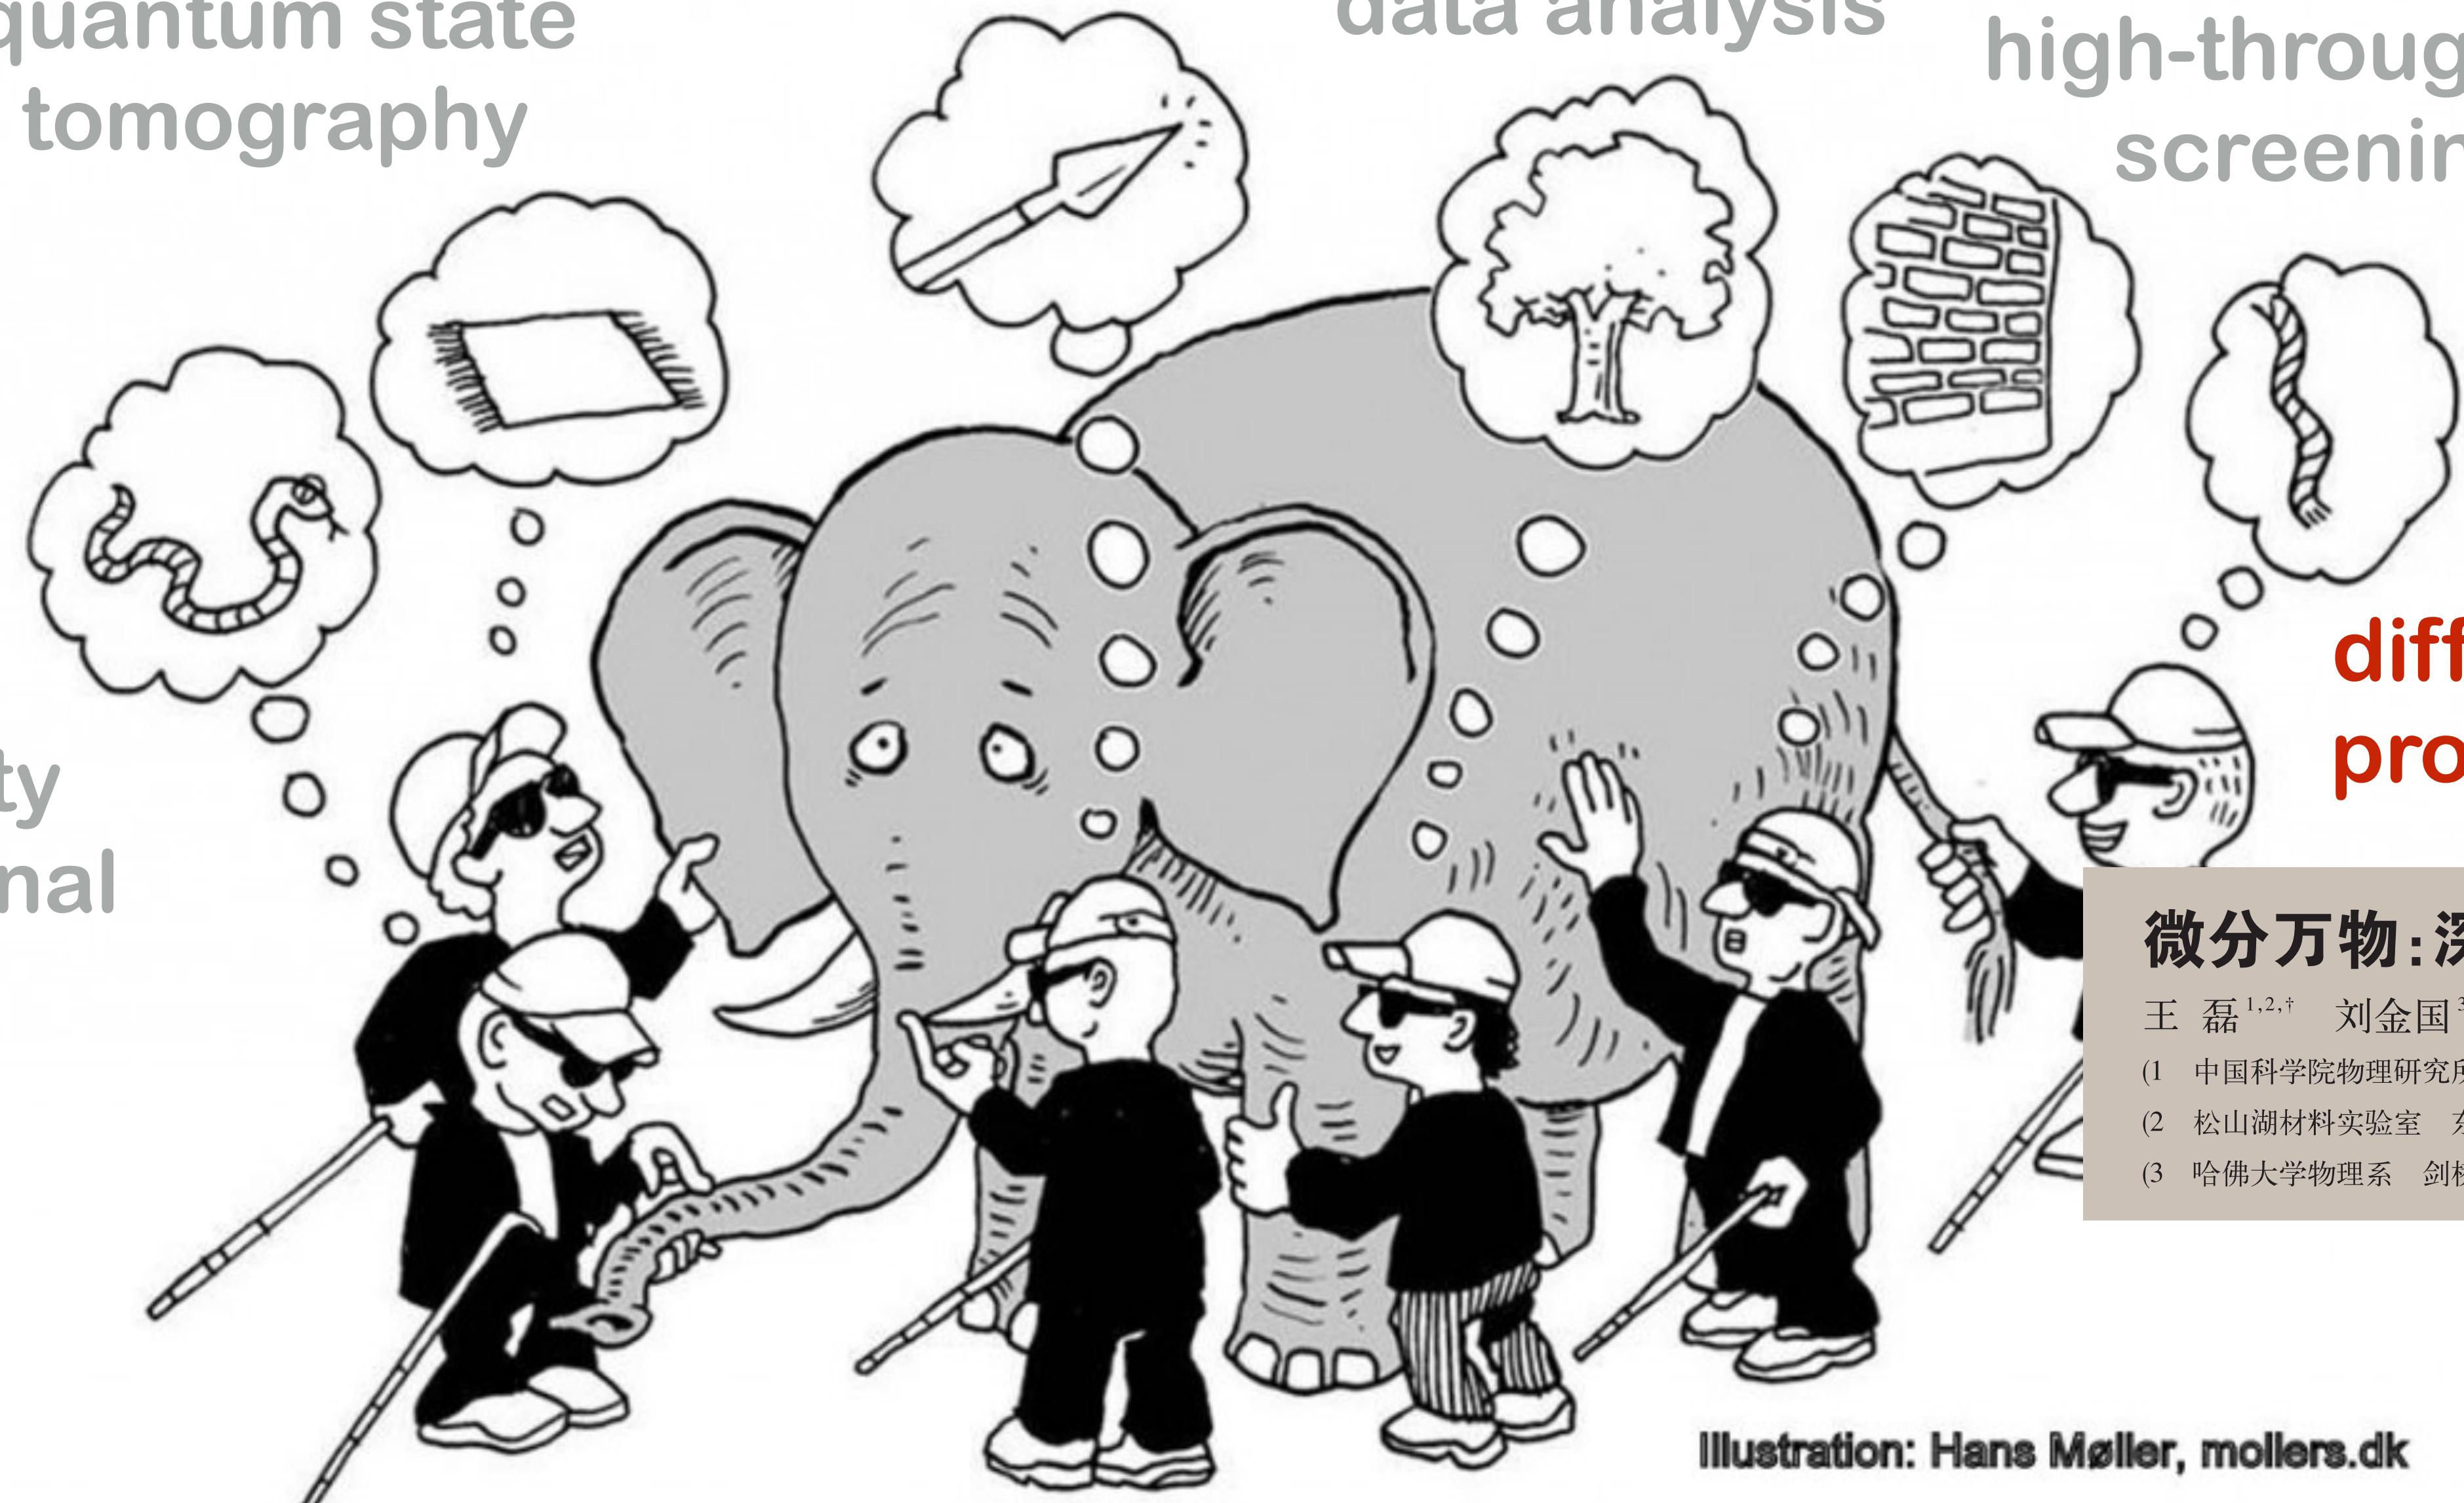
density
functional

experimental
data analysis

high-throughput
screening

inverse
design

**differentiable
programming**



微分万物：深度学习的启示*

王磊^{1,2,†} 刘金国³

(1) 中国科学院物理研究所 北京 100190)

(2) 松山湖材料实验室 东莞 523808)

(3) 哈佛大学物理系 剑桥 02138)

《物理》
2021年2月

Illustration: Hans Møller, moppers.dk

generative model for sampling (and more)

experimental
data analysis

quantum state
tomography

force
field

density
functional

high-throughput
screening

inverse
design

differentiable
programming

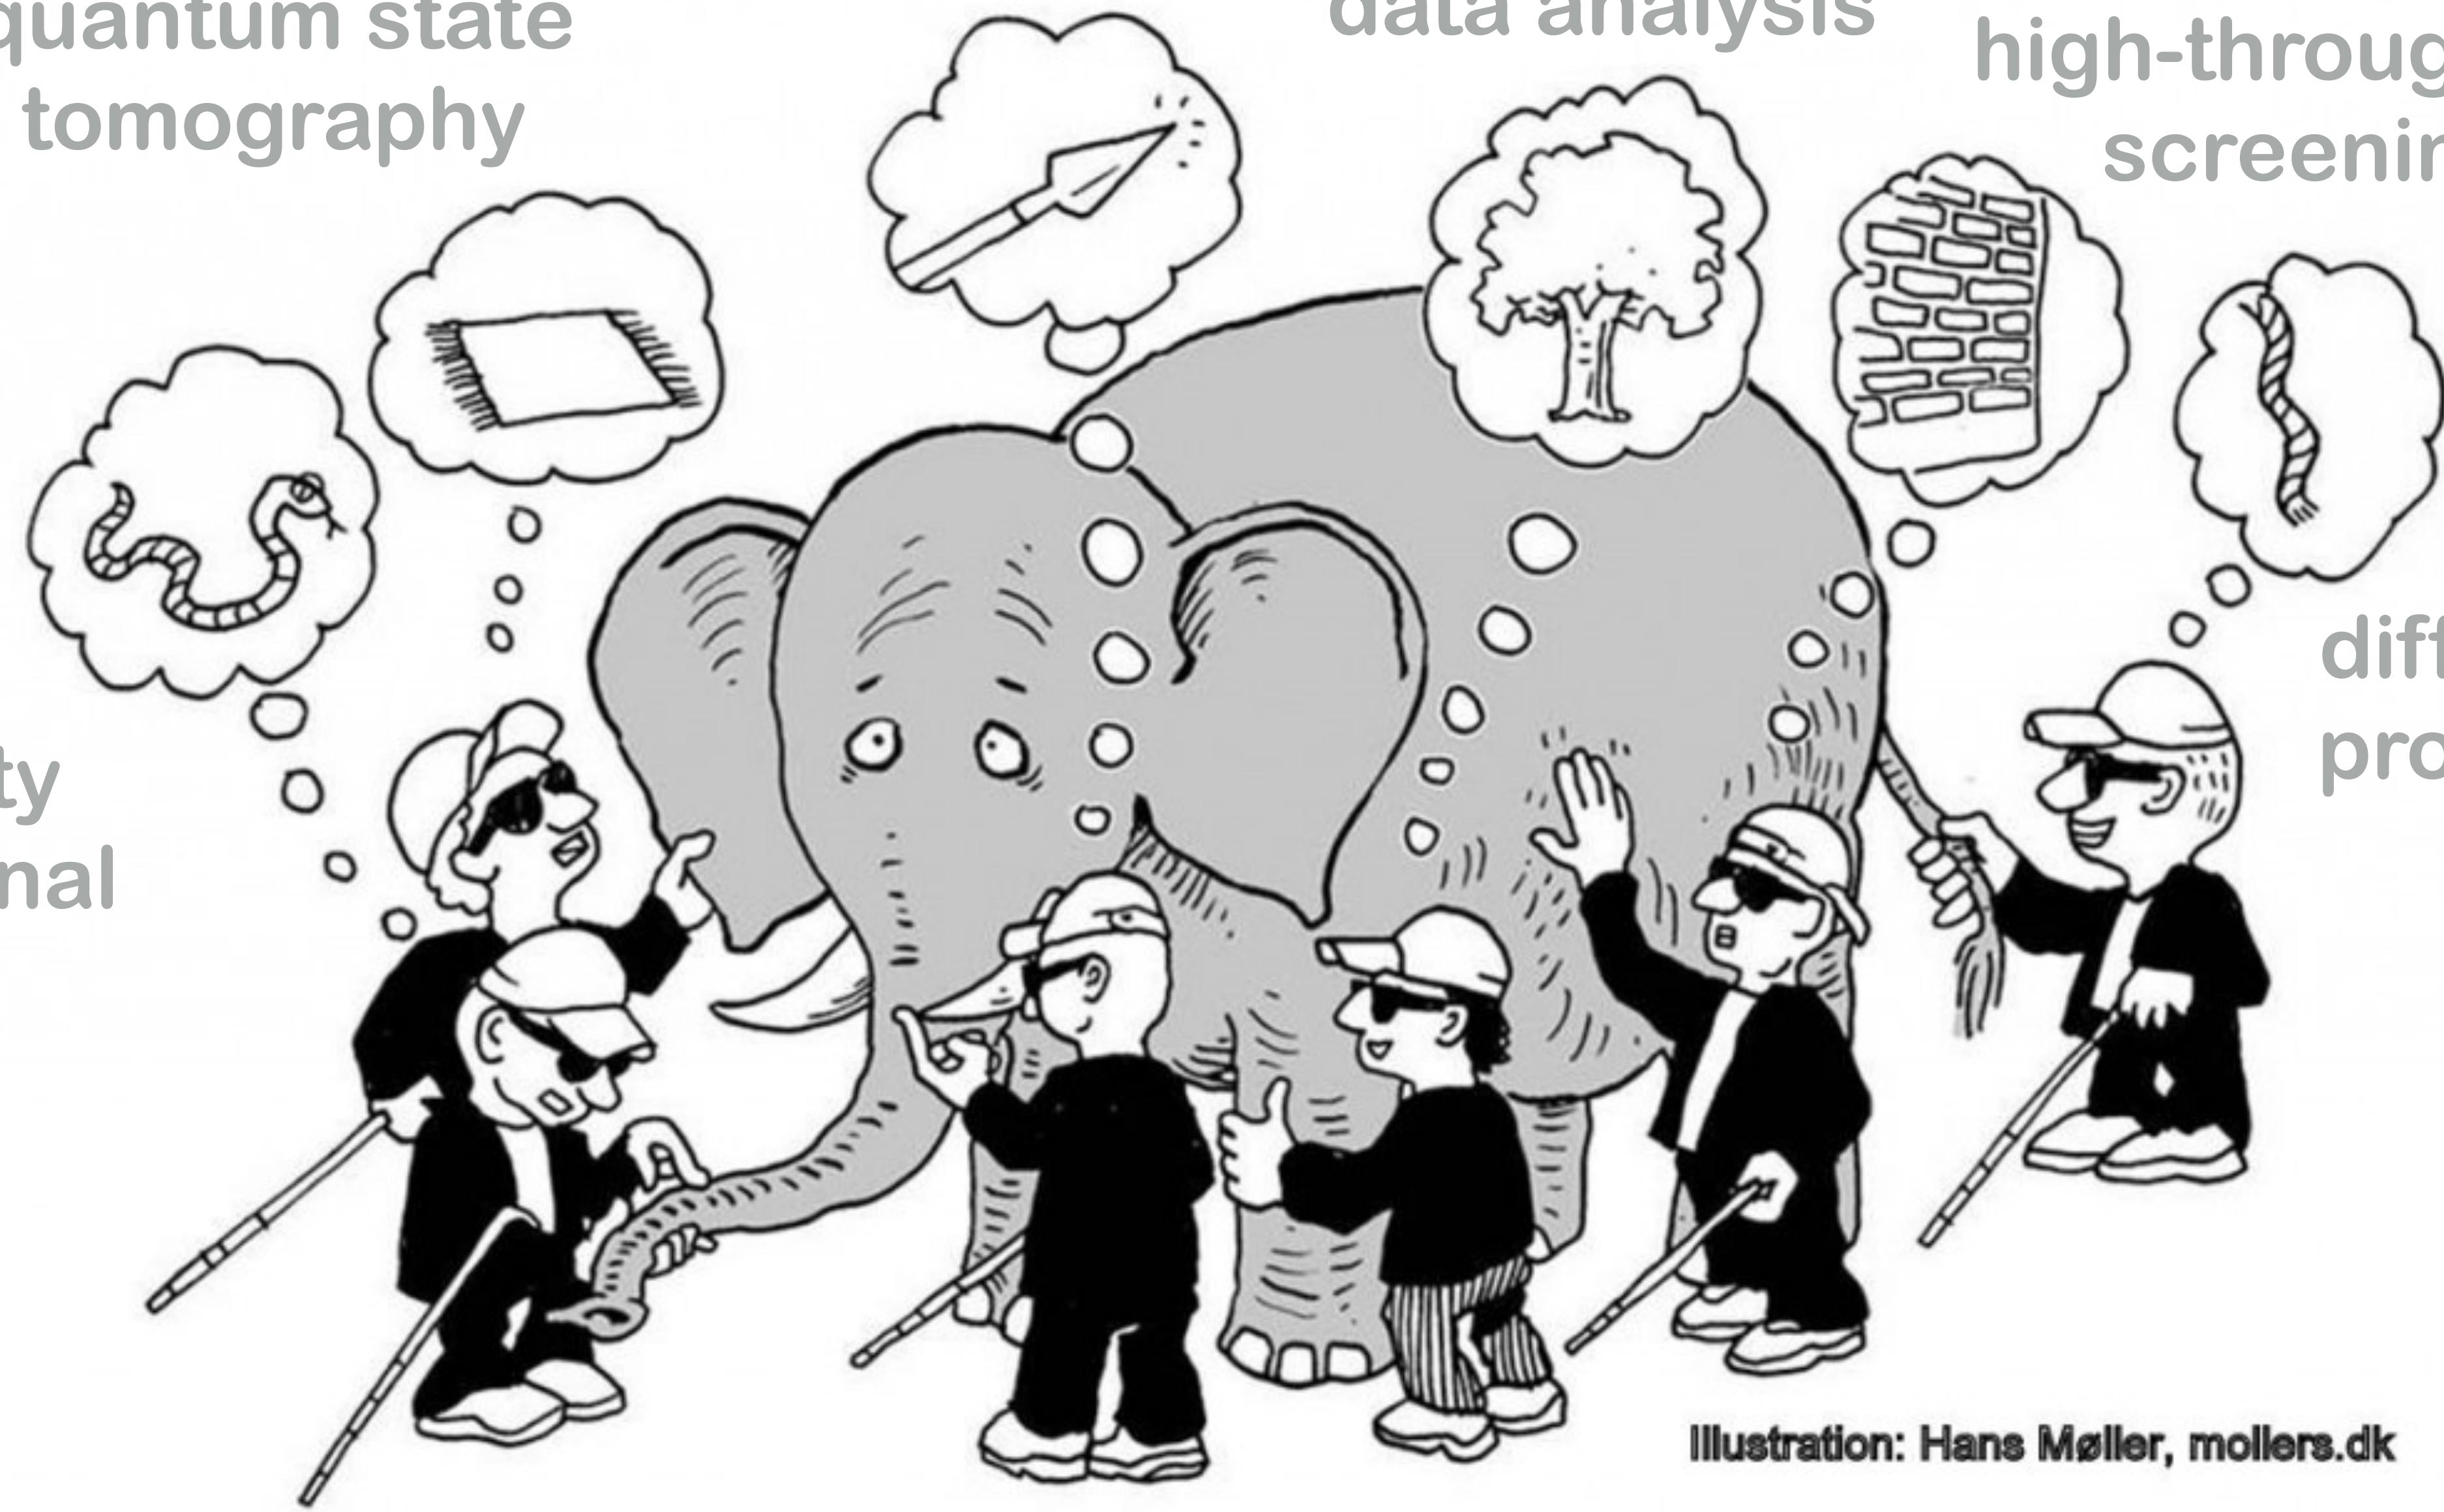
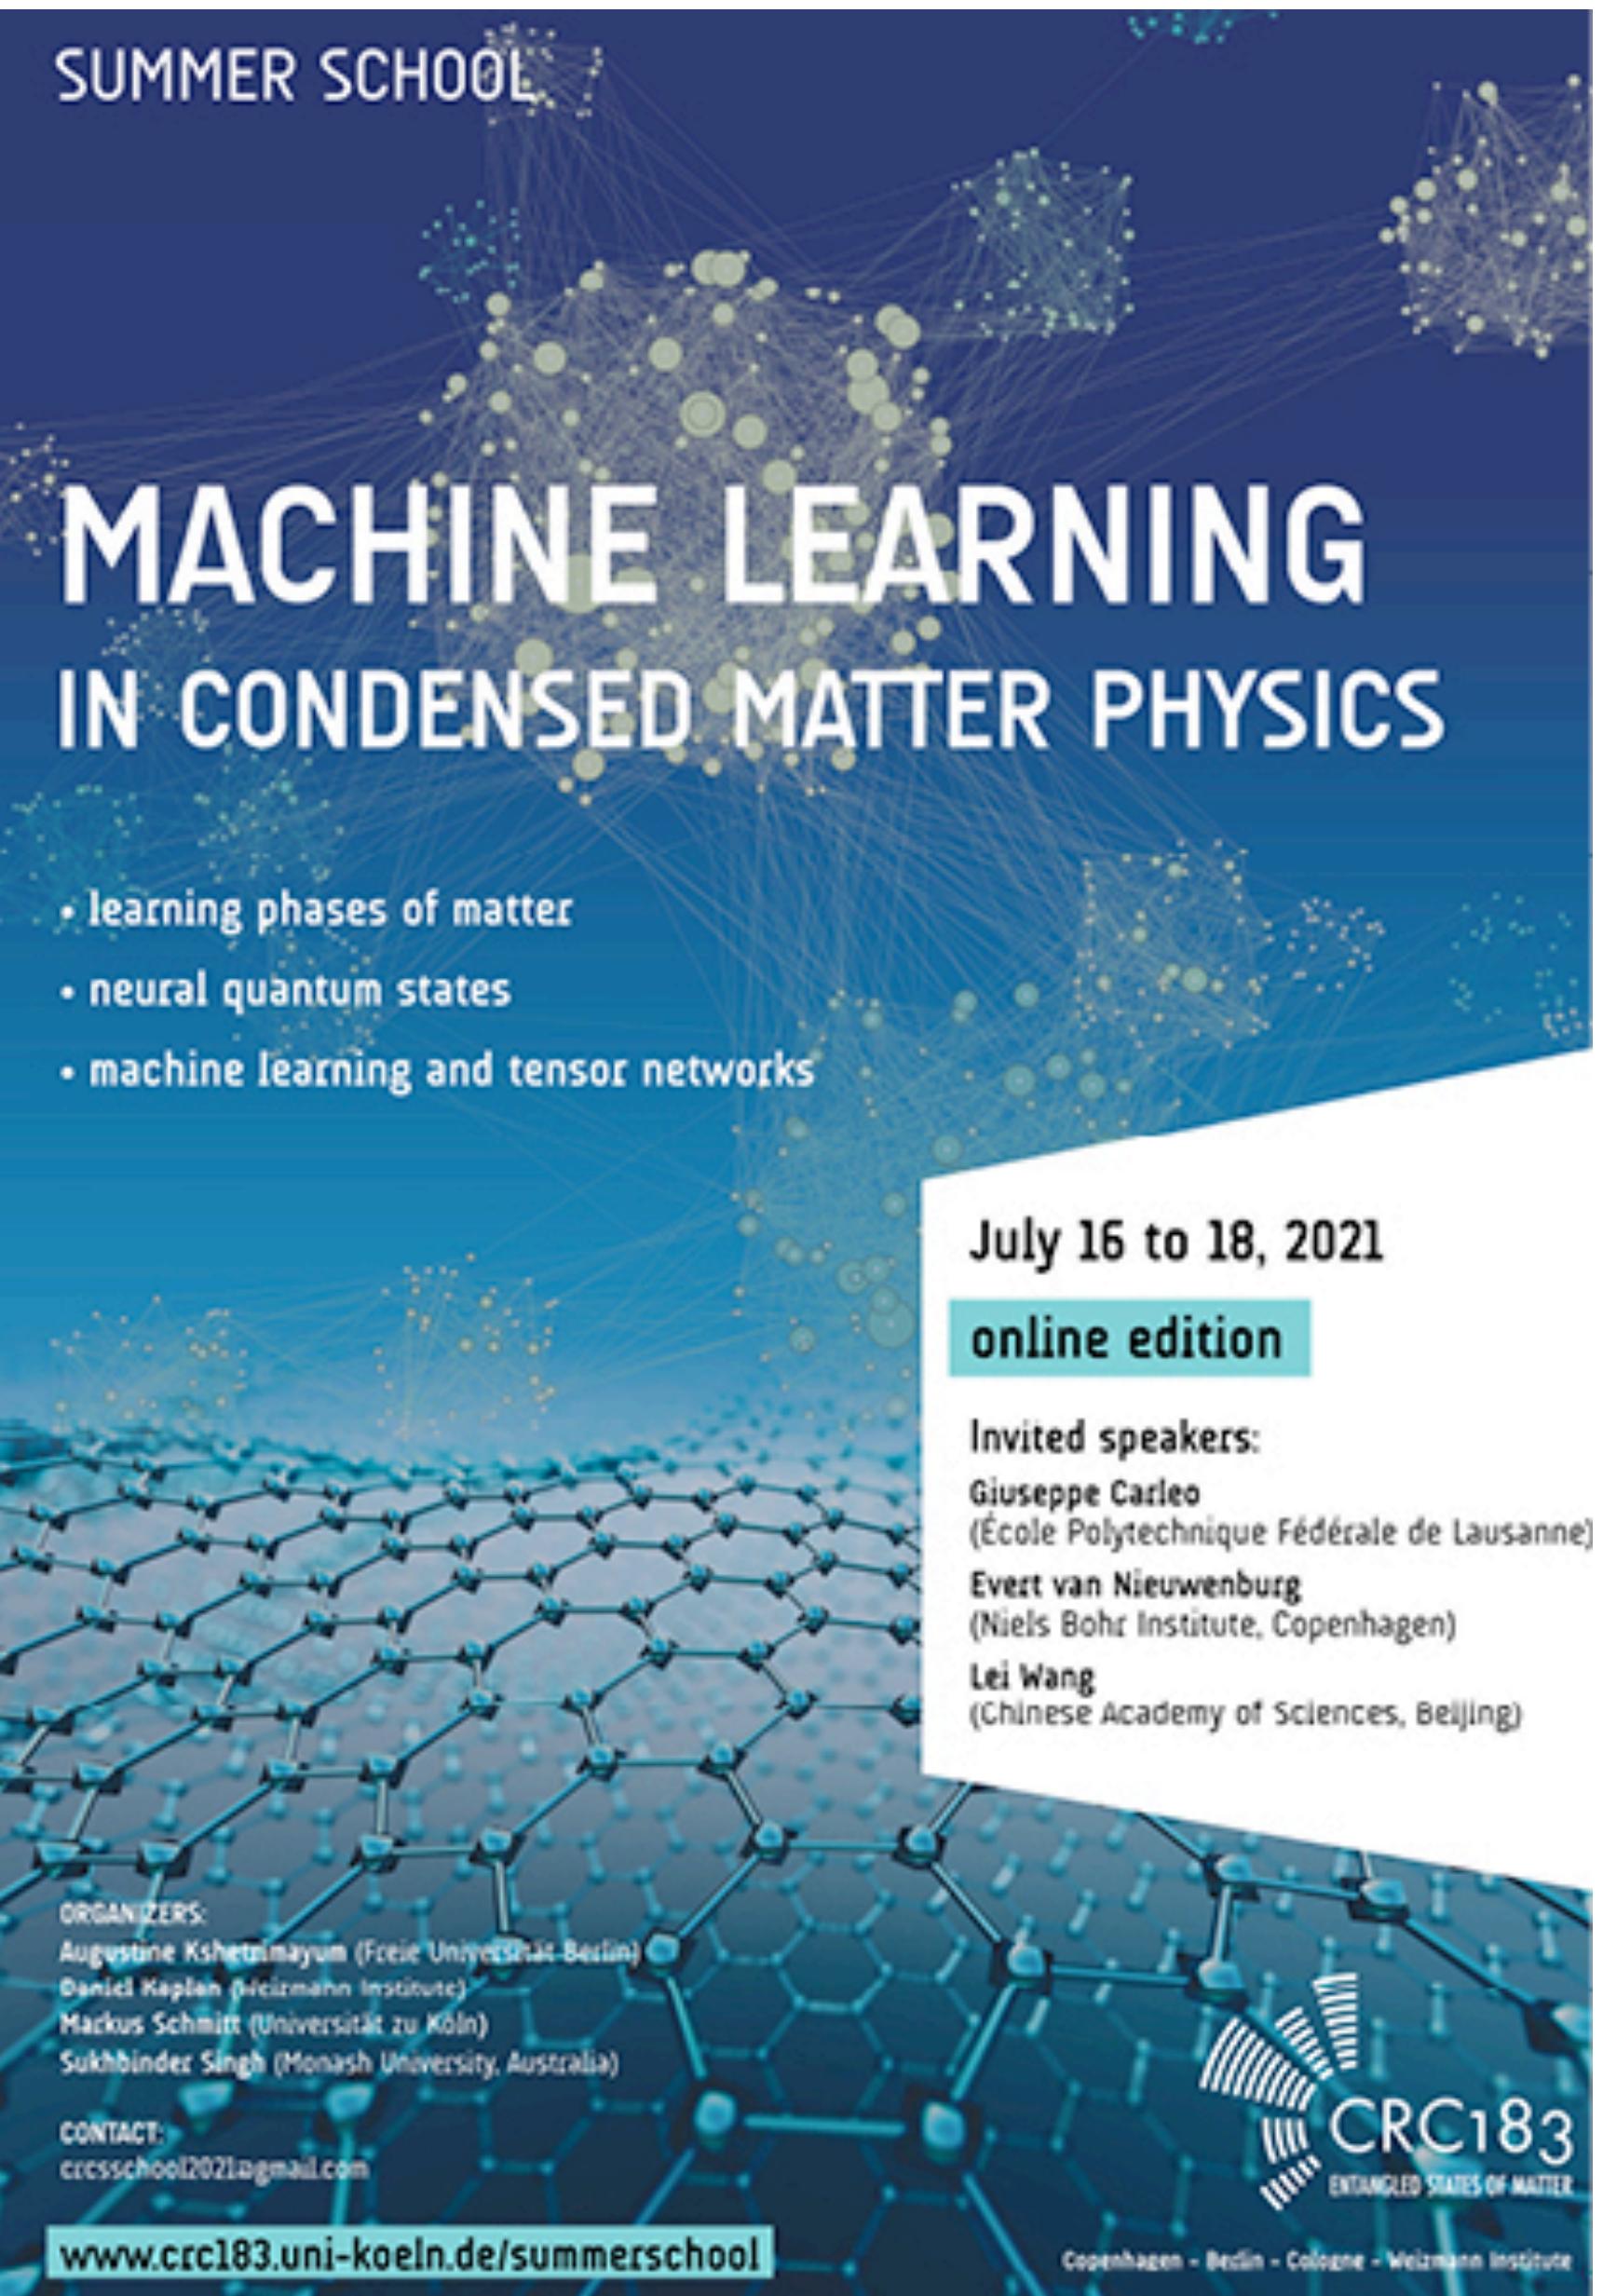


Illustration: Hans Møller, mollers.dk

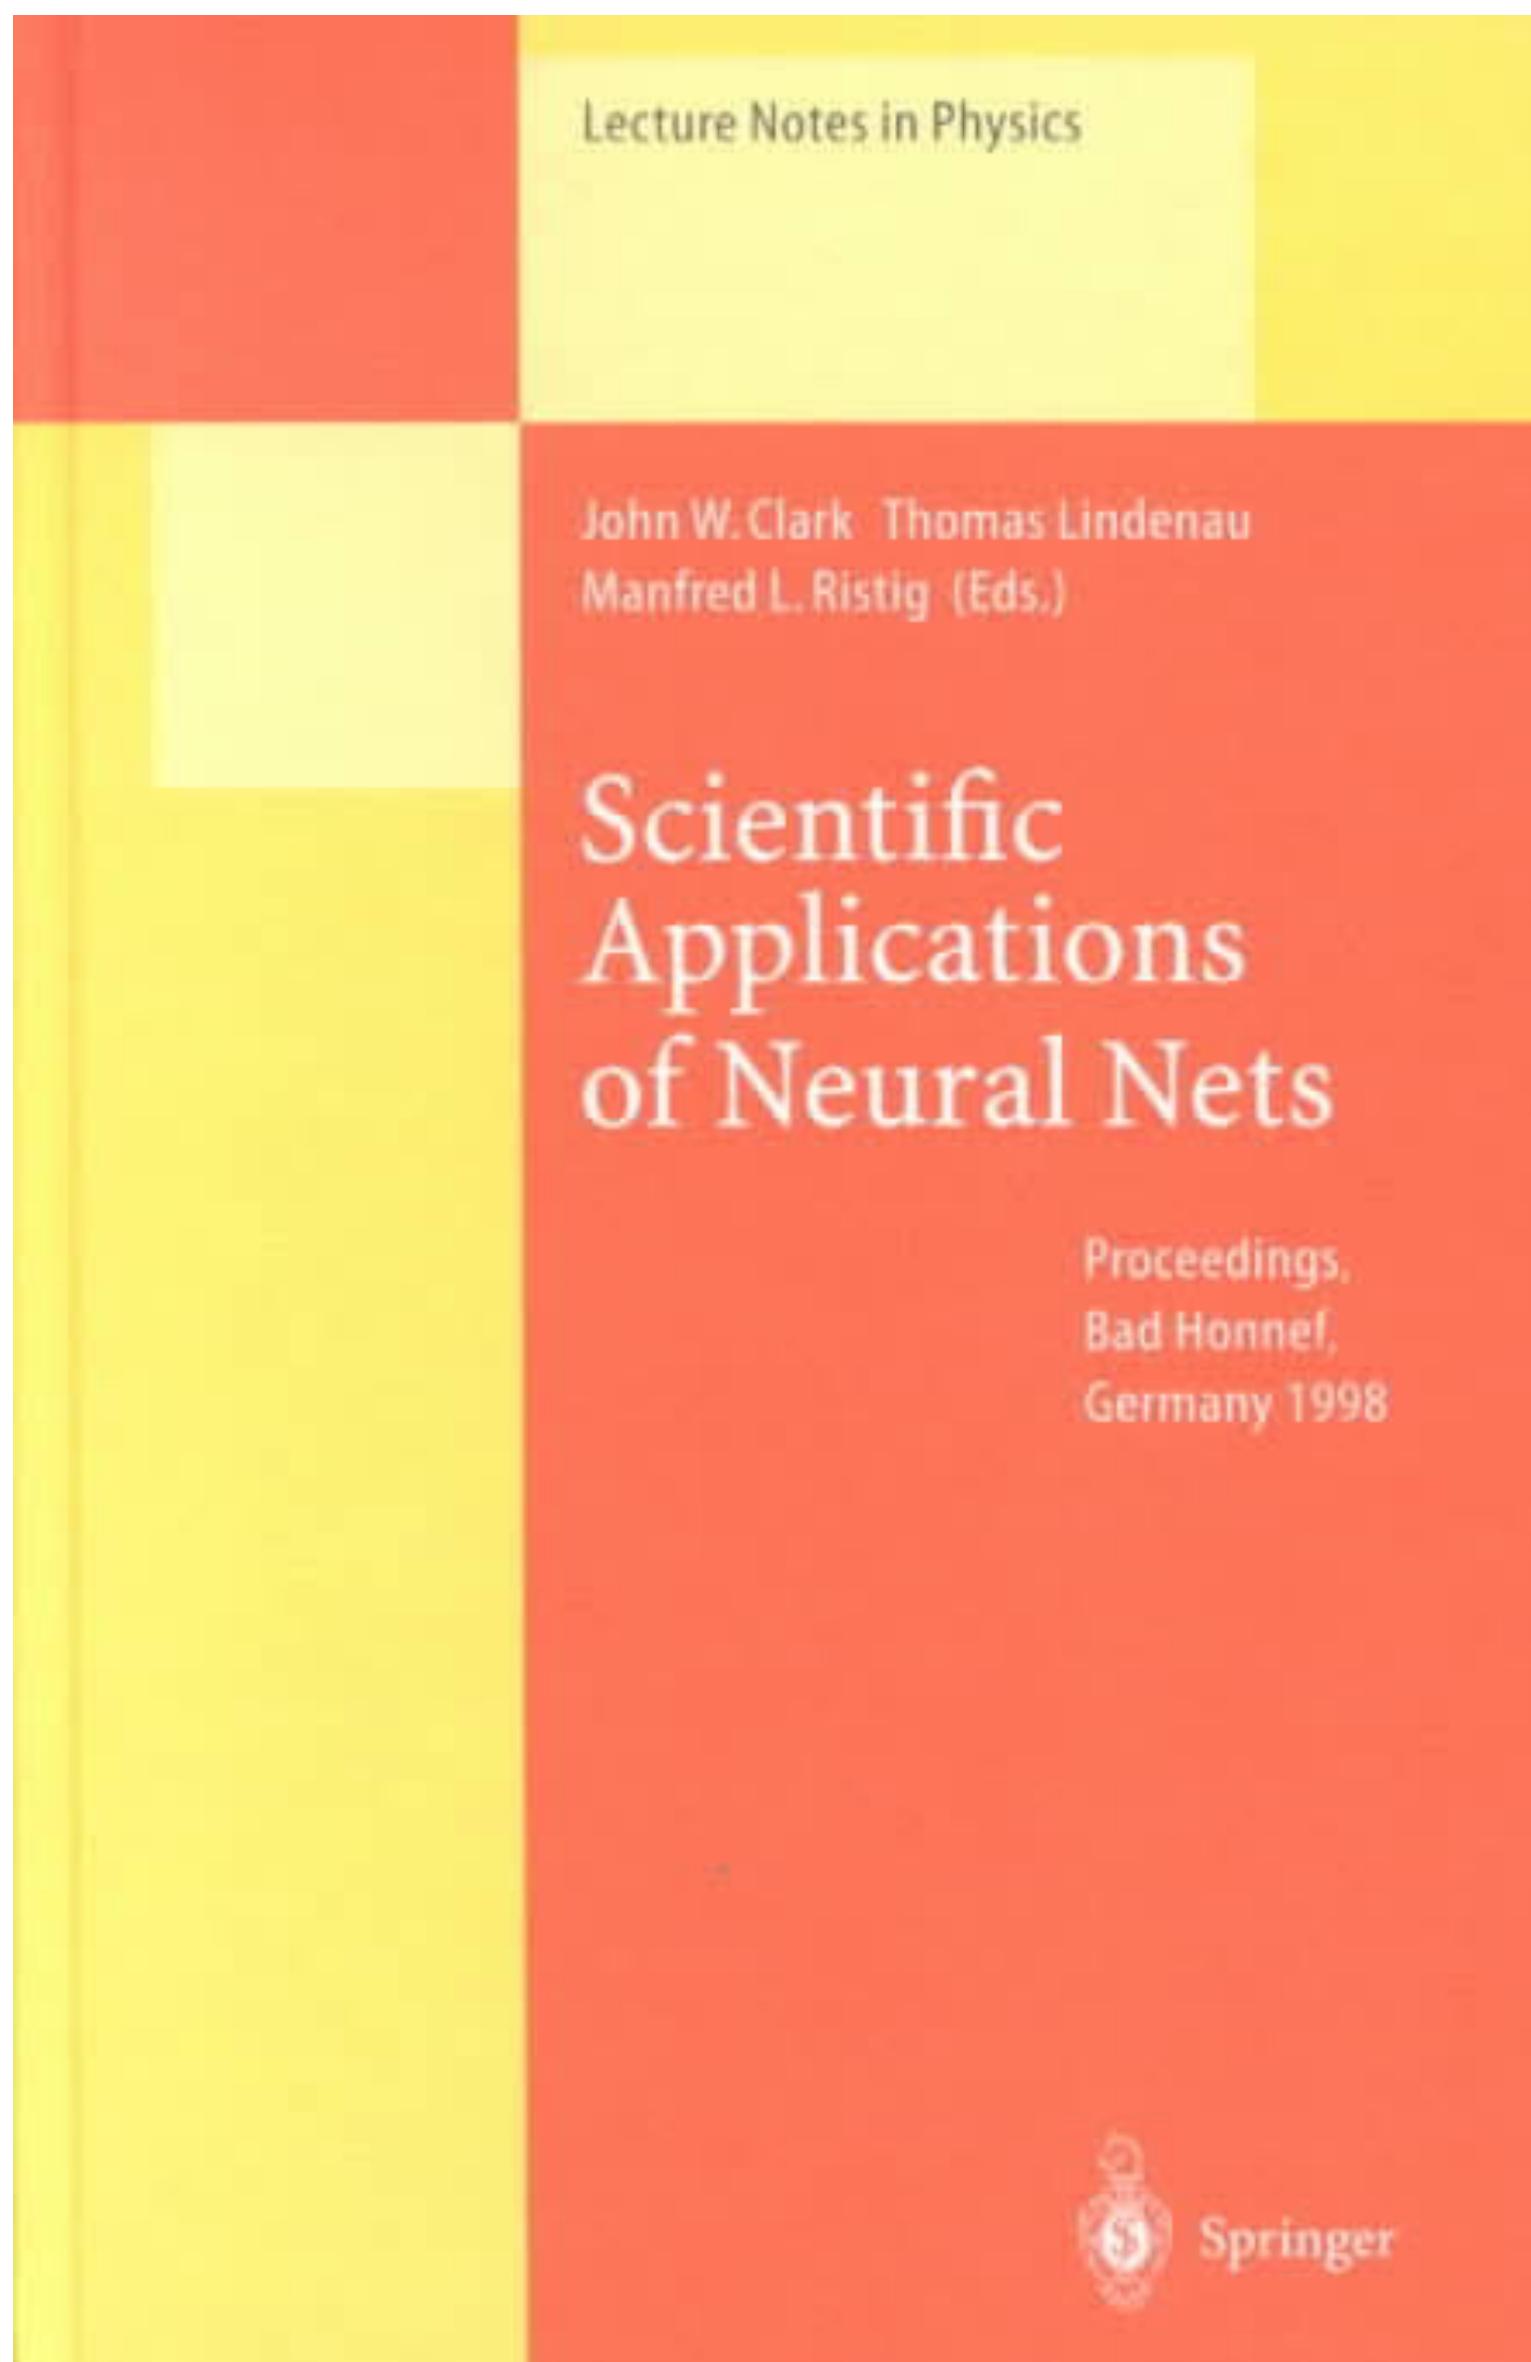


Lei Wang 1.5h x 3

1. Scientific machine learning with and without data (overview talk)
2. Generative models (mostly normalizing flows)
3. Differentiable programming (for tensor networks and quantum circuits)

Recordings <https://www.crc183.uni-koeln.de/summer-school-machine-learning/>

AI for science, 24 years ago

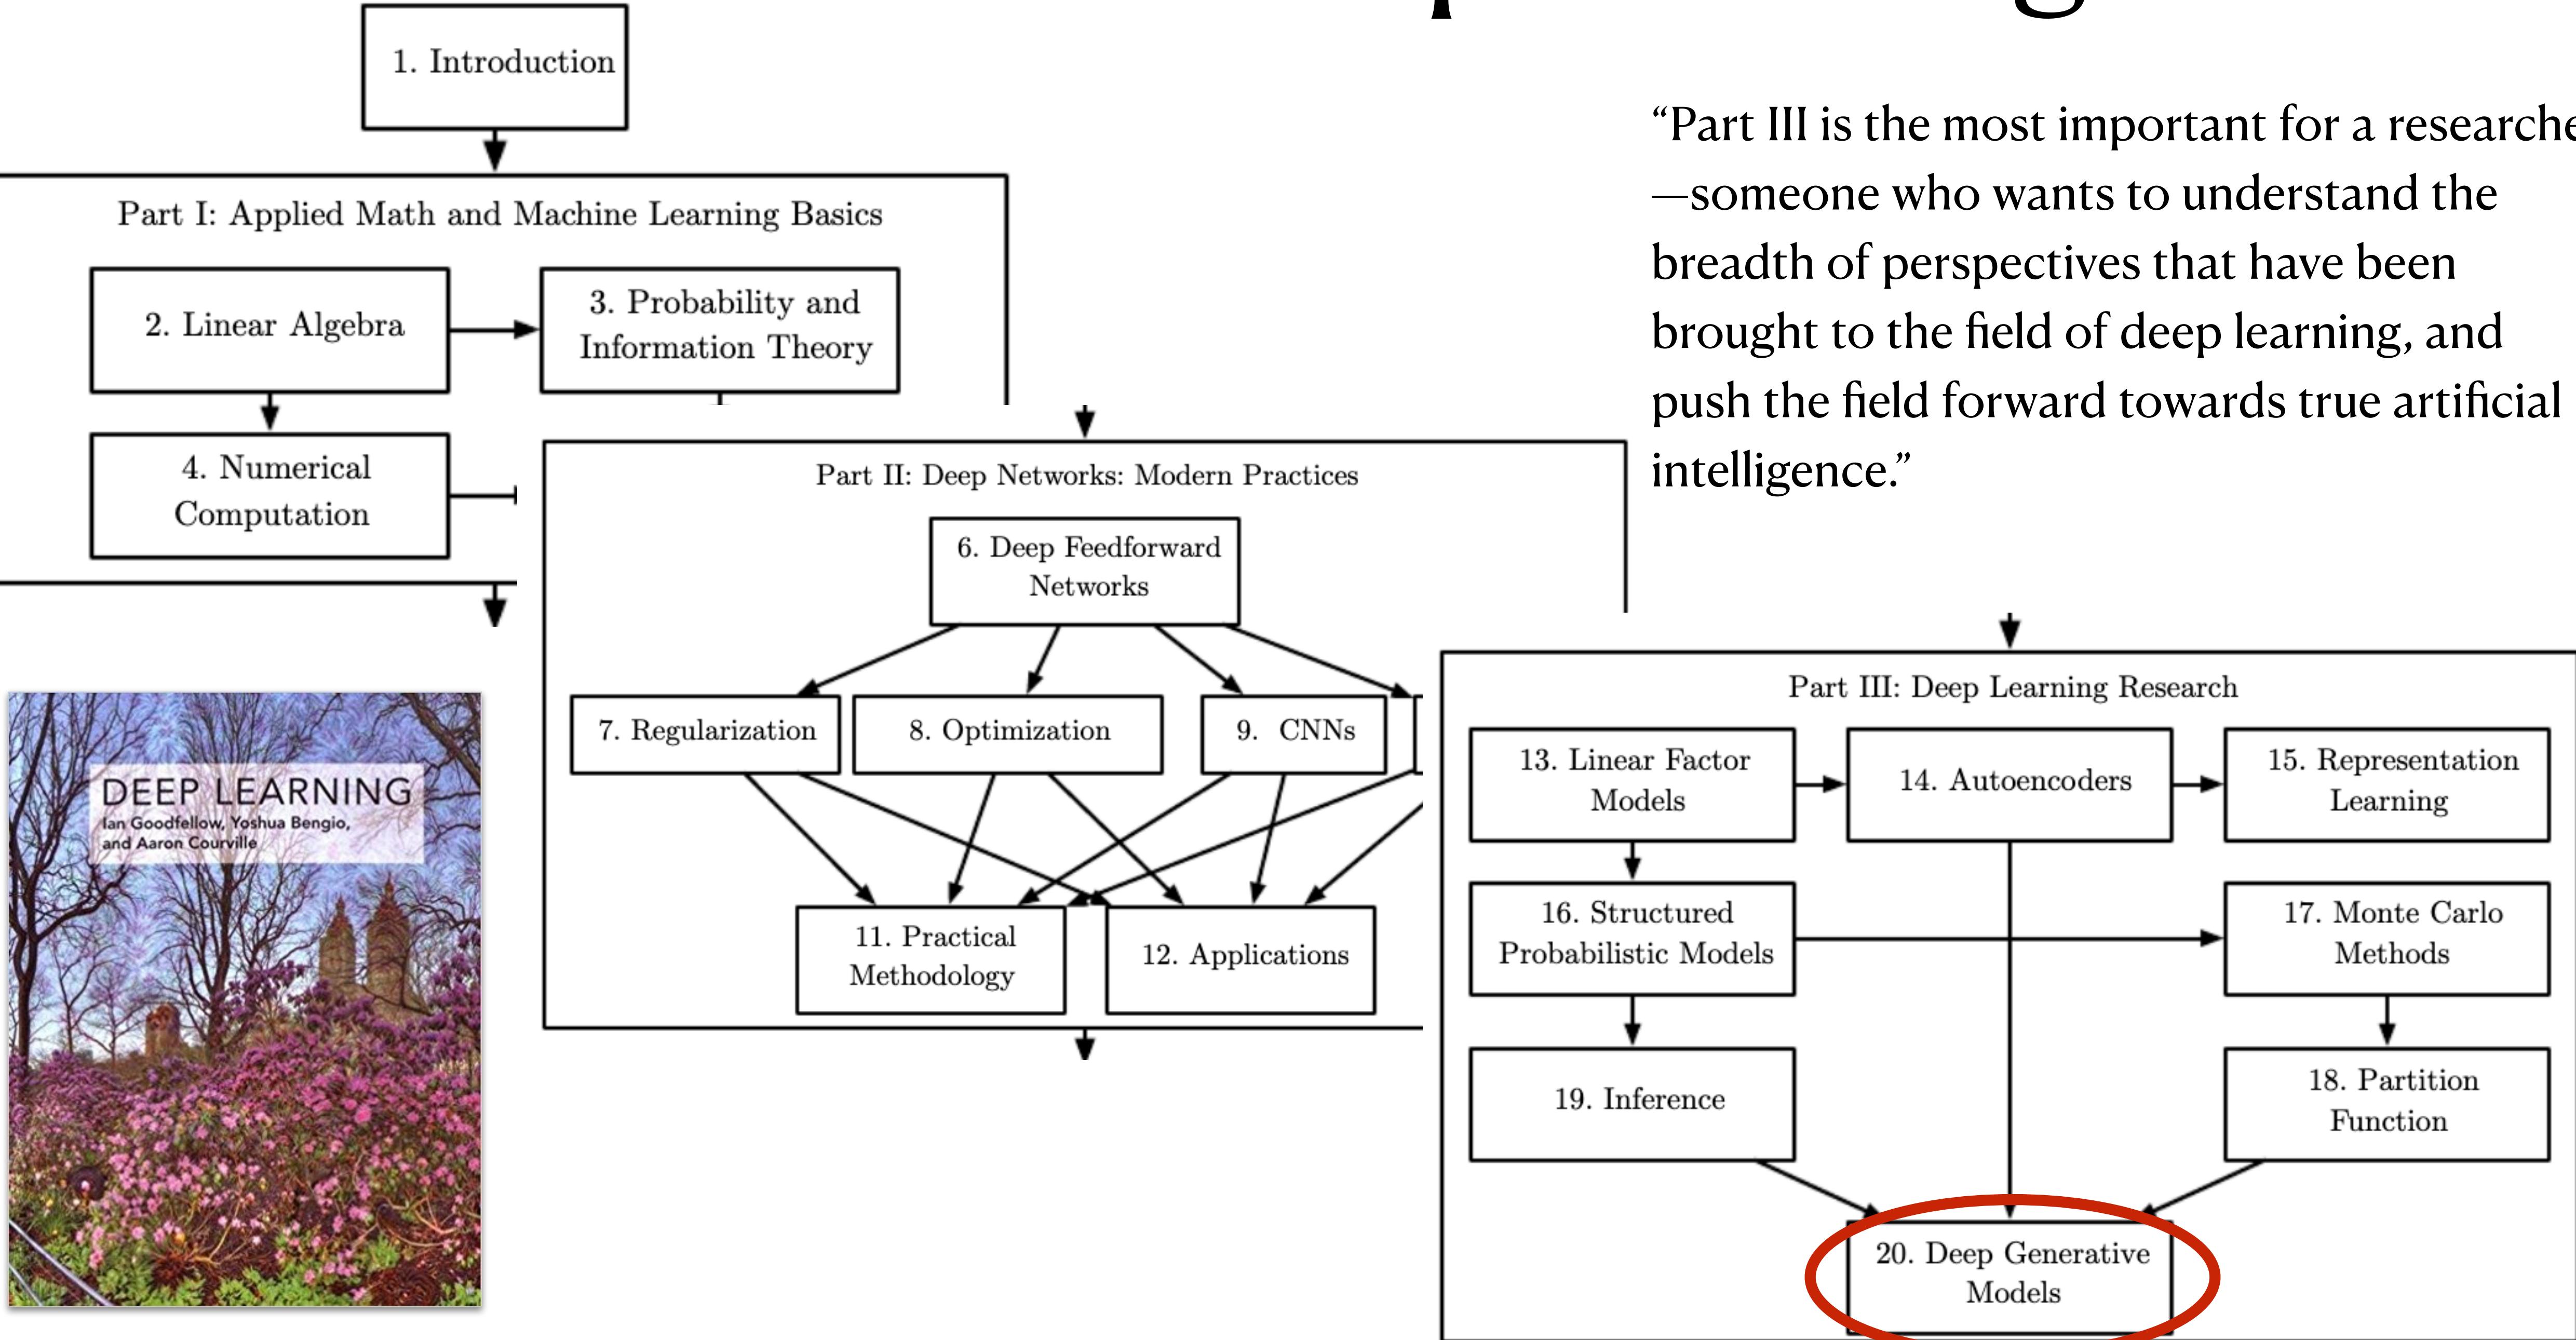


8 Doing Science With Neural Nets: Pride and Prejudice

When neural networks re-emerged on the scene in the mid-80s as a new and glamorous computational paradigm, the initial reaction in some sectors of the scientific community was perhaps too enthusiastic and not sufficiently critical. There was a tendency on the part of practitioners to oversell the powers of neural-network or “connectionist” solutions relative to conventional techniques – where conventional techniques can include both traditional theory-rich modeling and established statistical methods. The last five years have seen a correction phase, as some of the practical limitations of neural-network approaches have become apparent, and as scientists have become better acquainted with the wide array of advanced statistical tools that are currently available.

Why now, again ?
What has changed ?
What has not ?

A hint from the Deep Learning Book



Deep learning is more than fitting!

Discriminative learning



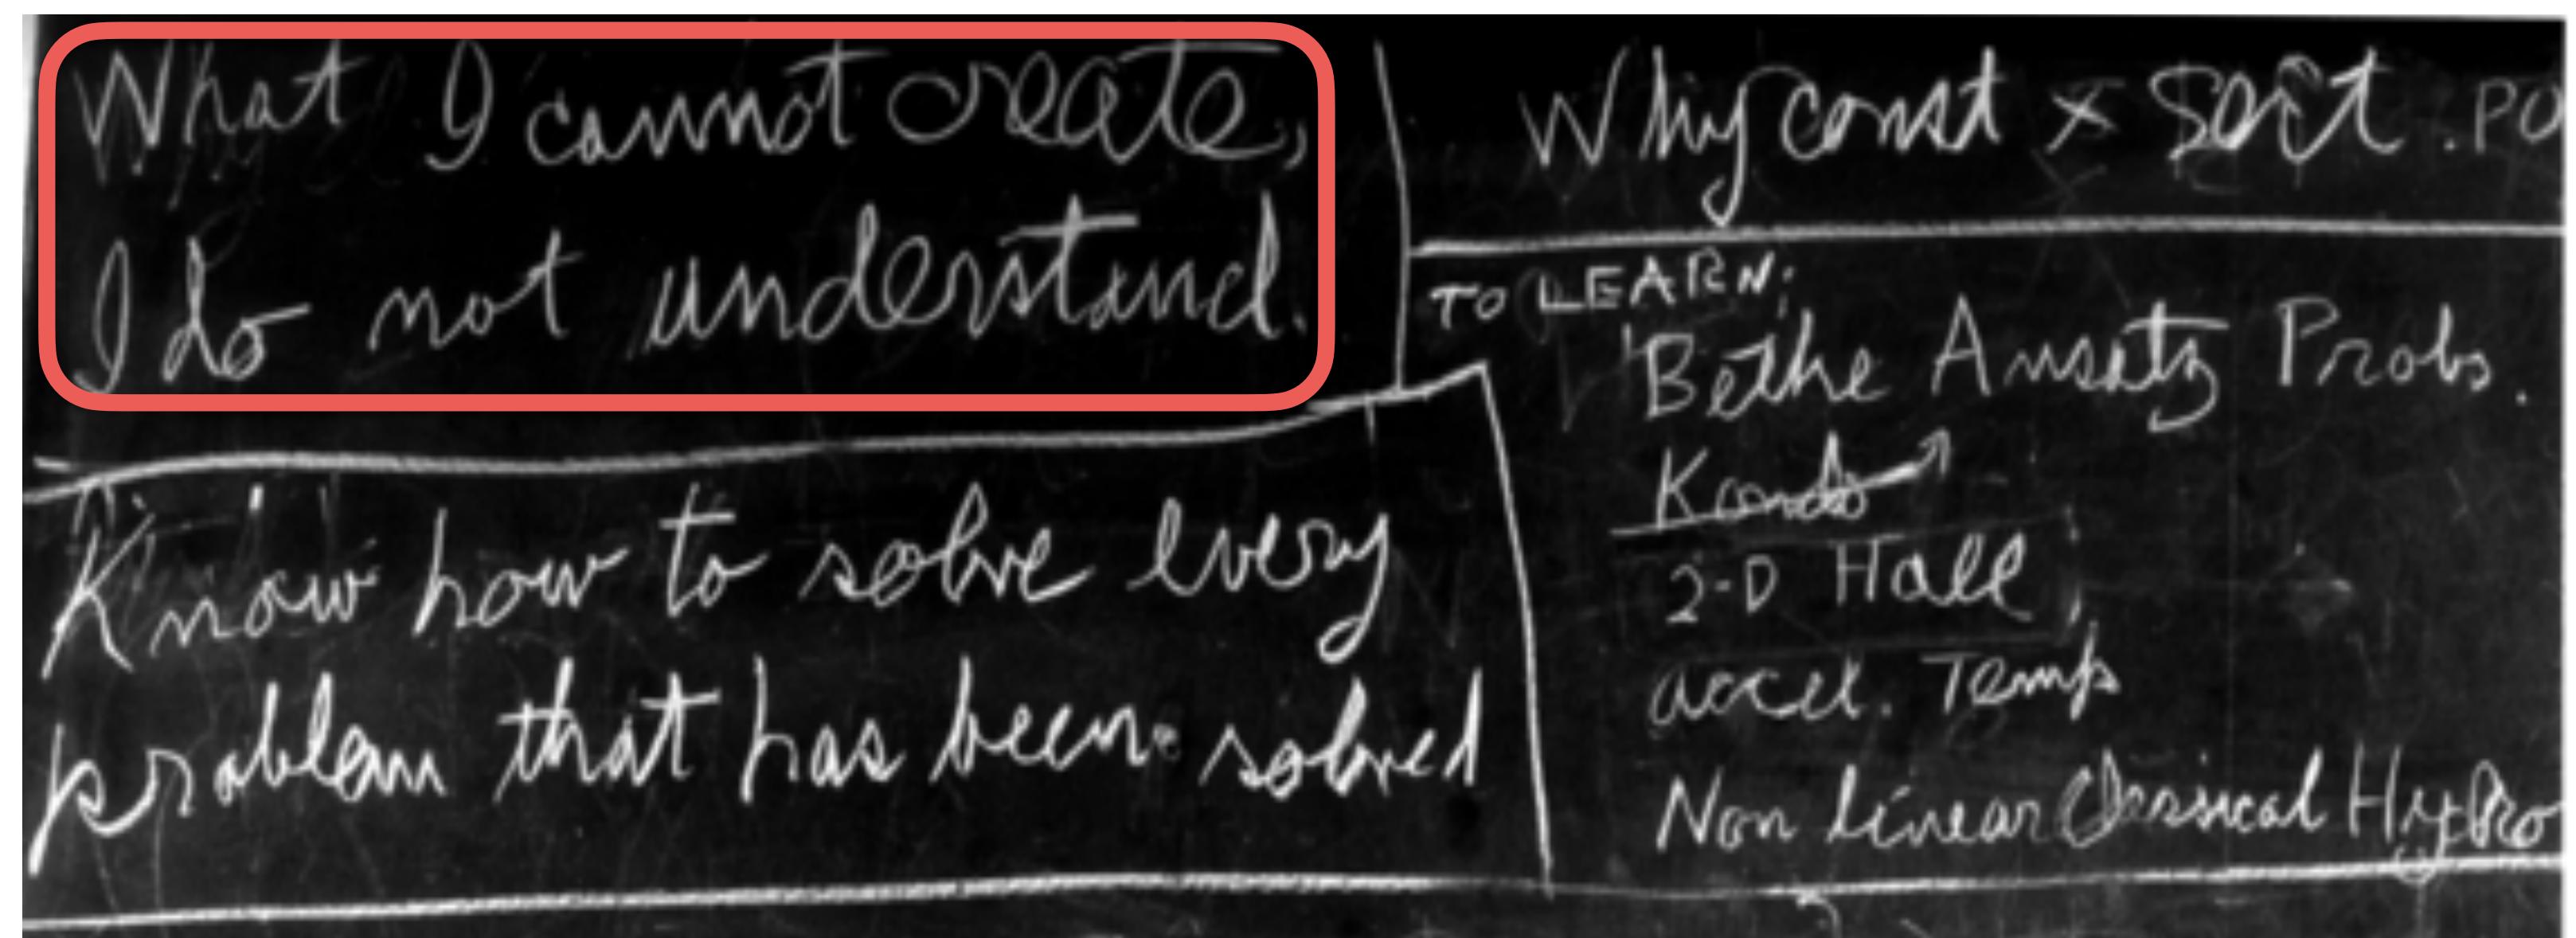
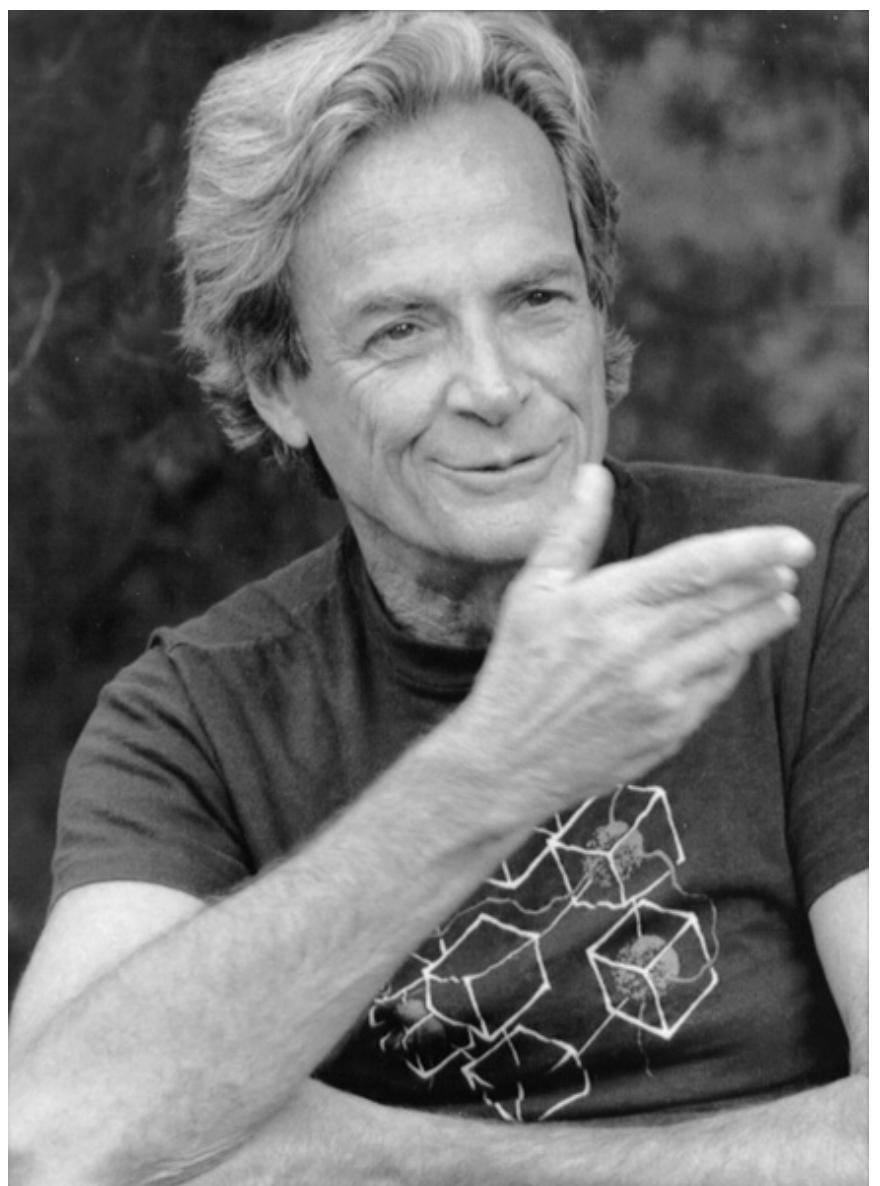
Generative learning



$$y = f(x)$$

or $p(y | x)$

$$p(x, y)$$

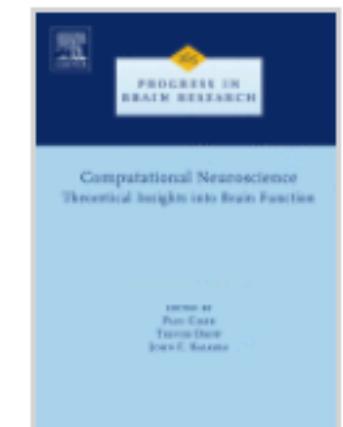


ELSEVIER

Progress in Brain Research

Volume 165, 2007, Pages 535–547

Computational Neuroscience: Theoretical Insights into Brain Function

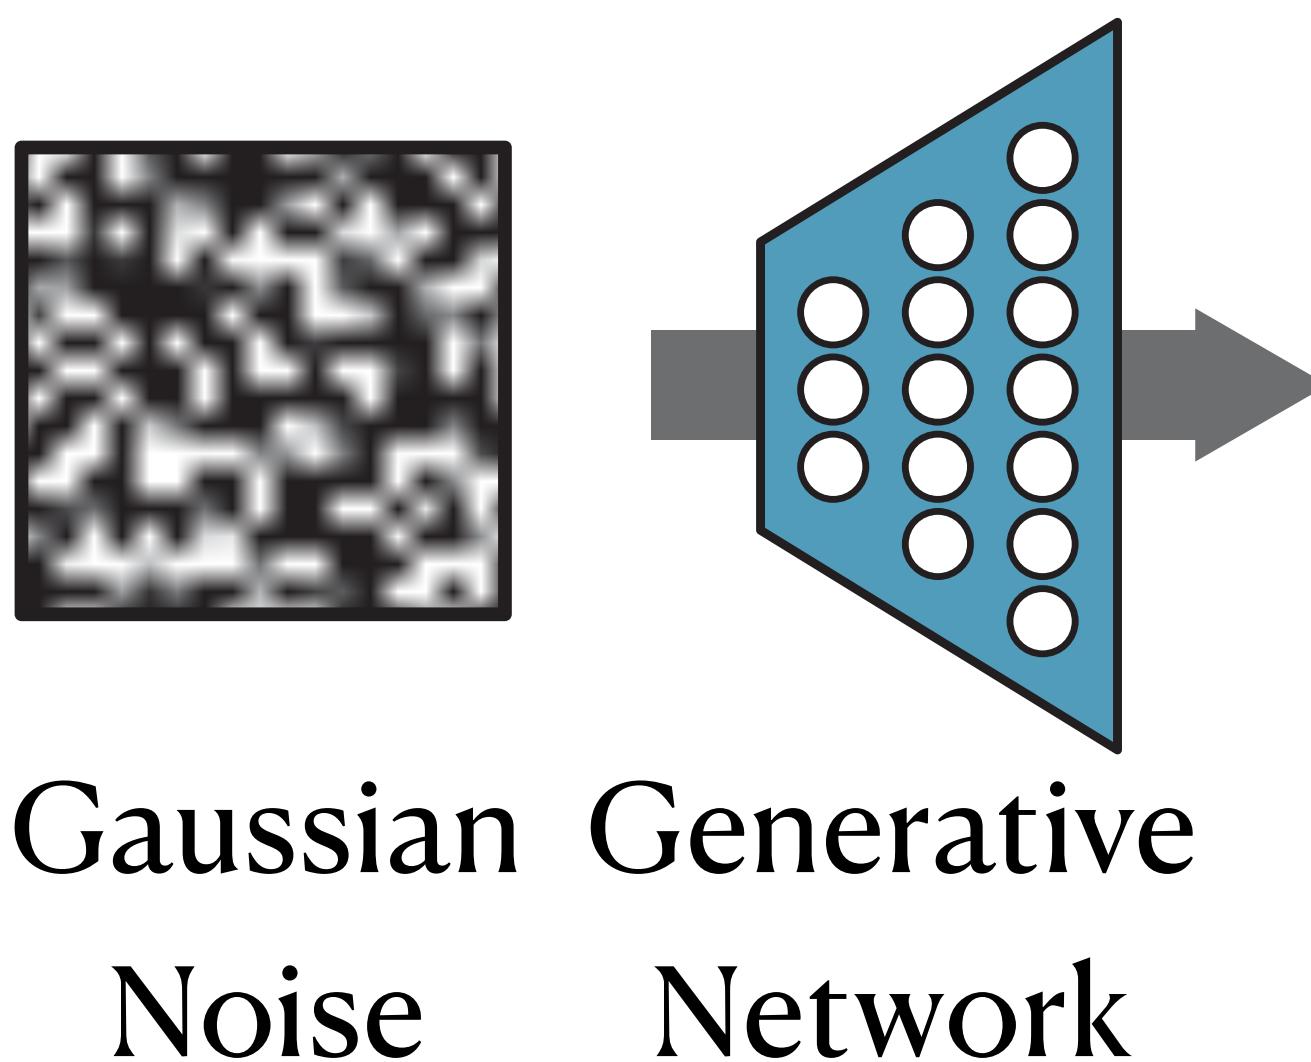


To recognize shapes, first learn to generate images

Geoffrey E. Hinton

Department of Computer Science, University of Toronto, 10 Kings College Road, Toronto, M5S 3G4
Canada

Generated arts

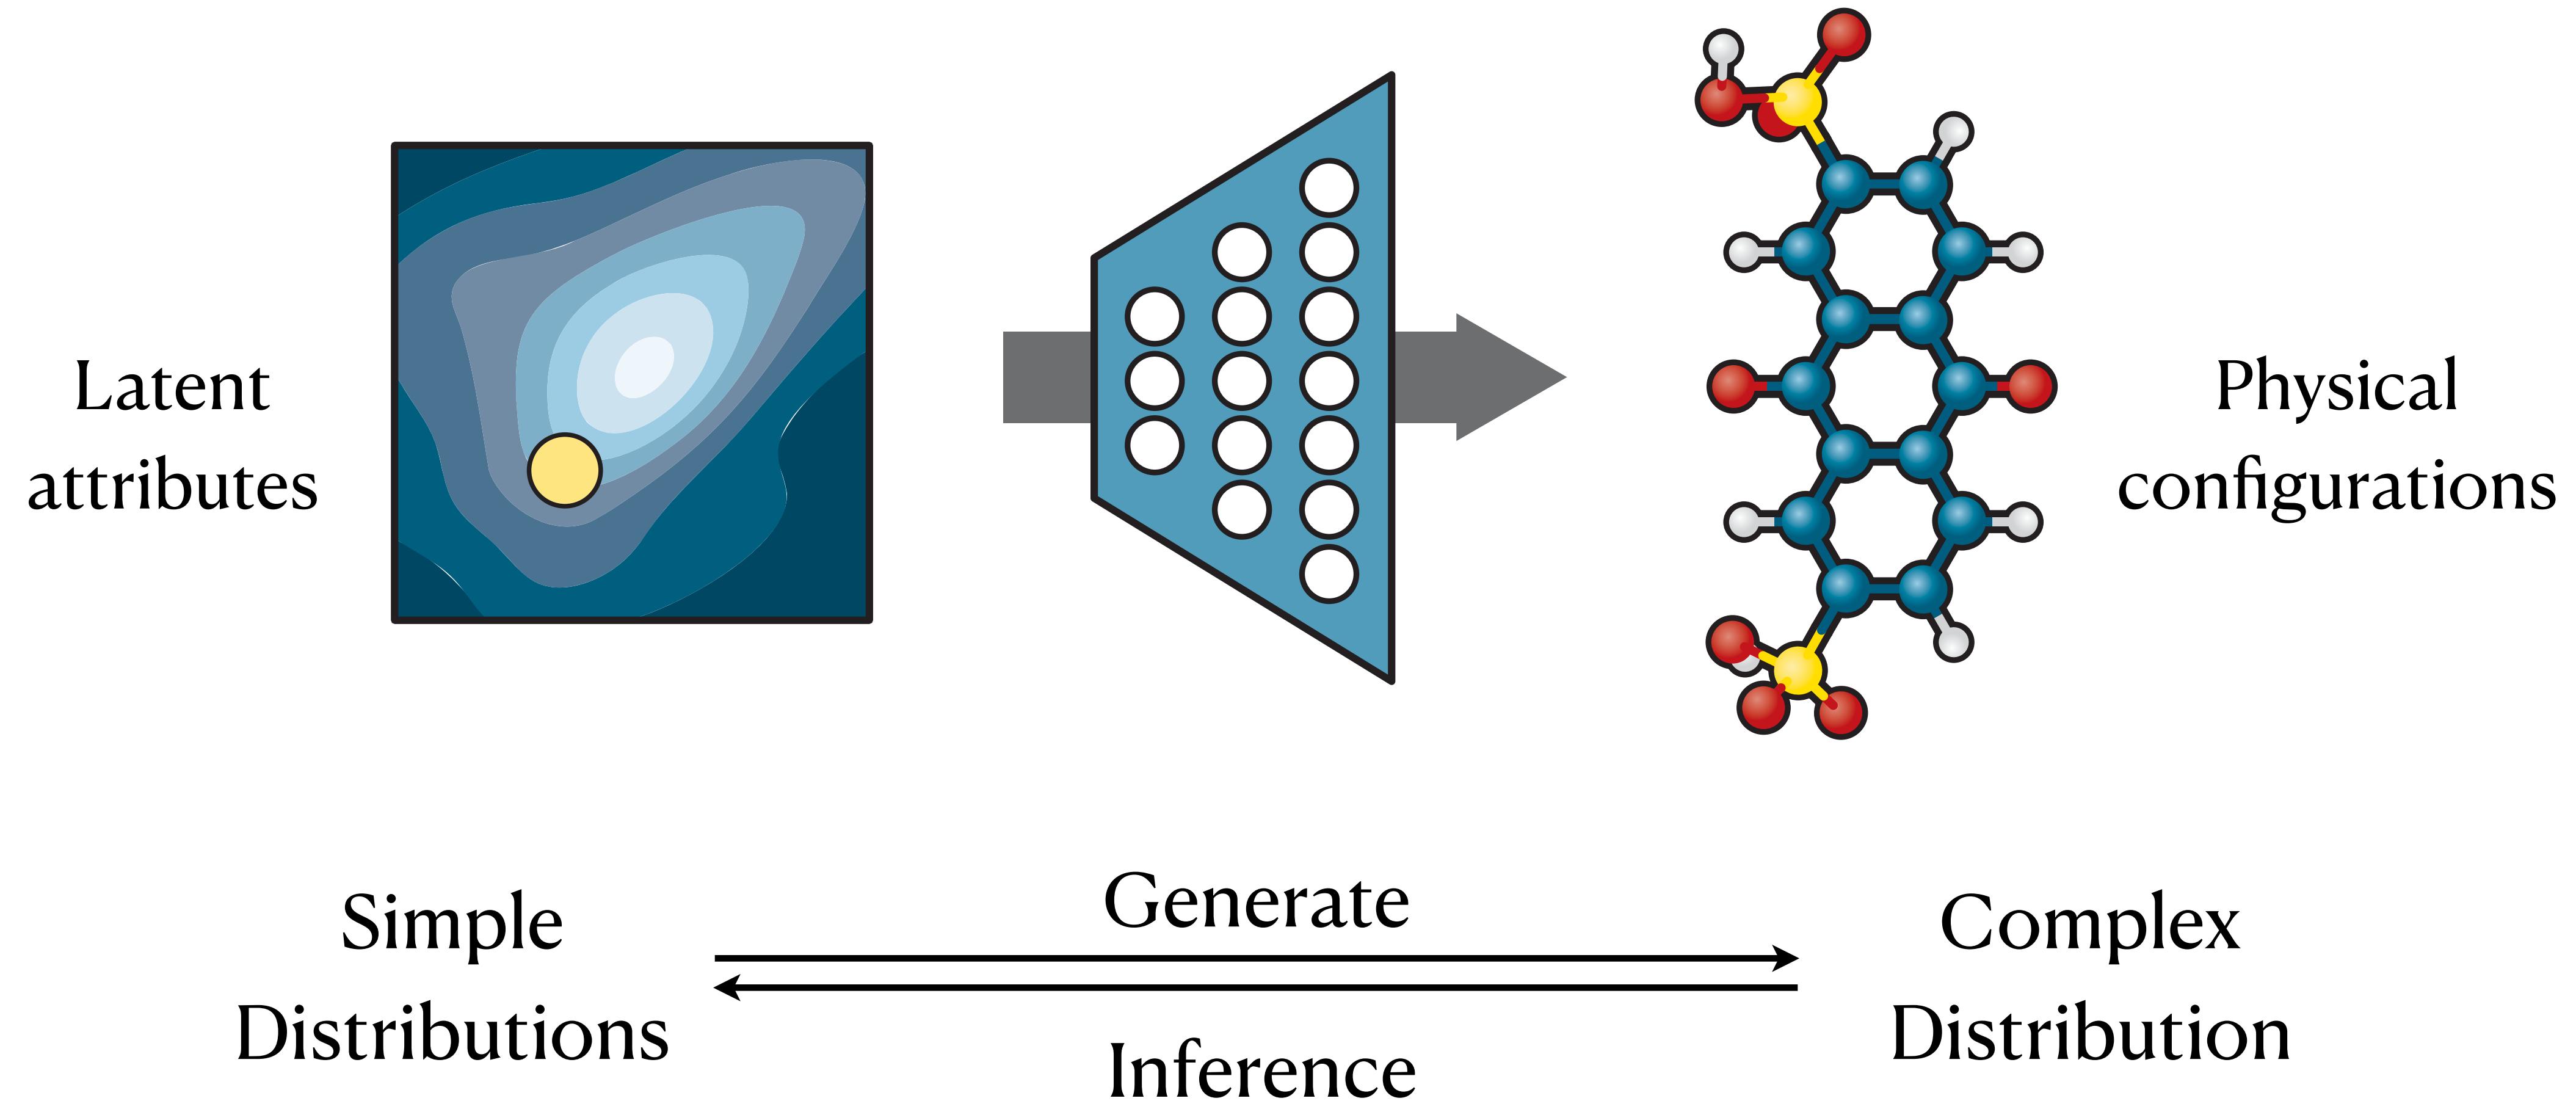


A painting of a man in a green jacket, framed in gold, with a mathematical equation overlaid. The equation is:

$$\min_{\mathcal{G}} \max_{\mathcal{D}} \mathbb{E}_{\mathbf{x}} [\log(\mathcal{D}(\mathbf{x}))] + \mathbb{E}_{\mathbf{z}} [\log(1 - \mathcal{D}(\mathcal{G}(\mathbf{z})))]$$

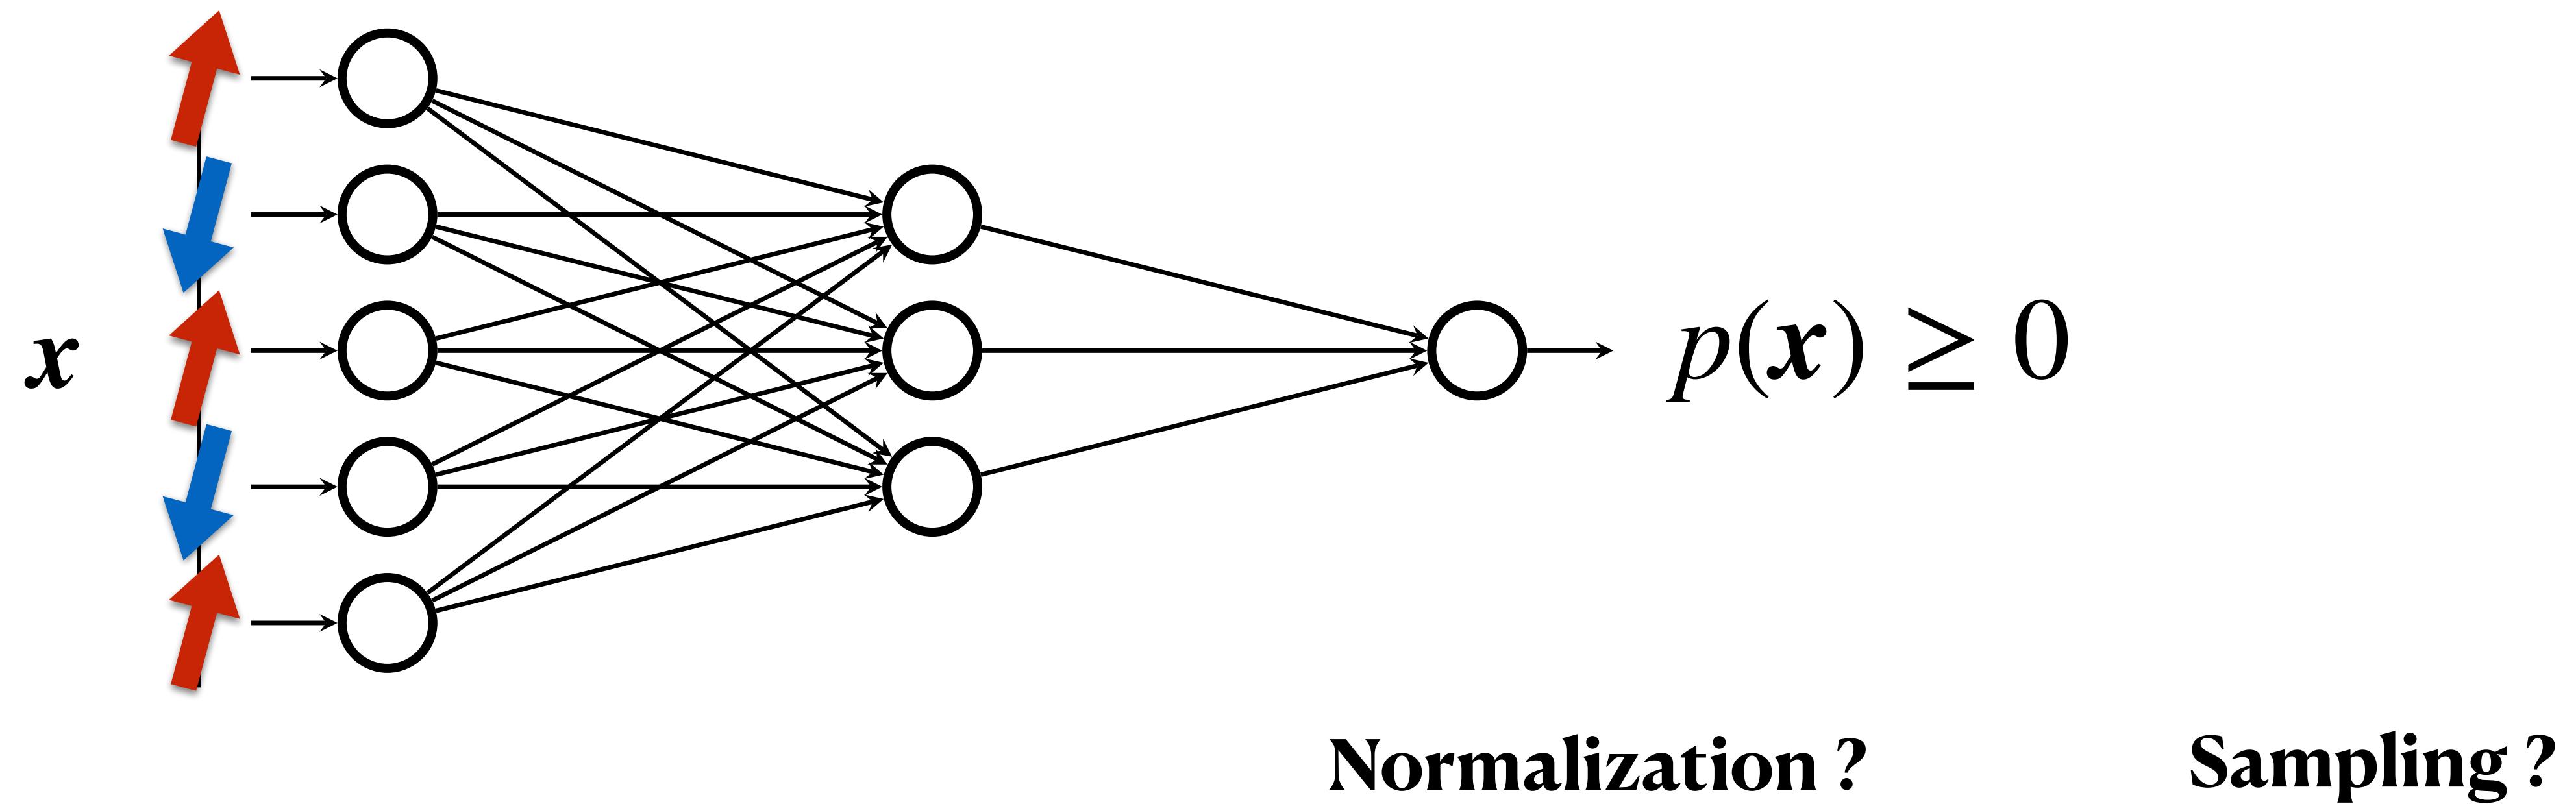
\$432,500
25 October 2018
Christie's New York

Generating molecules



Sanchez-Lengeling & Aspuru-Guzik,
Inverse molecular design using machine learning:
Generative models for matter engineering, Science '18

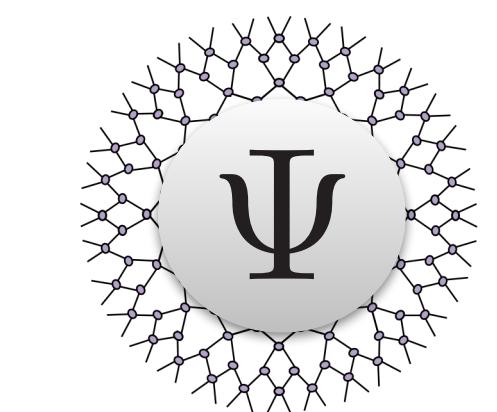
So, what is the fuss ?



$$\int dx p(x) = 1$$

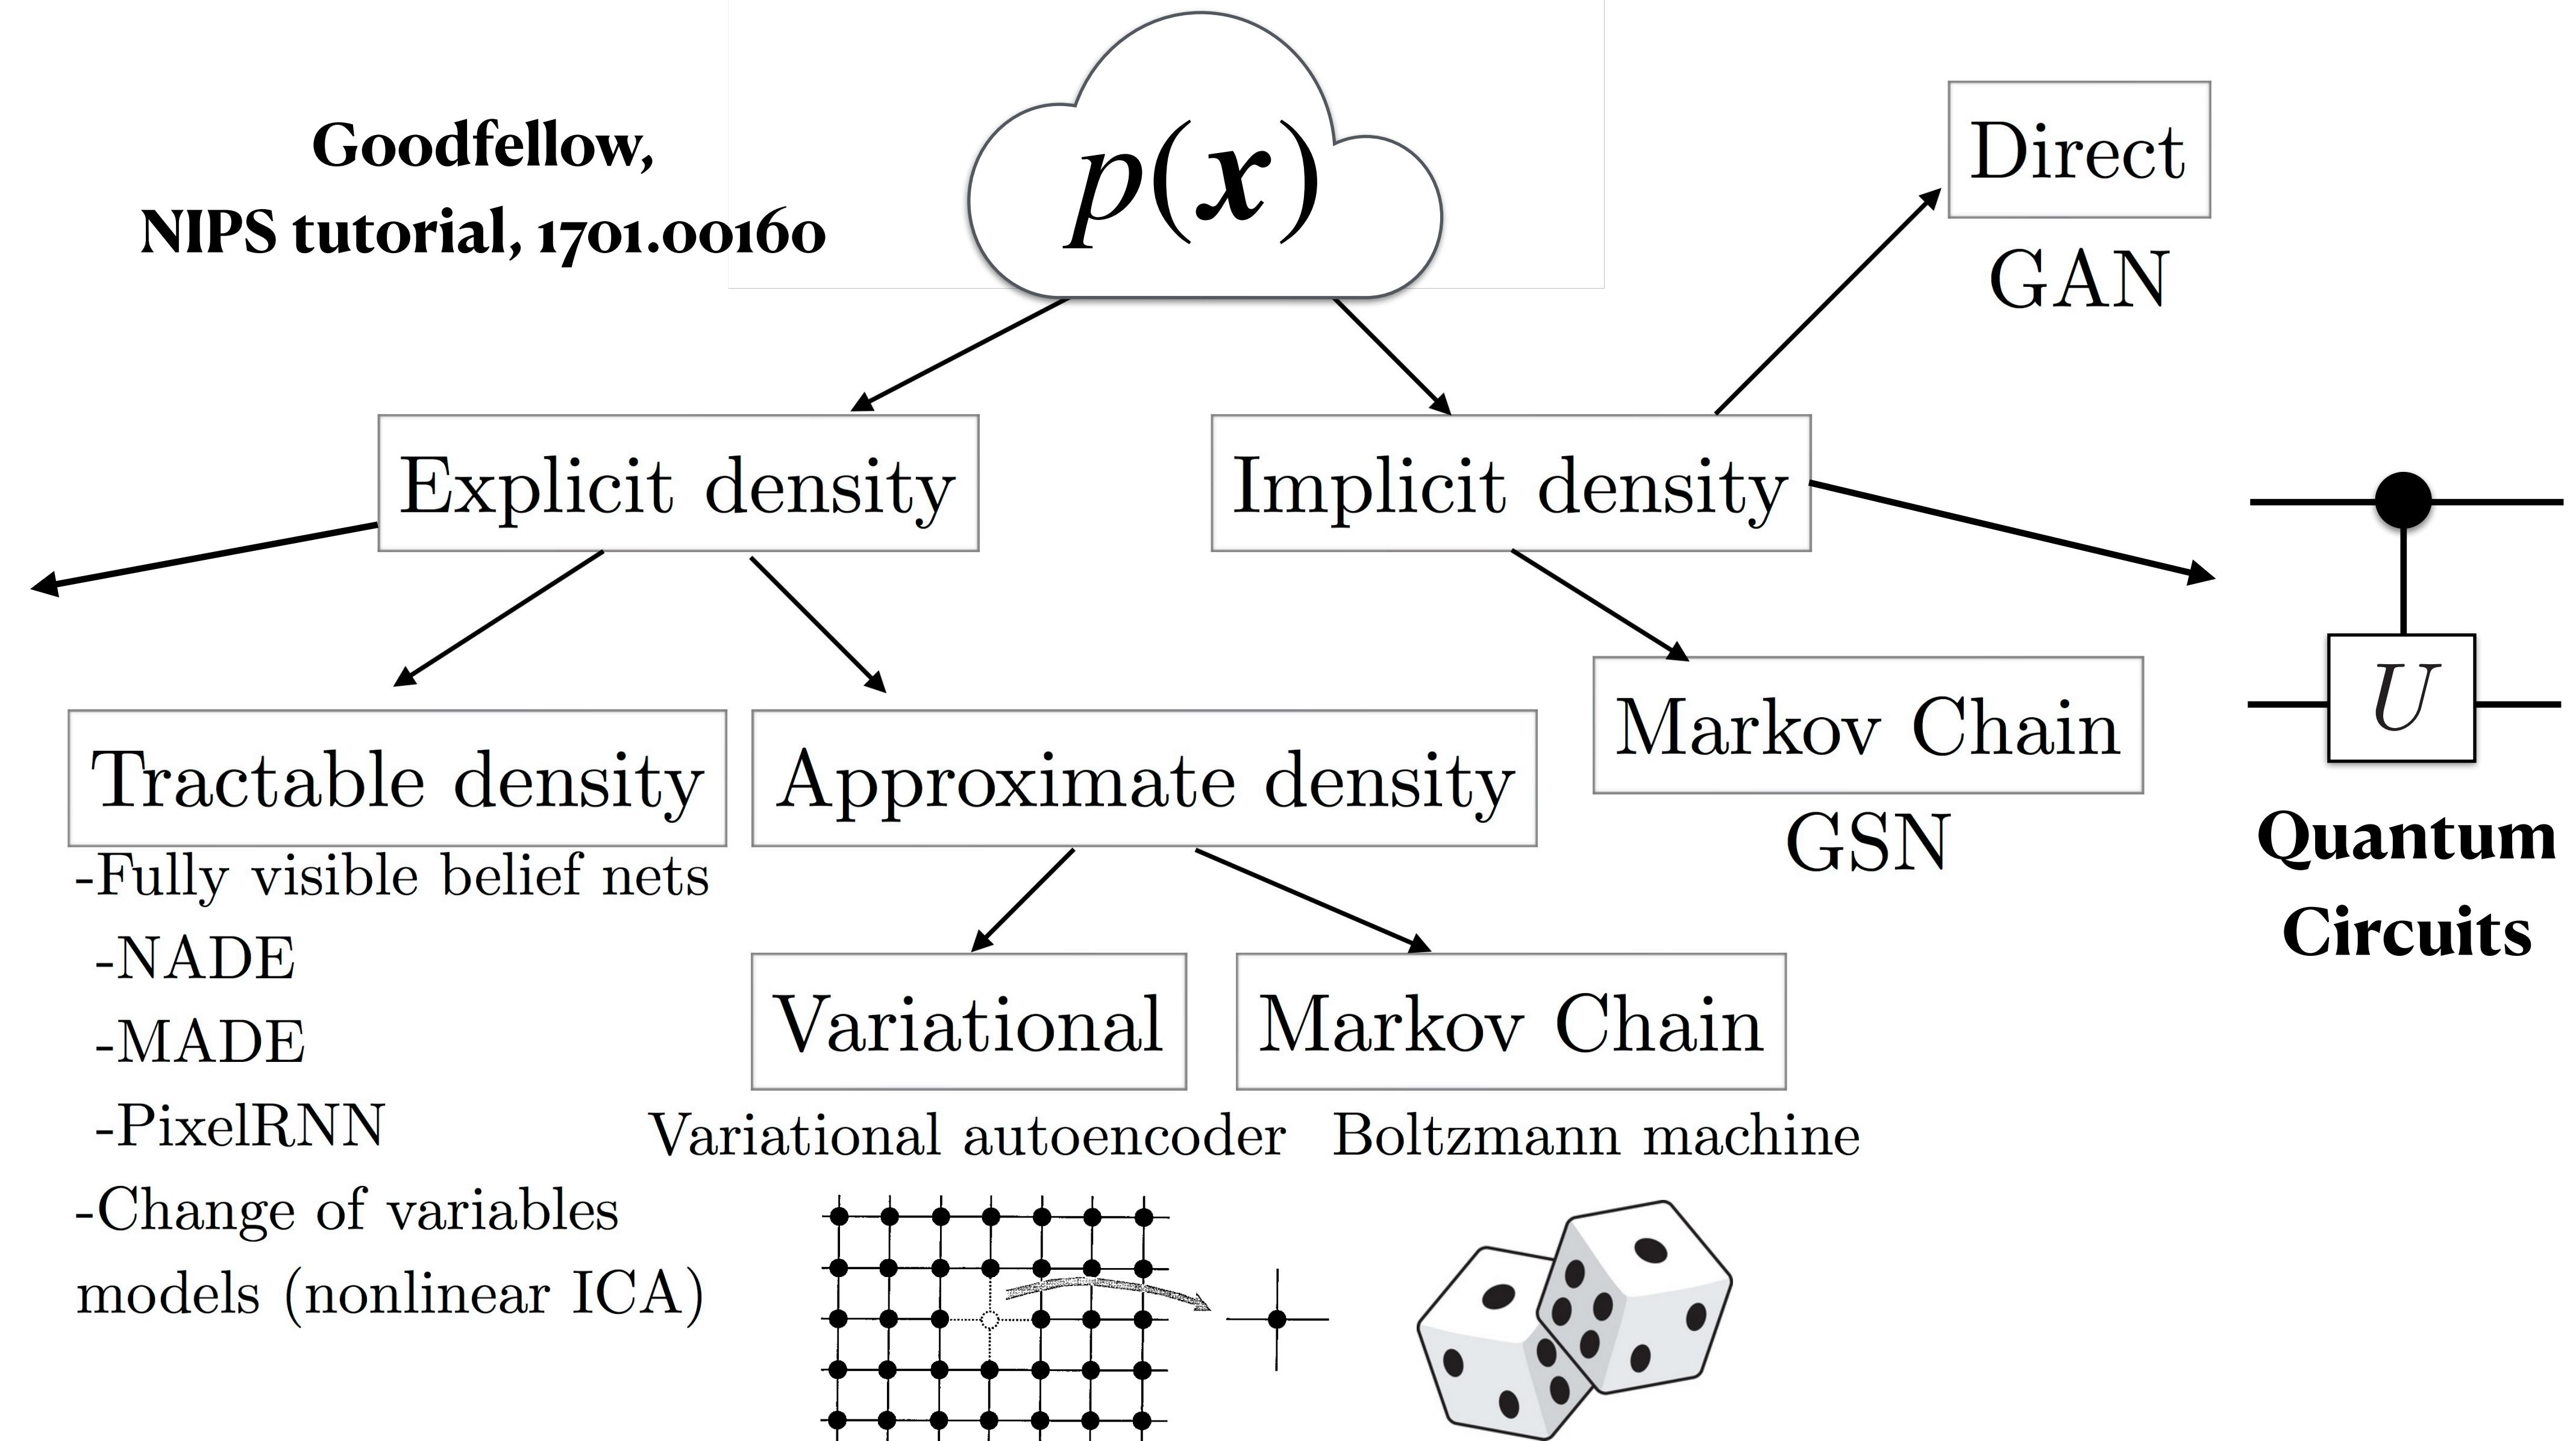
$$\mathbb{E}_{x \sim p(x)}$$

Generative models and their physics genes



**Tensor
Networks**

**Goodfellow,
NIPS tutorial, 1701.00160**



Lecture Note <http://wangleiphy.github.io/lectures/PILtutorial.pdf>

Generative Models for Physicists

Lei Wang*

Institute of Physics, Chinese Academy of Sciences
Beijing 100190, China

October 28, 2018

Abstract

Generative models generate unseen samples according to a learned joint probability distribution in the high-dimensional space. They find wide applications in density estimation, variational inference, representation learning and more. Deep generative models and associated techniques (such as differentiable programming and representation learning) are cutting-edge technologies physicists can learn from deep learning.

This note introduces the concept and principles of generative modeling, together with applications of modern generative models (autoregressive models, normalizing flows, variational autoencoders etc) as well as the old ones (Boltzmann machines) to physics problems. As a bonus, this note puts some emphasize on physics-inspired generative models which take insights from statistical, quantum, and fluid mechanics.

The latest version of the note is at <http://wangleiphy.github.io/>. Please send comments, suggestions and corrections to the email address in below.

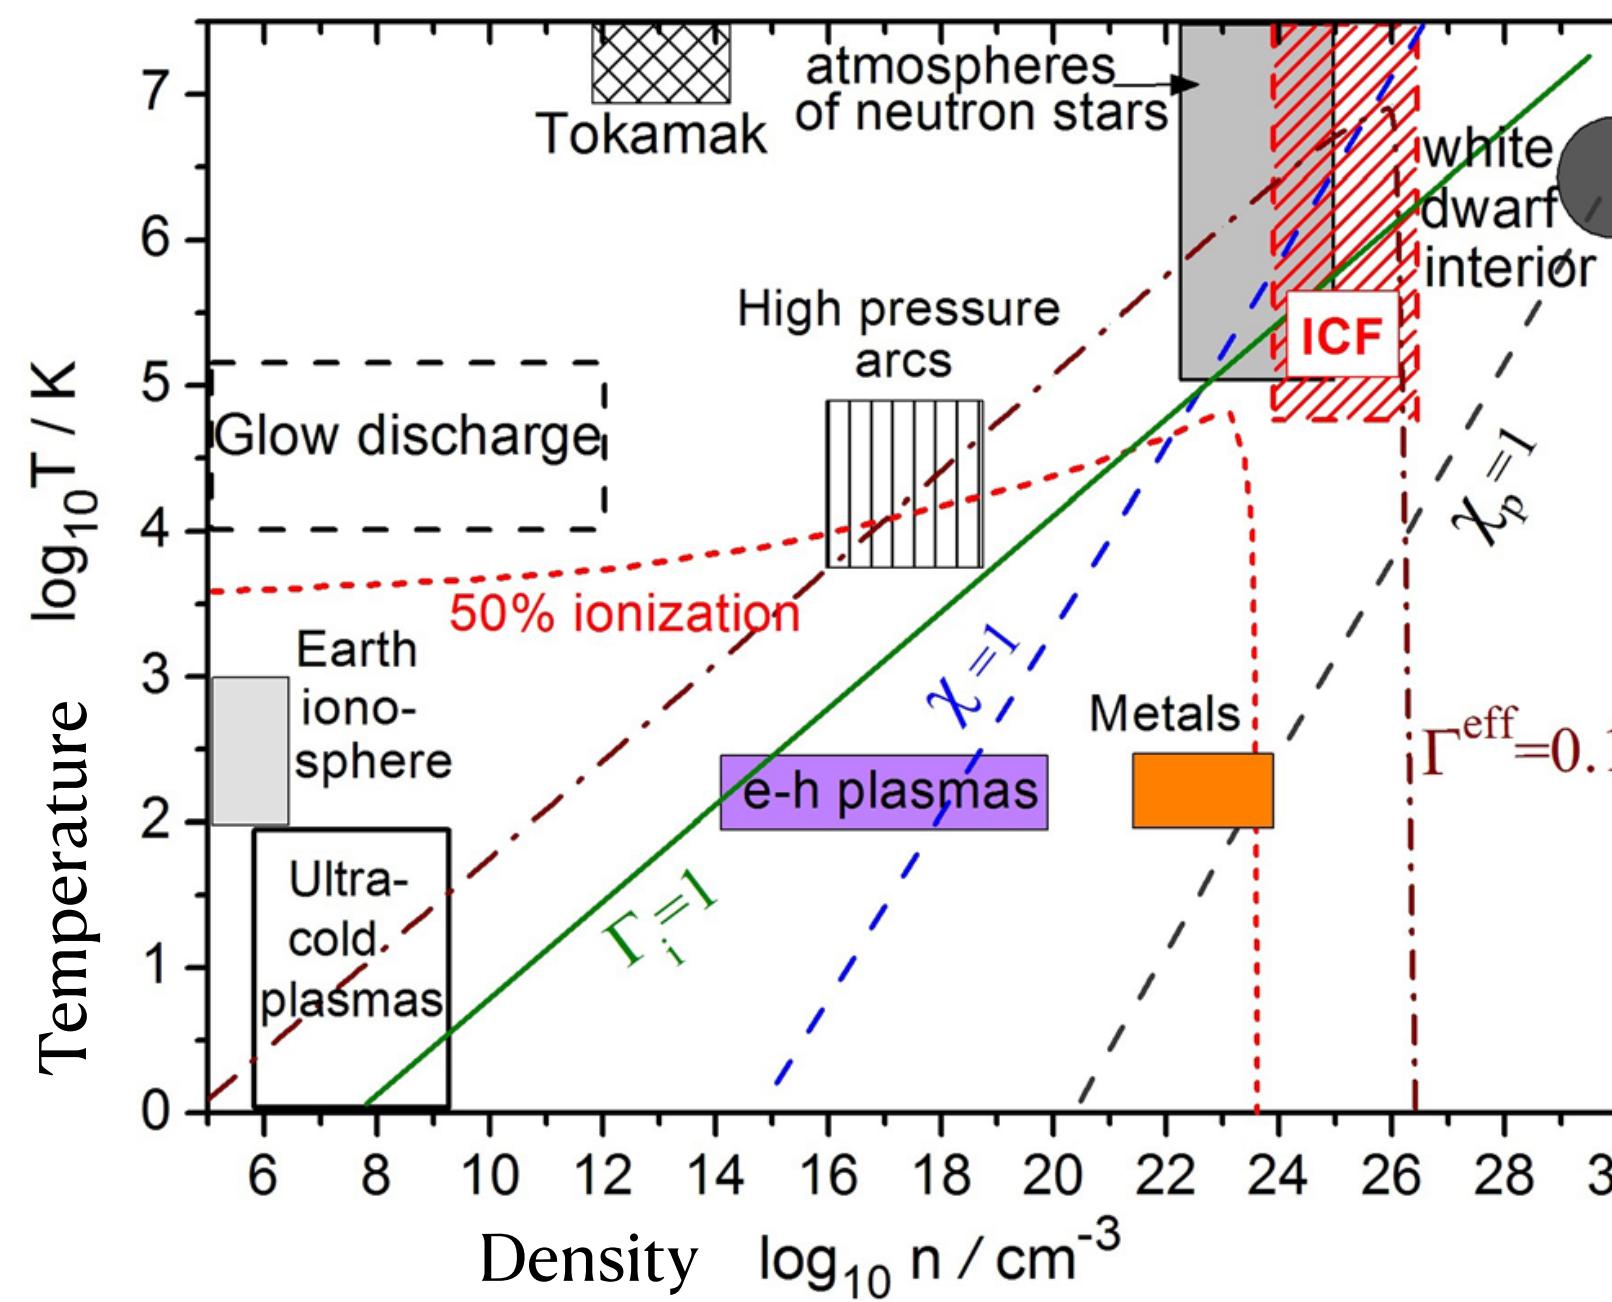
CONTENTS

1 GENERATIVE MODELING	2
1.1 Probabilistic Generative Modeling	2
1.2 Generative Model Zoo	4
1.2.1 Boltzmann Machines	5
1.2.2 Autoregressive Models	8
1.2.3 Normalizing Flow	9
1.2.4 Variational Autoencoders	13
1.2.5 Tensor Networks	15
1.2.6 Generative Adversarial Networks	17
1.2.7 Generative Moment Matching Networks	18
1.3 Summary	20
2 PHYSICS APPLICATIONS	21
2.1 Variational Ansatz	21
2.2 Renormalization Group	22
2.3 Monte Carlo Update Proposals	22
2.4 Chemical and Material Design	23
2.5 Quantum Information Science and Beyond	24
3 RESOURCES	25
BIBLIOGRAPHY	26

Ab-initio study of quantum matters at finite T

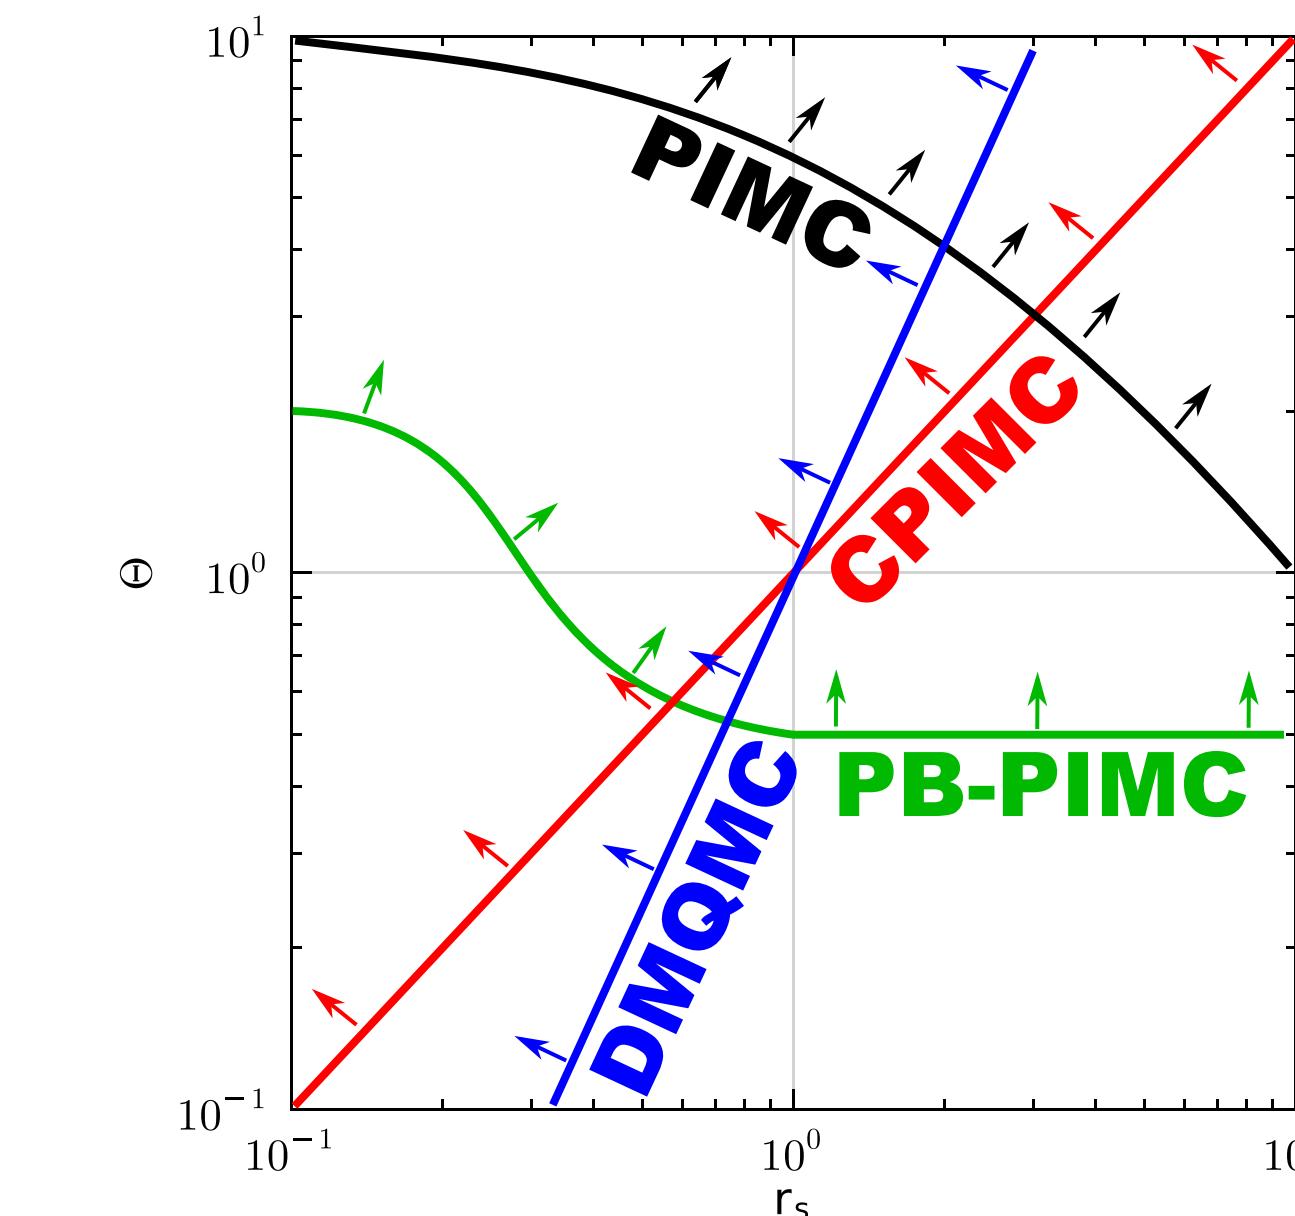
$$H = - \sum_i \frac{\hbar^2}{2m_e} \nabla_i^2 - \sum_{I,i} \frac{Z_I e^2}{|R_I - r_i|} + \frac{1}{2} \sum_{i \neq j} \frac{e^2}{|r_i - r_j|} - \sum_I \frac{\hbar^2}{2m_I} \nabla_I^2 + \frac{1}{2} \sum_{I \neq J} \frac{Z_I Z_J e^2}{|R_I - R_J|}$$

$$Z = \text{Tr}(e^{-\beta H})$$



Bonitz et al, Phys. Plasmas '20

Application range
of the workhorse
quantum Monte Carlo



Dornheim et al, Phys. Plasmas '17

The classical variational free-energy approach

Gibbs–Bogolyubov–Feynman variational principle

$$F = \int d\mathbf{x} p(\mathbf{x}) [k_B T \ln p(\mathbf{x}) + H(\mathbf{x})] \geq -k_B T \ln Z$$

↓ ↓
 entropy energy

**Difficulties in Applying the Variational
Principle to Quantum Field Theories¹**

Richard P. Feynman

¹transcript of Professor Feynman's talk in 1987

deep
generative
models !

Generative modeling



Known: samples

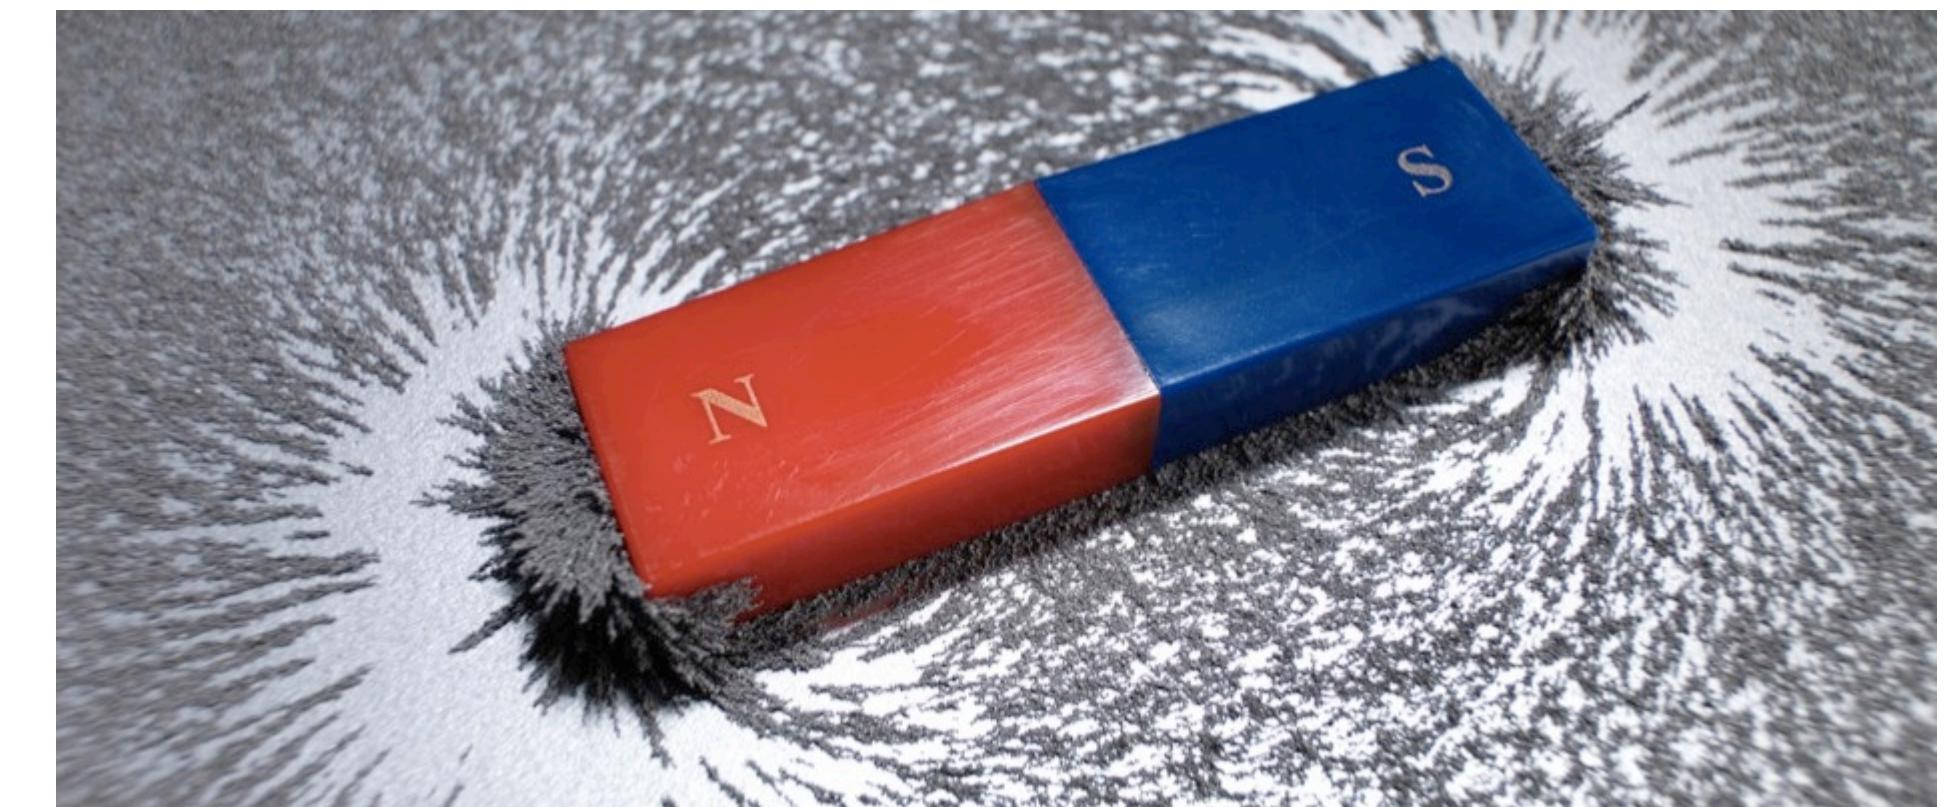
Unknown: generating distribution

Maximum likelihood estimation

“learn from data”

$$\mathcal{L} = - \mathbb{E}_{x \sim \text{dataset}} [\ln p(x)]$$

Statistical physics



Known: energy function

Unknown: samples, partition function

Variational free energy

“learn from Hamiltonian”

$$F = \mathbb{E}_{x \sim p(x)} [k_B T \ln p(x) + H(x)]$$

Deep variational free-energy approach

Use deep generative models as the variational density

$$F = \mathbb{E}_{x \sim p(x)} [k_B T \ln p(x) + H(x)]$$

↓ ↓
😊 entropy energy



Tractable entropy



Direct sampling



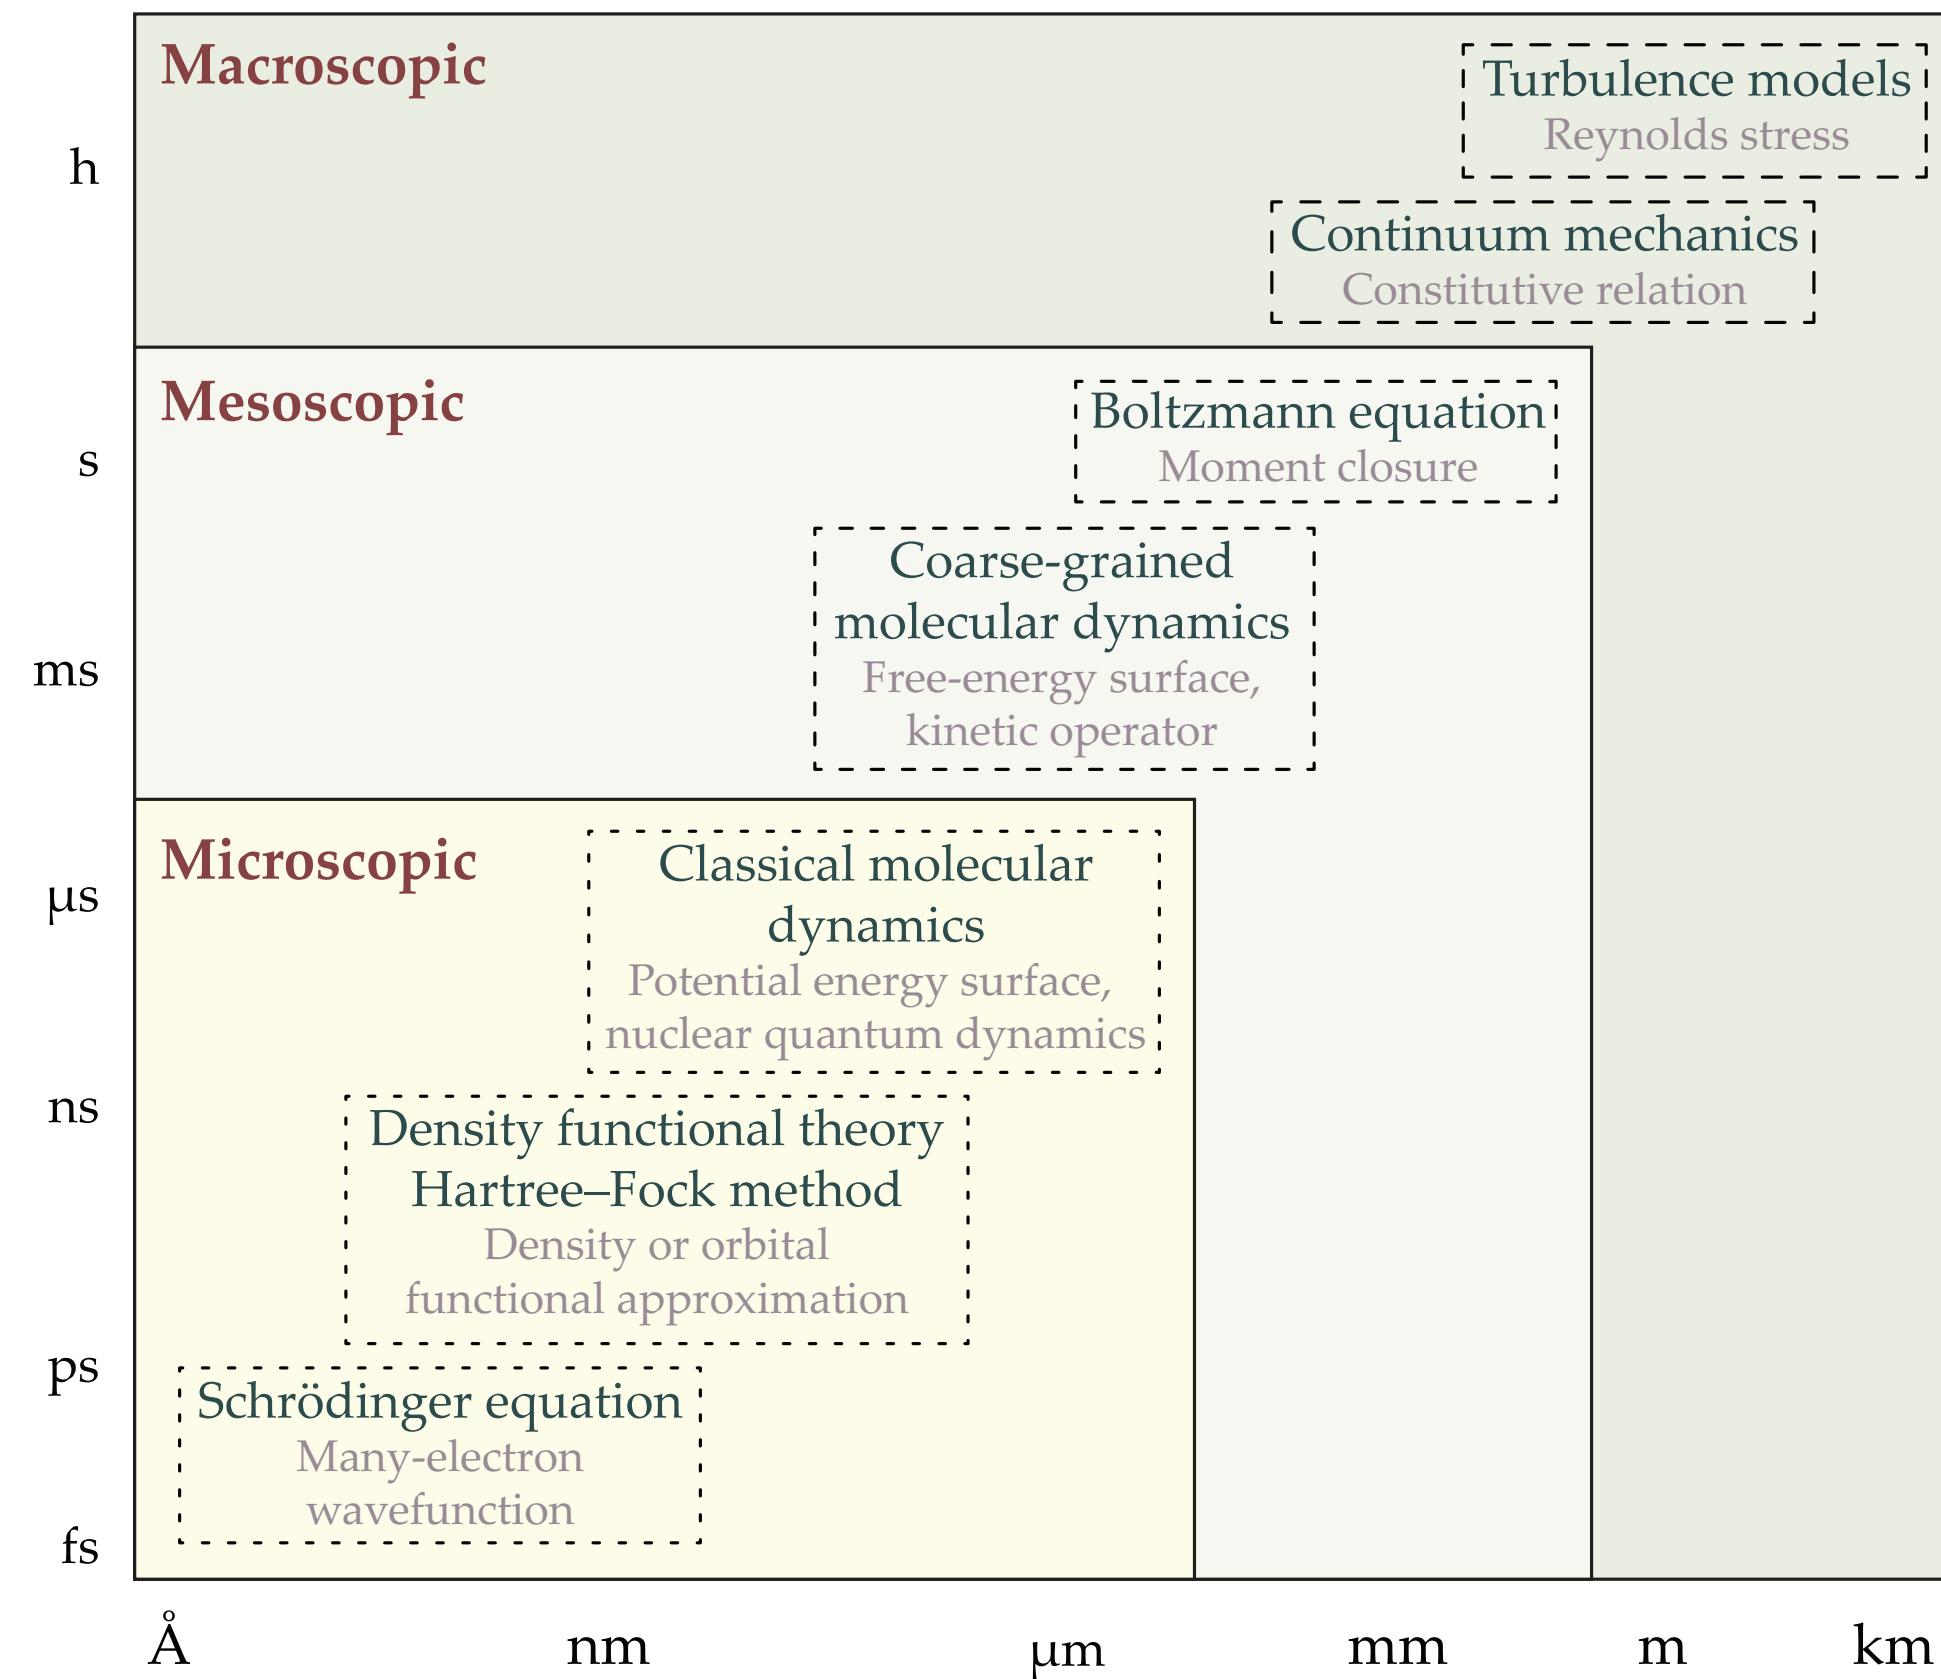
Turning sampling problem to an optimization problem

leverages the deep learning engine:



Variational free-energy in the context

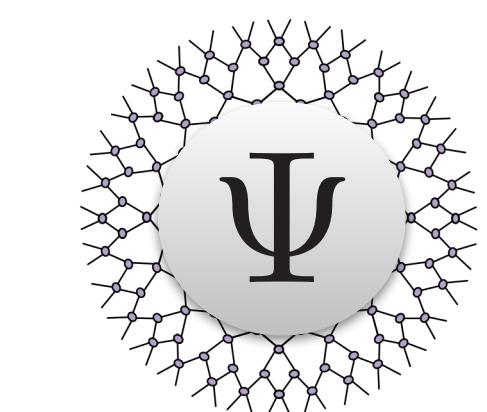
E, Han,Zhang, Physics Today 2020



Application	Model	Data	Objective
MD potential energy surface	3N-dim function	DFT energy/ force	Generalization
DFT xc functional	3-dim functional	QMC/ CCSD/...	
Variational free-energy	3N-dim functional	No	Optimization

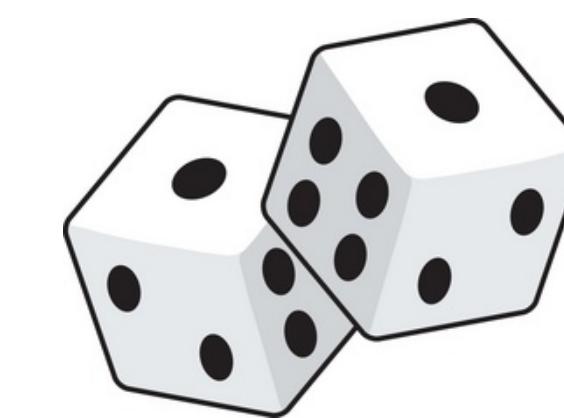
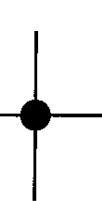
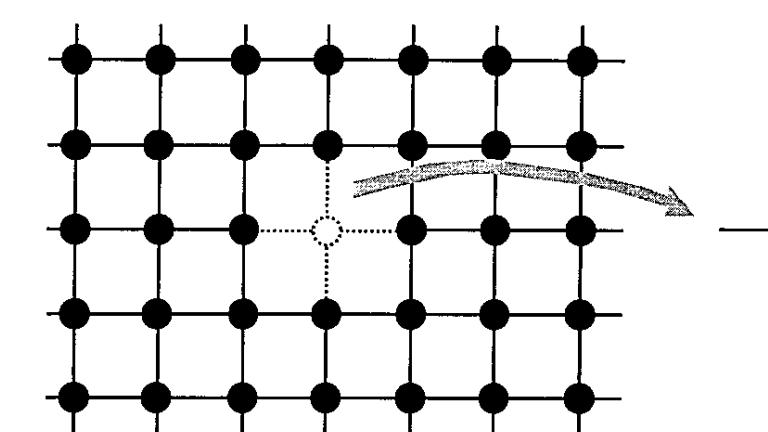
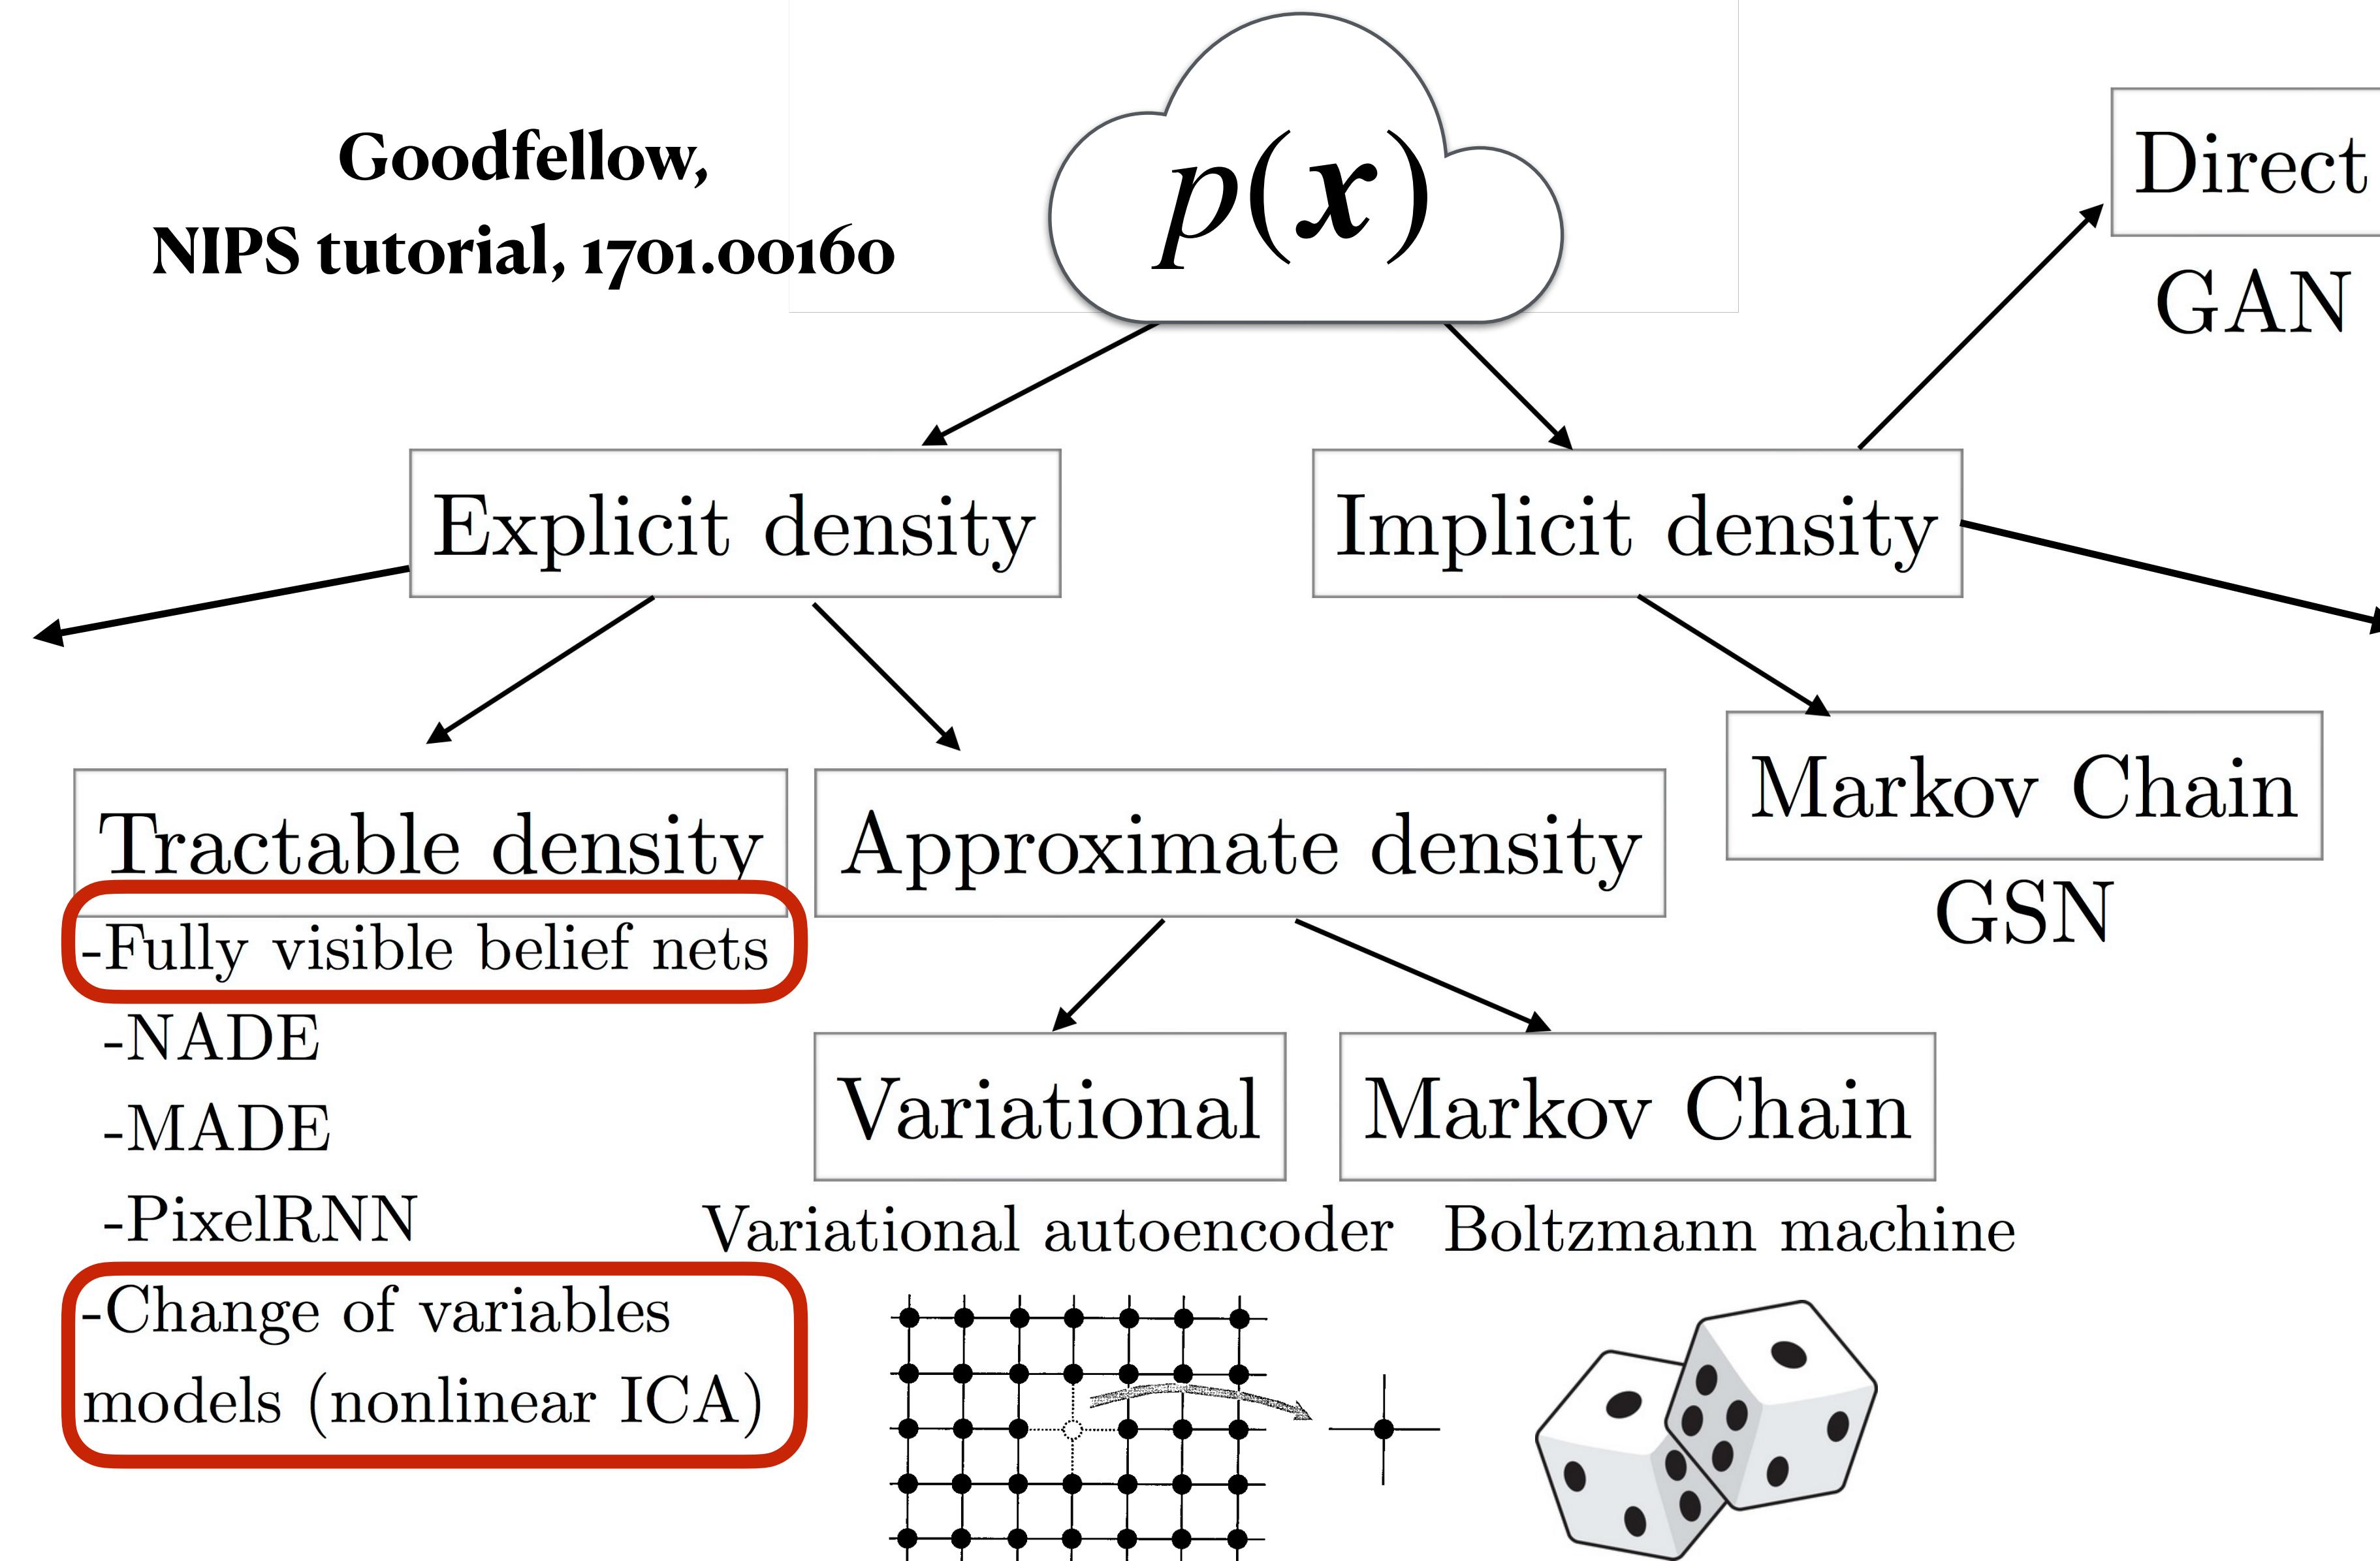
more fundamental, more difficult, more limited

Generative models and their physics genes



**Tensor
Networks**

**Goodfellow,
NIPS tutorial, 1701.00160**



Generative modeling with normalizing flows



WaveNet 1609.03499 1711.10433

<https://deepmind.com/blog/wavenet-generative-model-raw-audio/>

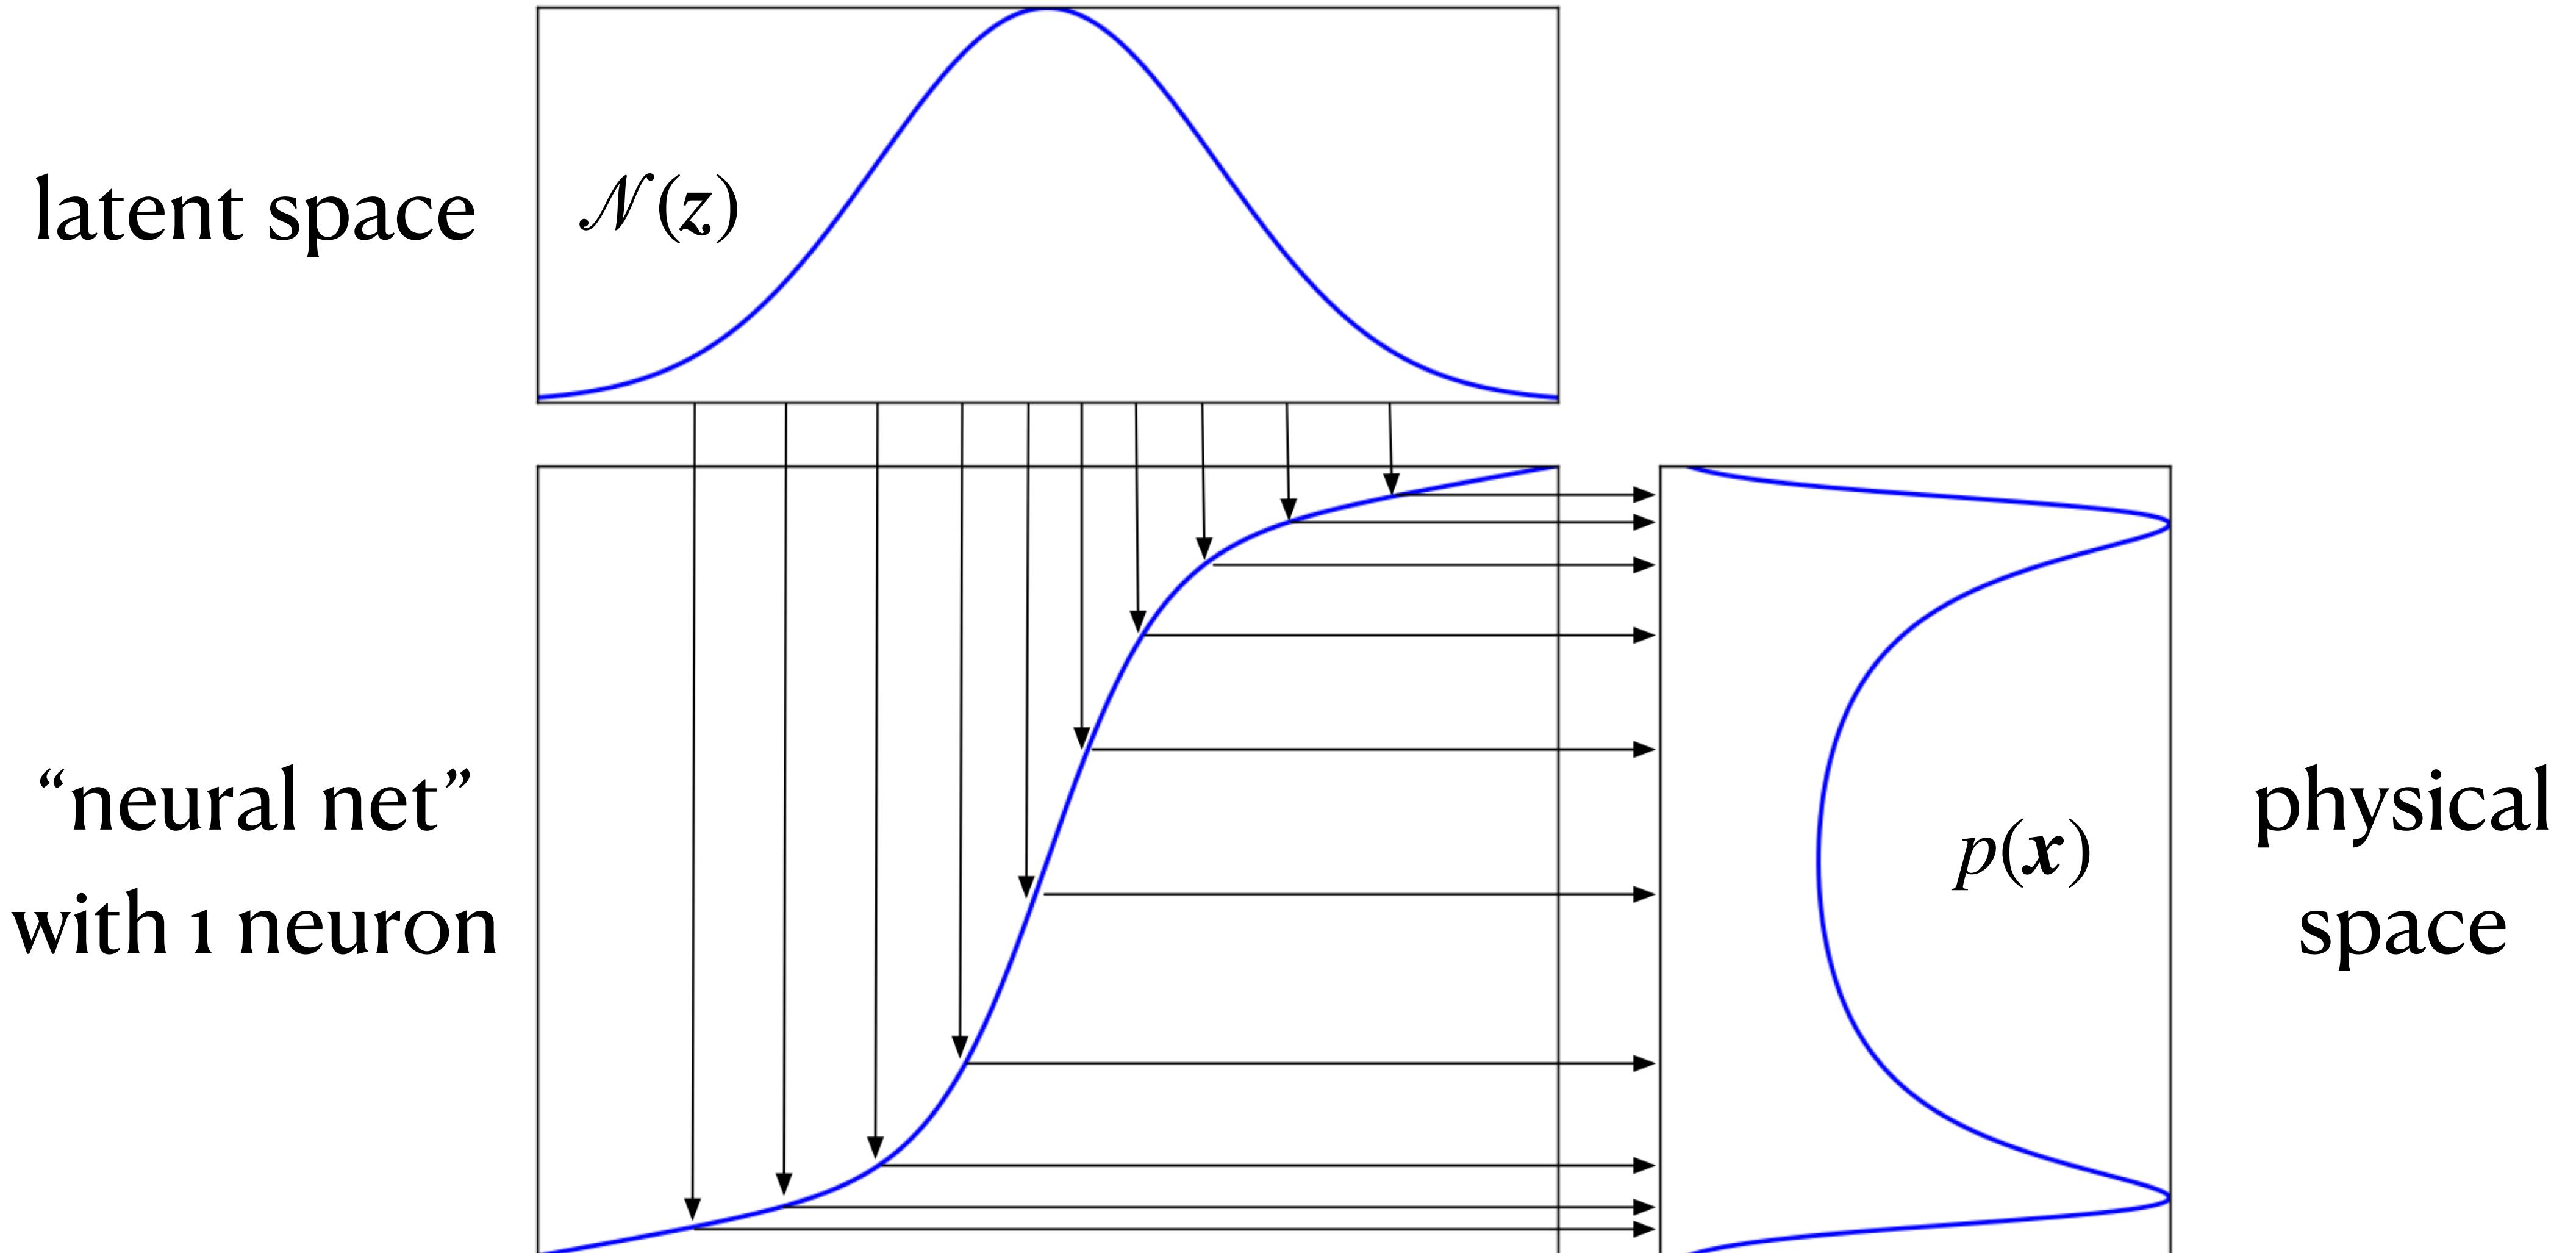
<https://deepmind.com/blog/high-fidelity-speech-synthesis-wavenet/>



Glow 1807.03039

<https://blog.openai.com/glow/>

Normalizing flow in a nutshell



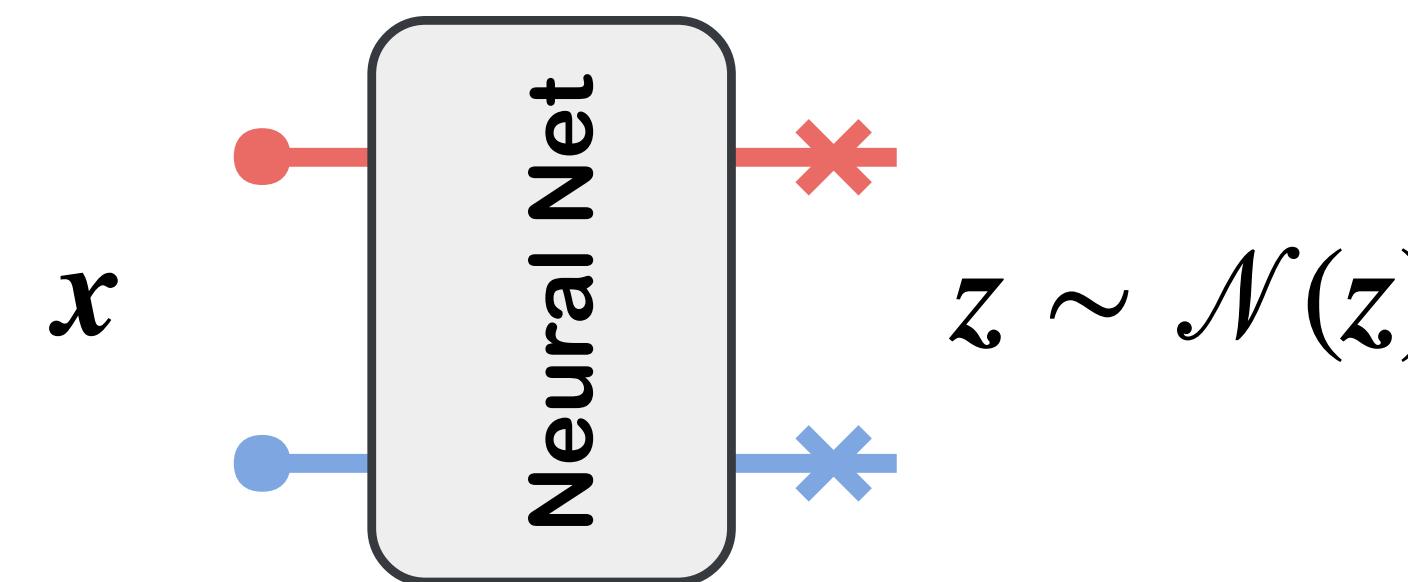
Normalizing Flows

Change of variables $x \leftrightarrow z$ with deep neural nets

$$p(x) = \mathcal{N}(z) \left| \det \left(\frac{\partial z}{\partial x} \right) \right|$$

Review article 1912.02762
Tutorial https://iclr.cc/virtual_2020/speaker_4.html

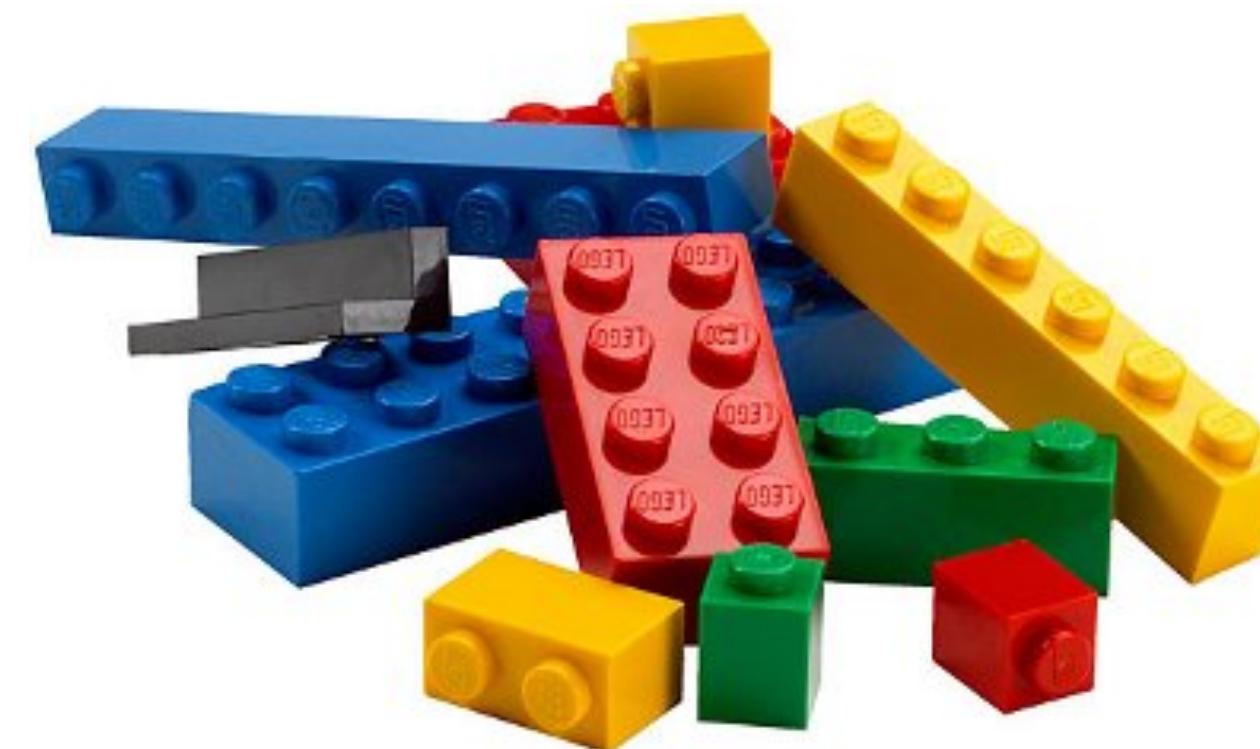
composable, differentiable, and invertible mapping between manifolds



Learn probability transformations with normalizing flows

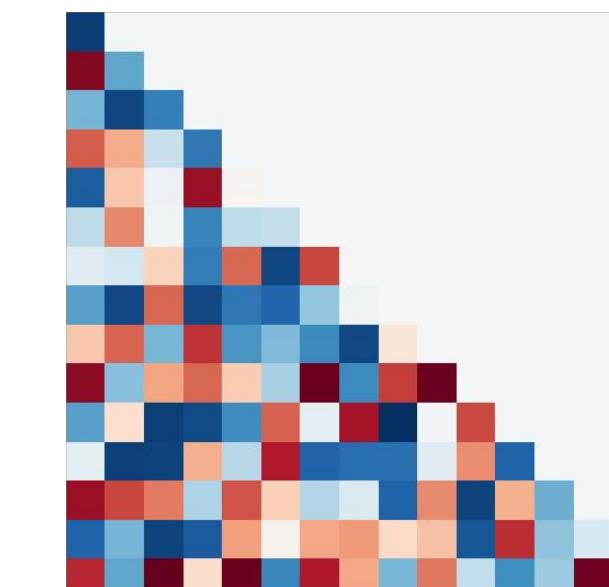
Architecture design principle

Composability

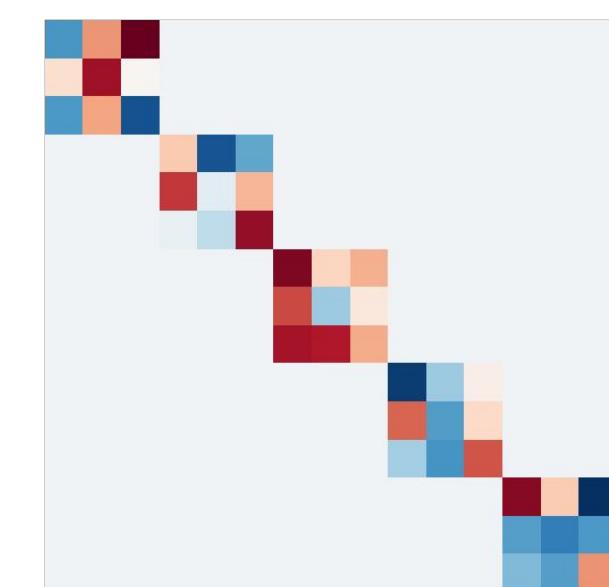


Balanced
efficiency &
inductive bias

$$\left| \det \left(\frac{\partial z}{\partial x} \right) \right|$$



Autoregressive



Neural RG

$$z = \mathcal{T}(x)$$
$$\mathcal{T} = \mathcal{T}_1 \circ \mathcal{T}_2 \circ \mathcal{T}_3 \circ \dots$$

$$\frac{\partial \rho(x, t)}{\partial t} + \nabla \cdot [\rho(x, t)v] = 0$$

Continuous flow

Example of a building block

Forward

$$\begin{cases} \mathbf{x}_< = \mathbf{z}_< \\ \mathbf{x}_> = \mathbf{z}_> \odot e^{s(\mathbf{z}_<)} + t(\mathbf{z}_<) \end{cases}$$

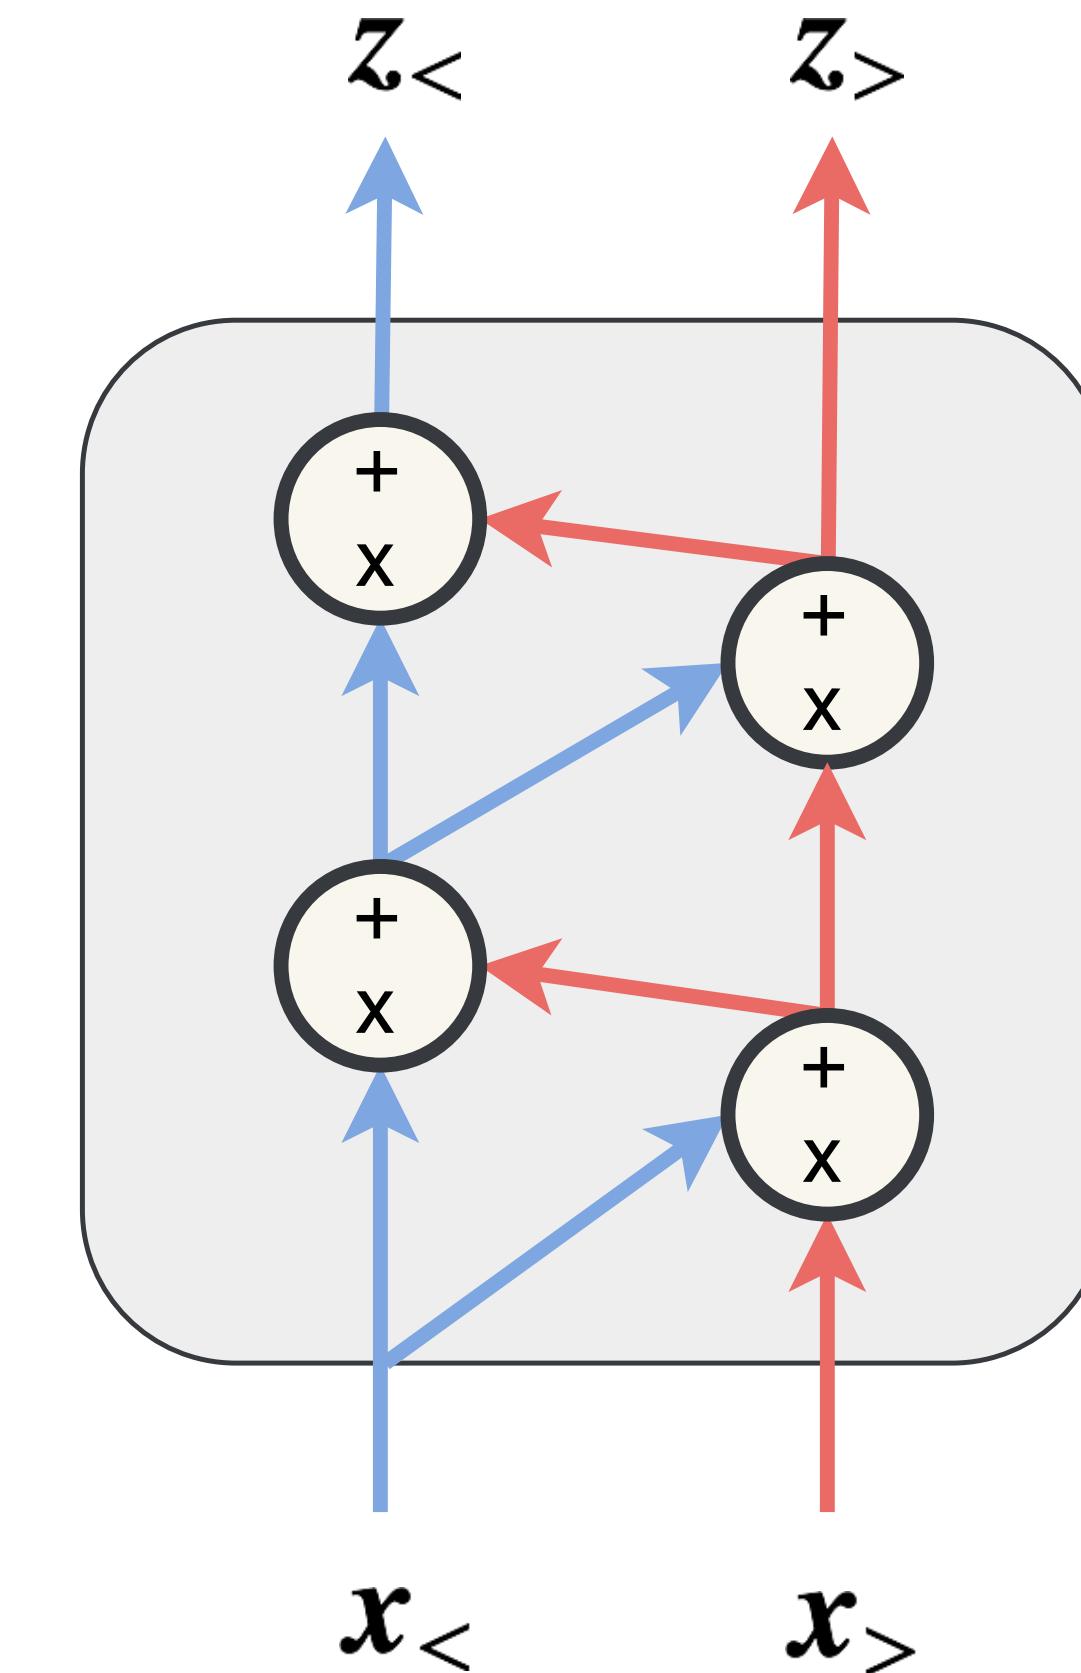
arbitrary
neural nets

Inverse

$$\begin{cases} \mathbf{z}_< = \mathbf{x}_< \\ \mathbf{z}_> = (\mathbf{x}_> - t(\mathbf{x}_<)) \odot e^{-s(\mathbf{x}_<)} \end{cases}$$

Log-Abs-Jacobian-Det

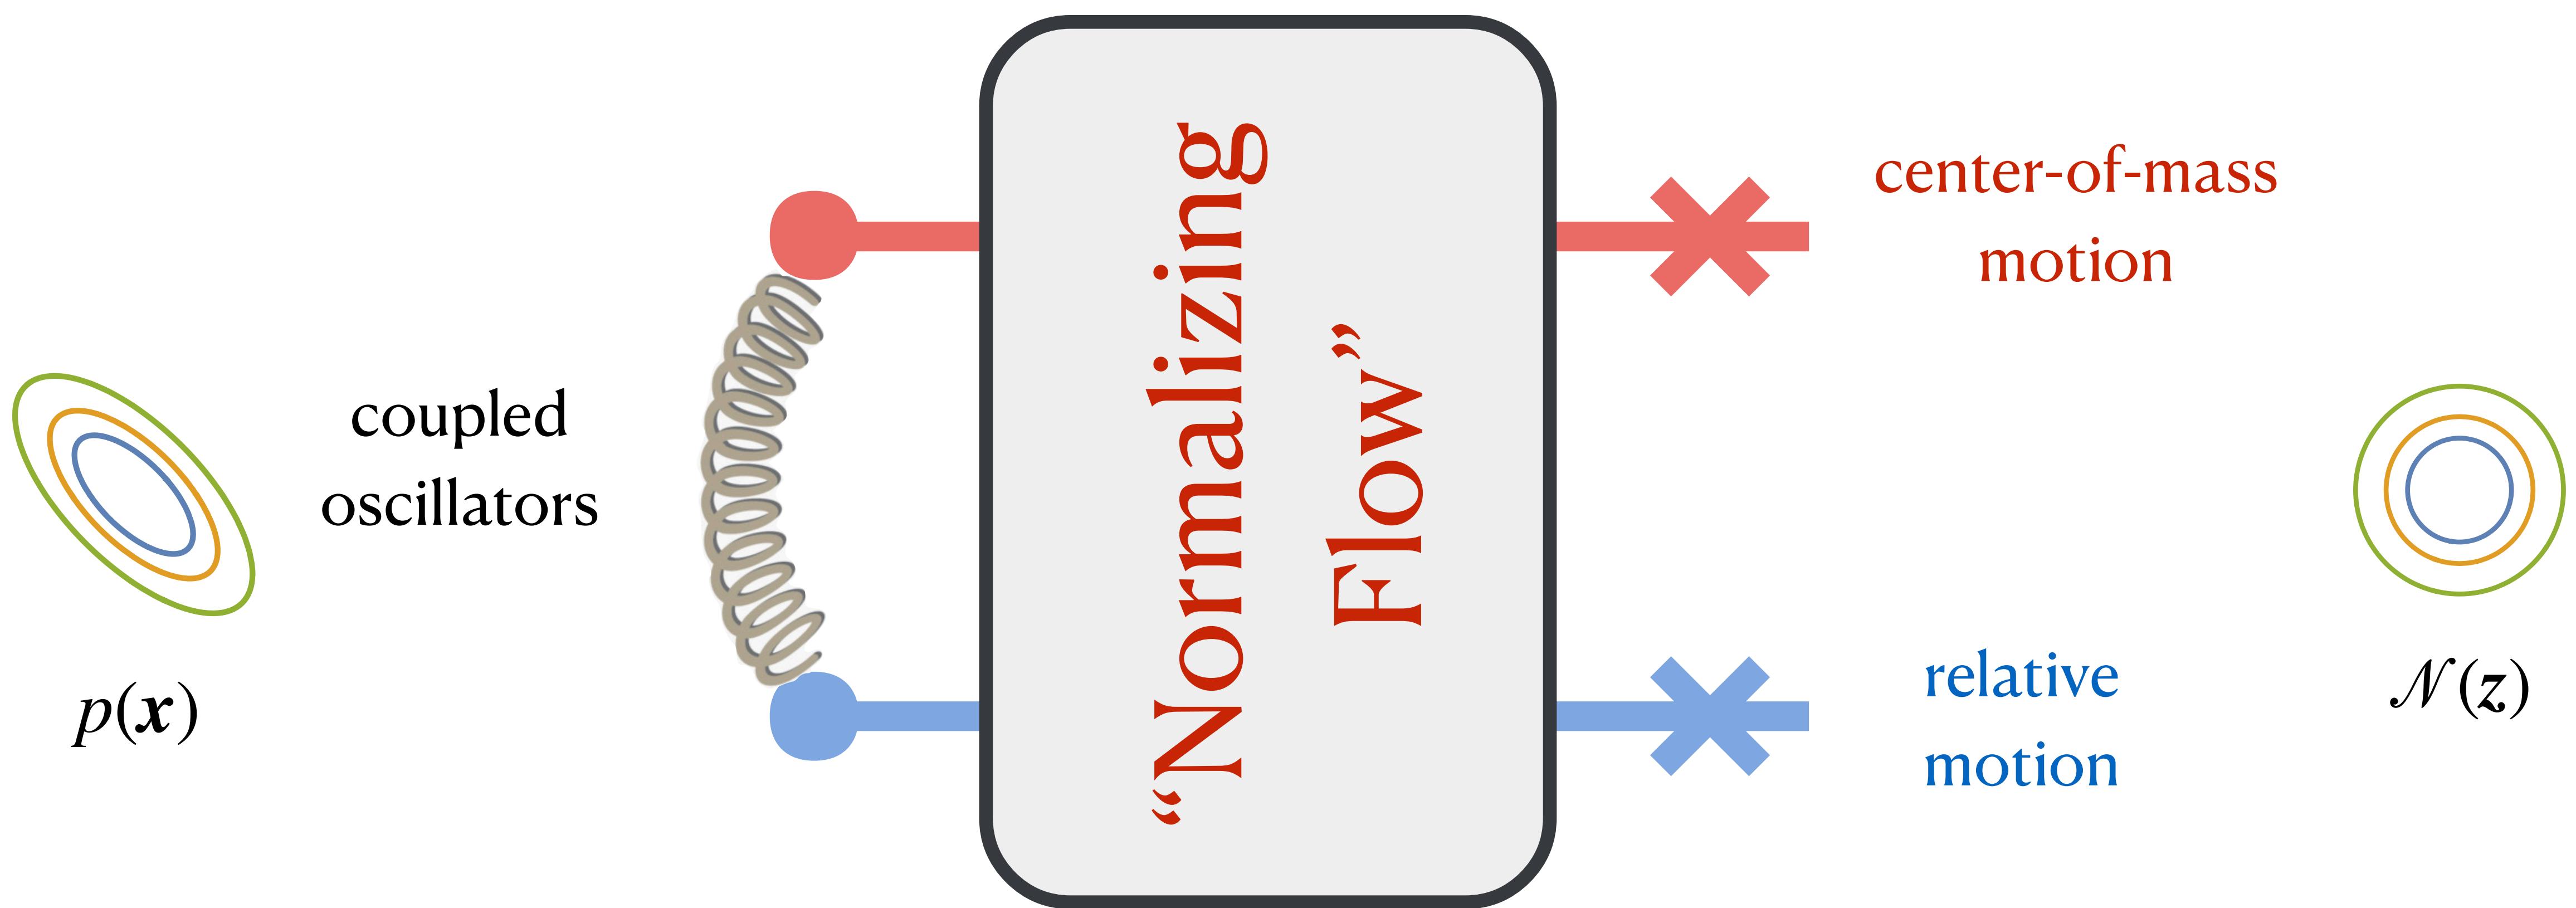
$$\ln \left| \det \left(\frac{\partial \mathbf{x}}{\partial \mathbf{z}} \right) \right| = \sum_i [s(\mathbf{z}_<)]_i$$



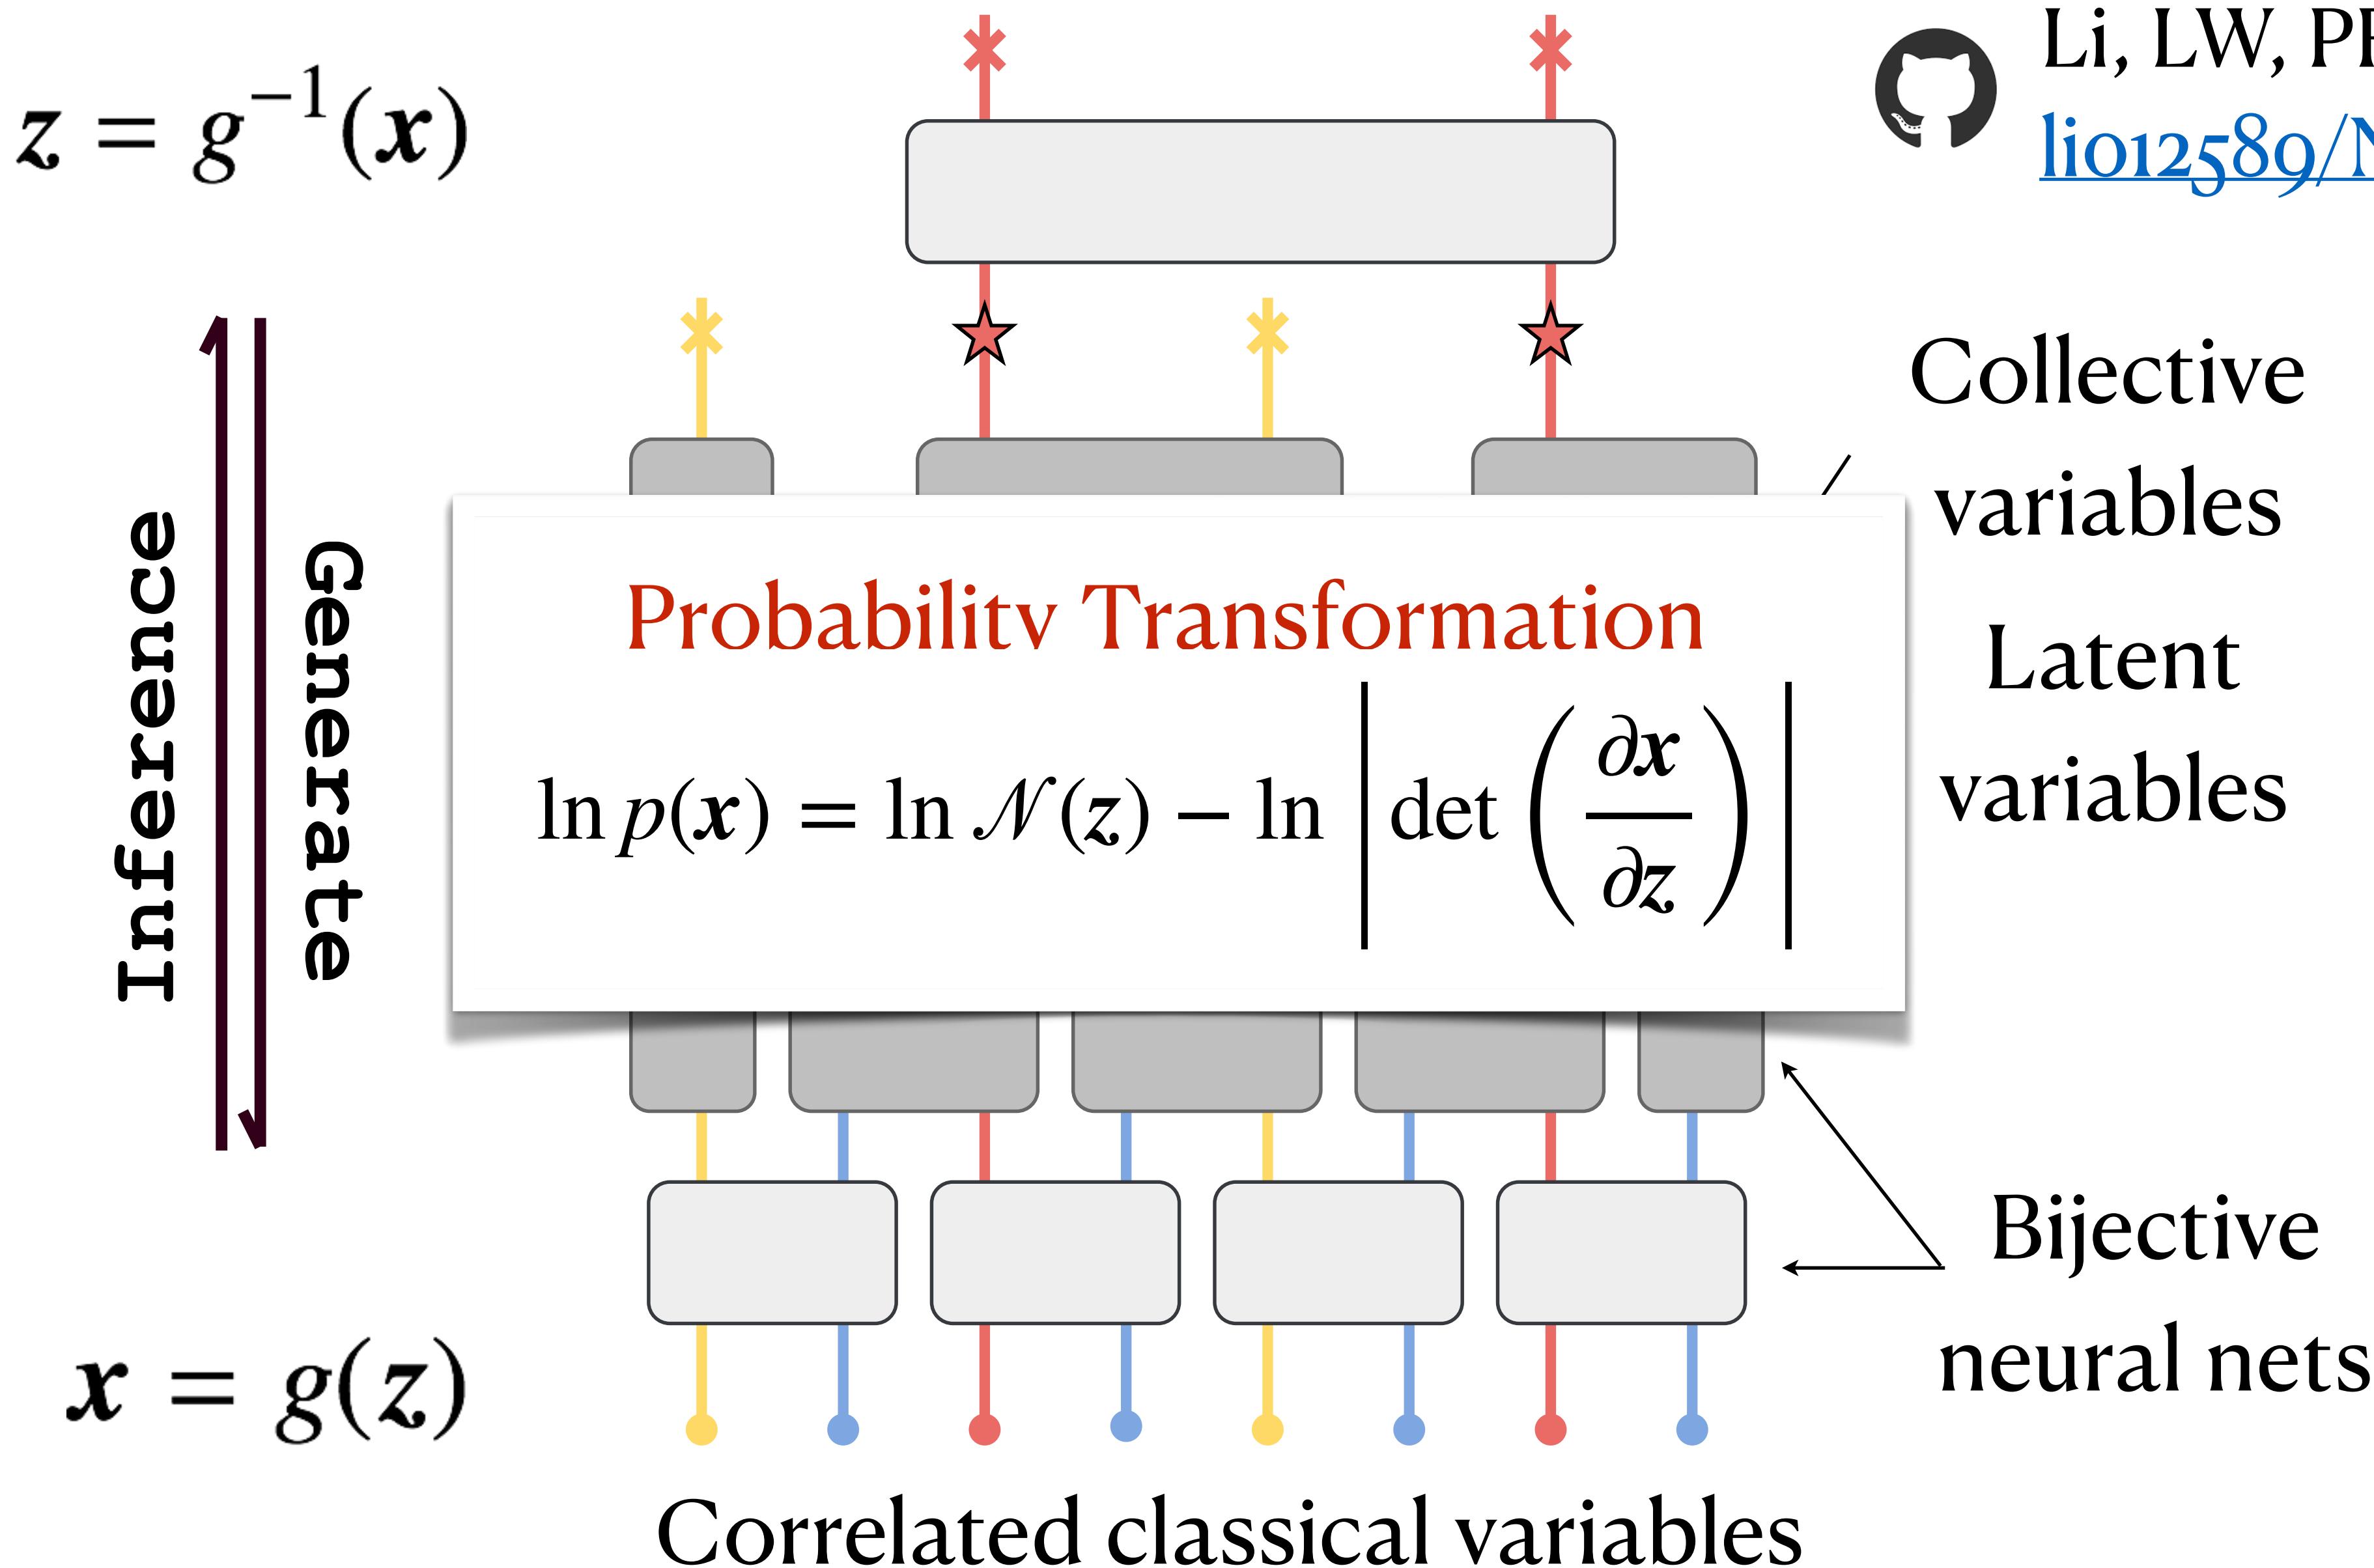
Real NVP, Dinh et al, 1605.08803

Turns out to have surprising connection Störmer–Verlet integration (later)

Normalizing flow in physics



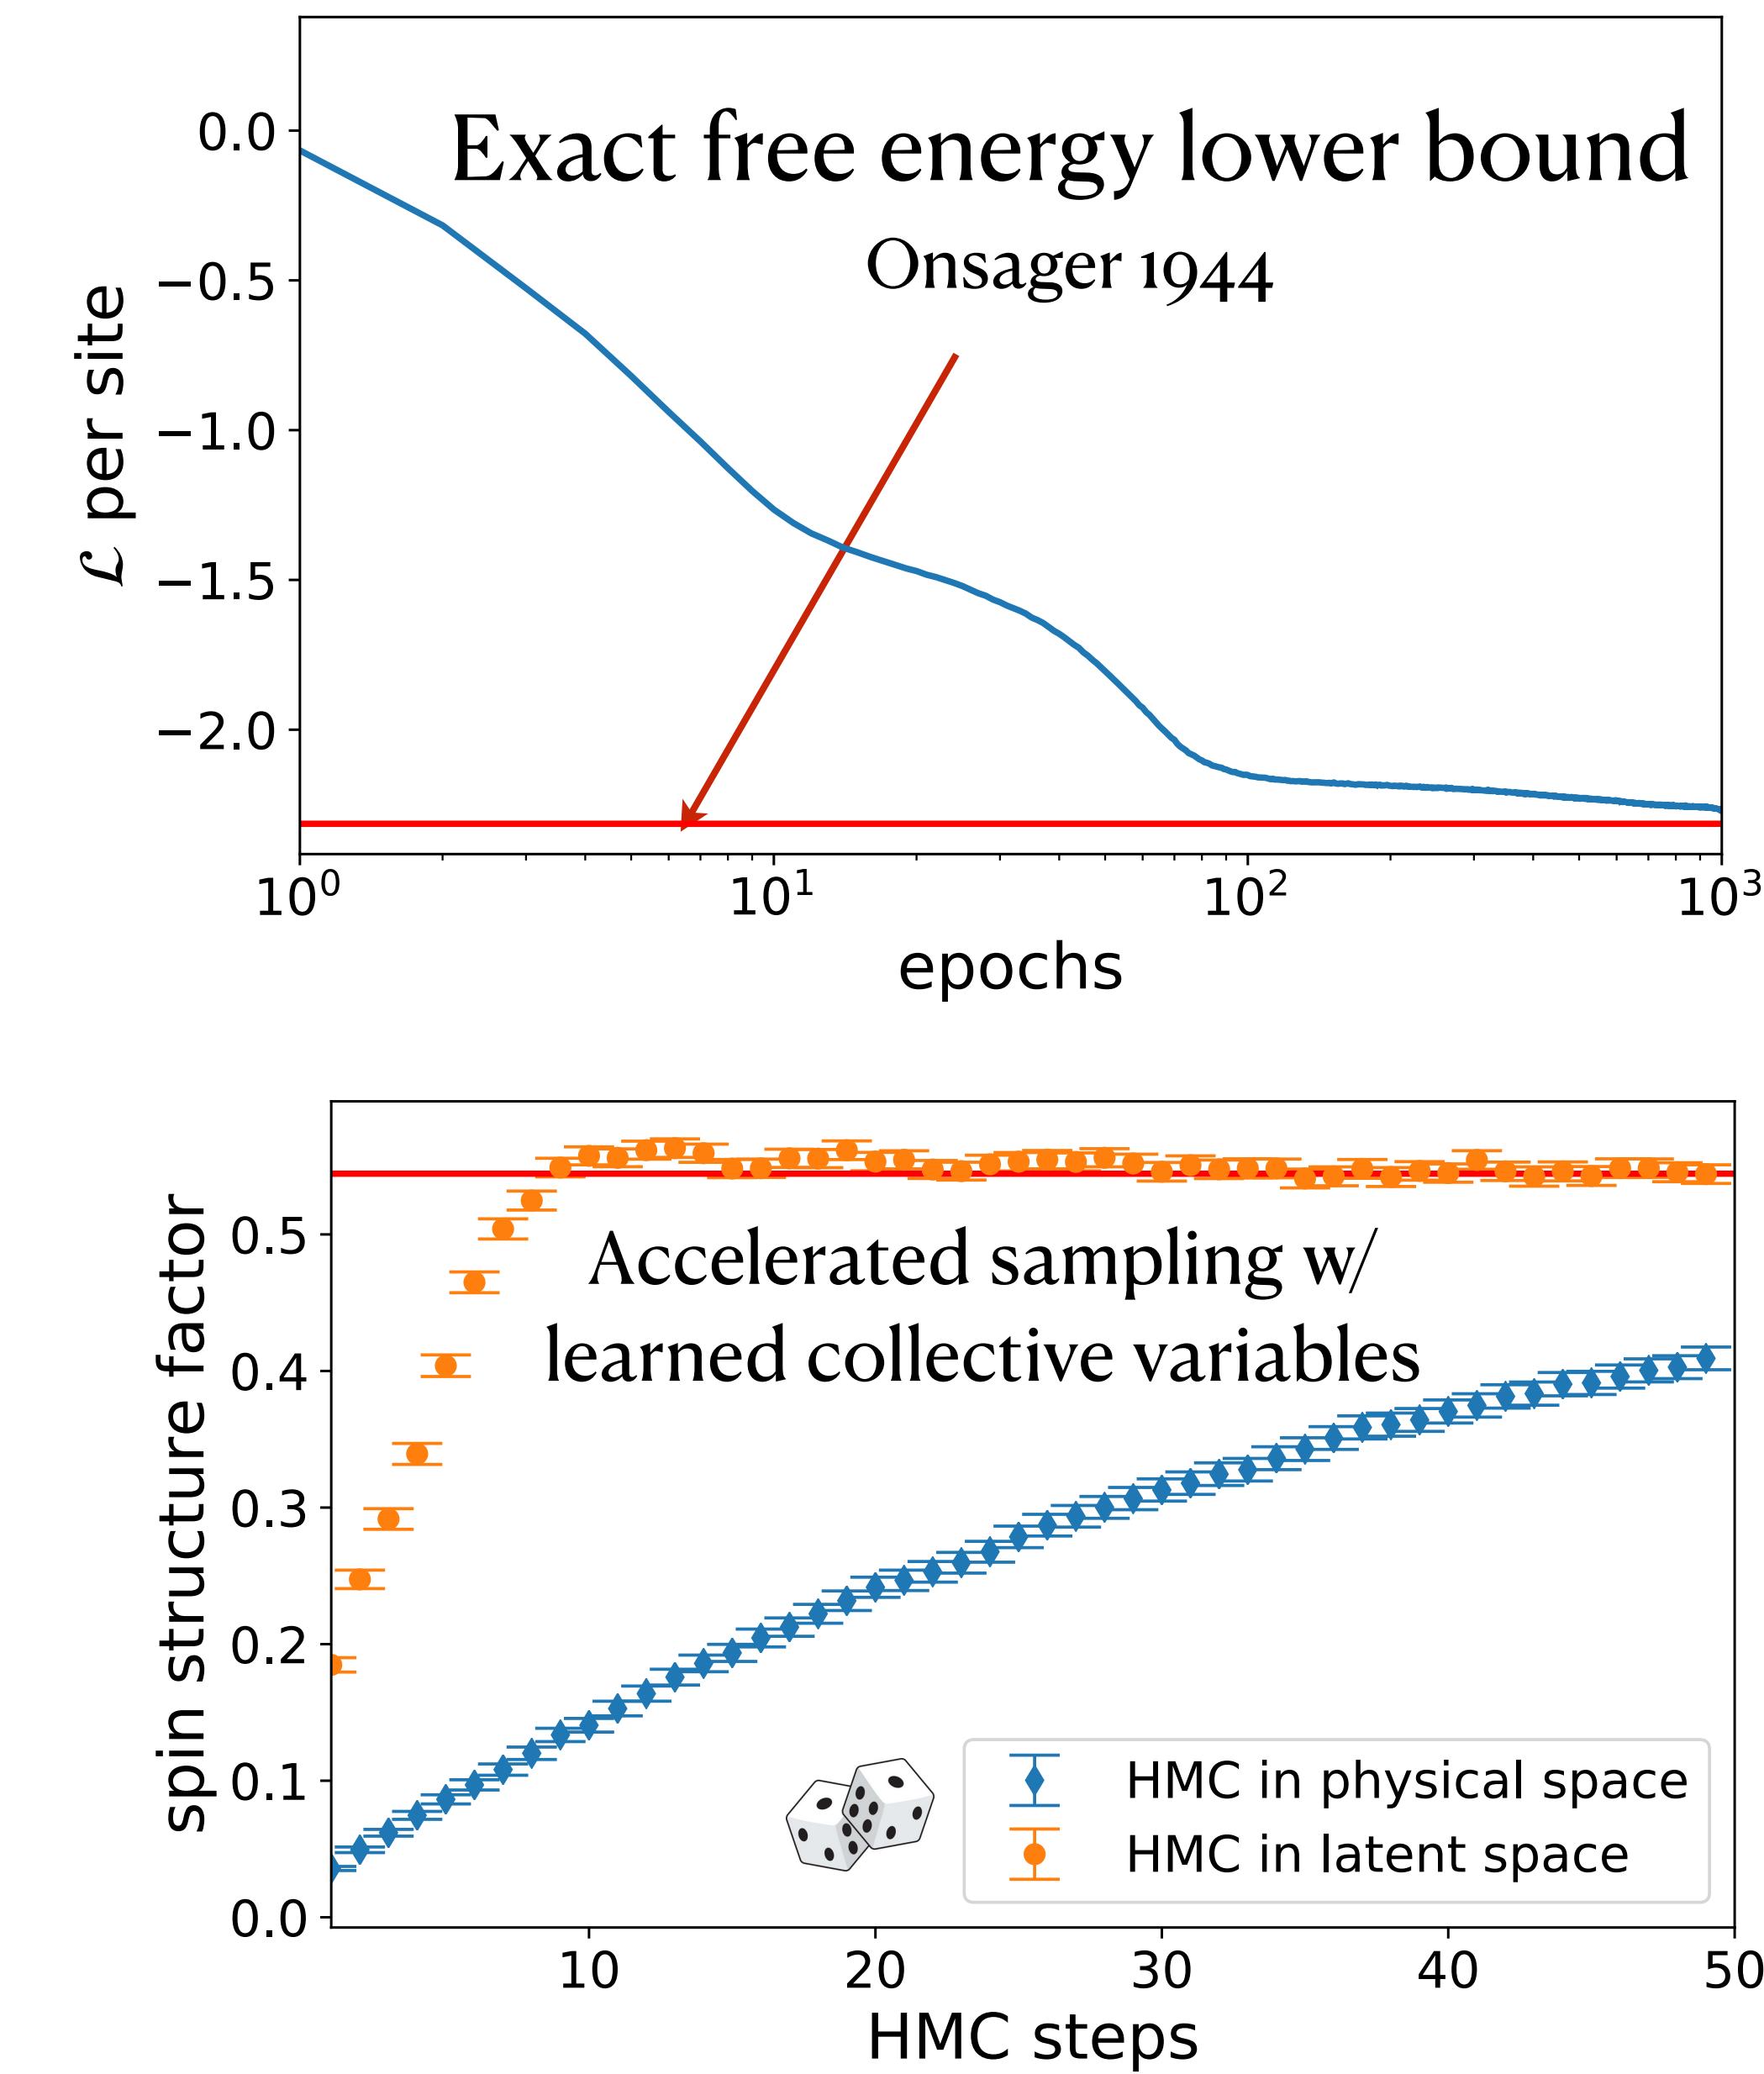
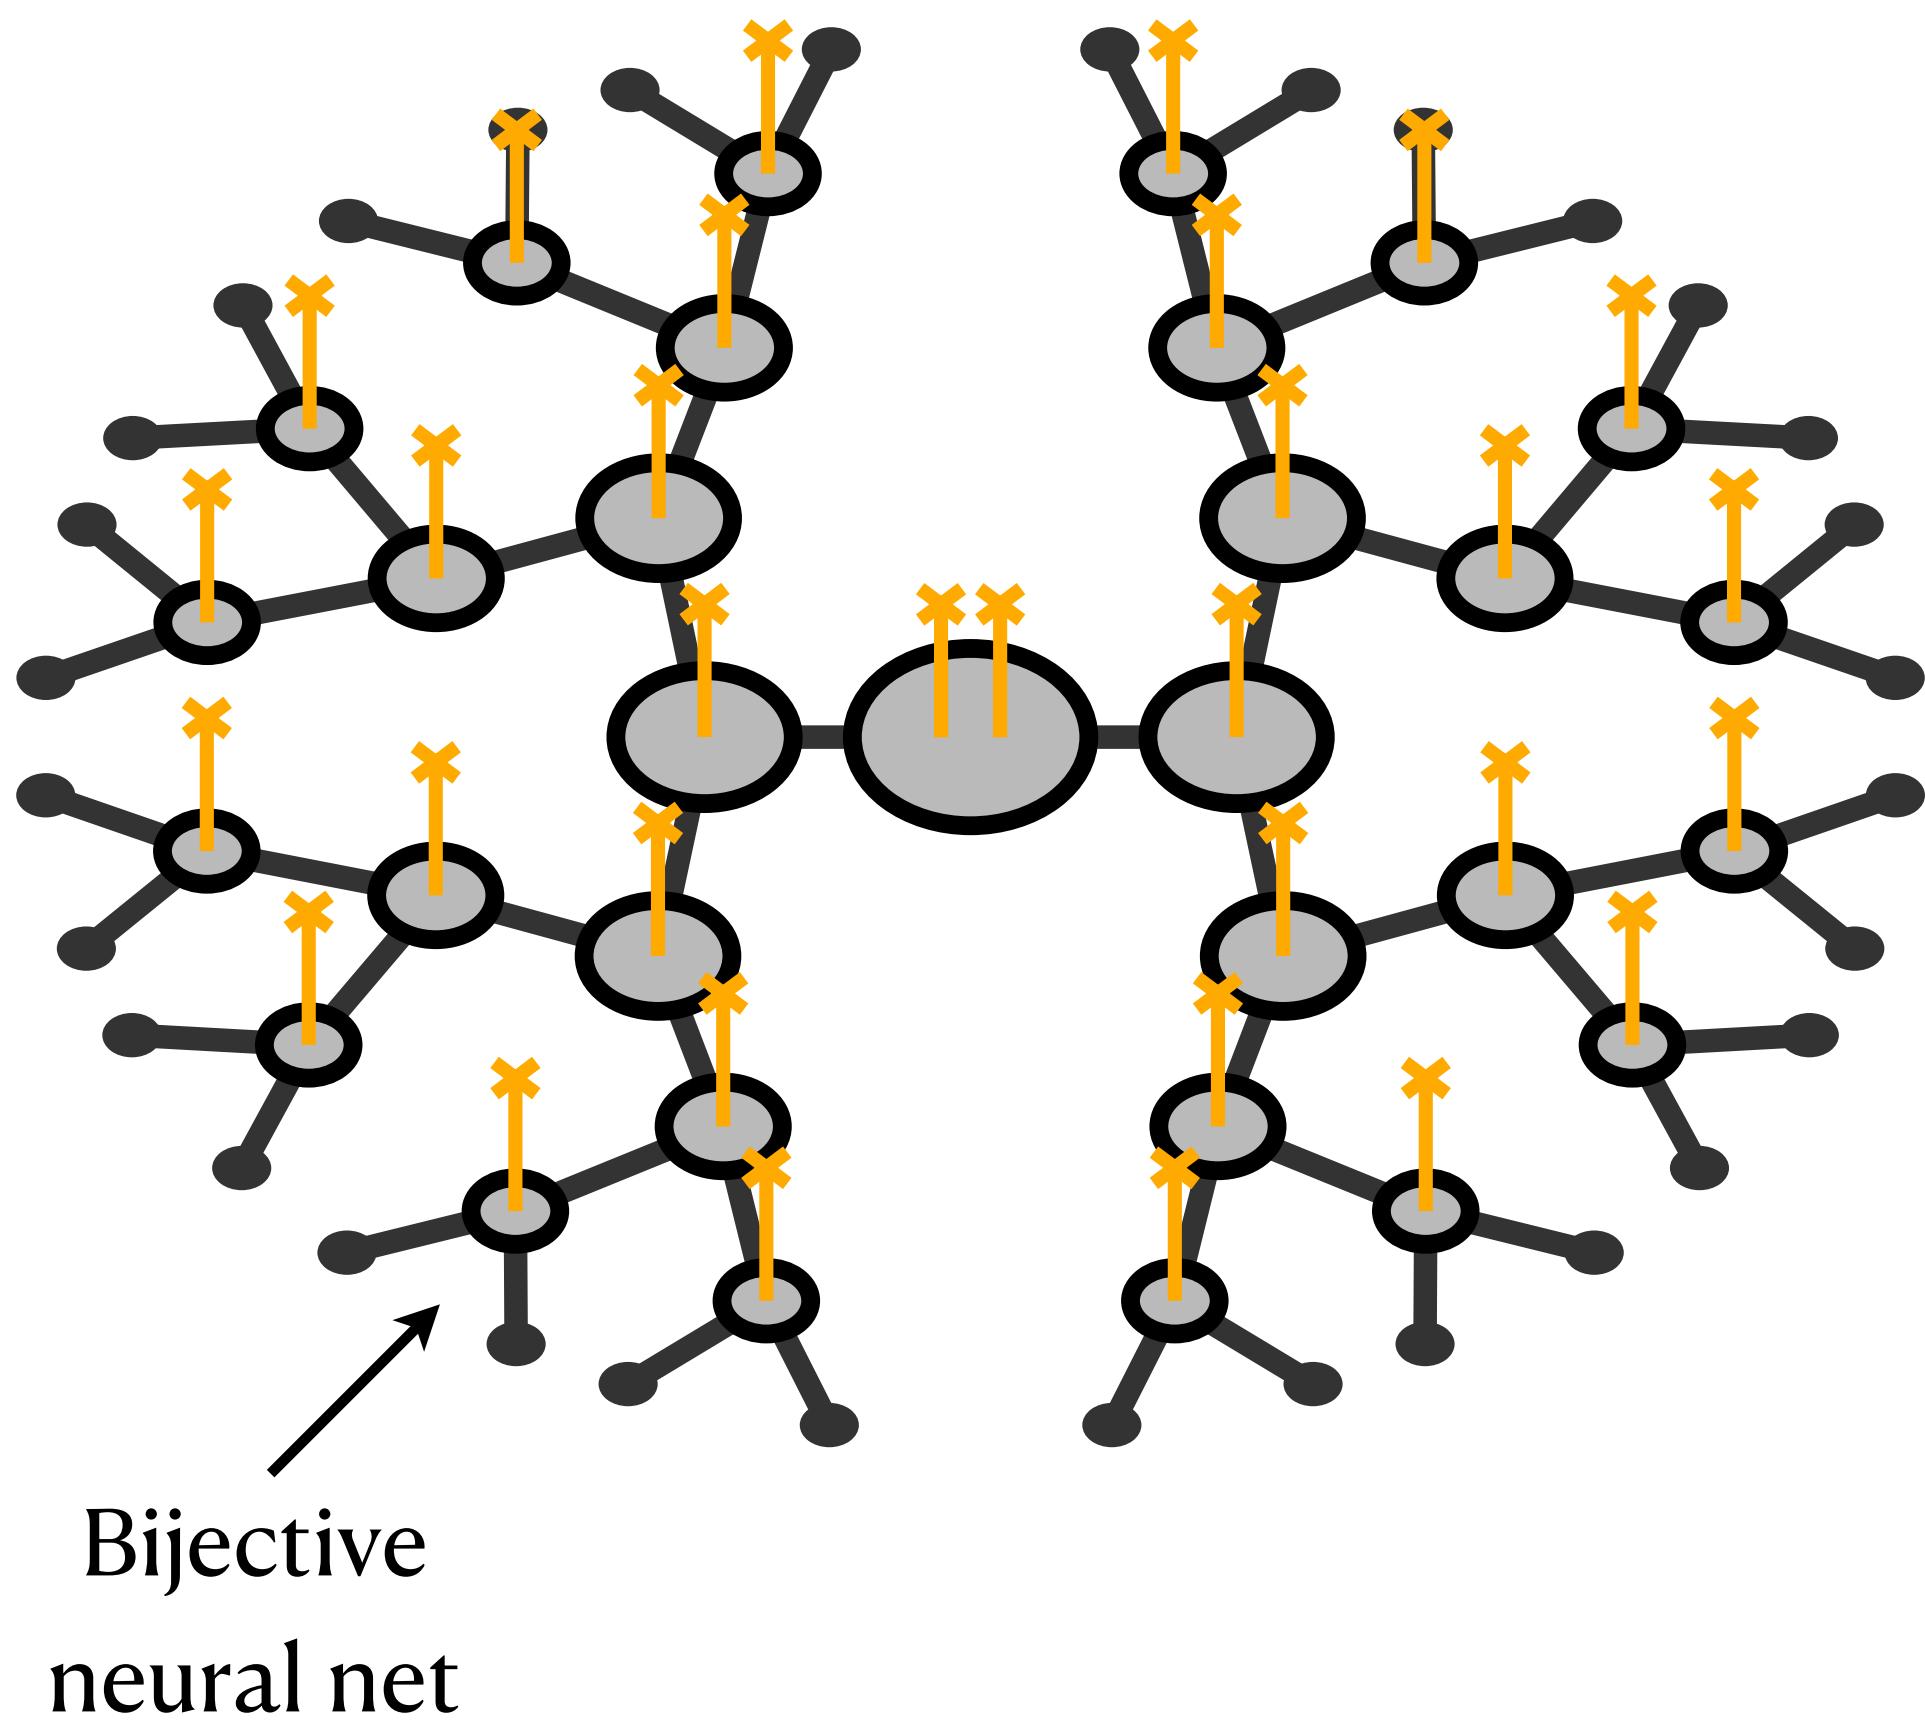
Neural Network Renormalization Group



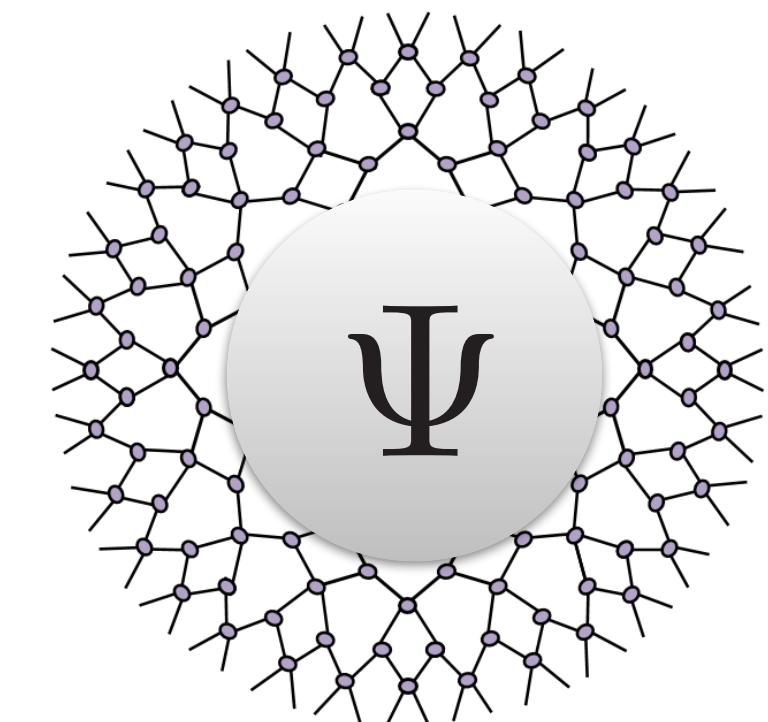
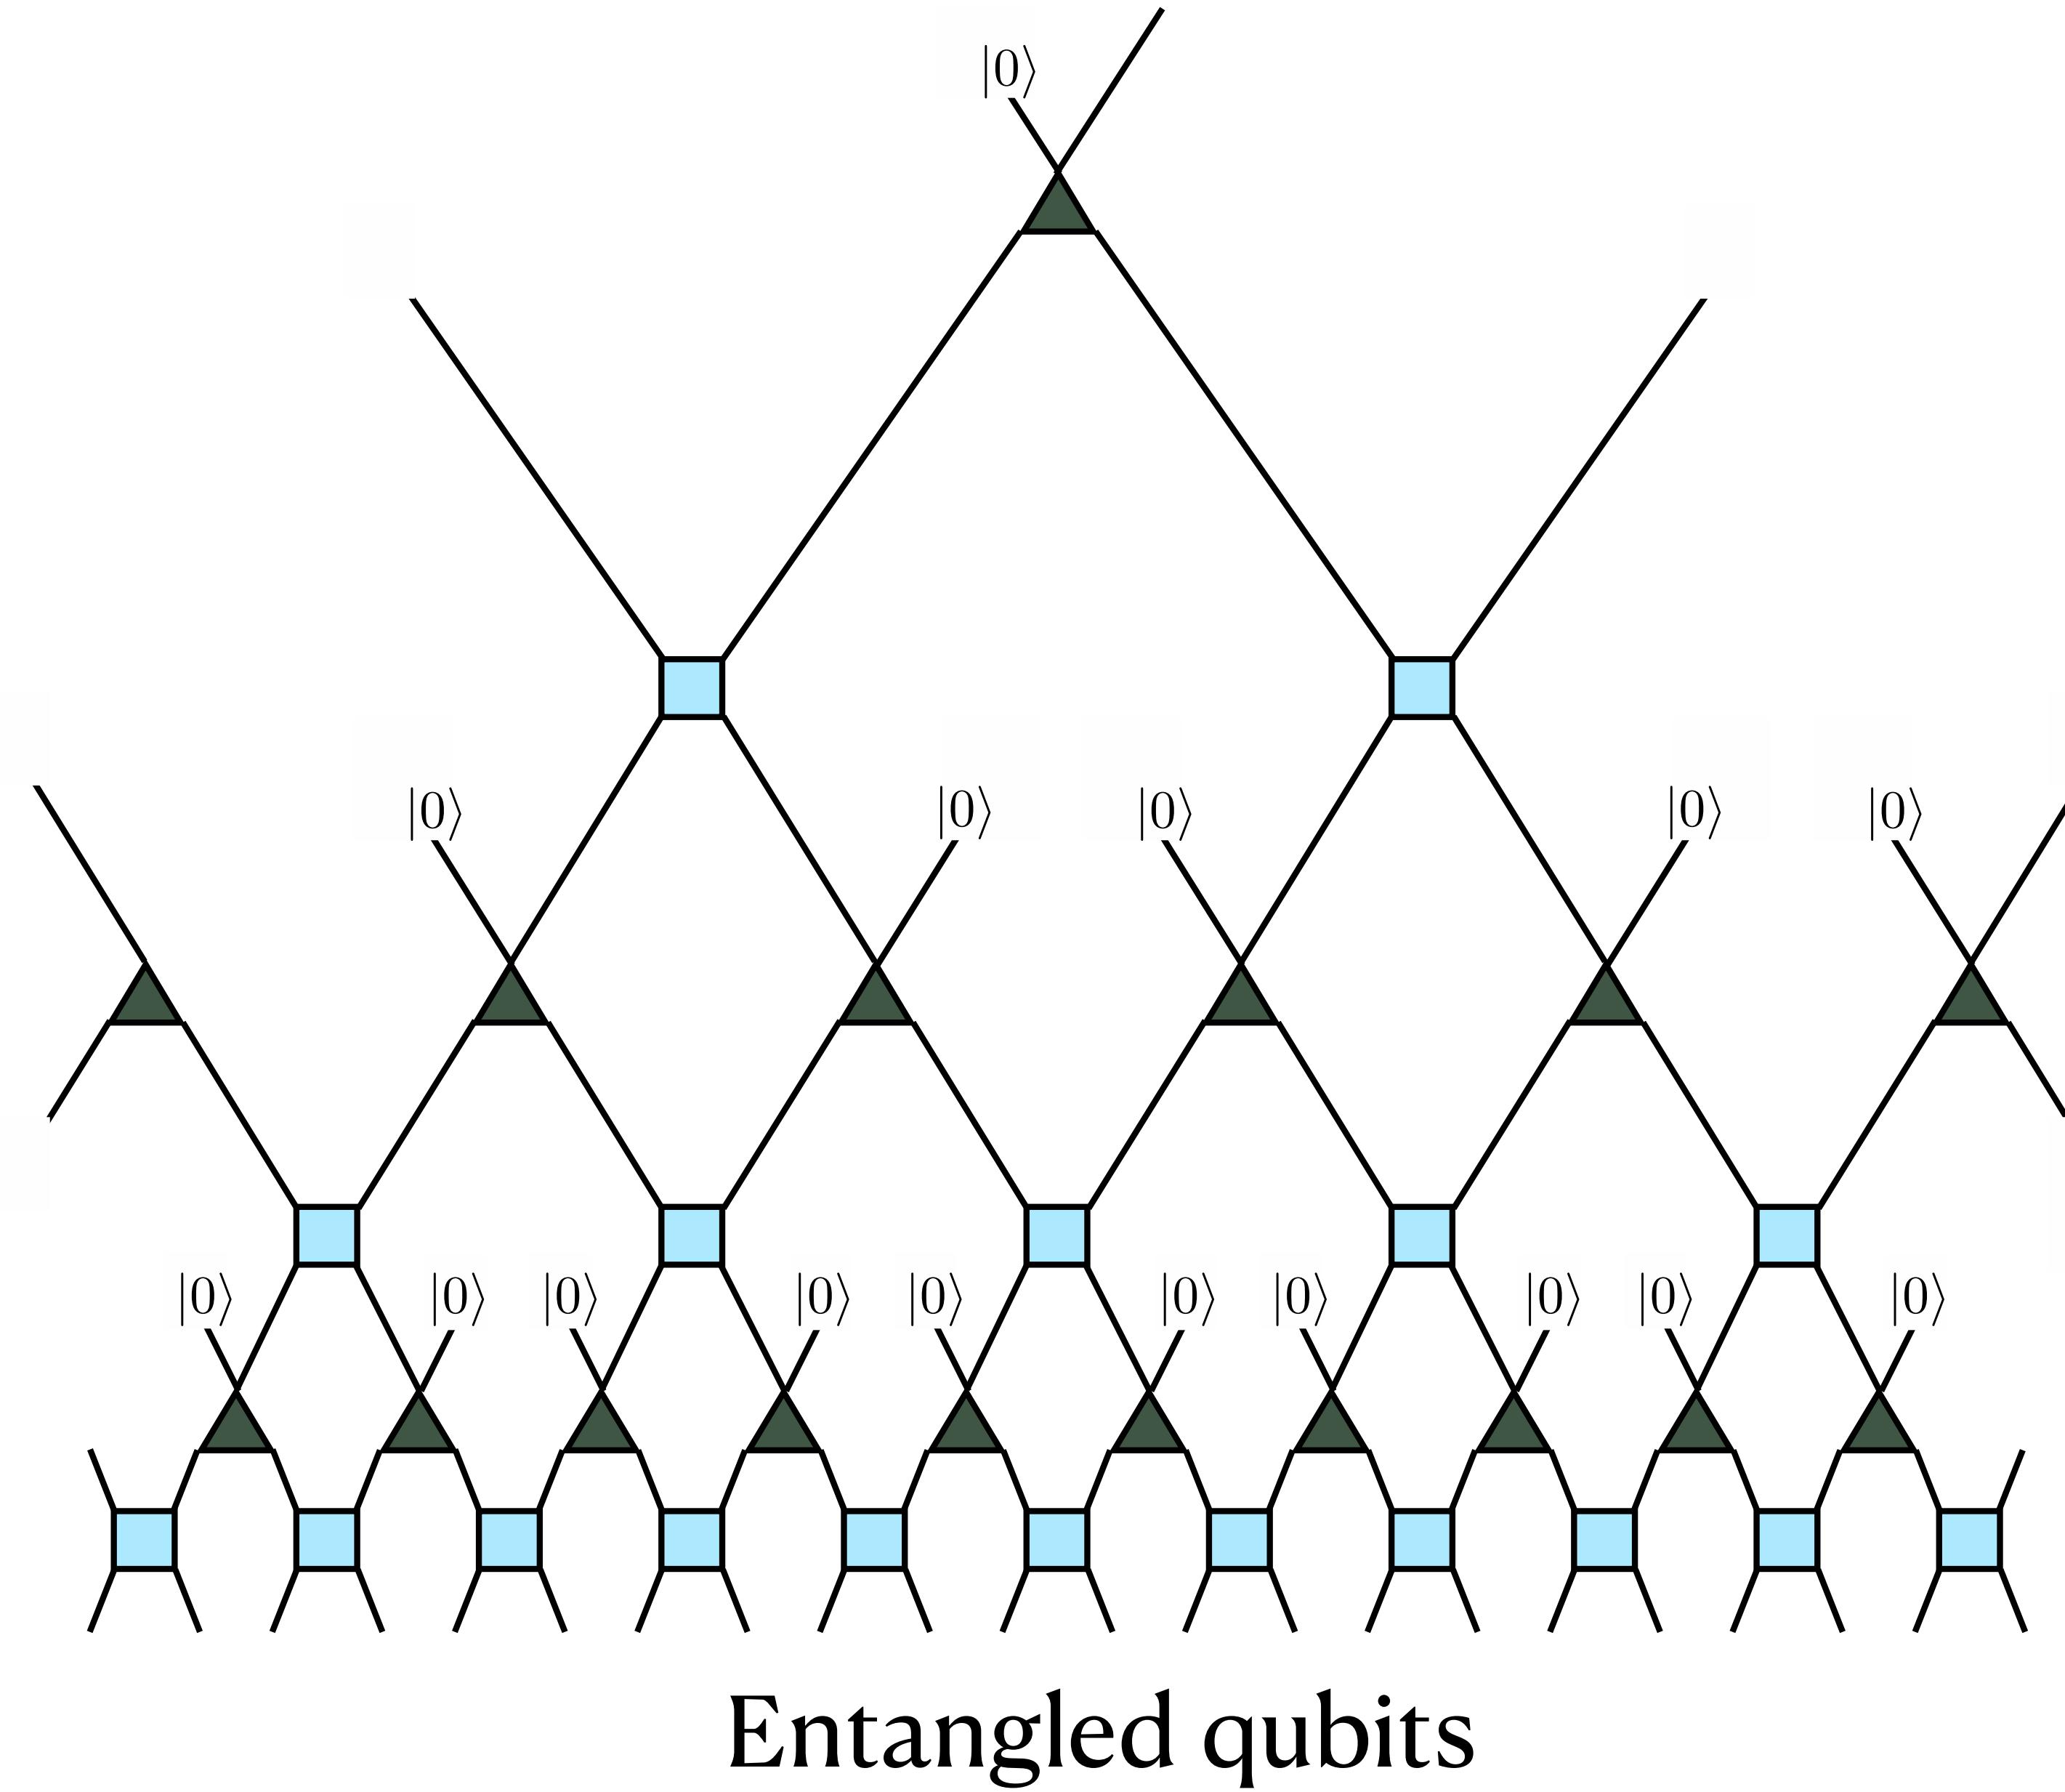
Li, LW, PRL '18
[lio12589/NeuralRG](https://github.com/lio12589/NeuralRG)

Neural network renormalization group

Li, LW, PRL '18
github.com/lio12589/NeuralRG

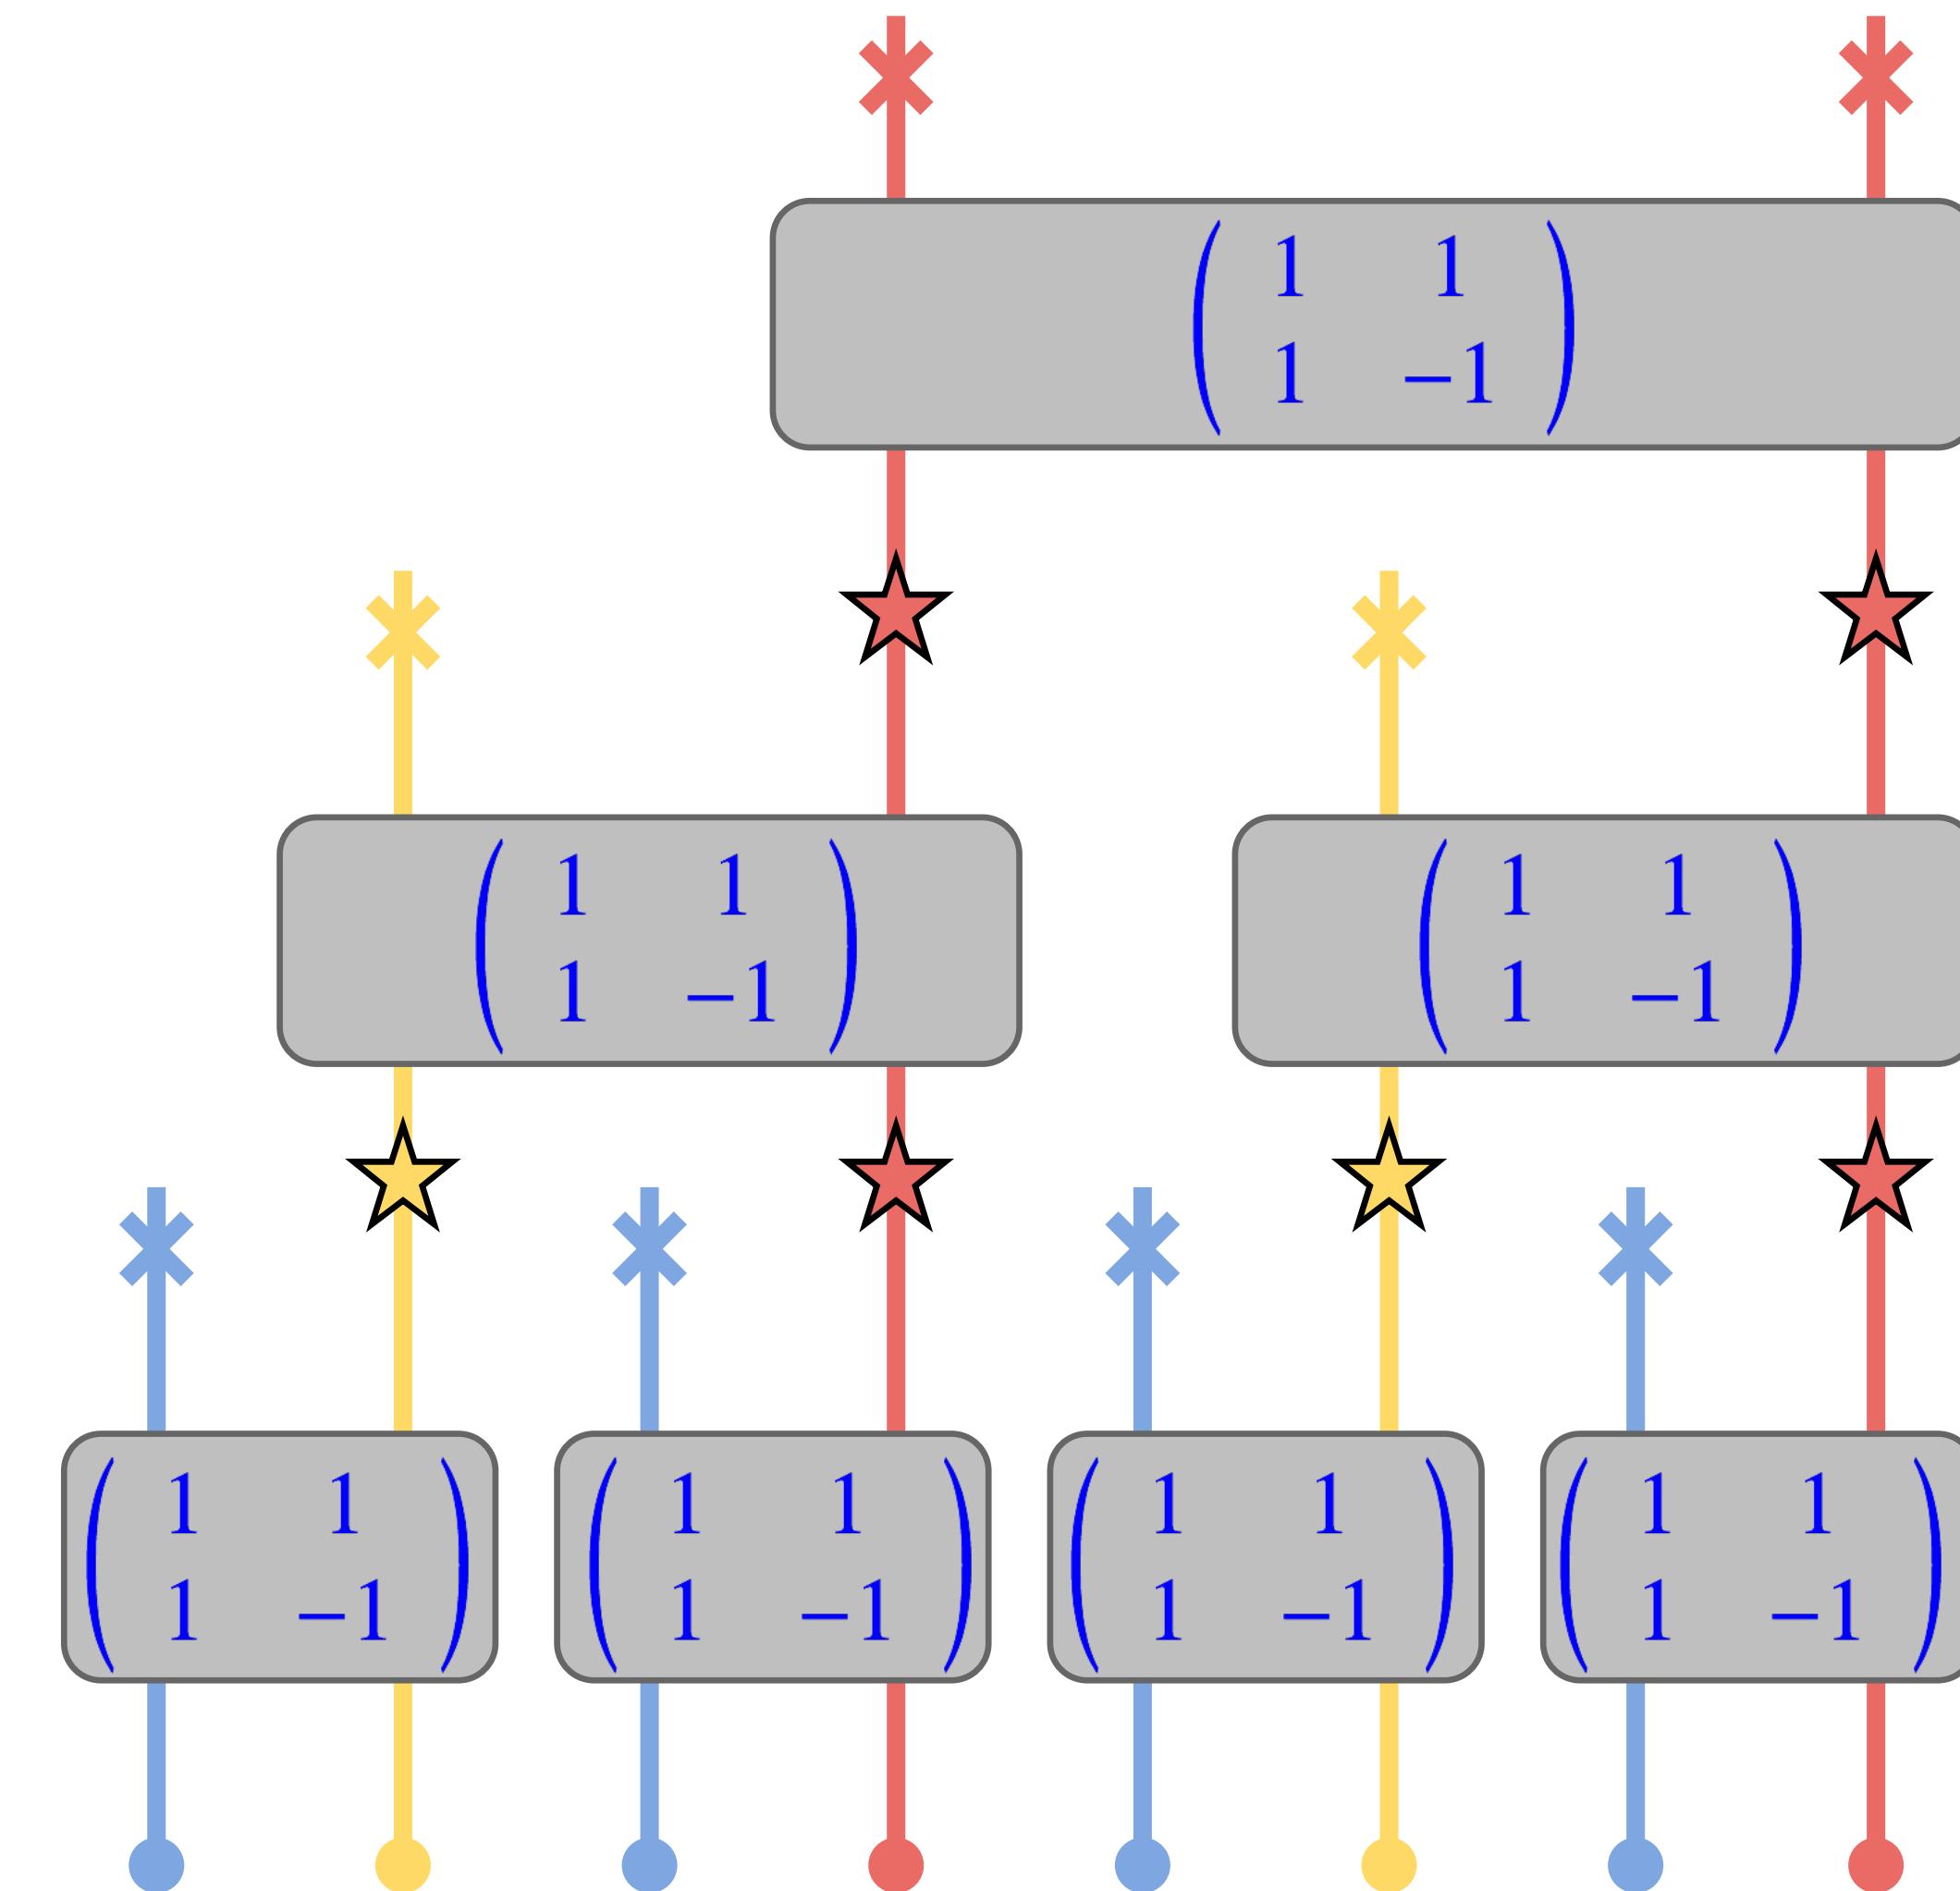
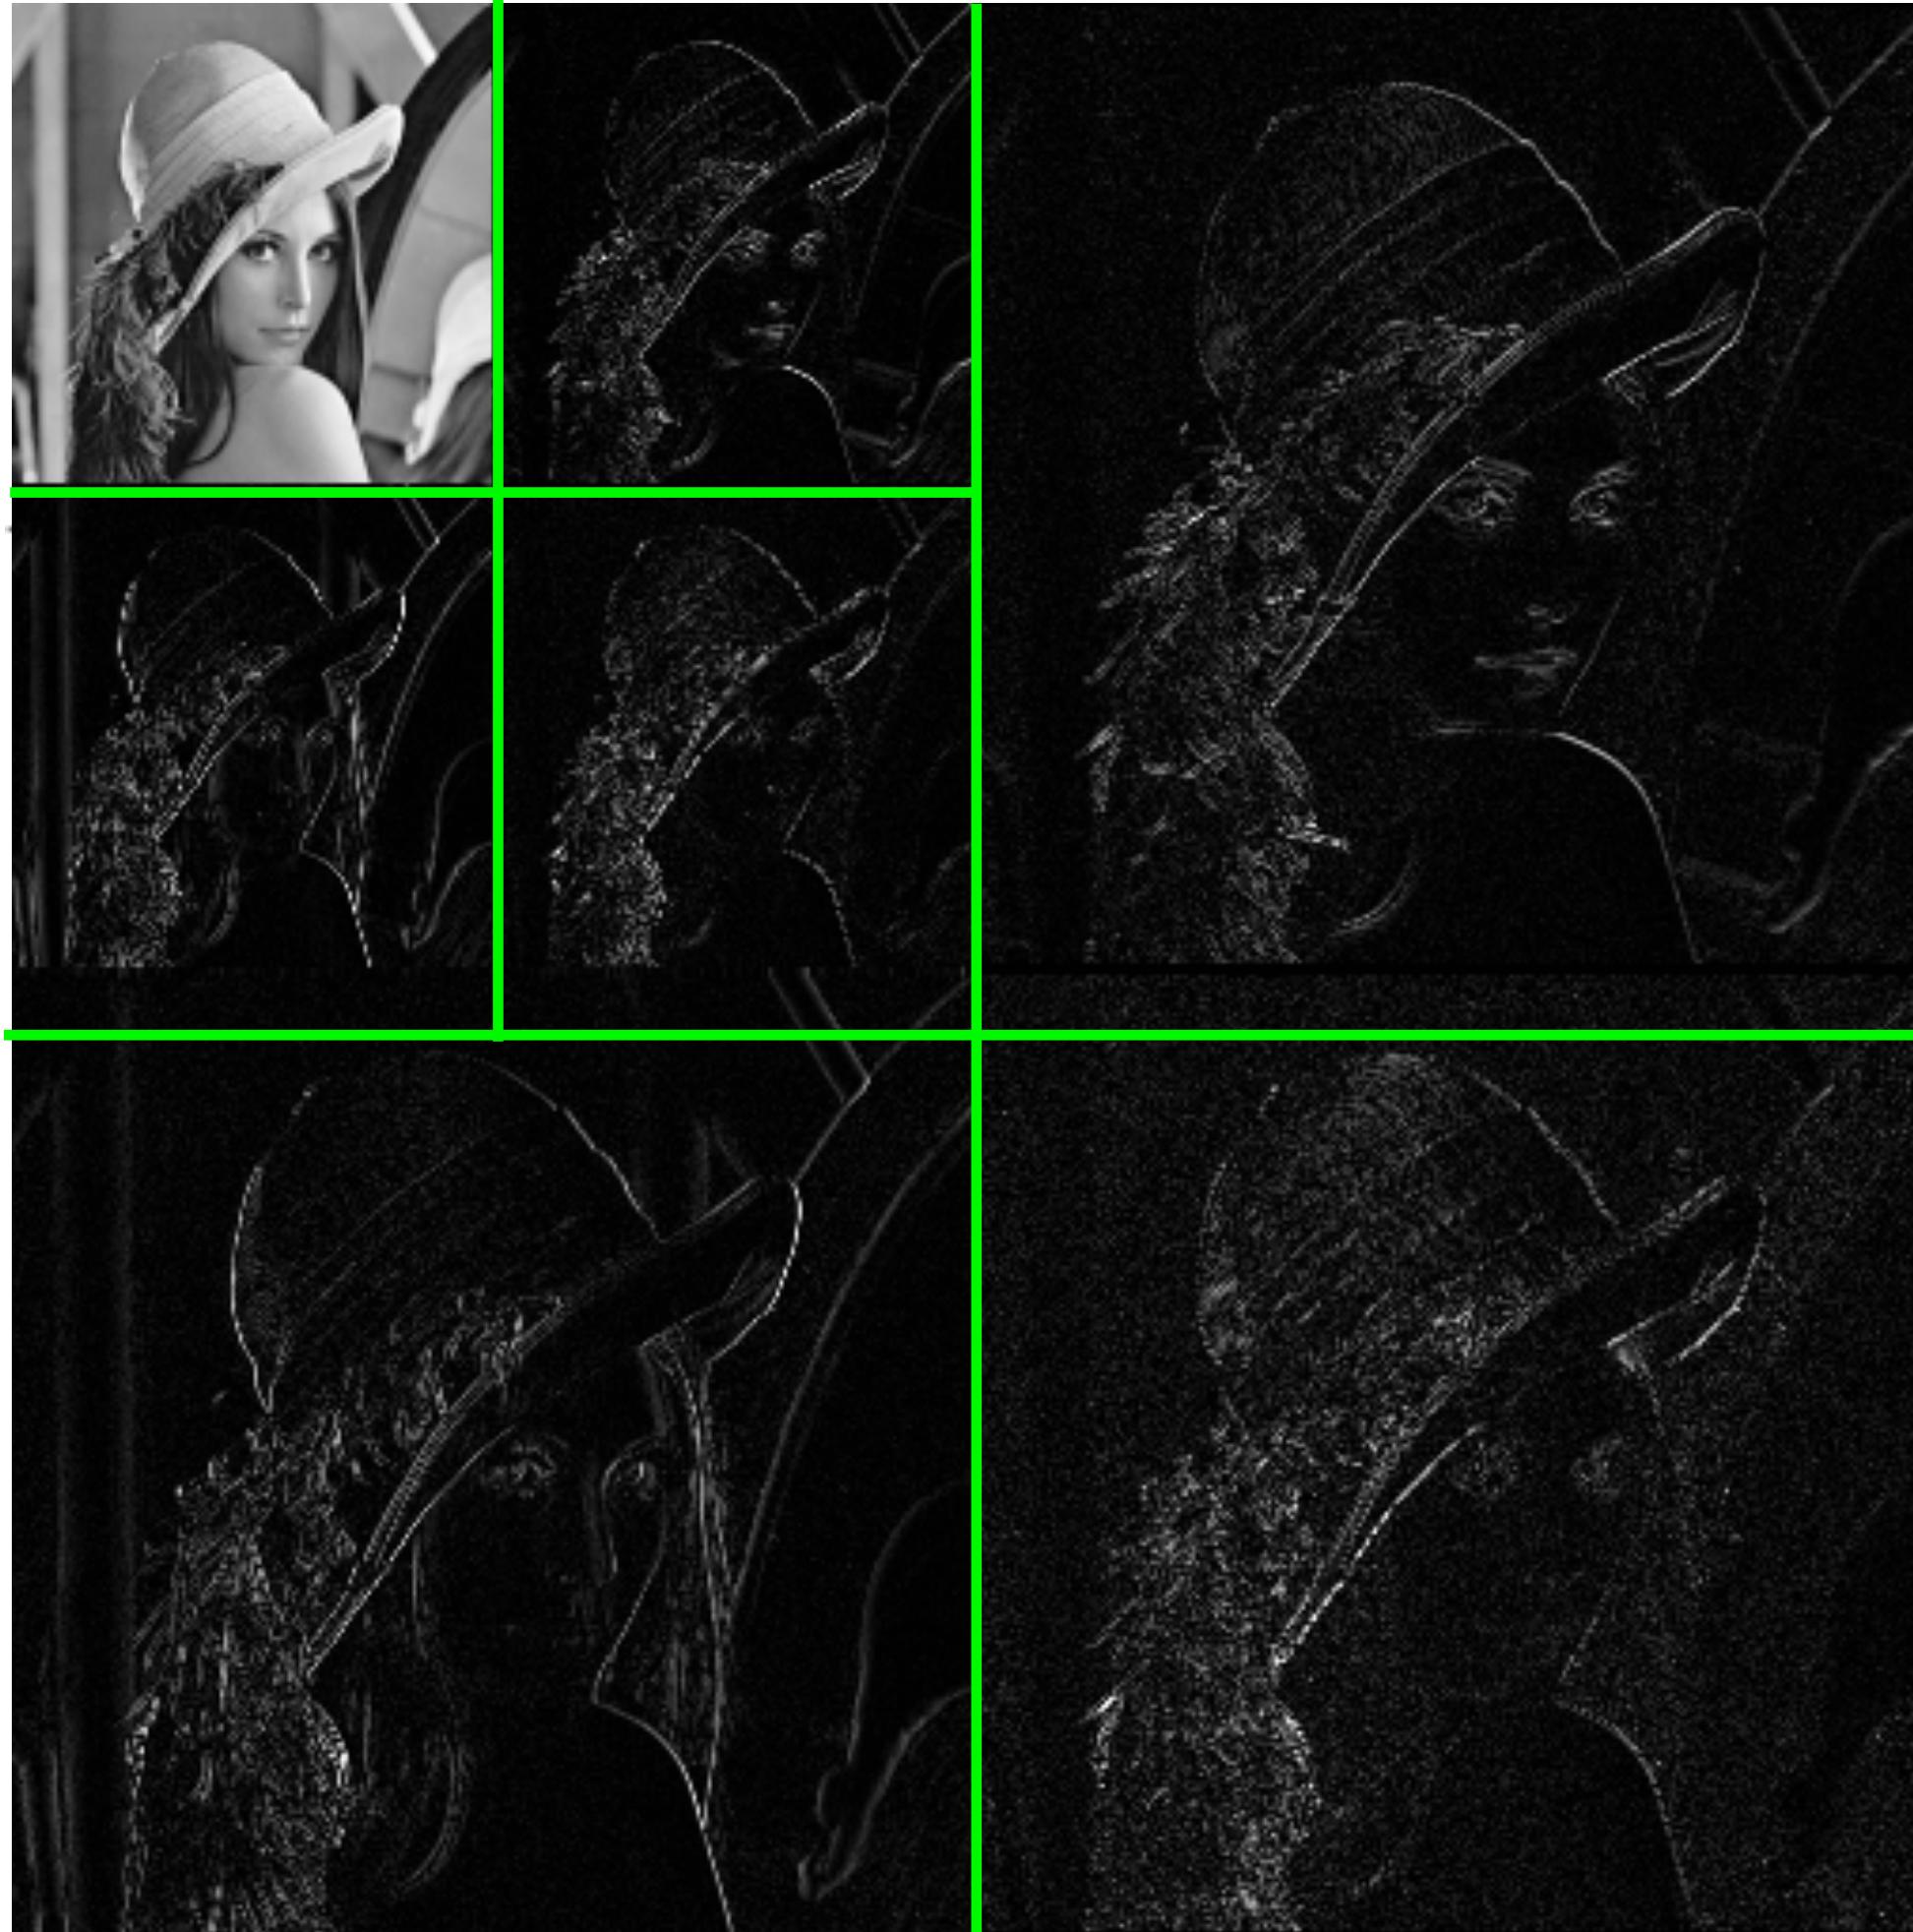


Quantum origin of the architecture



**Multi-Scale
Entanglement
Renormalization
Ansatz**

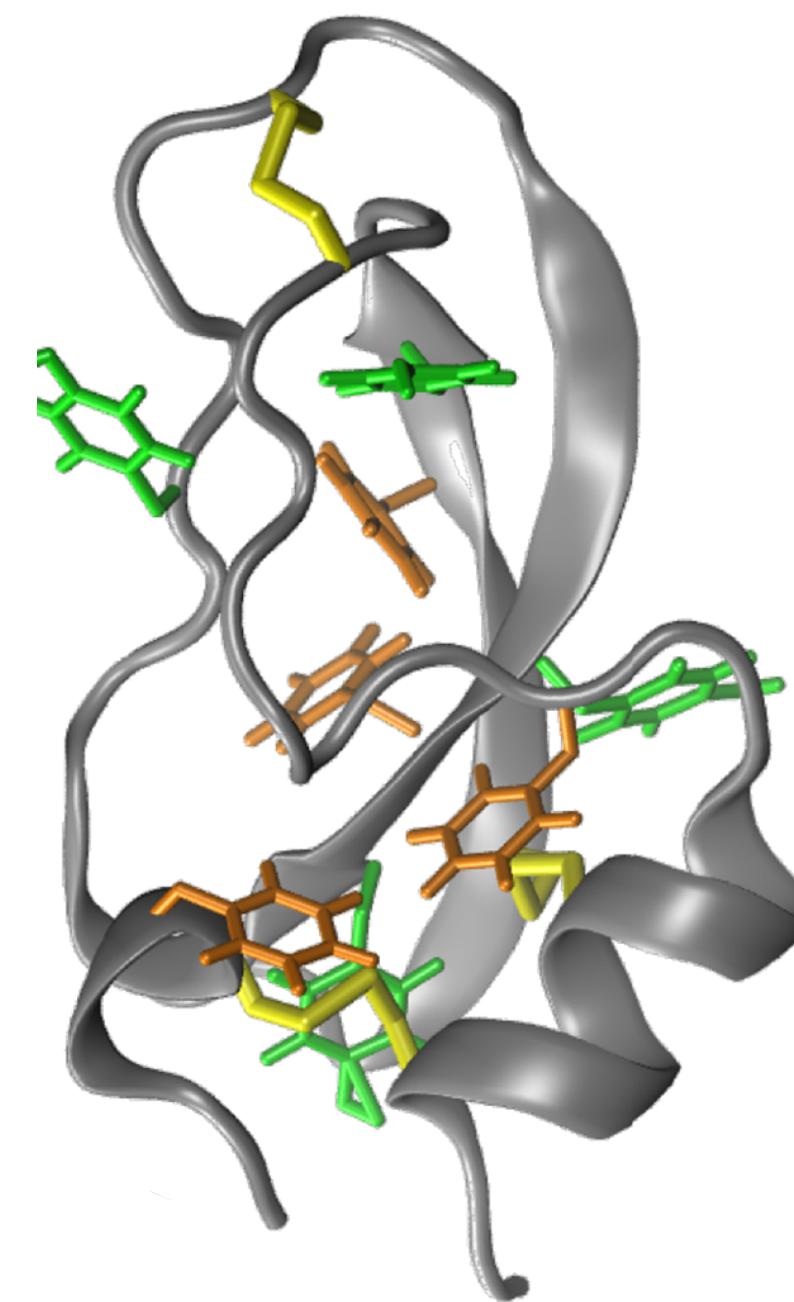
Connection to wavelets



Nonlinear & adaptive generalizations of wavelets

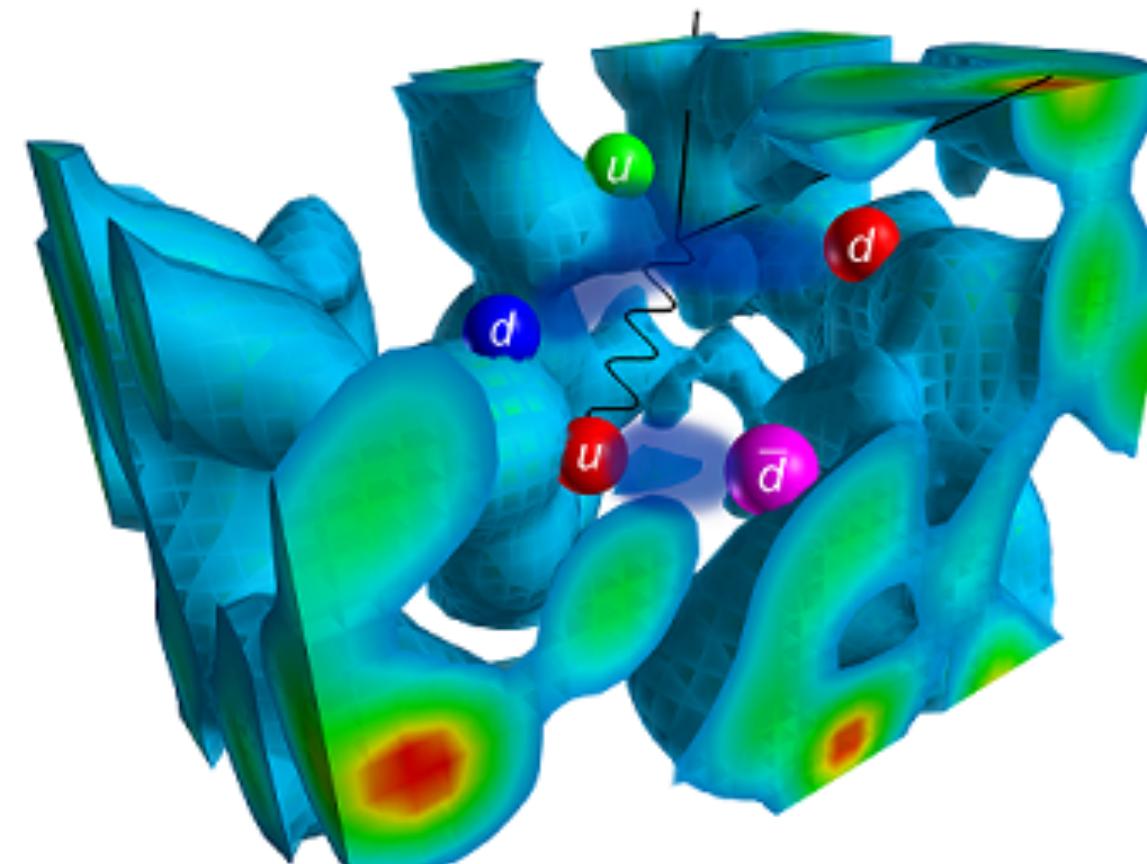
Normalizing flow in physics

Molecular simulation



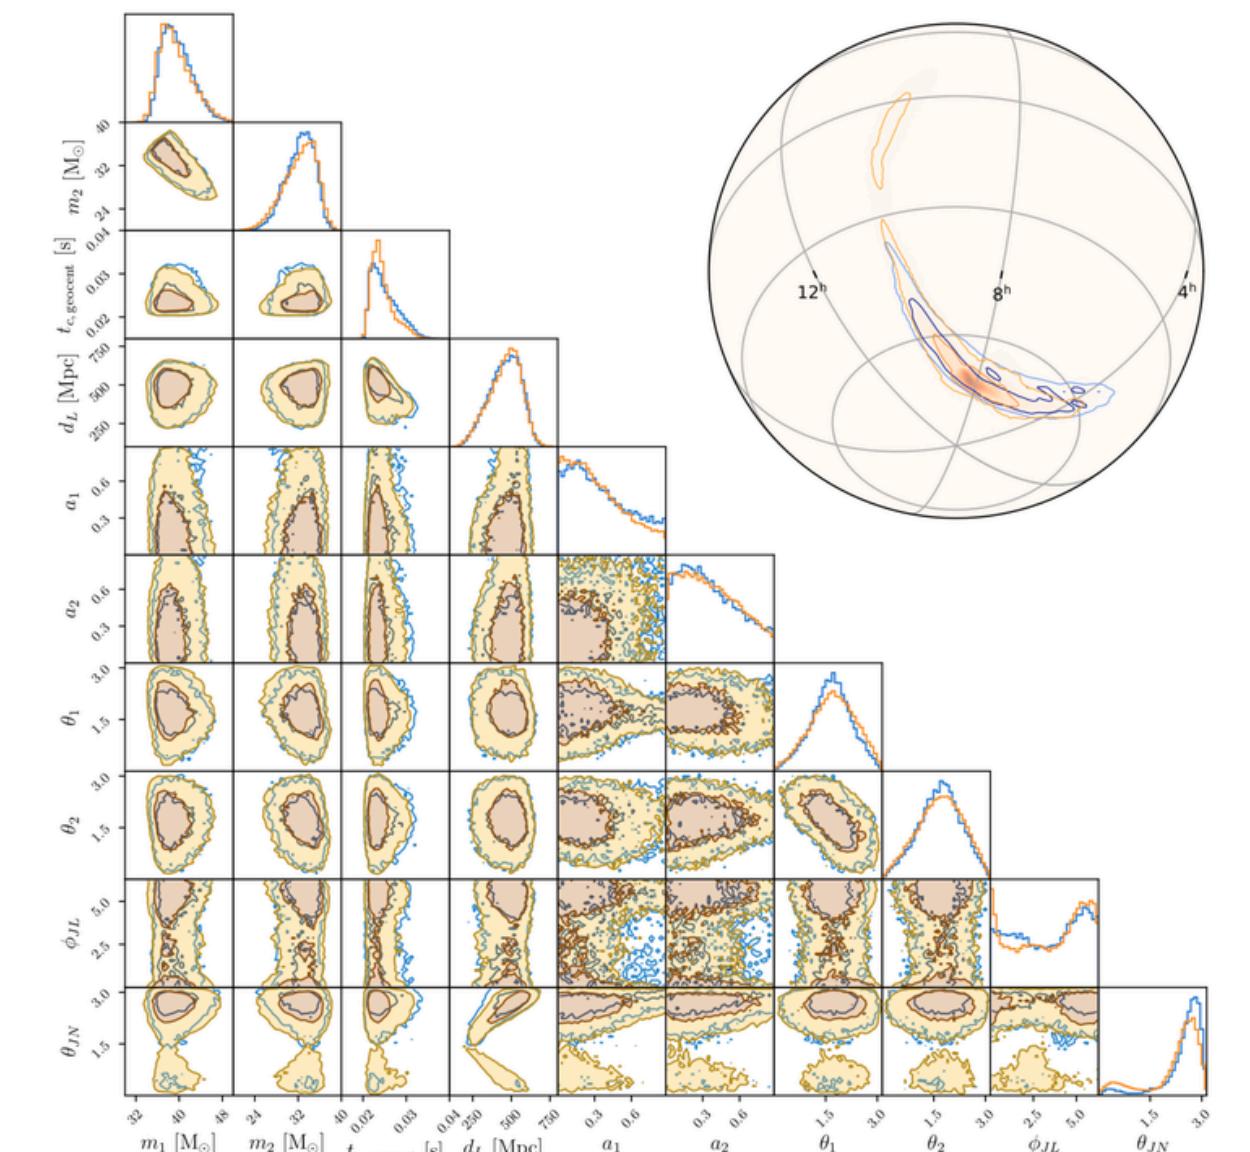
Noe et al, Science '19
Wirnsberger et al, JCP '20

Lattice field theory



Albergo et al, PRD '19
Kanwar et al, PRL '20

Gravitational wave detection



Green et al, MLST '21
Dex et al, PRL '21

Continuous normalizing flows

$$\ln p(\mathbf{x}) = \ln \mathcal{N}(\mathbf{z}) - \ln \left| \det \left(\frac{\partial \mathbf{x}}{\partial \mathbf{z}} \right) \right|$$

Consider infinitesimal change-of-variables Chen et al 1806.07366

$$\mathbf{x} = \mathbf{z} + \varepsilon \boldsymbol{\nu}$$

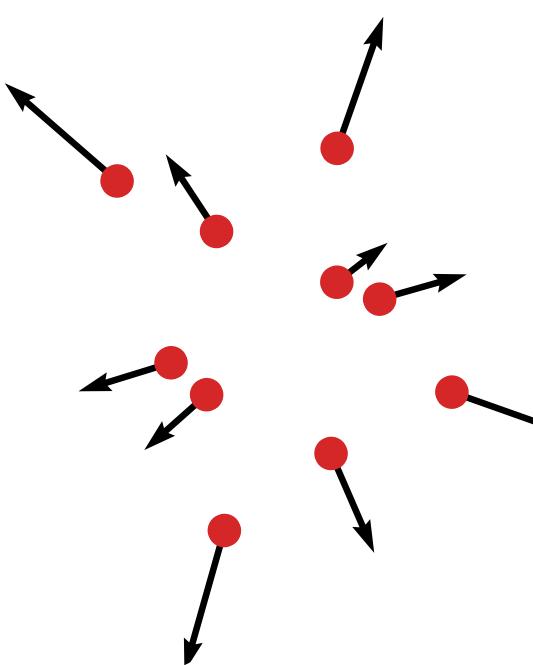
$$\varepsilon \rightarrow 0$$

$$\frac{d\mathbf{x}}{dt} = \boldsymbol{\nu}$$

$$\ln p(\mathbf{x}) - \ln \mathcal{N}(\mathbf{z}) = - \ln \left| \det \left(1 + \varepsilon \frac{\partial \boldsymbol{\nu}}{\partial \mathbf{z}} \right) \right|$$
$$\frac{d \ln \rho(\mathbf{x}, t)}{dt} = - \nabla \cdot \boldsymbol{\nu}$$

Fluid physics behind flows

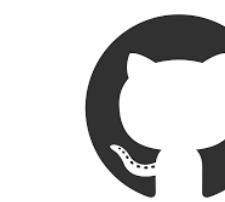
$$\frac{dx}{dt} = v$$



$$\frac{d \ln \rho(x, t)}{dt} = - \nabla \cdot v$$

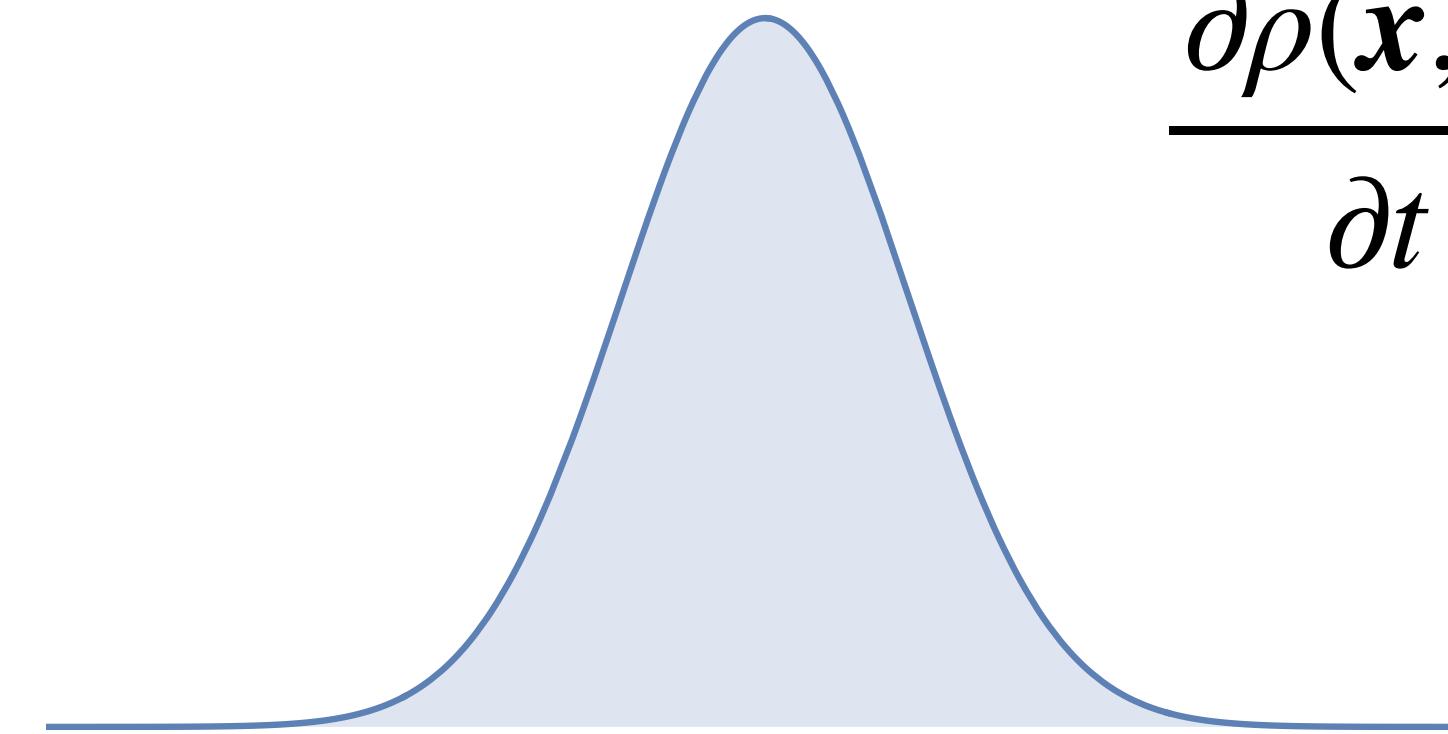
$$\frac{d}{dt} = \frac{\partial}{\partial t} + v \cdot \nabla$$

“material derivative”



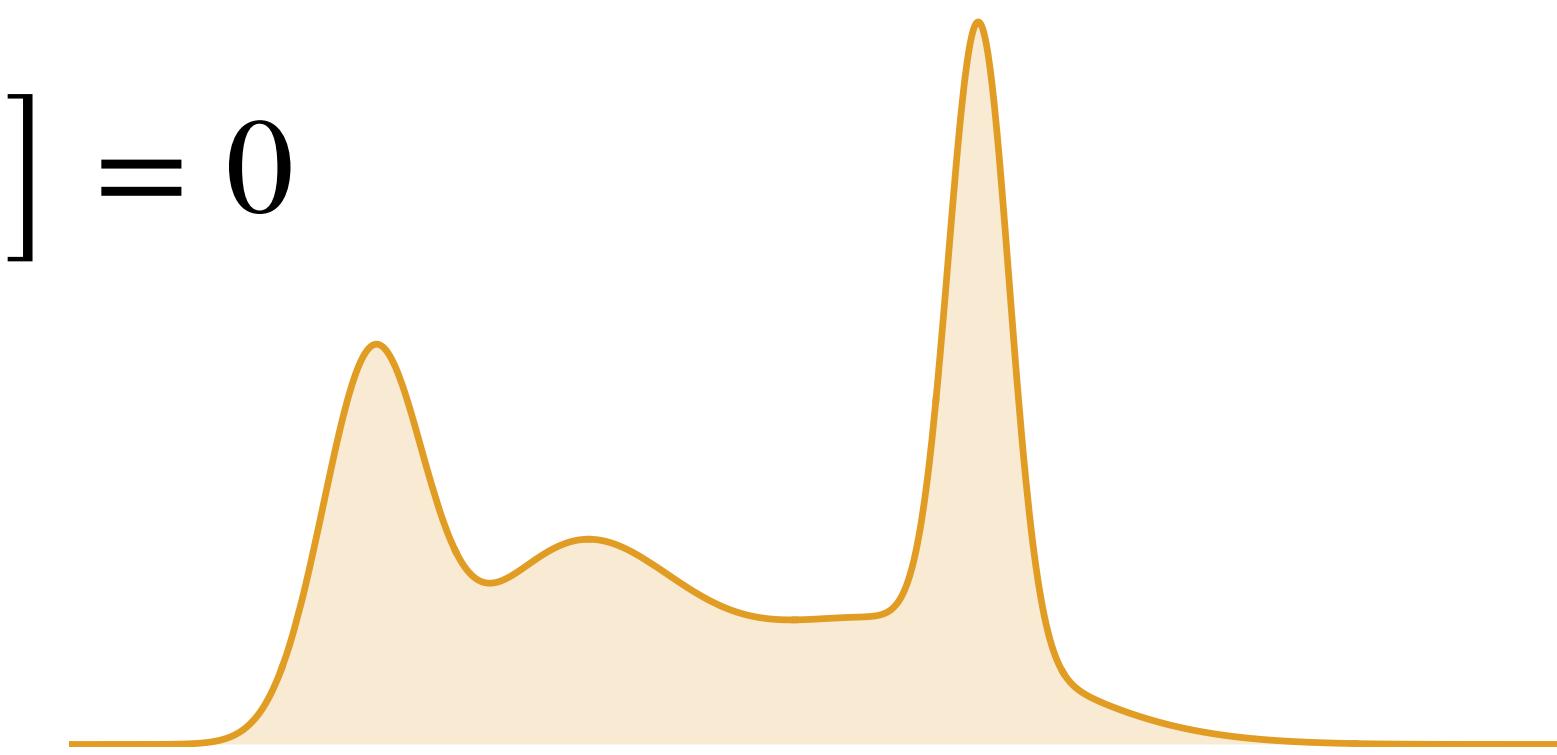
Zhang, E, LW 1809.10188

[wangleiphy/MongeAmpereFlow](https://github.com/wangleiphy/MongeAmpereFlow)



Simple density

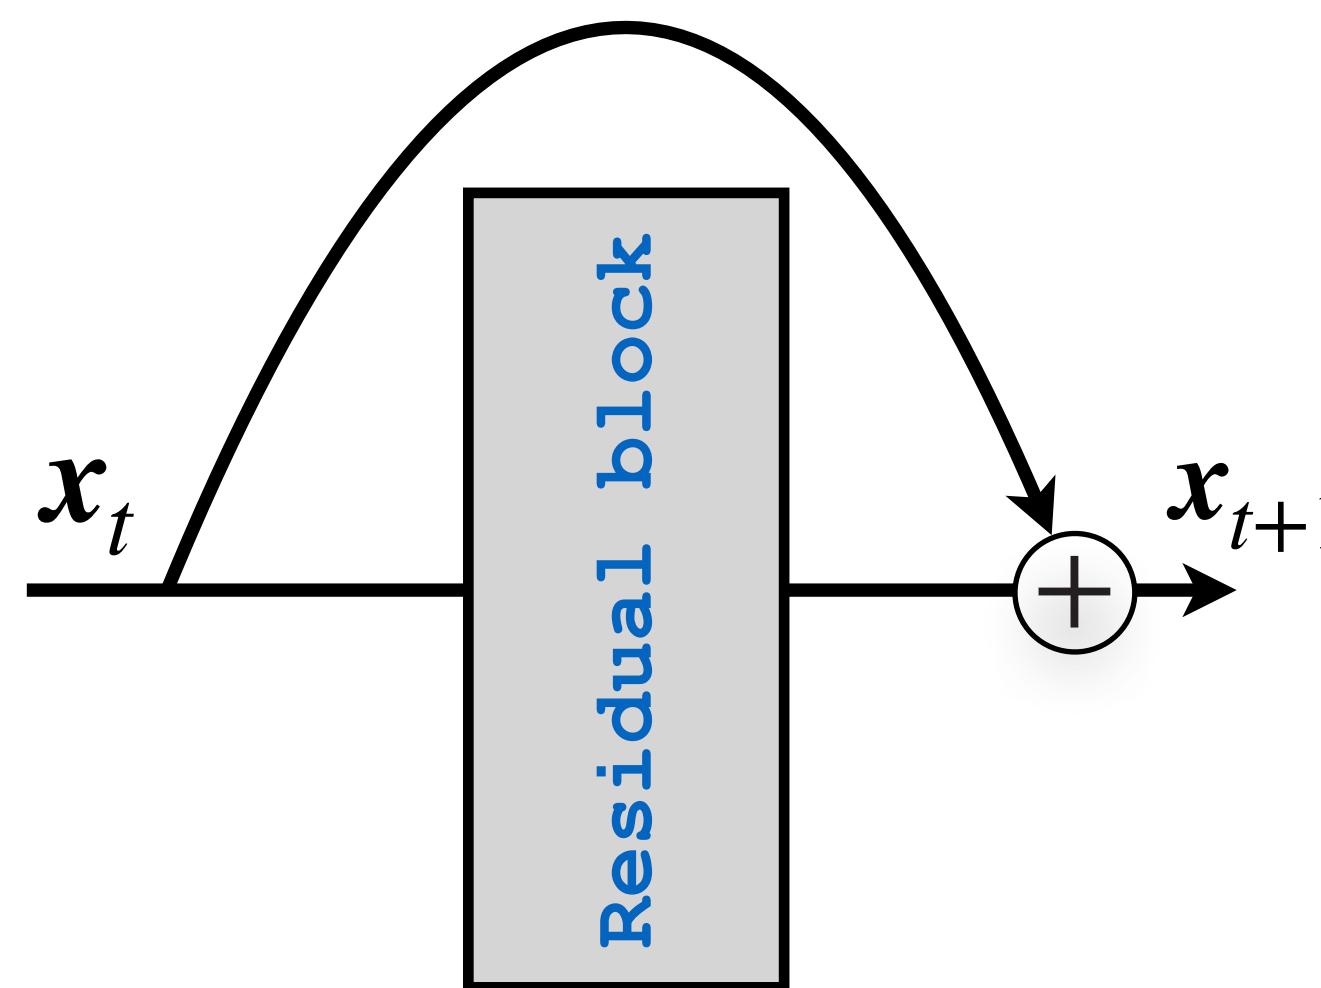
$$\frac{\partial \rho(x, t)}{\partial t} + \nabla \cdot [\rho(x, t)v] = 0$$



Complex density

Neural Ordinary Differential Equations

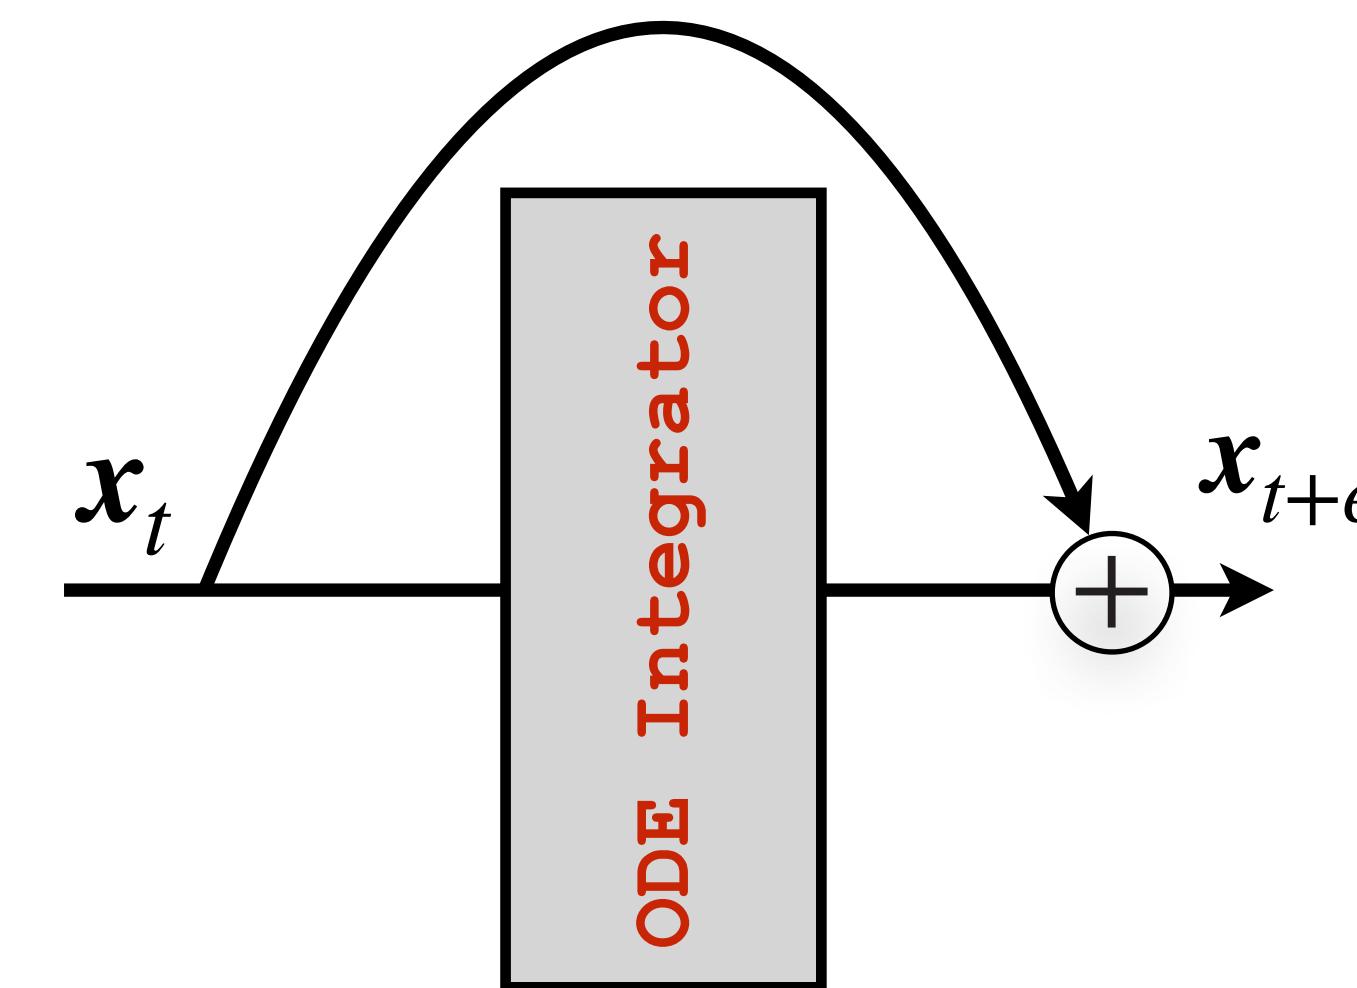
Residual network



$$x_{t+1} = x_t + v(x_t)$$

Chen et al, 1806.07366

ODE integration

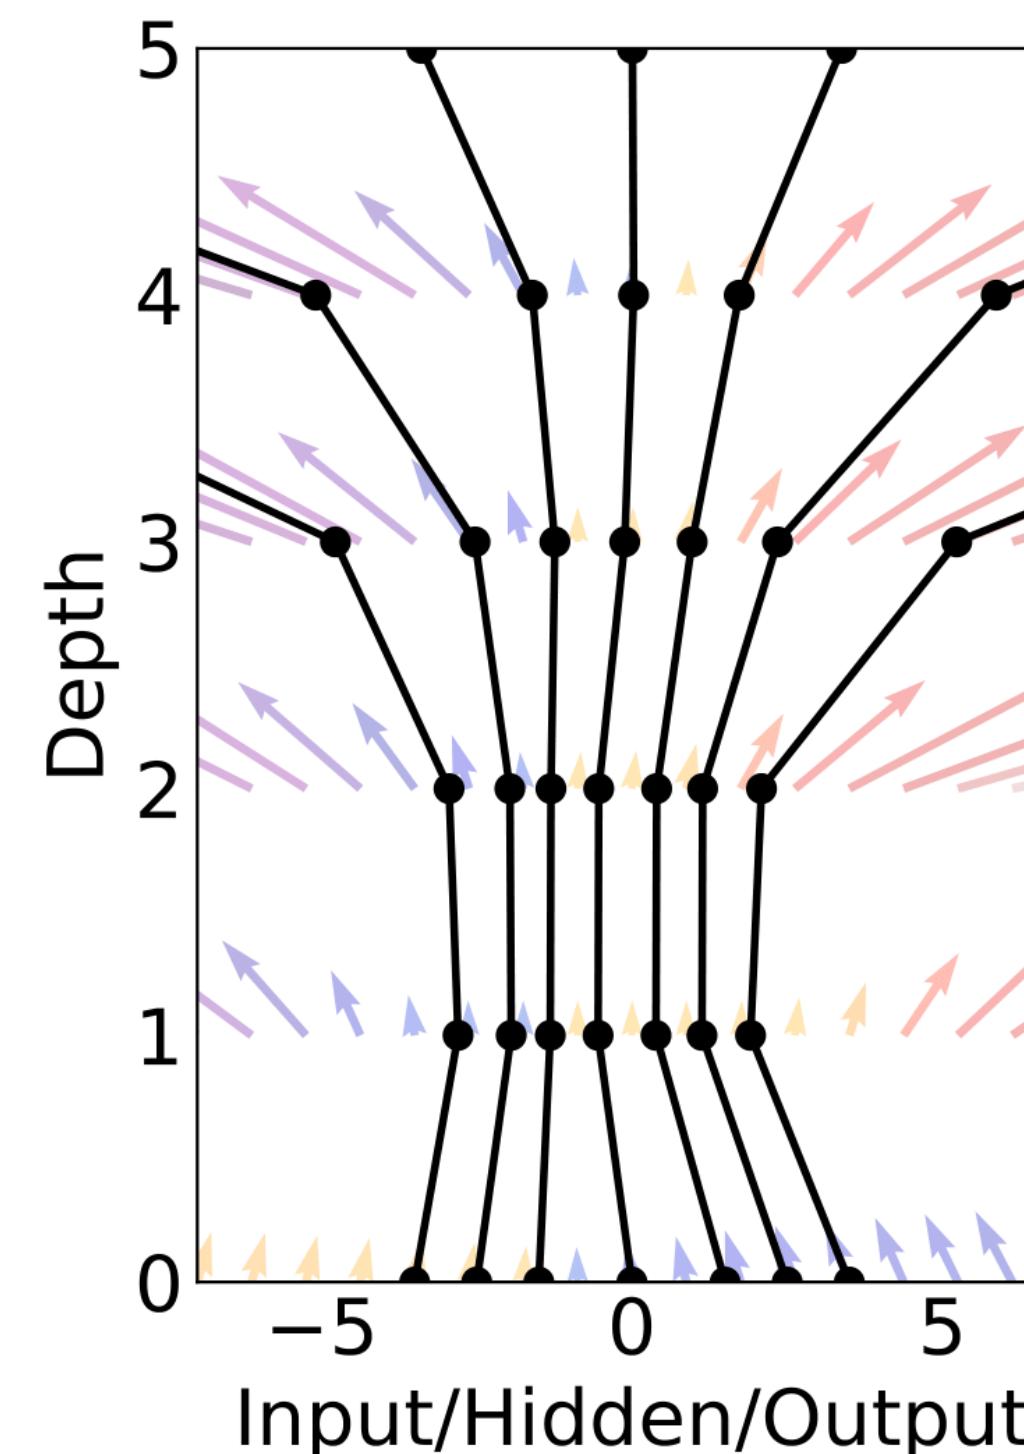


$$dx/dt = v(x)$$

Harbor et al 1705.03341
Lu et al 1710.10121,
E Commun. Math. Stat 17'

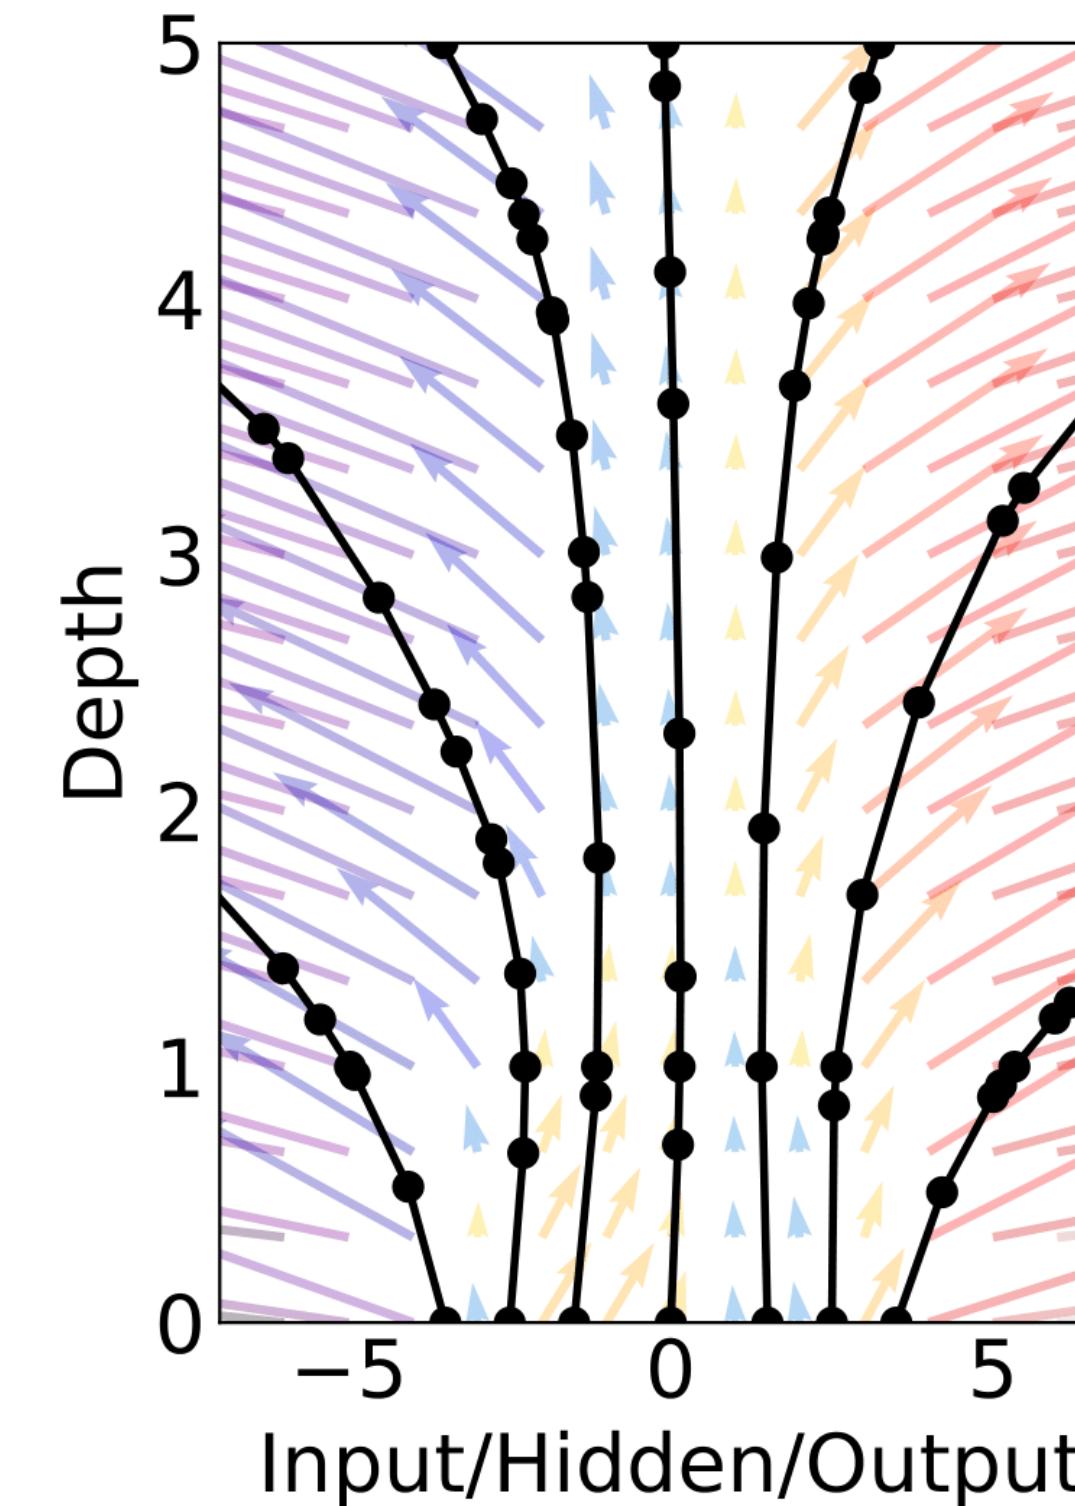
Neural Ordinary Differential Equations

Residual network



$$x_{t+1} = x_t + v(x_t)$$

ODE integration



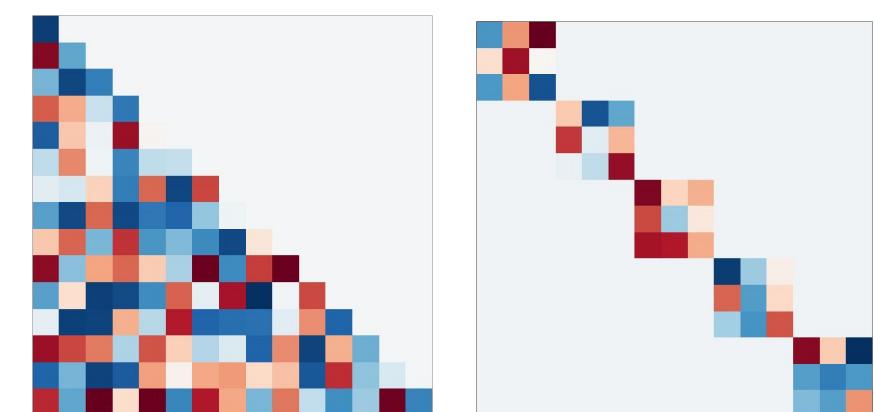
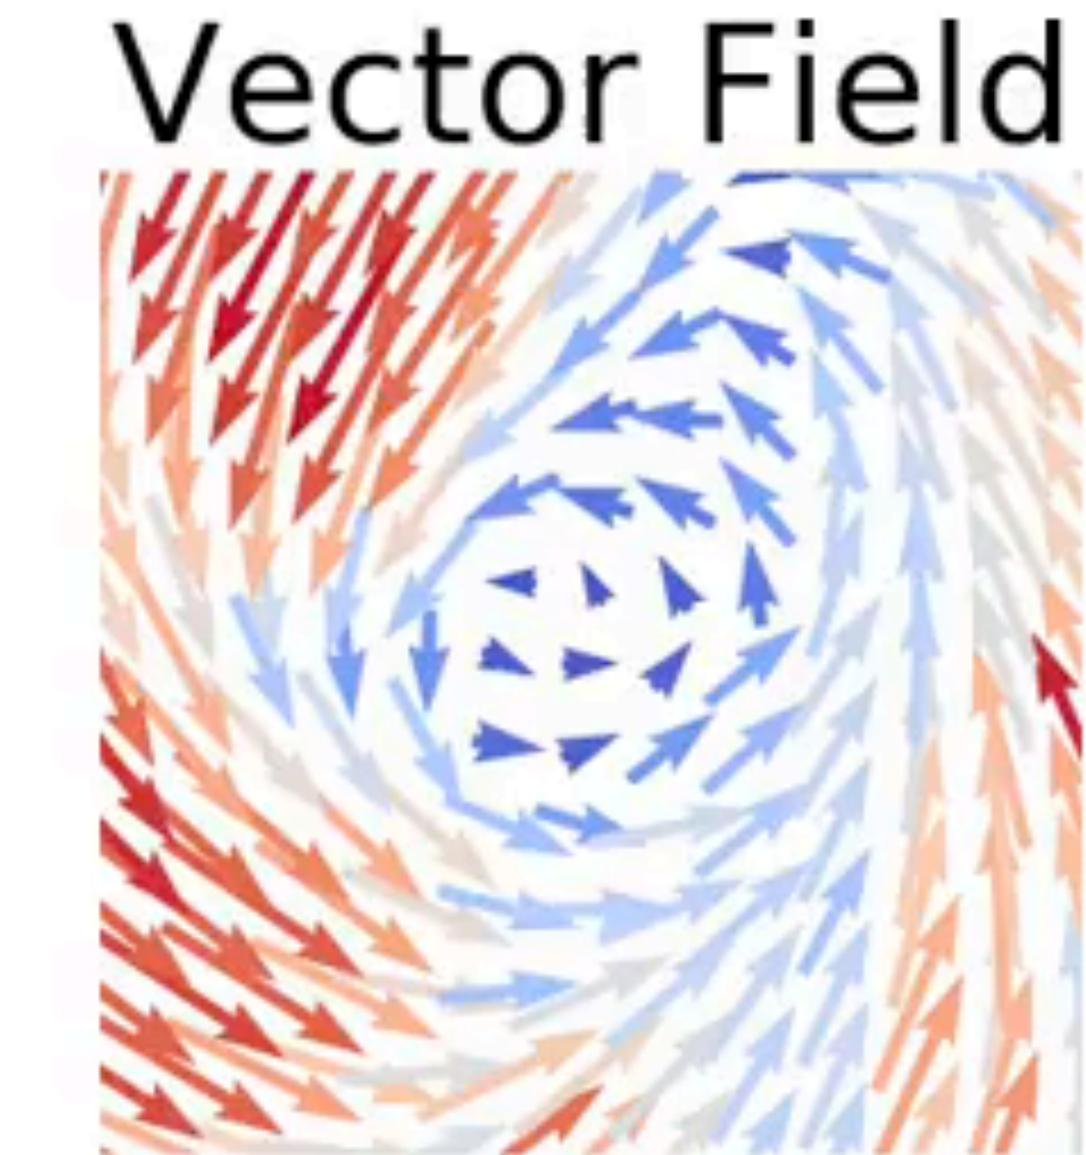
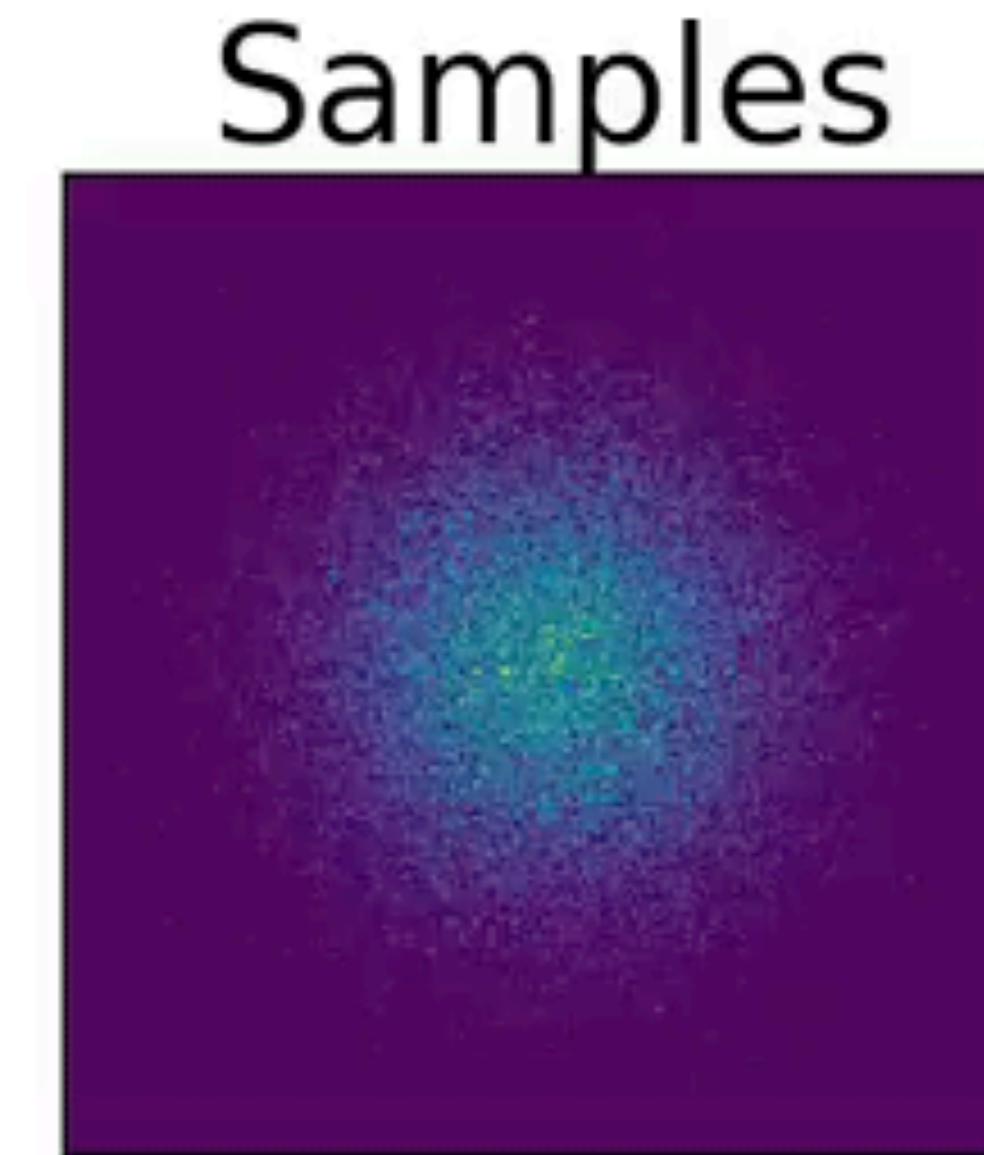
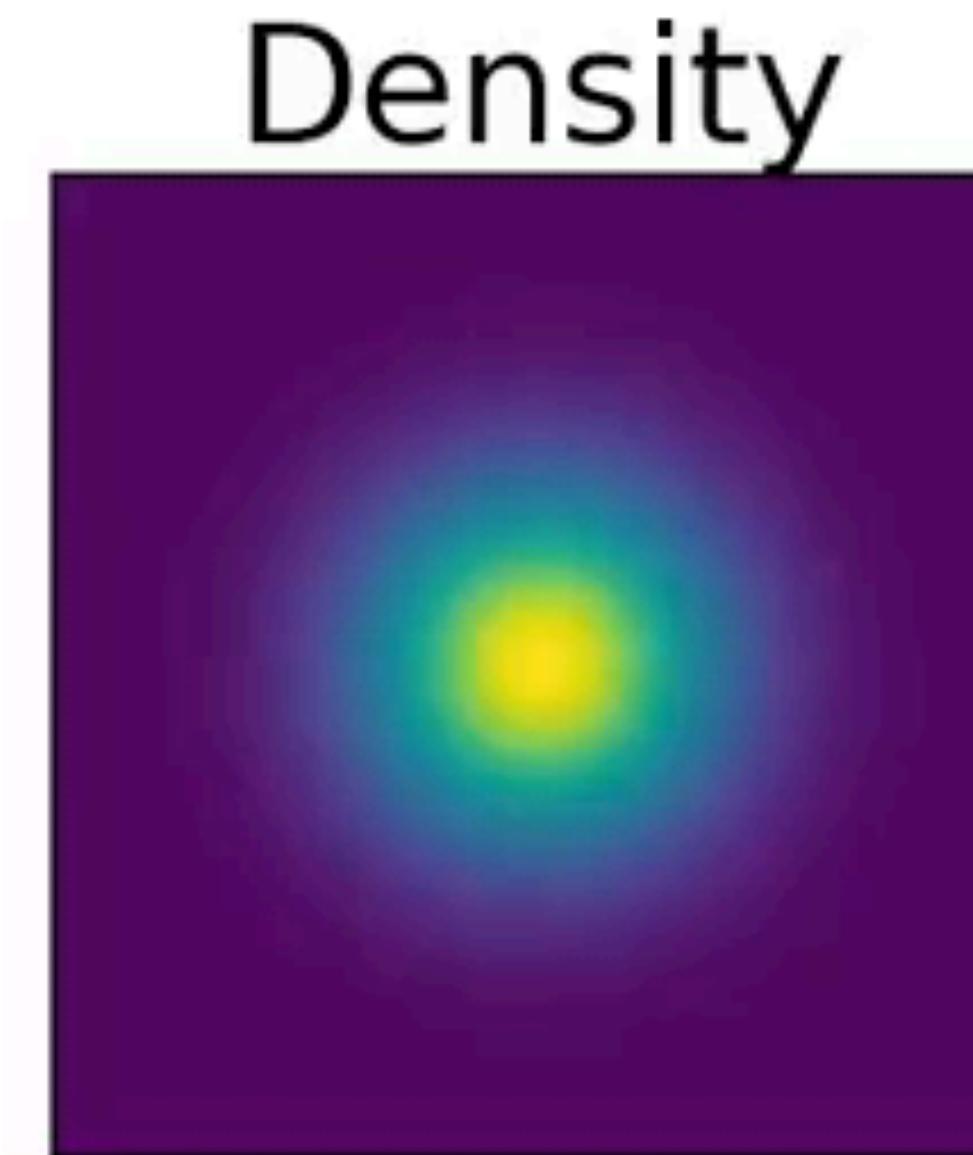
$$dx/dt = v(x)$$

Chen et al, 1806.07366

Harbor et al 1705.03341
Lu et al 1710.10121,
E Commun. Math. Stat 17'...

Continuous normalizing flows implemented with NeuralODE

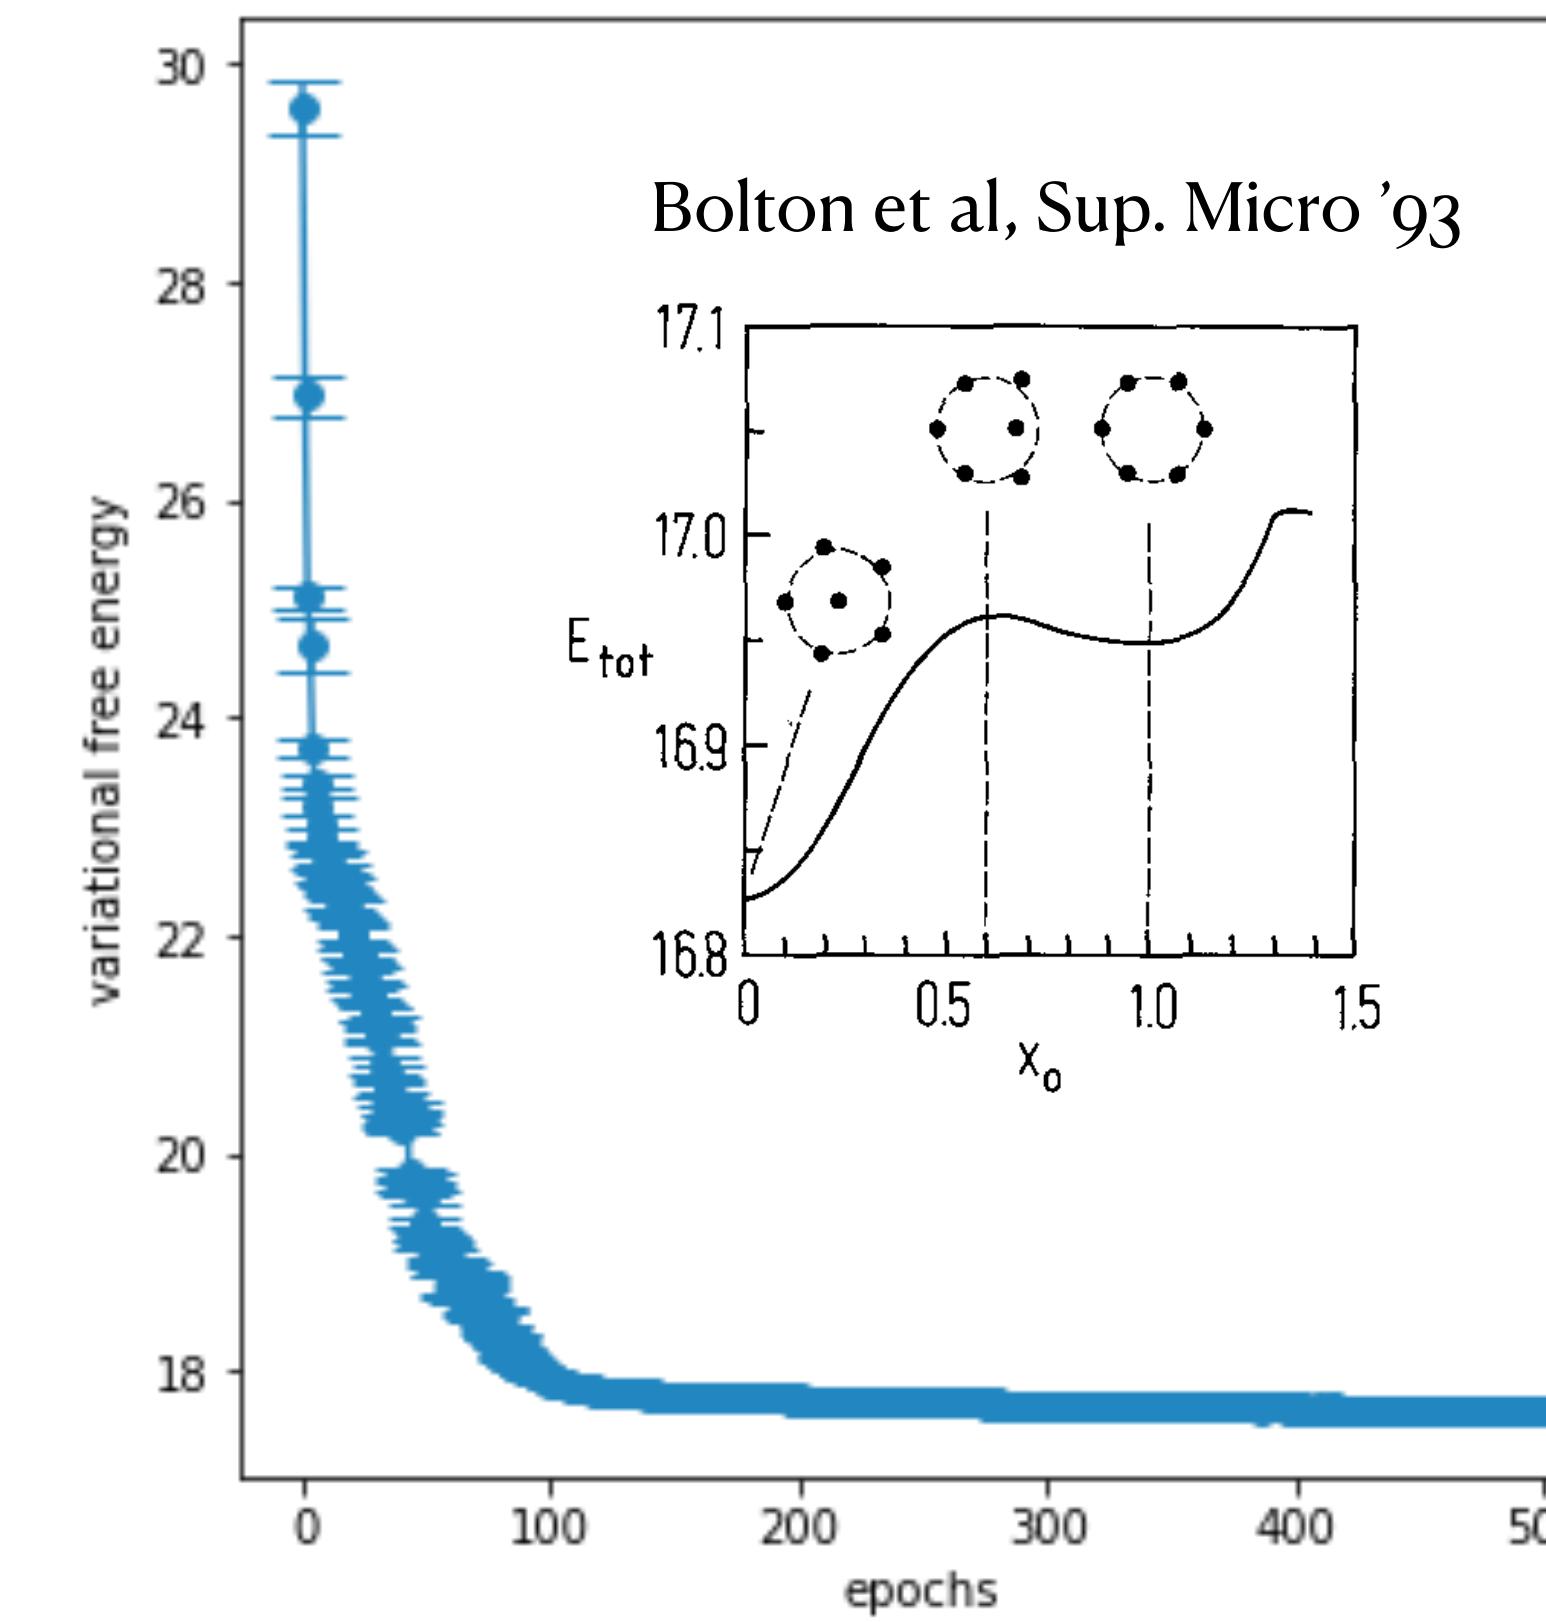
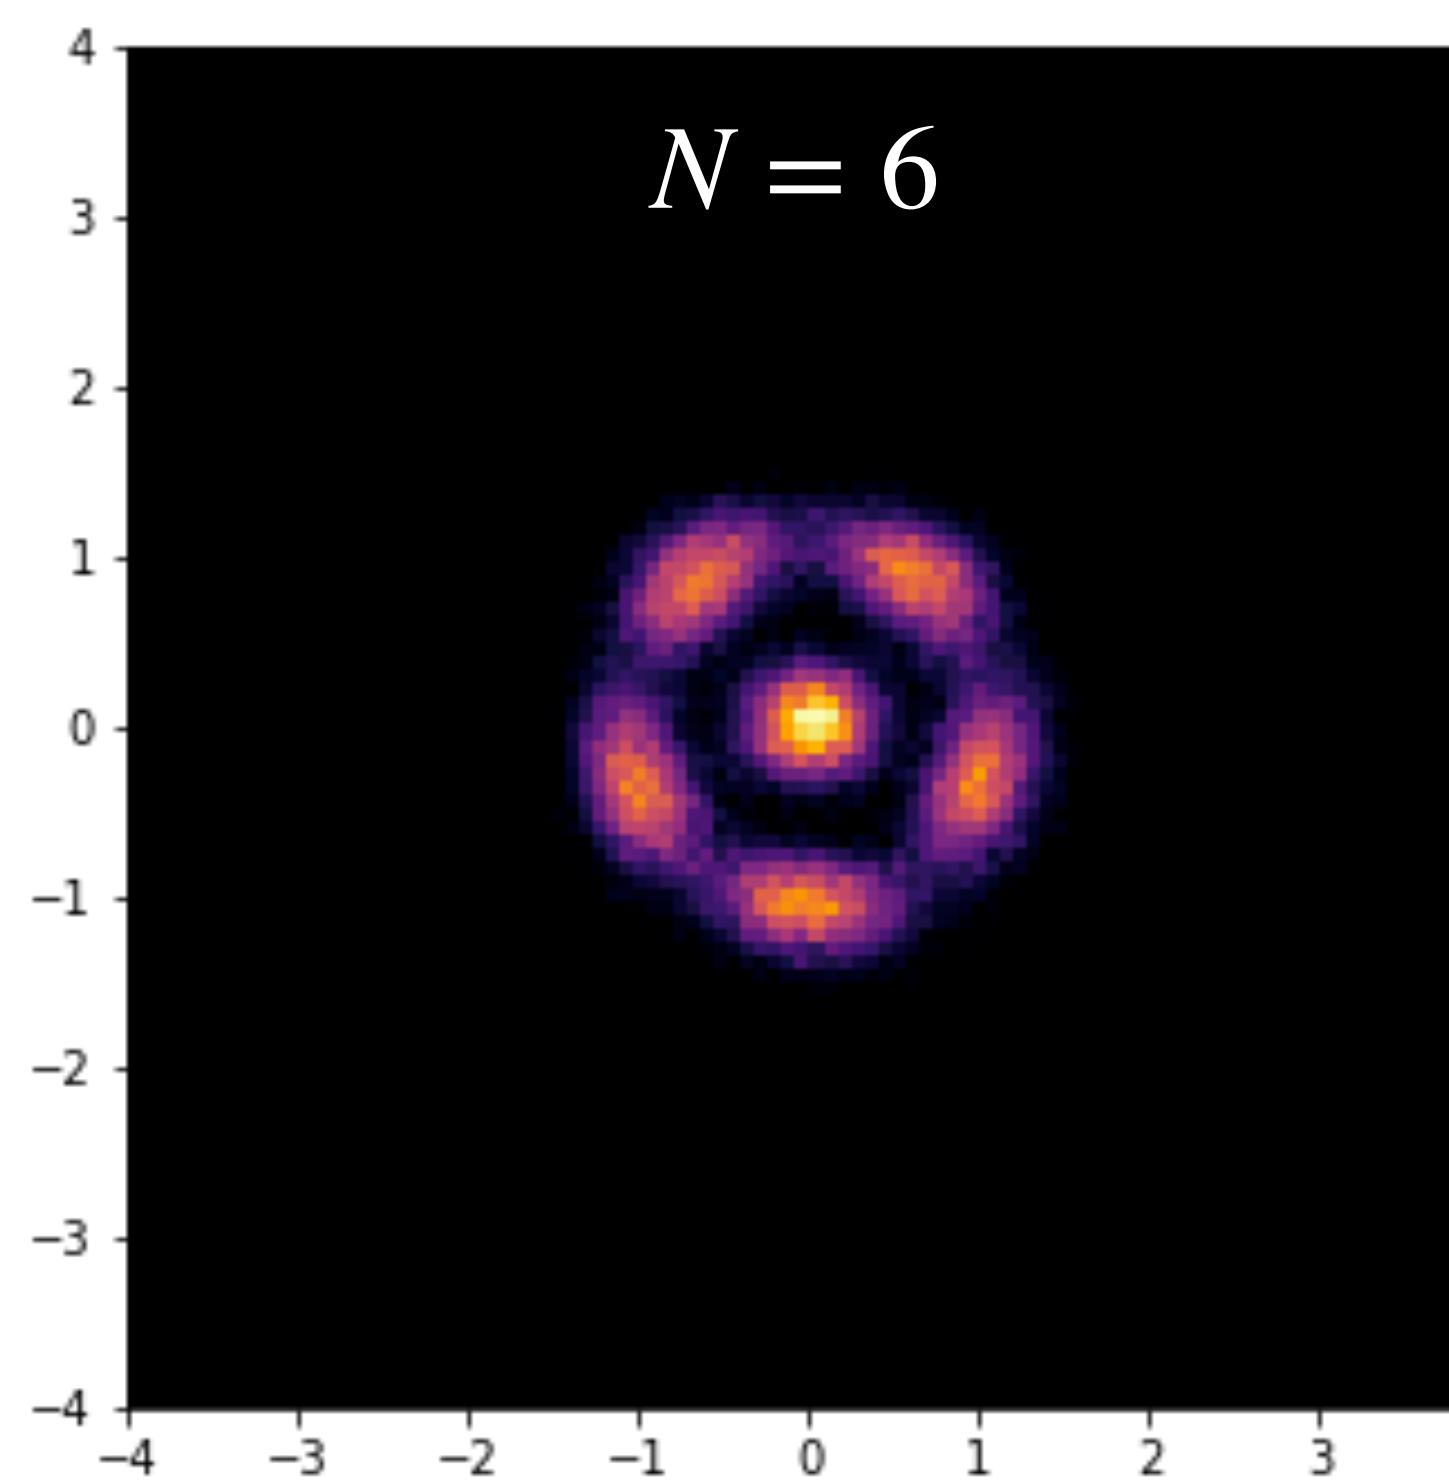
Chen et al, 1806.07366, Grathwohl et al 1810.01367



Continuous normalizing flow have no structural
constraints on the transformation Jacobian

Tutorial: Classical Coulomb gas in a harmonic trap

$$H = \sum_{i < j} \frac{1}{|x_i - x_j|} + \sum_i x_i^2$$



Training: Monte Carlo Gradient Estimators

Review: 1906.10652

$$\nabla_{\theta} \mathbb{E}_{x \sim p_{\theta}} [f(x)]$$

Reinforcement learning

Variational inference

Variational Monte Carlo

Variational quantum algorithms

...

$$\nabla_{\theta} \mathbb{E}_{x \sim p_{\theta}} [f(x)] = \mathbb{E}_{x \sim p_{\theta}} [f(x) \nabla_{\theta} \ln p_{\theta}(x)]$$

Score function estimator (REINFORCE)

Pathwise estimator (Reparametrization trick) $x = g_{\theta}(z)$

$$\nabla_{\theta} \mathbb{E}_{x \sim p_{\theta}} [f(x)] = \mathbb{E}_{z \sim \mathcal{N}(z)} [\nabla_{\theta} f(g_{\theta}(z))]$$

10.1 Guidance in Choosing Gradient Estimators

With so many competing approaches, we offer our rules of thumb in choosing an estimator, which follow the intuition we developed throughout the paper:

- If our estimation problem involves continuous functions and measures that are continuous in the domain, then using the pathwise estimator is a good default. It is relatively easy to implement and a default implementation, one without other variance reduction, will typically have variance that is low enough so as not to interfere with the optimisation.
- If the cost function is not differentiable or a black-box function then the score-function or the measure-valued gradients are available. If the number of parameters is low, then the measure-valued gradient will typically have lower variance and would be preferred. But if we have a high-dimensional parameter set, then the score function estimator should be used.
- If we have no control over the number of times we can evaluate a black-box cost function, effectively only allowing a single evaluation of it, then the score function is the only estimator of the three we reviewed that is applicable.
- The score function estimator should, by default, always be implemented with at least a basic variance reduction. The simplest option is to use a baseline control variate estimated with a running average of the cost value.
- When using the score-function estimator, some attention should be paid to the dynamic range of the cost function and its variance, and to find ways to keep its value bounded within a reasonable range, e.g., transforming the cost so that it is zero mean, or using a baseline.
- For all estimators, track the variance of the gradients if possible and address high variance by using a larger number of samples from the measure, decreasing the learning rate, or clipping the gradient values. It may also be useful to restrict the range of some parameters to avoid extreme values, e.g., by clipping them to a desired interval.
- The measure-valued gradient should be used with some coupling method for variance reduction. Coupling strategies that exploit relationships between the positive and negative components of the density decomposition, and which have shared sampling paths, are known for the commonly-used distributions.
- If we have several unbiased gradient estimators, a convex combination of them might have lower variance than any of the individual estimators.
- If the measure is discrete on its domain then the score-function or measure-valued gradient are available. The choice will again depend on the dimensionality of the parameter space.
- In all cases, we strongly recommend having a broad set of tests to verify the unbiasedness of the gradient estimator when implemented.

1906.10652

$$\nabla_{\theta} \mathbb{E}_{x \sim p_{\theta}} [f(x)]$$

When to use which ?

More discussions

Roeder et al, 1703.09194

Vaitl et al 2206.09016, 2207.08219

A few words about tooling

HIPS/autograd

theano



[M]^s
MindSpore

PyTorch

TensorFlow

The SciML logo consists of three overlapping colored shapes: red, green, and purple, forming a stylized 'S' or swirl pattern.

SciML

The Taichi logo features a blue stylized 'T' character with a circular motion effect around it.

The JAX logo is a stylized 'X' composed of colored 3D cubes in blue, green, and purple.

K Keras

The NiLang logo features a colorful, abstract circular pattern with radiating lines and dots in various colors like red, green, and purple.

NiLang

The Taichi logo features a black Tai Chi symbol (yin-yang) with the word "Taichi" written next to it.

Differentiable programming frameworks

Autoregressive model

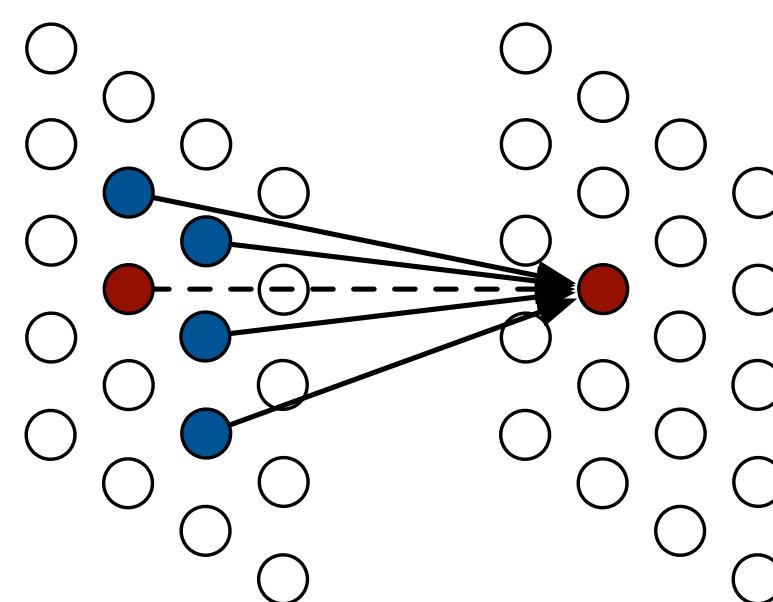
$$p(\mathbf{x}) = p(x_1)p(x_2 | x_1)p(x_3 | x_1, x_2)\dots$$

Language: Casual transformer 1706.03762

Speech: WaveNet 1609.03499



Image: PixelCNN 1601.06759



Exercise

Wait, Isn't WaveNet a normalizing flow?

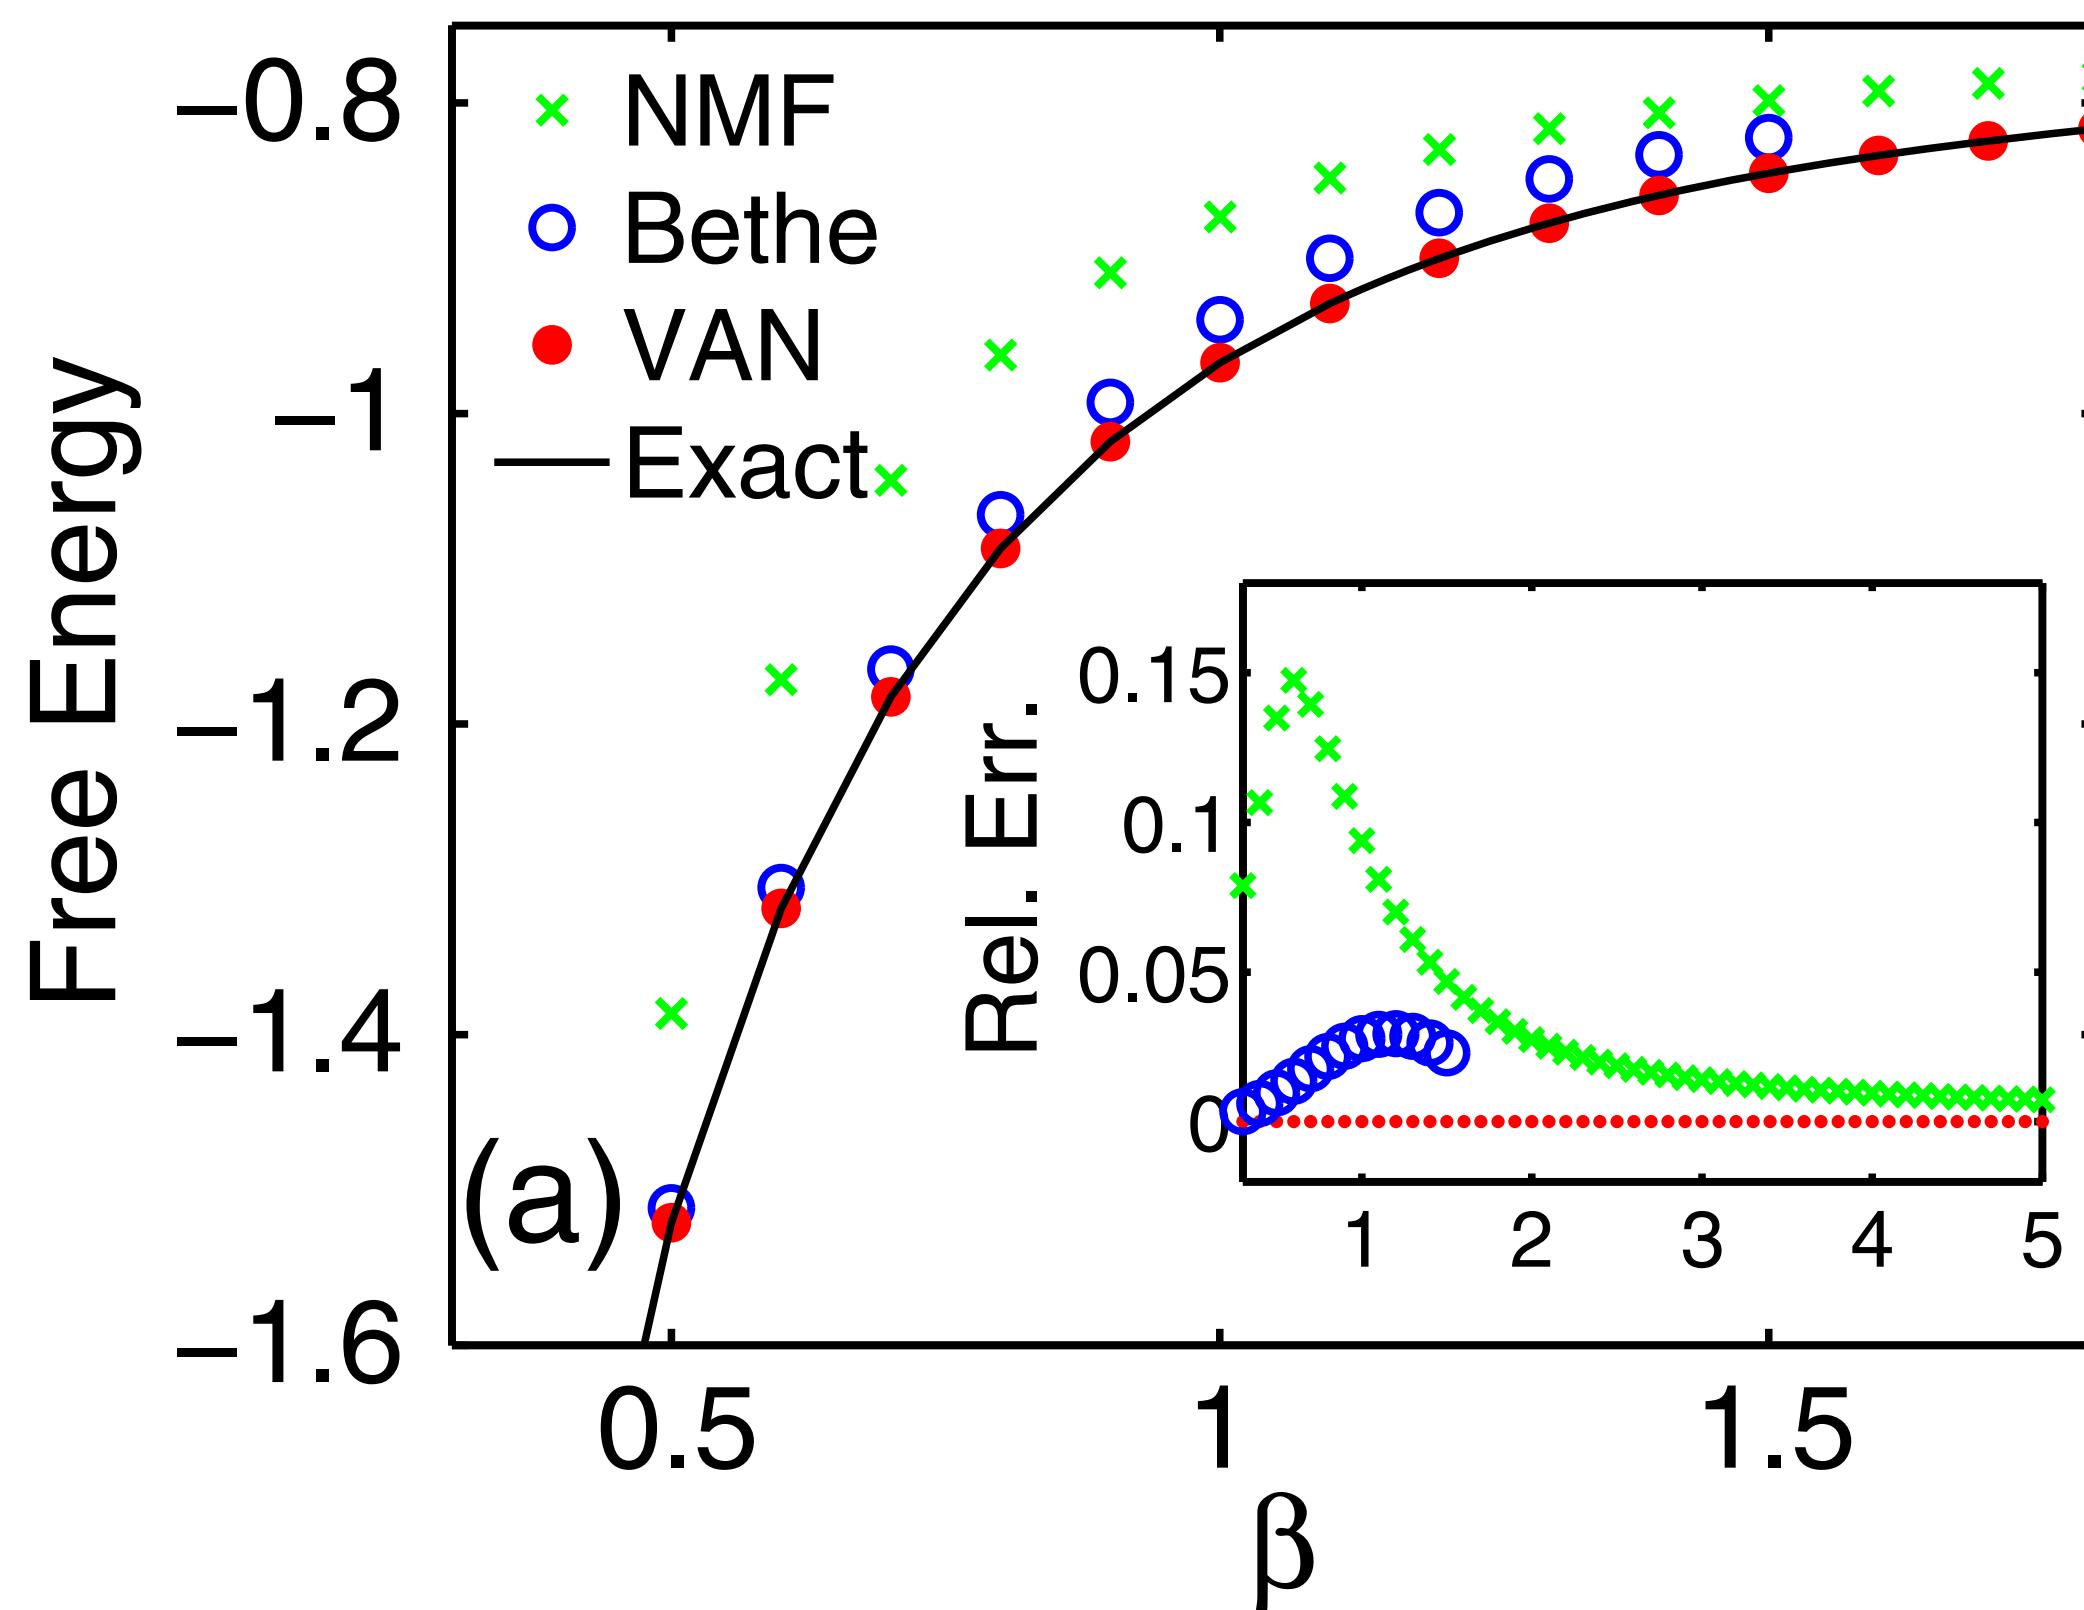


Hint: read Papamakarios et al,[1705.07057](#)
and van den Oord et al,[1711.10433](#)

Variational autoregressive networks

Conventional approaches

Sherrington-Kirkpatrick spin glass



Naive mean-field
factorized probability

$$p(\mathbf{x}) = \prod_i p(x_i)$$

Bethe approximation
pairwise interaction

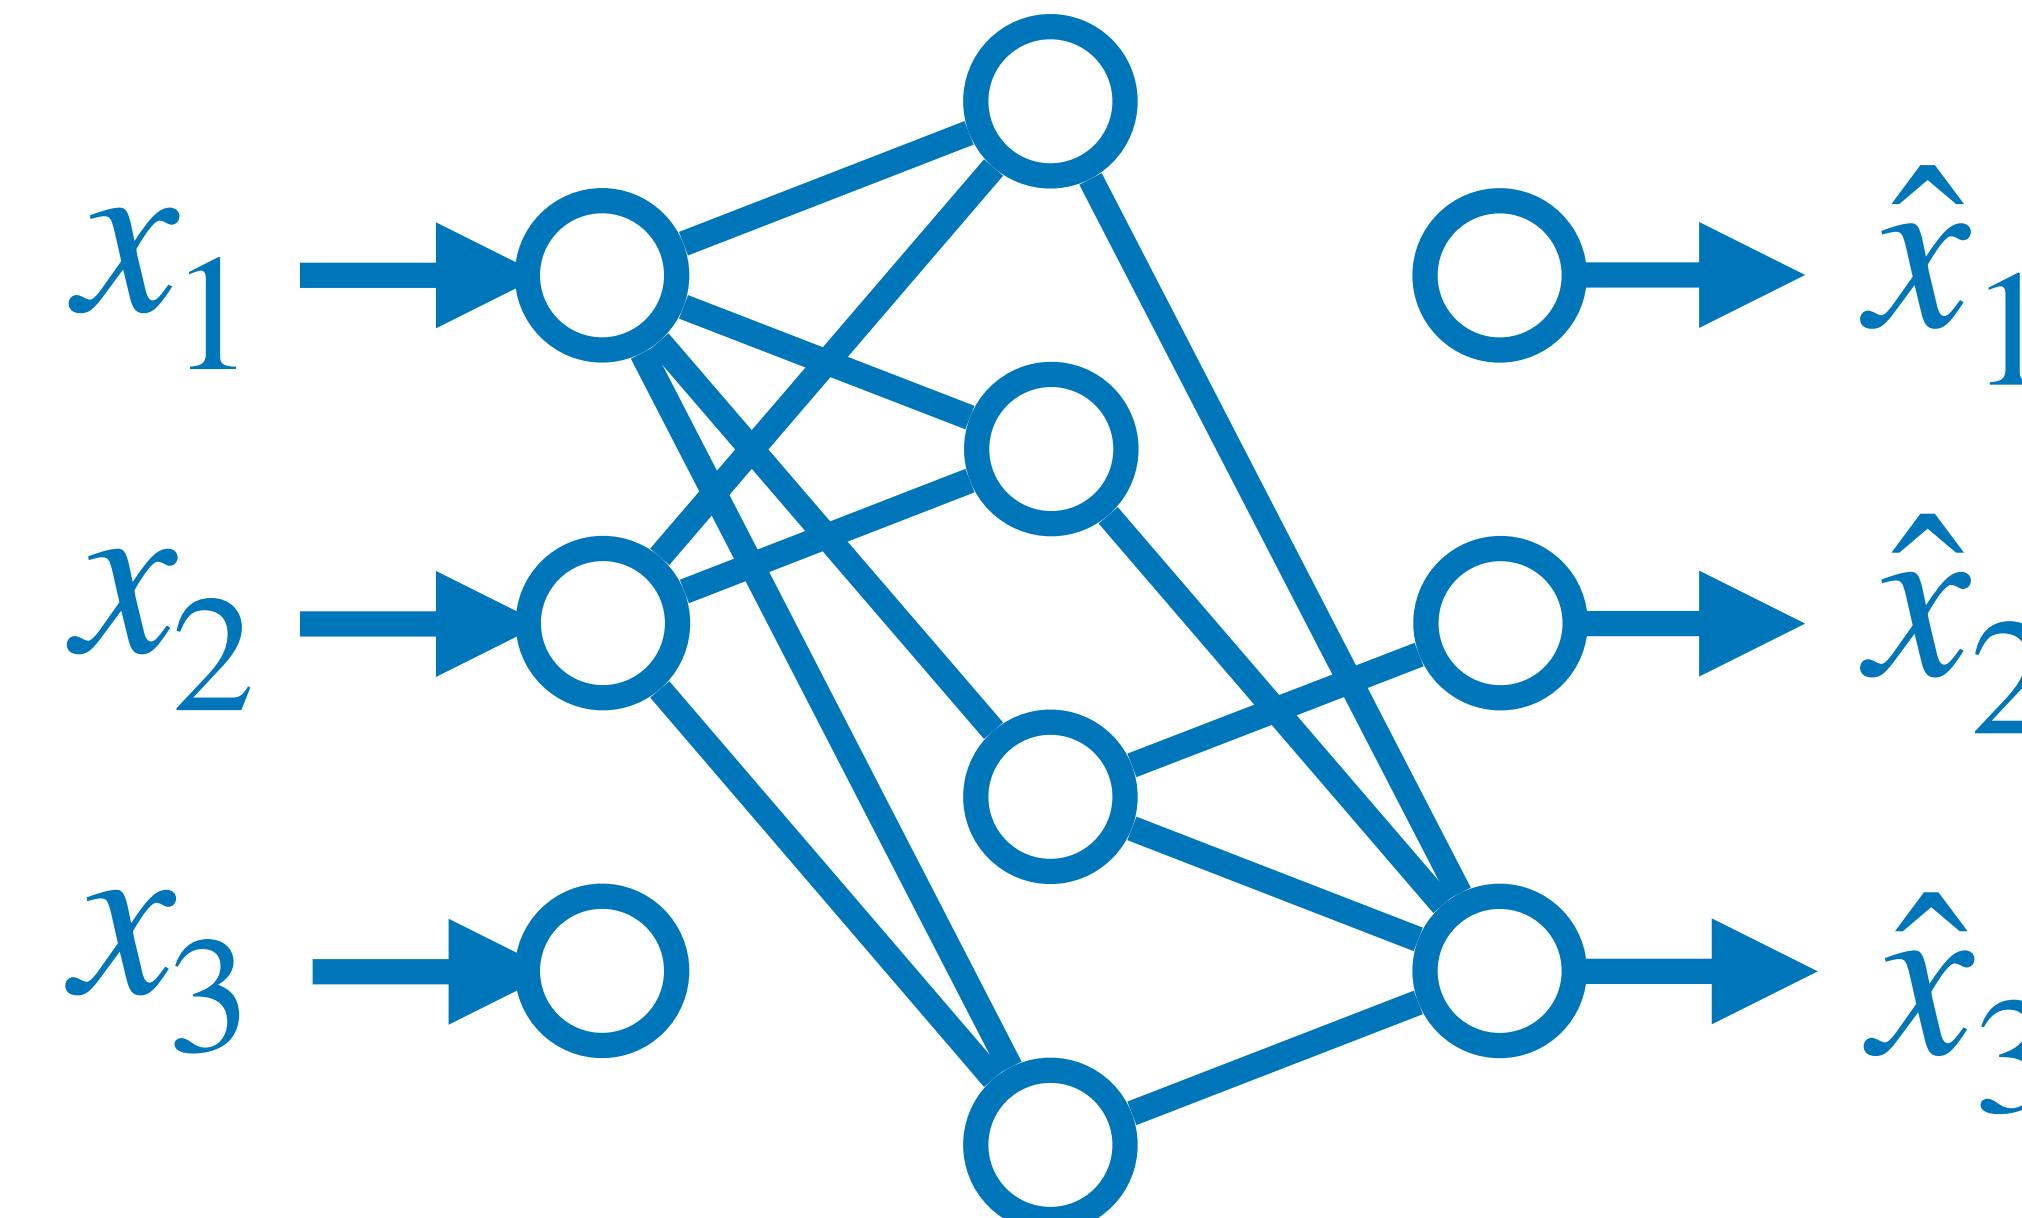
$$p(\mathbf{x}) = \prod_i p(x_i) \prod_{(i,j) \in E} \frac{p(x_i, x_j)}{p(x_i)p(x_j)}$$

Variational autoregressive network

$$p(\mathbf{x}) = \prod_i p(x_i | \mathbf{x}_{<i})$$

Wu, LW, Zhang, PRL '19
github.com/wdphy16/stat-mech-van

Implementation: autoregressive masks



Masked Autoencoder
Germain et al, 1502.03509

$$p(x_1) = \text{Bern}(\hat{x}_1)$$

$$p(x_2 | x_1) = \text{Bern}(\hat{x}_2)$$

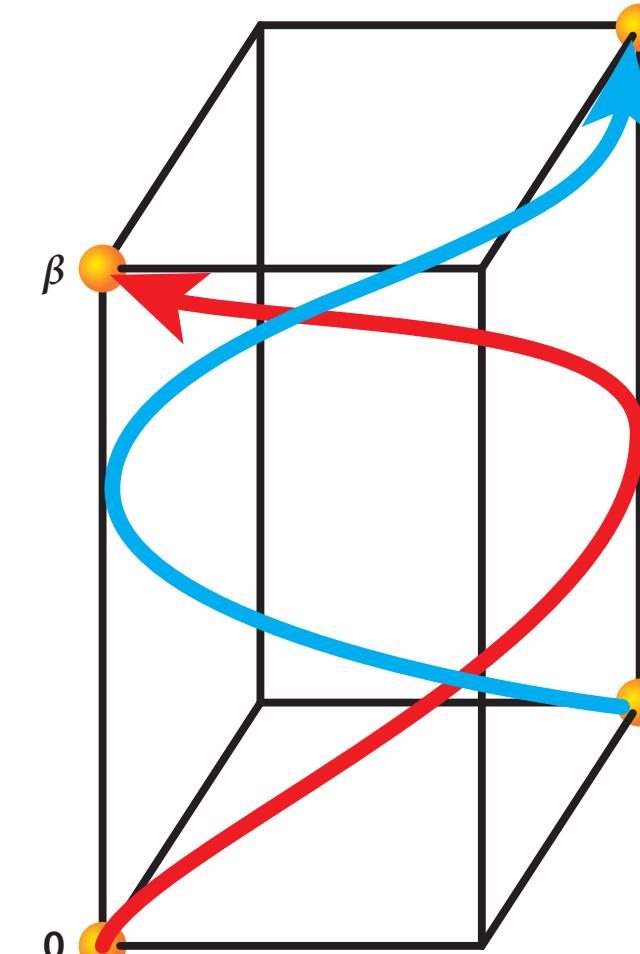
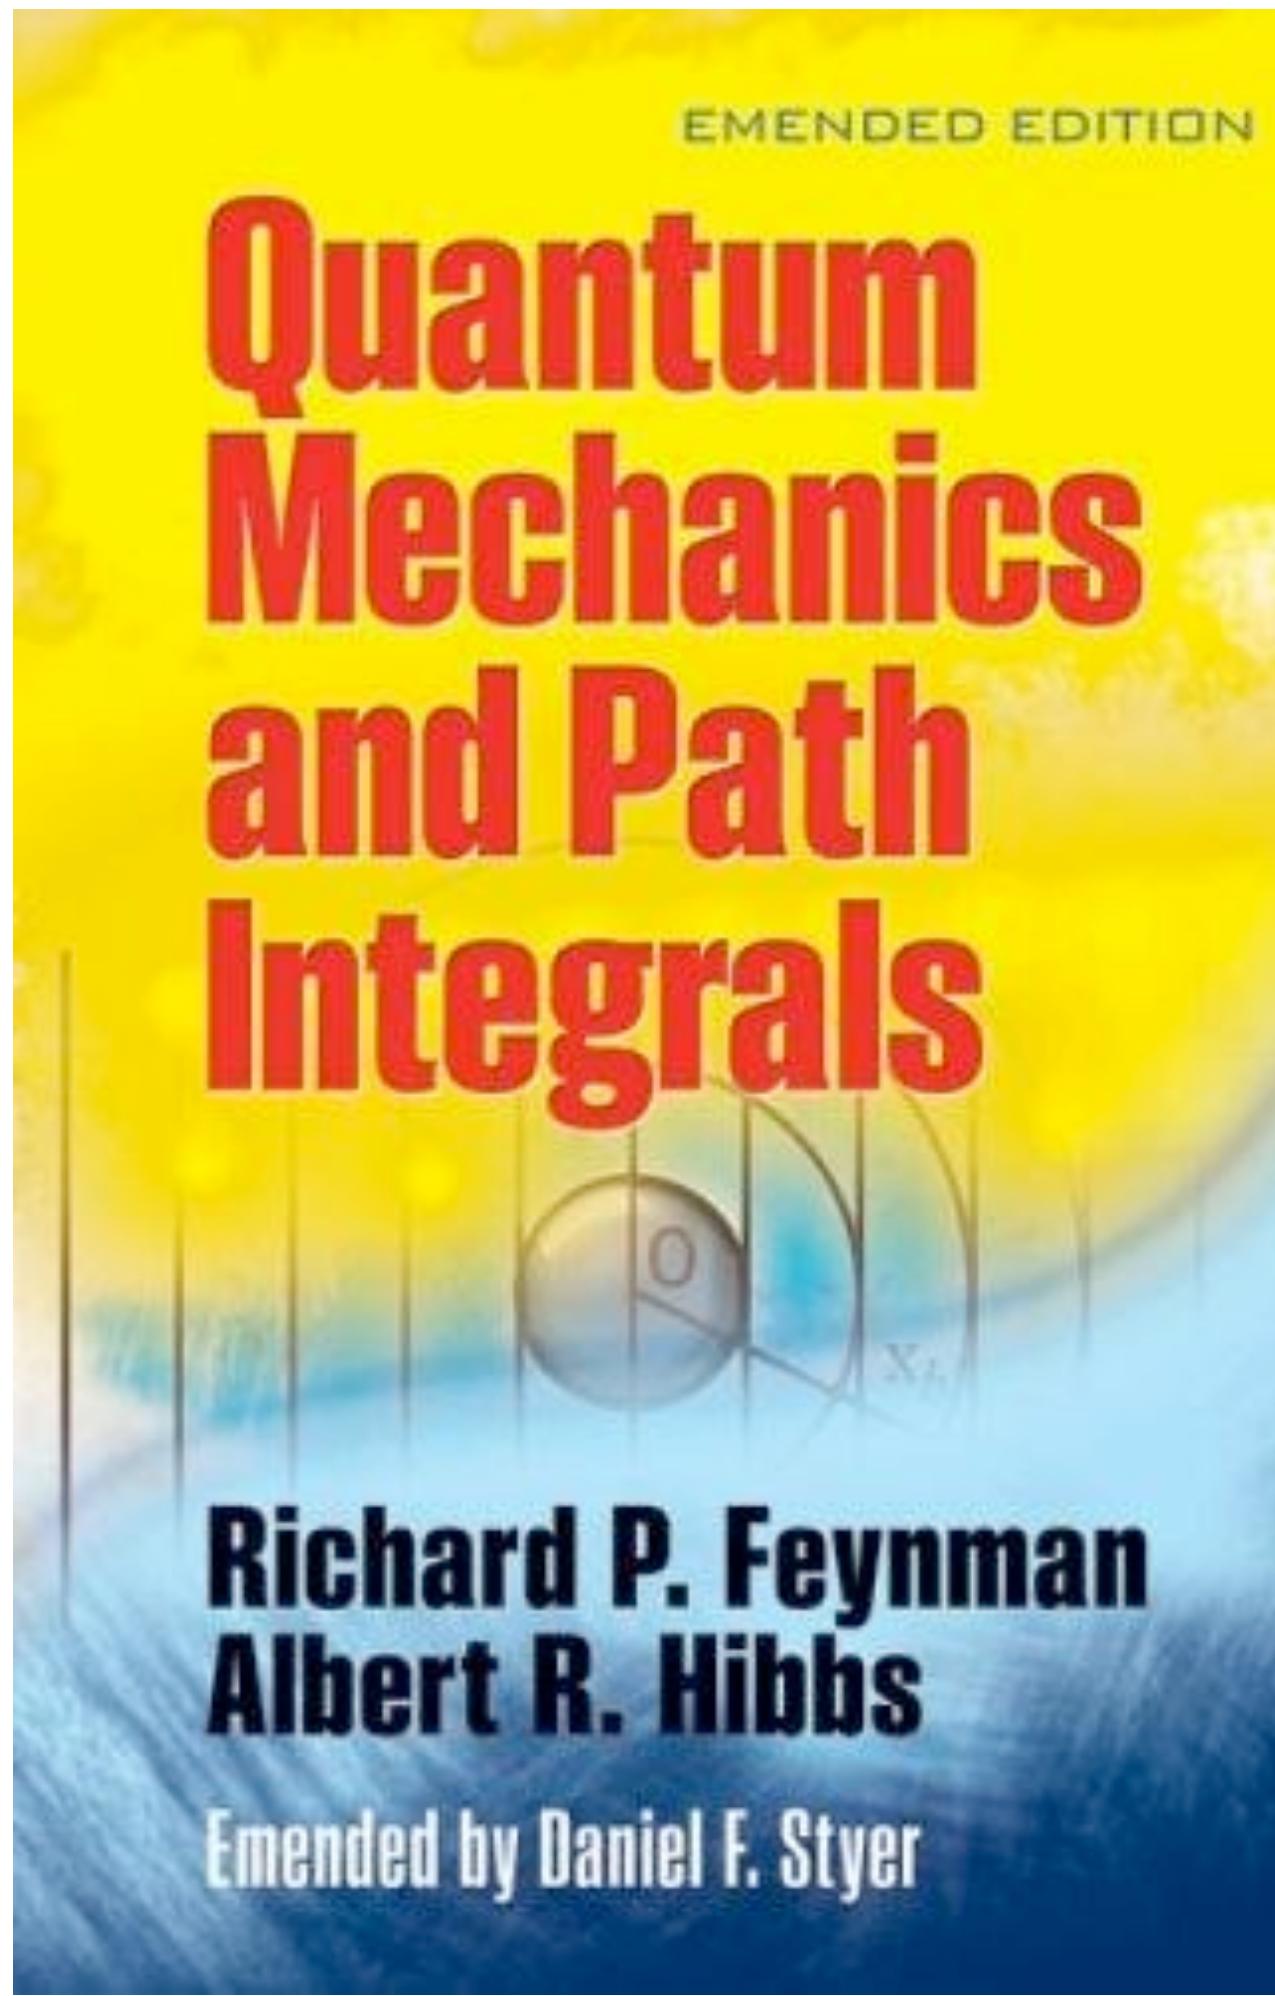
$$p(x_3 | x_1, x_2) = \text{Bern}(\hat{x}_3)$$

Other examples: PixelCNN, van den Oord et al, 1601.06759 Casual transformer, 1706.03762
Other ways to implement autoregressive models: recurrent networks



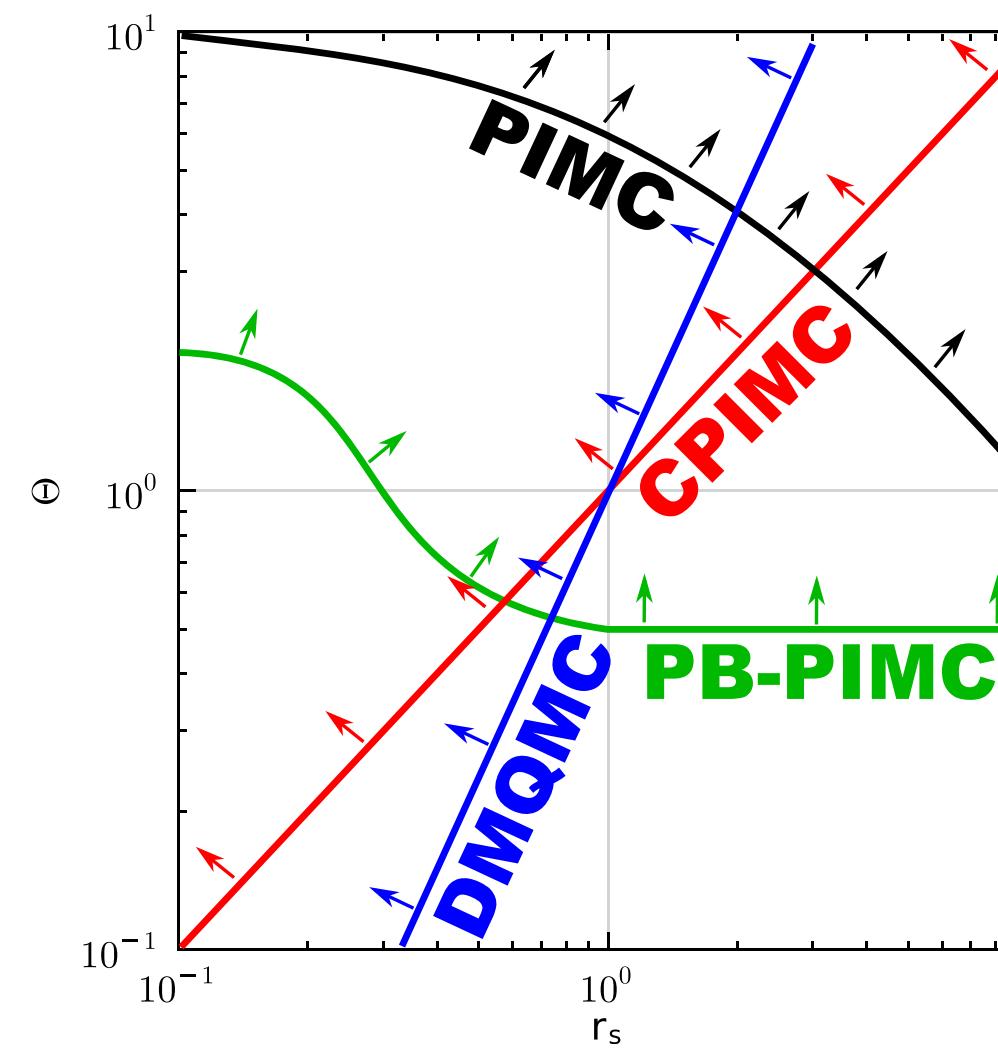
How about quantum systems?

Quantum-to-classical mapping



$d+1$ -dim
spacetime integral

$$Z = \text{Tr}(e^{-\beta H}) = \int d^{d+1}x \dots$$



However, the “weight”
may not be positive definite.
Sign problem!

The quantum variational free-energy approach

Gibbs–Bogolyubov–Feynman–Delbrück–Molière variational principle

$$\min F[\rho] = k_B T \operatorname{Tr}(\rho \ln \rho) + \operatorname{Tr}(H\rho)$$



$$\text{s.t. } \operatorname{Tr}\rho = 1 \quad \rho > 0 \quad \rho^\dagger = \rho \quad \langle x | \rho | x' \rangle = (-)^{\mathcal{P}} \langle \mathcal{P}x | \rho | x' \rangle$$

Exercise

Prove $F[\rho] \geq -k_B T \ln Z$

where $Z = \operatorname{Tr}(e^{-H/k_B T})$

Search “Quantum relative entropy” on wikipedia

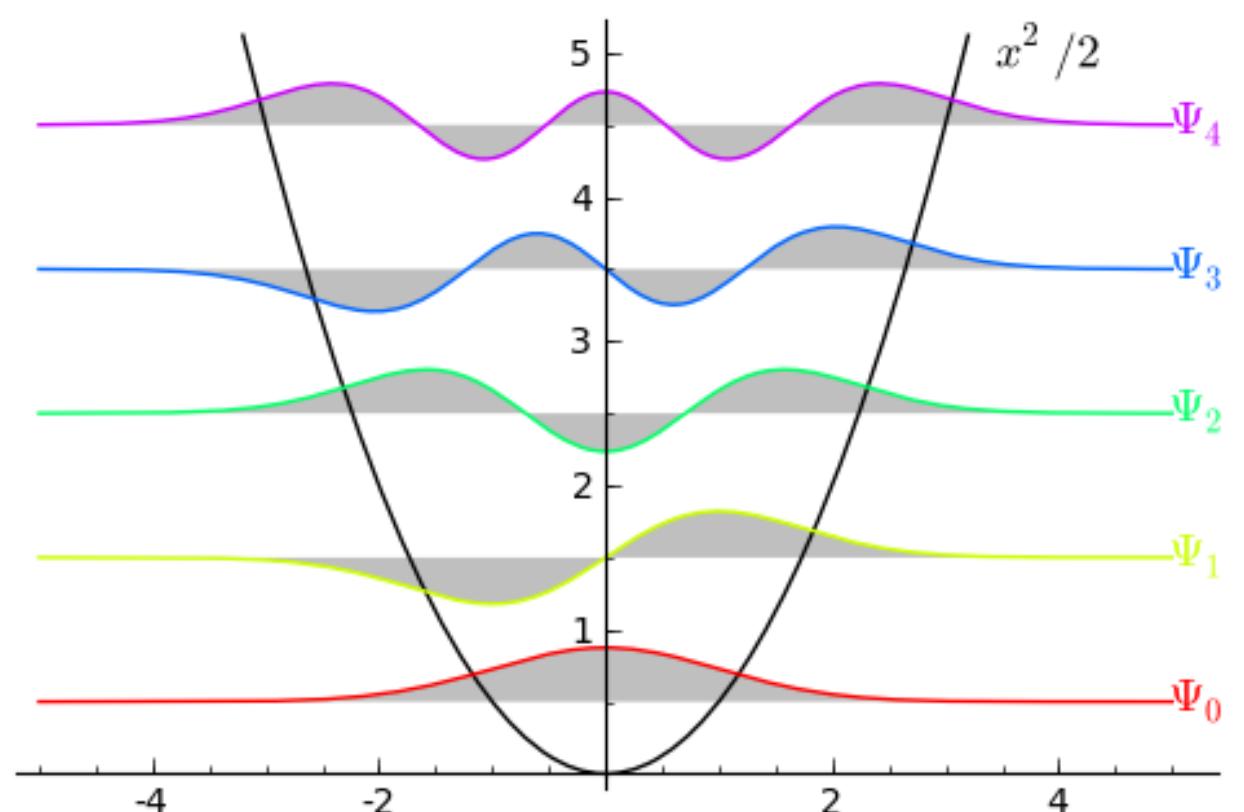
Exercise

Think about how to solve the
quantum Coulomb gas problem
using this principle.

Density matrix

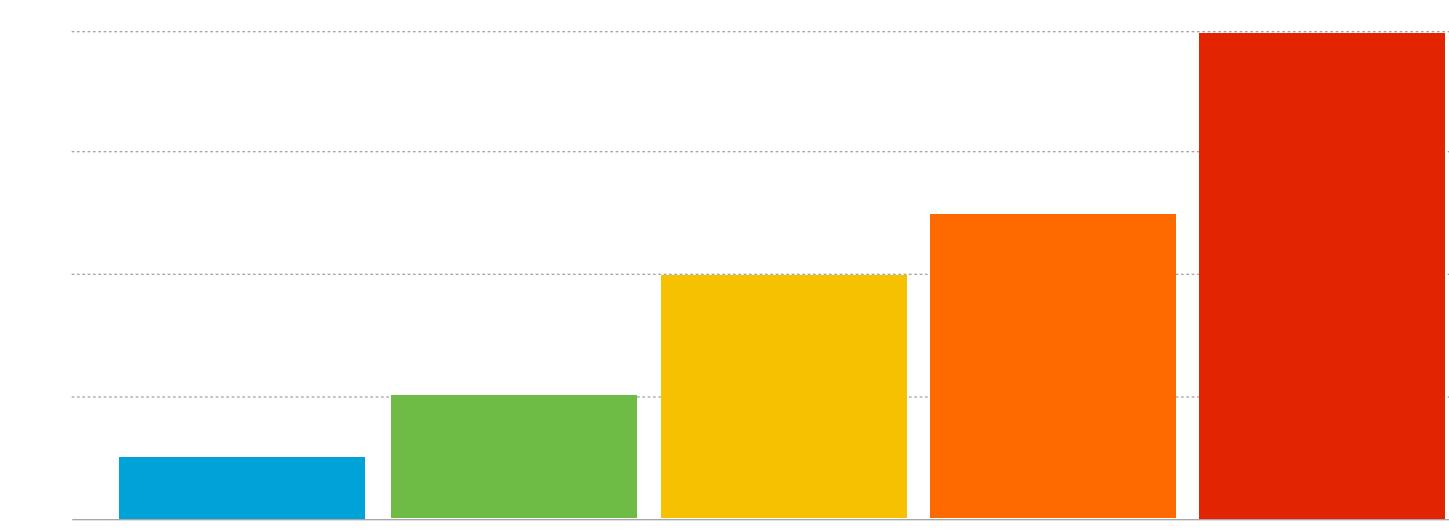
$$\rho = \sum_n \mu_n |\Psi_n\rangle\langle\Psi_n|$$

Quantum states $\Psi_n(x) = \langle x | \Psi_n \rangle$



$$\langle \Psi_m | \Psi_n \rangle = \delta_{mn}$$

Classical probability $0 < \mu_n < 1$

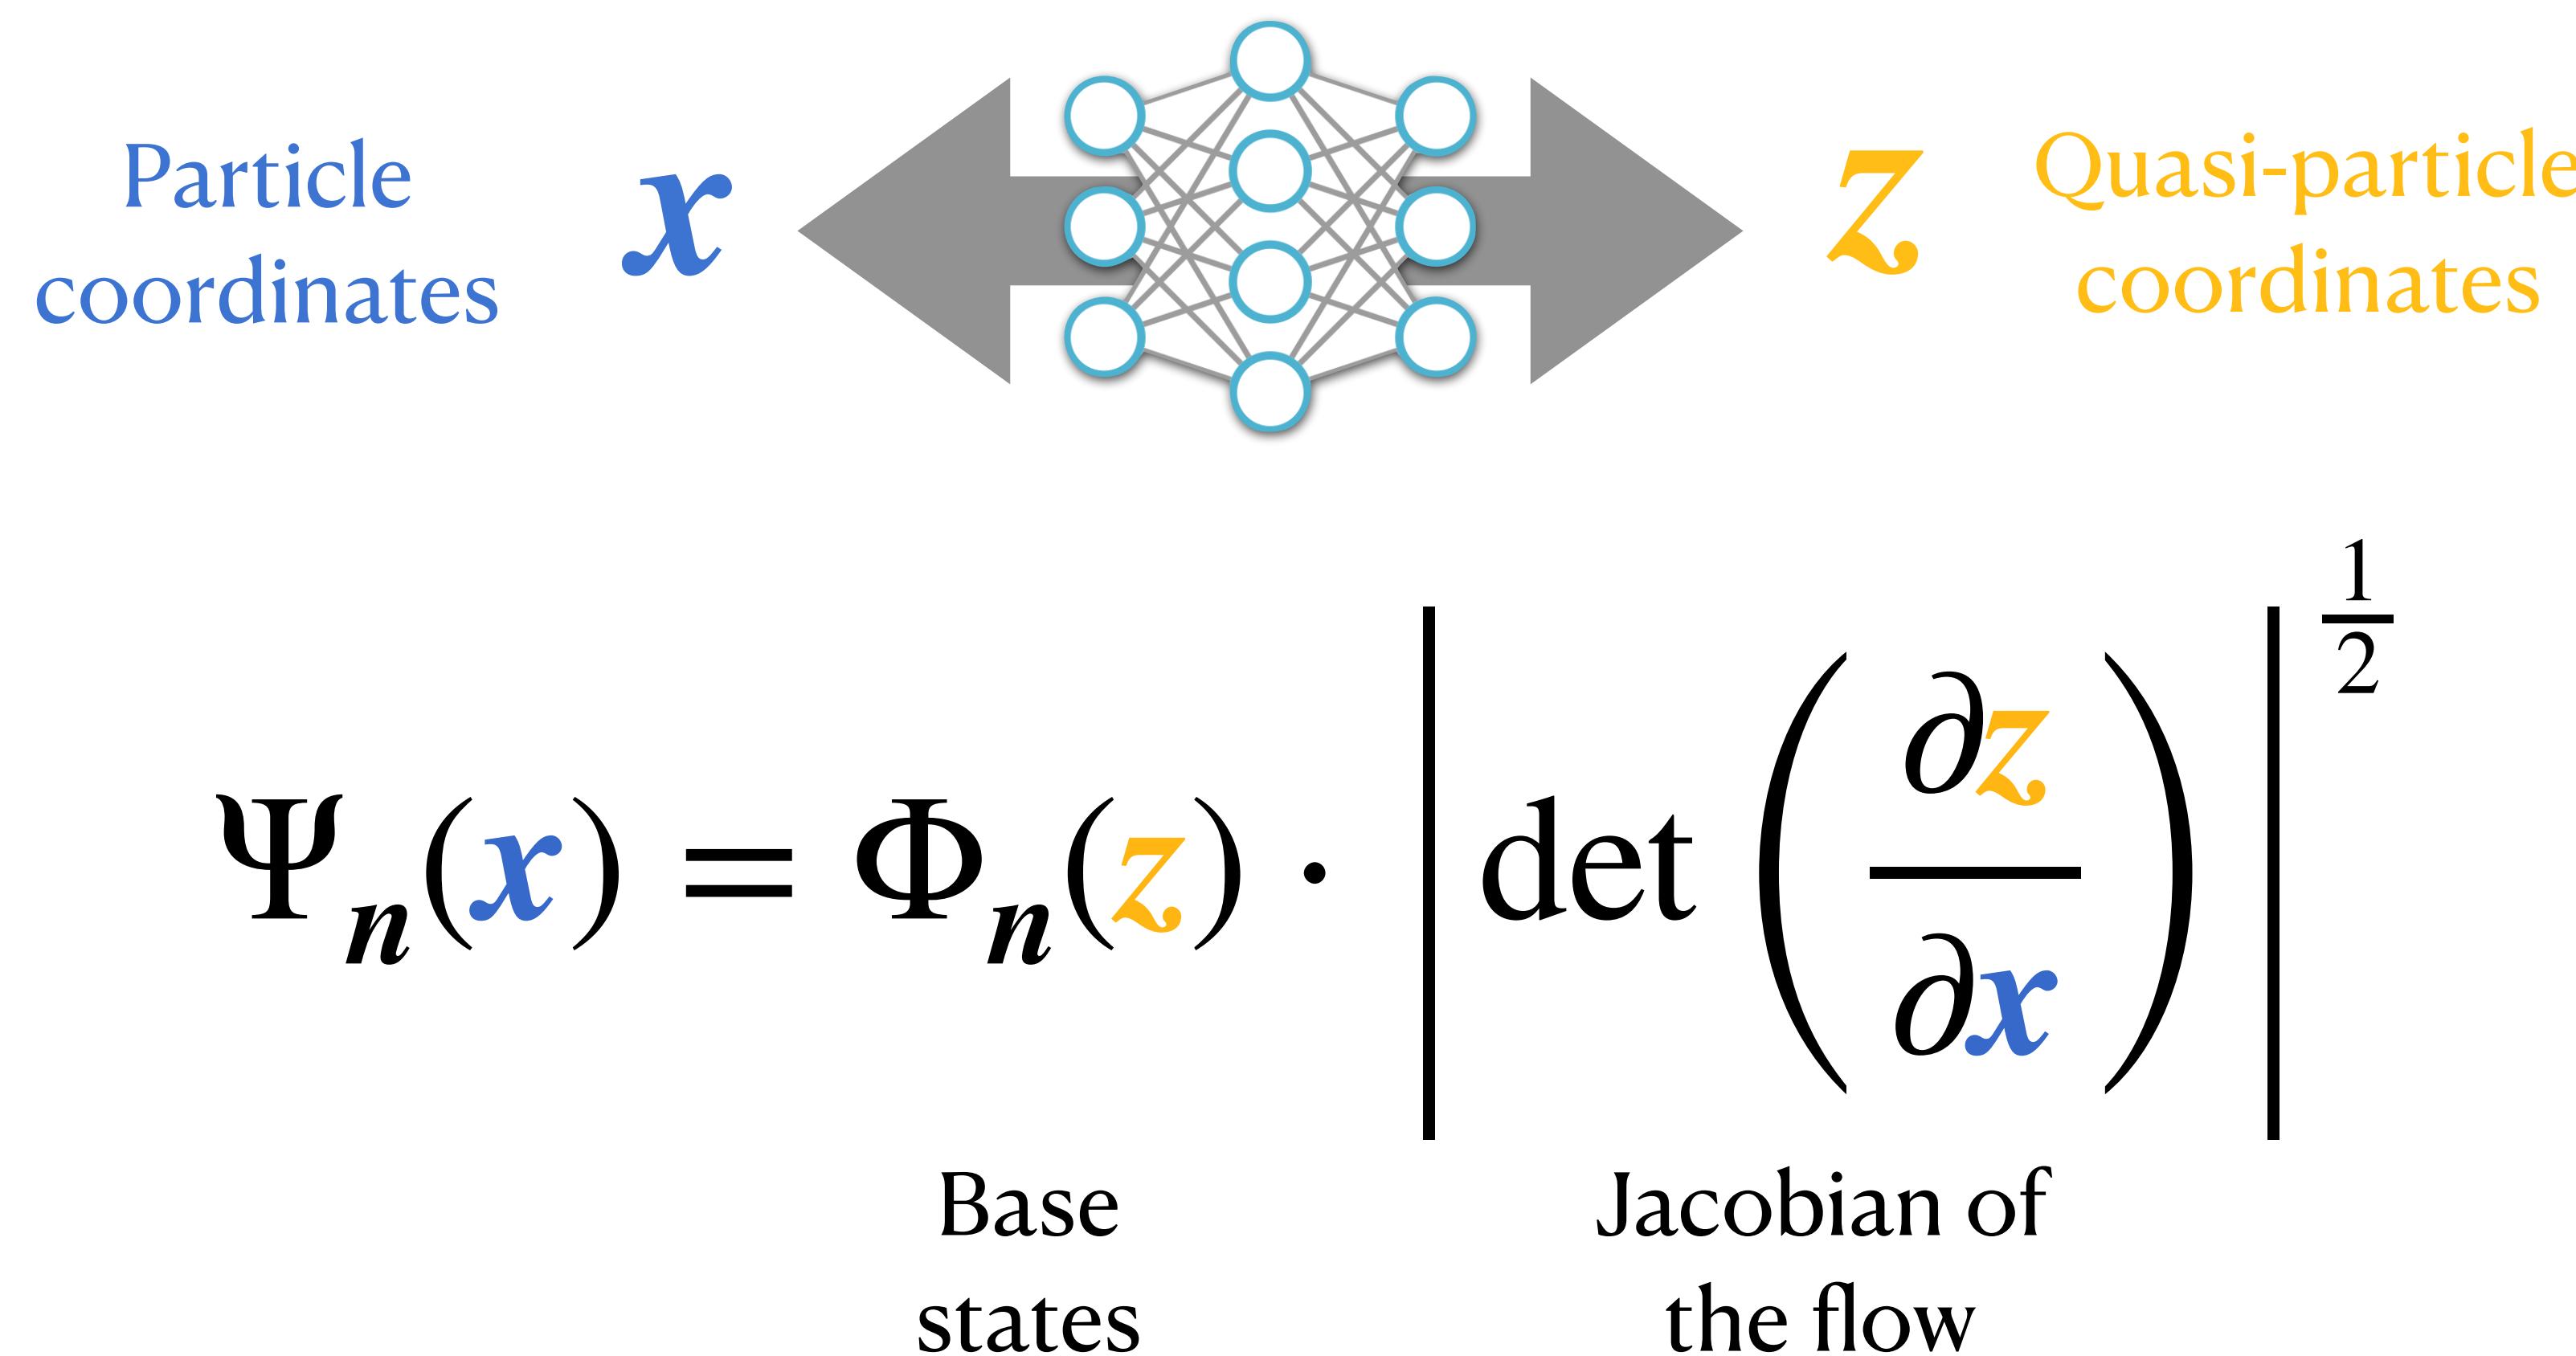


$$\sum_n \mu_n = 1$$

How to represent them ??

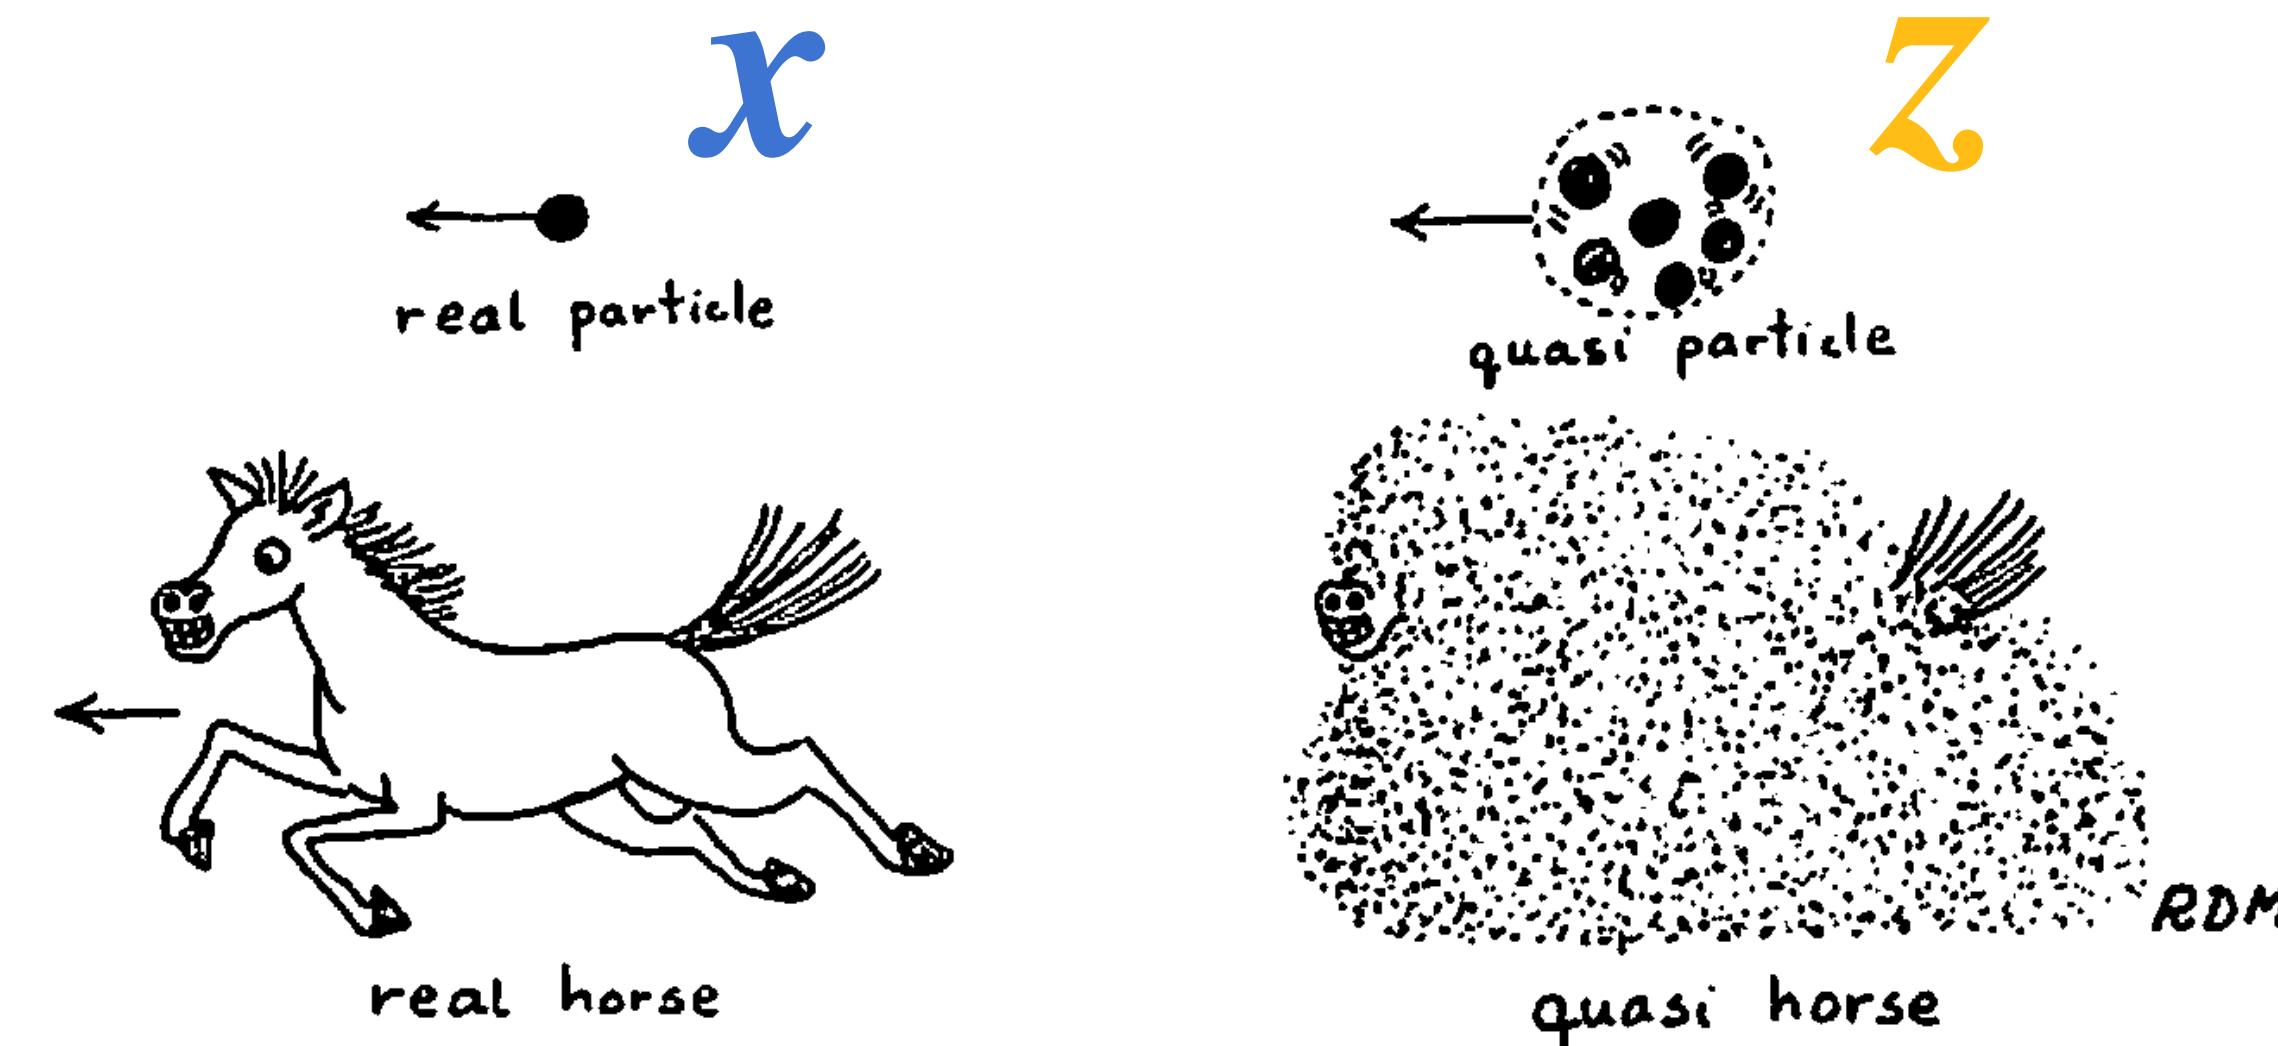
Use two deep generative models !!

“Square root” of a normalizing flow



The flow implements a *learnable* many-body unitary transformation
hence the name “neural canonical transformation” a classical generalization of Li, Dong, Zhang, LW, PRX ‘20

Feynman's backflow in the deep learning era



$$z_i = x_i + \sum_{j \neq i} \eta(|x_i - x_j|) (x_j - x_i)$$

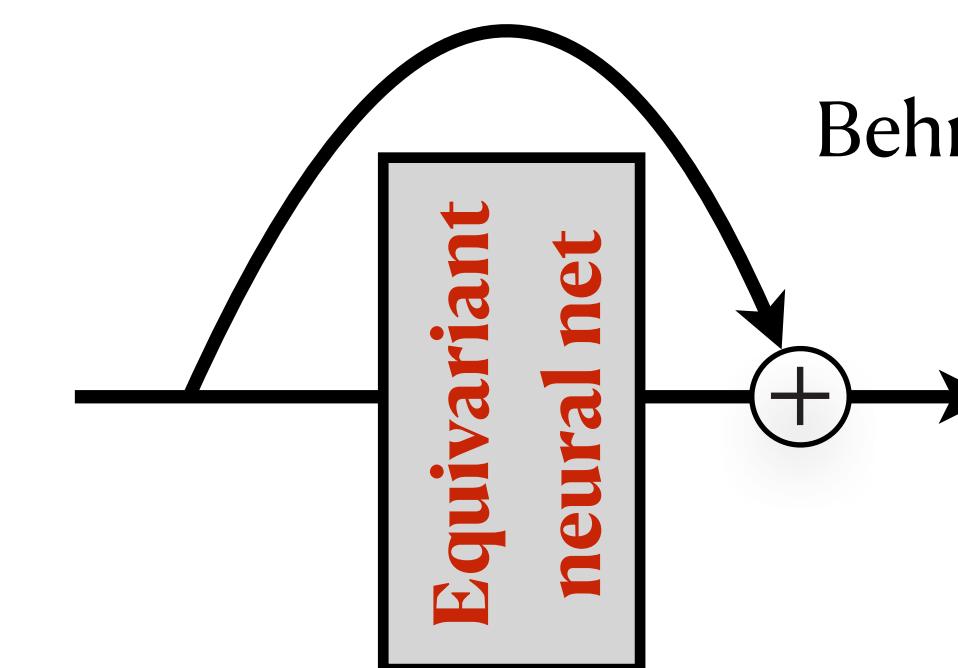
Feynman & Cohen 1956
wavefunction for liquid Helium

1

Backflow can be made unitary (if we track its Jacobian)

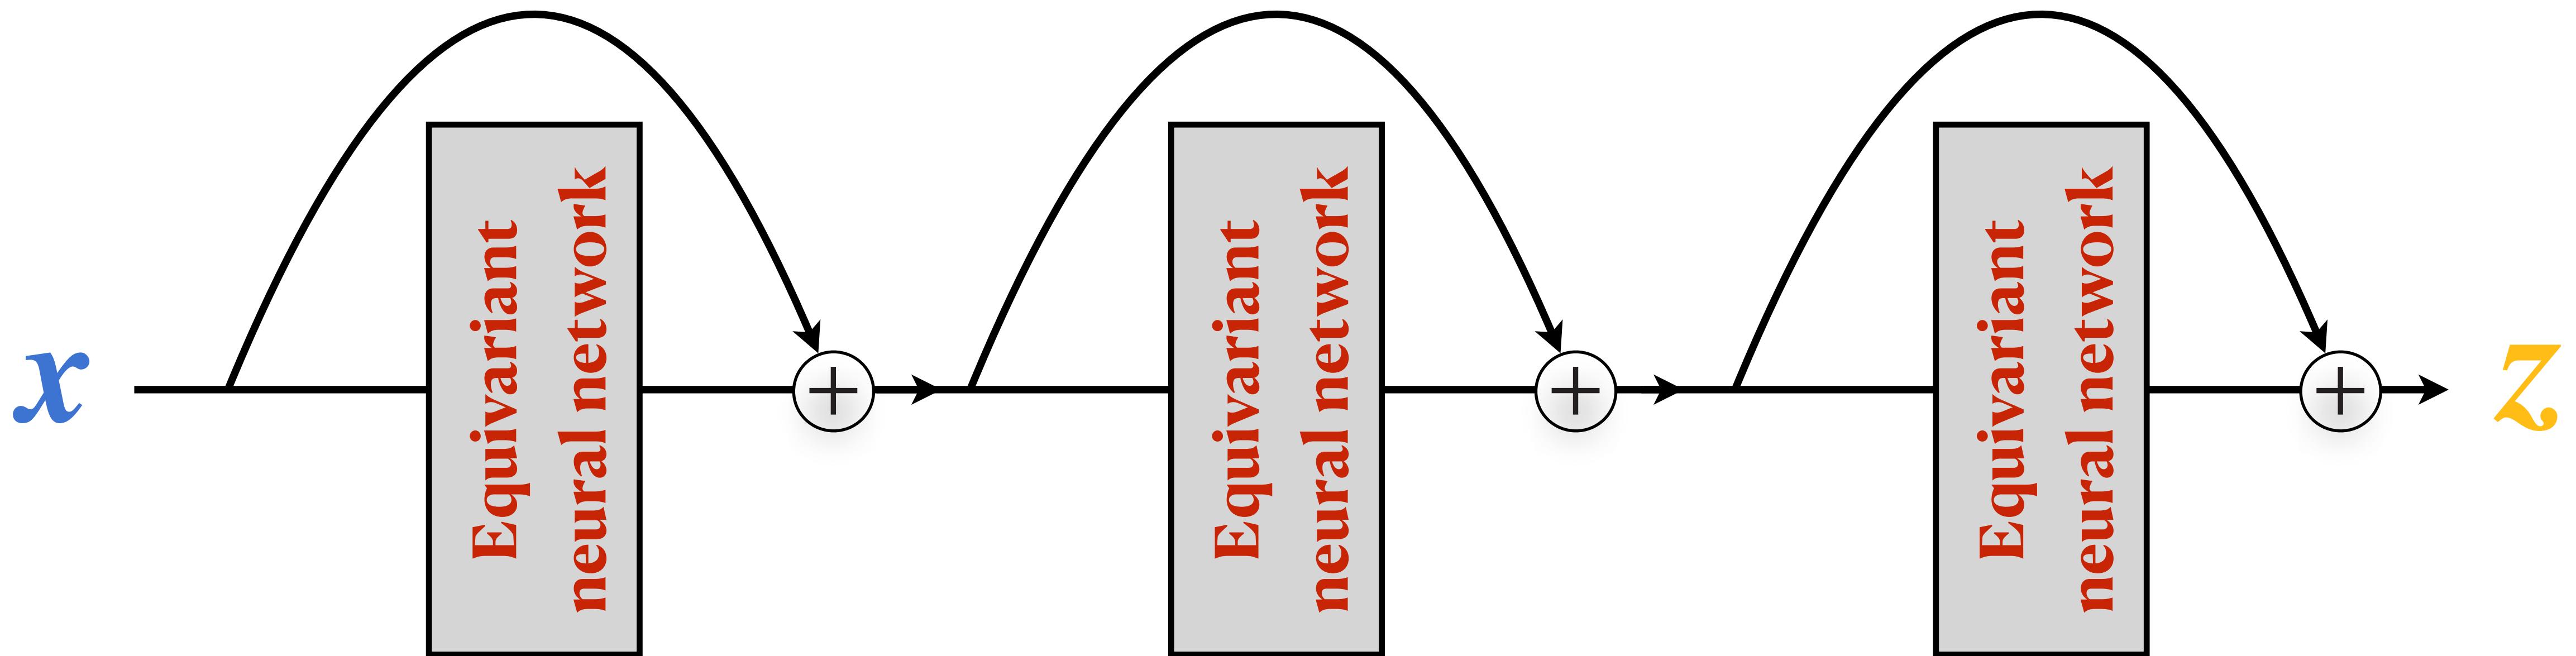
2

Backflow is an equivariant residual flow



Behrmann et al, 1811.00995
Chen et al, 1906.02735

Feynman's backflow in the deep learning era



Deep residual networks can be regarded as
discretization of a continuous dynamics



Fermi Flow

Xie, Zhang, LW, 2105.08644, JML '22

github.com/fermiflow

Continuous flow of electron density in a quantum dot

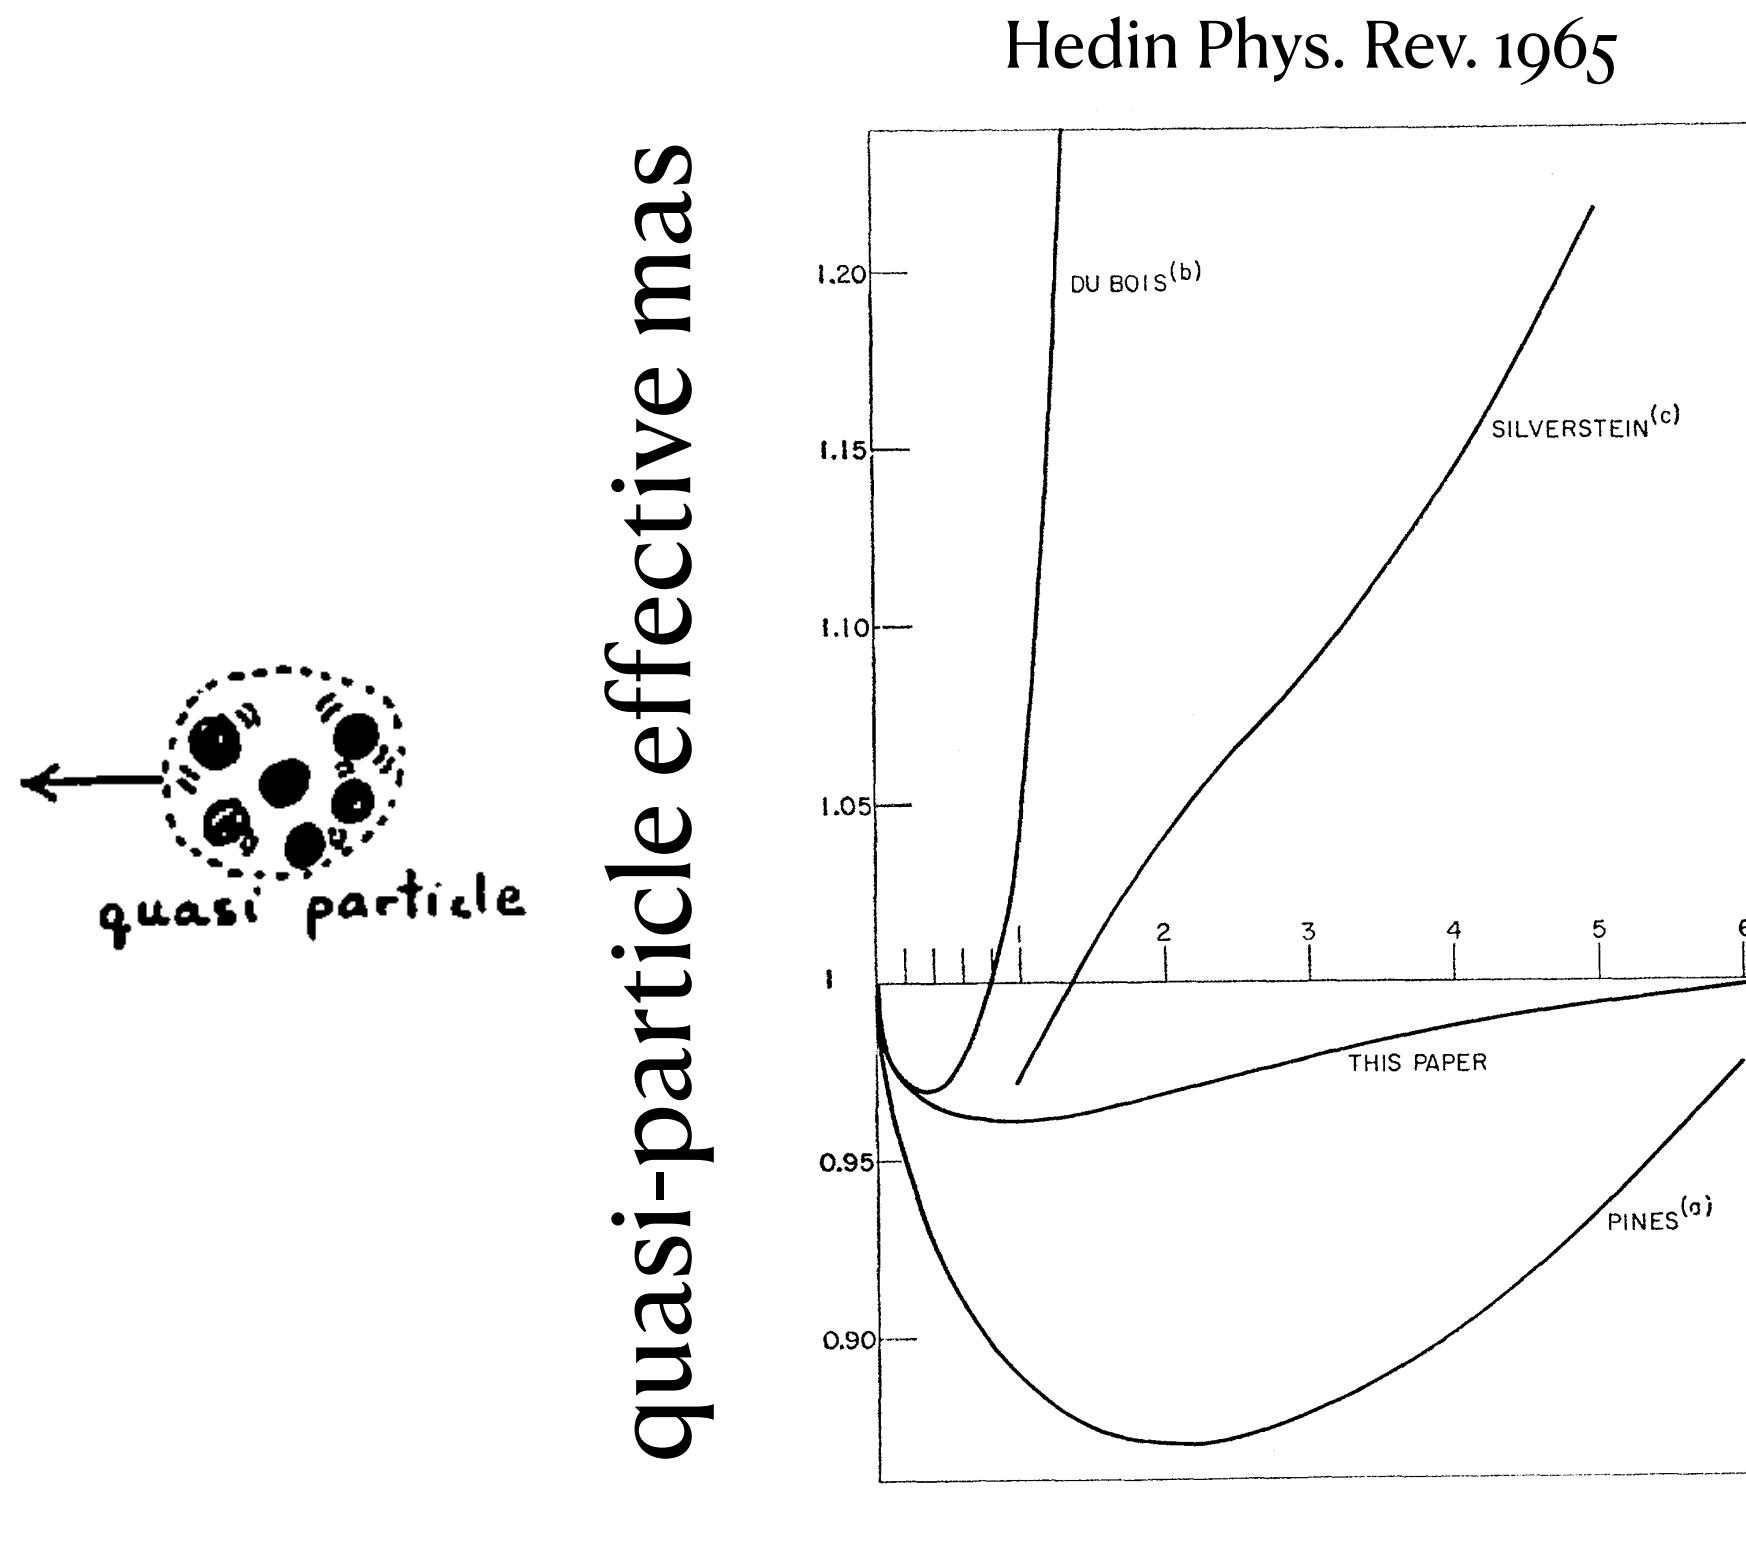
Exercise

Recall that $\rho = \sum_n \mu_n |\Psi_n\rangle\langle\Psi_n|$, prove

$$\text{Tr}(\rho \ln \rho) = \mathbb{E}_{n \sim \mu_n} [\ln \mu_n]$$

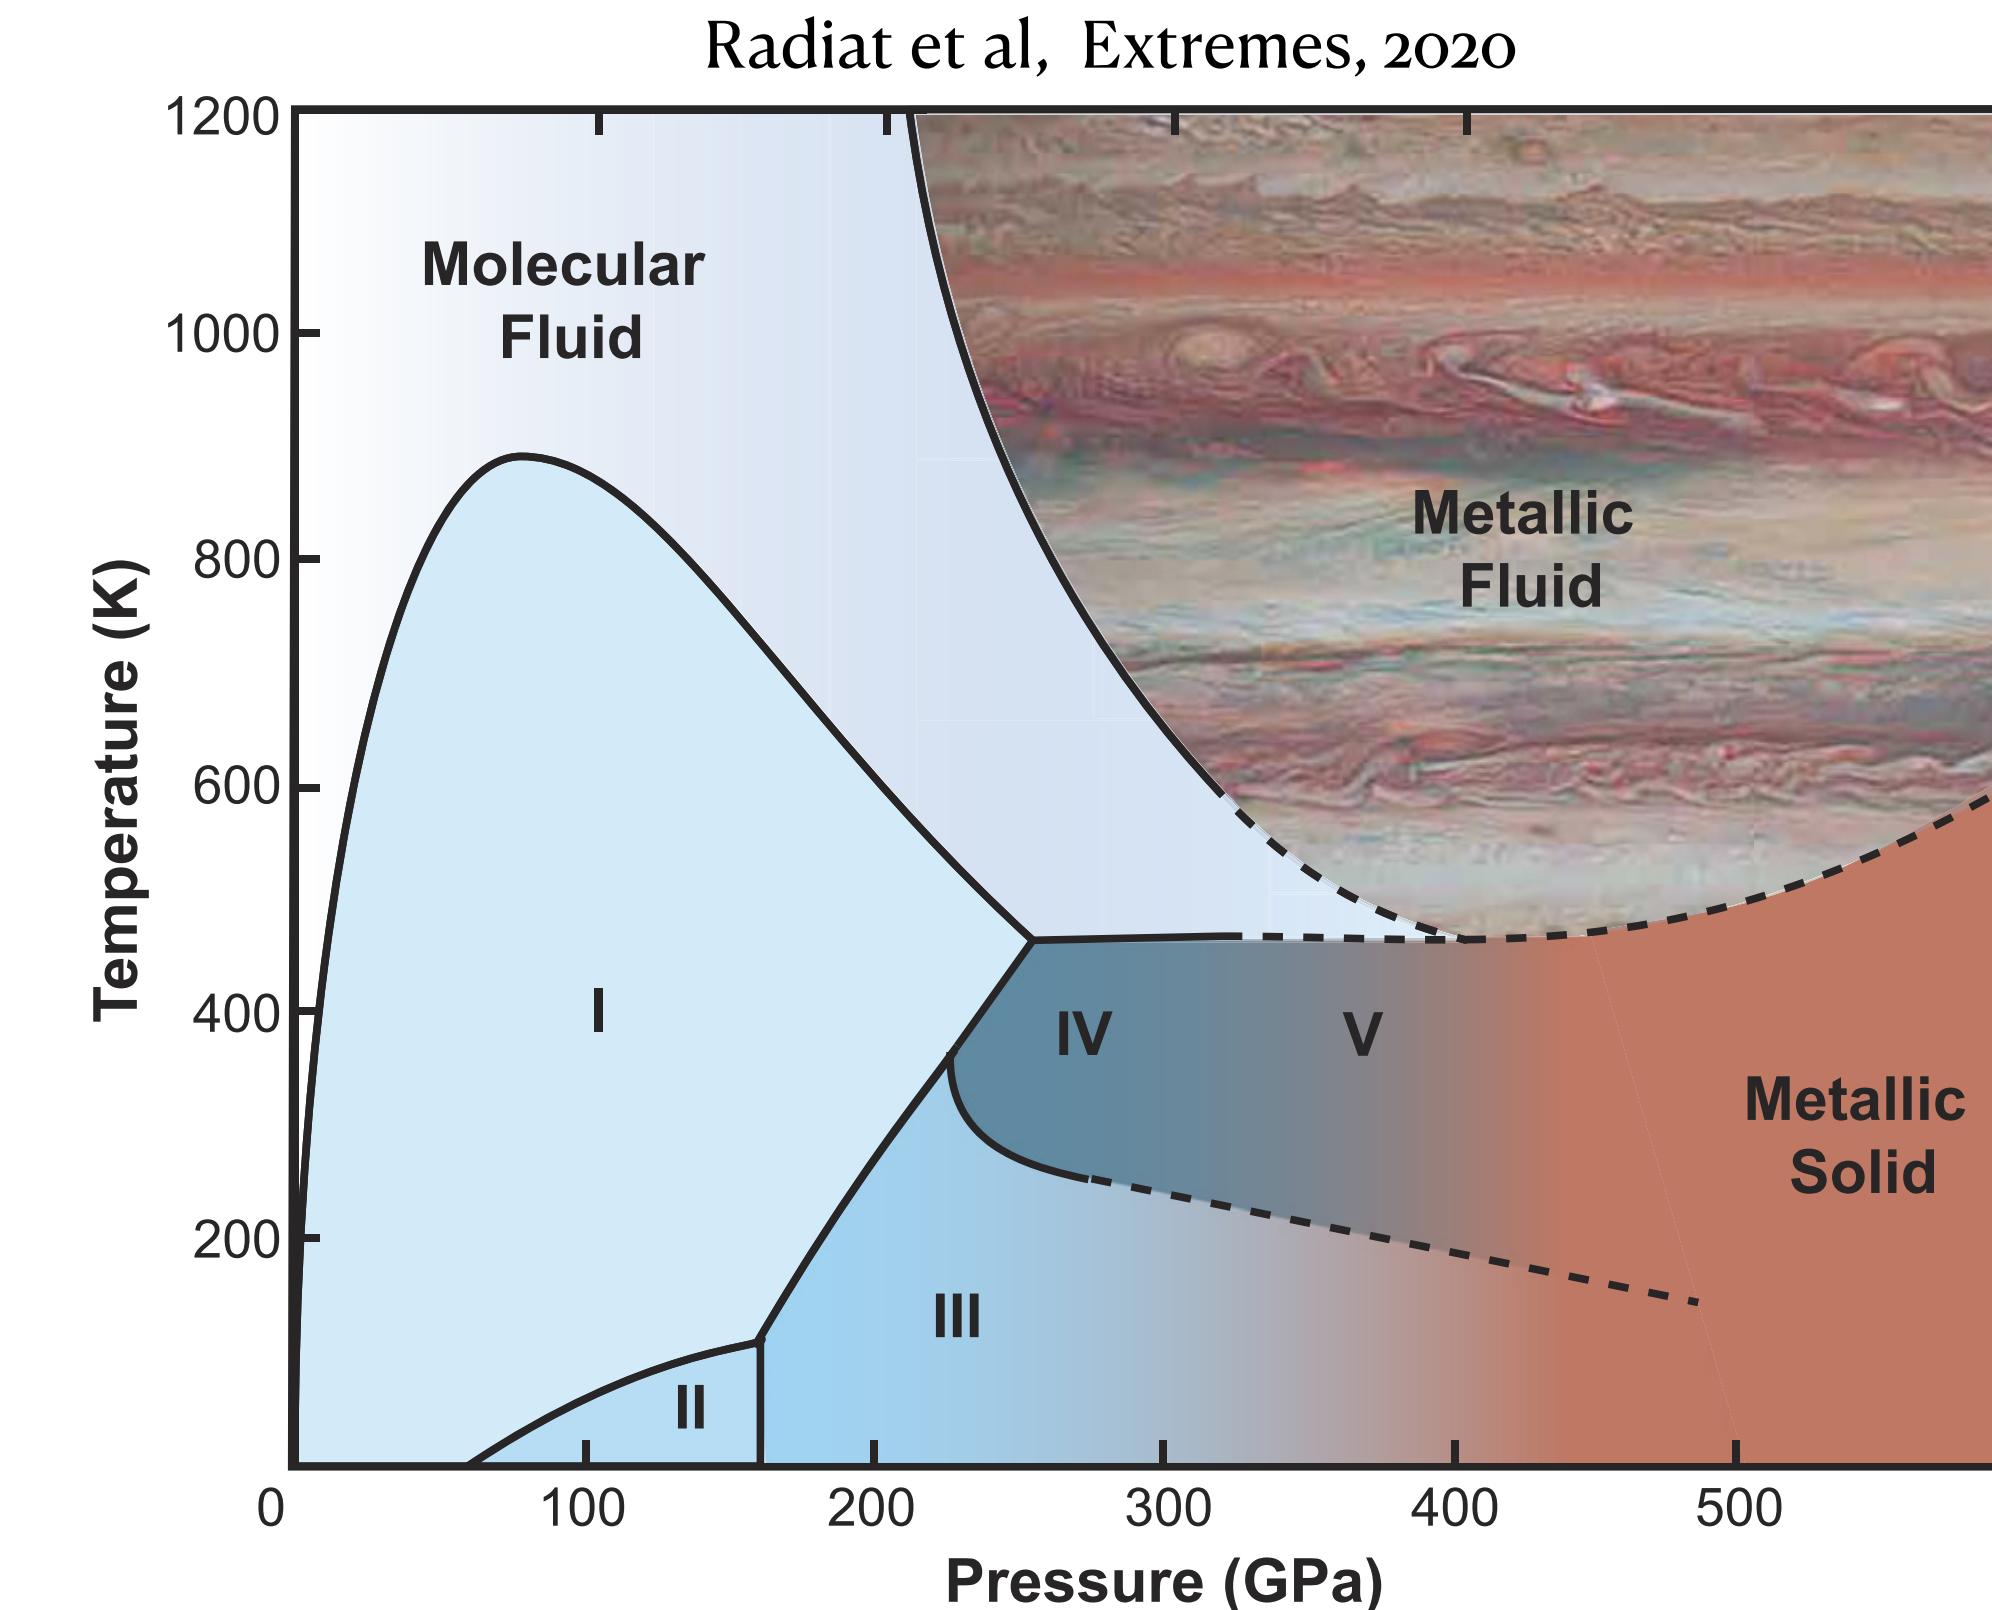
Applications

Uniform electron gas



Xie, Zhang, LW, arXiv '22

Dense hydrogen



1 to 1000: model architecture based on physics, pretraining, large scale optimization...

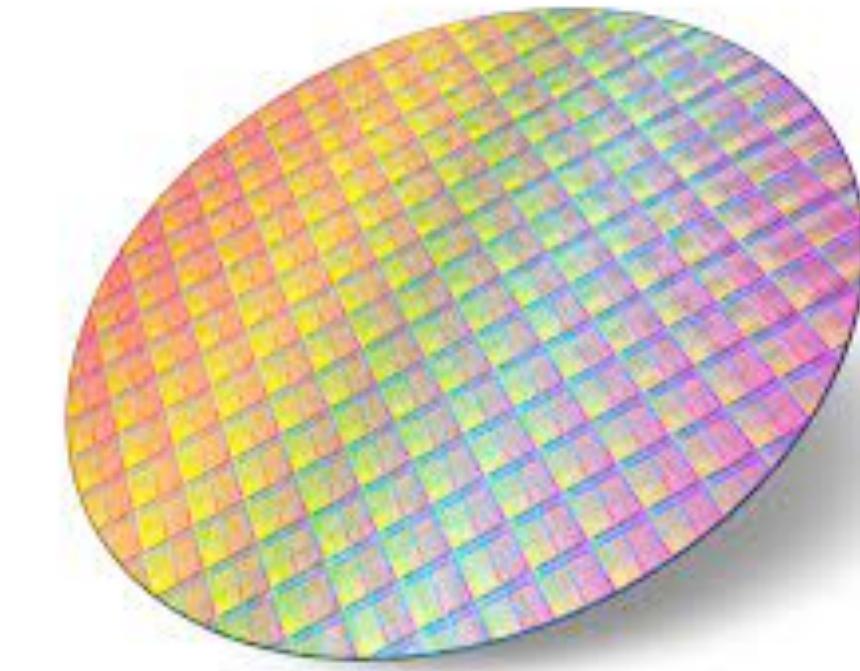
Triumph of condensed matter physics



Insulators



Metal



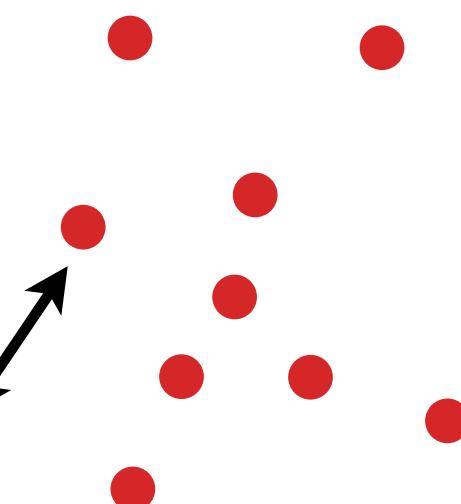
Semiconductors

Why metal is metal?



Uniform electron gas

$$H = - \sum_{i=1}^N \frac{\hbar^2 \nabla_i^2}{2m} + \sum_{i < j} \frac{e^2}{|x_i - x_j|}$$



r_s of typical metals, Richard Martin, *Electronic structure*

Z = 1	Z = 2	Z = 1	Z = 2	Z = 3	Z = 4
Li 3.23	Be 1.88		B	C 1.31	
Na 3.93	Mg 2.65		Al 2.07	Si 2.00	
K 4.86	Ca 3.27	Cu 2.67	Zn 2.31	Ga 2.19	Ge 2.08
Rb 5.20	Sr 3.56	Ag 3.02	Cd 2.59	In 2.41	Sn 2.39
Cs 5.63	Ba 3.69	Au 3.01	Hg 2.15	Tl	Pb 2.30

Metal density $2 < r_s < 6$: Coulomb repulsion
is nonperturbative compared to kinetic energy

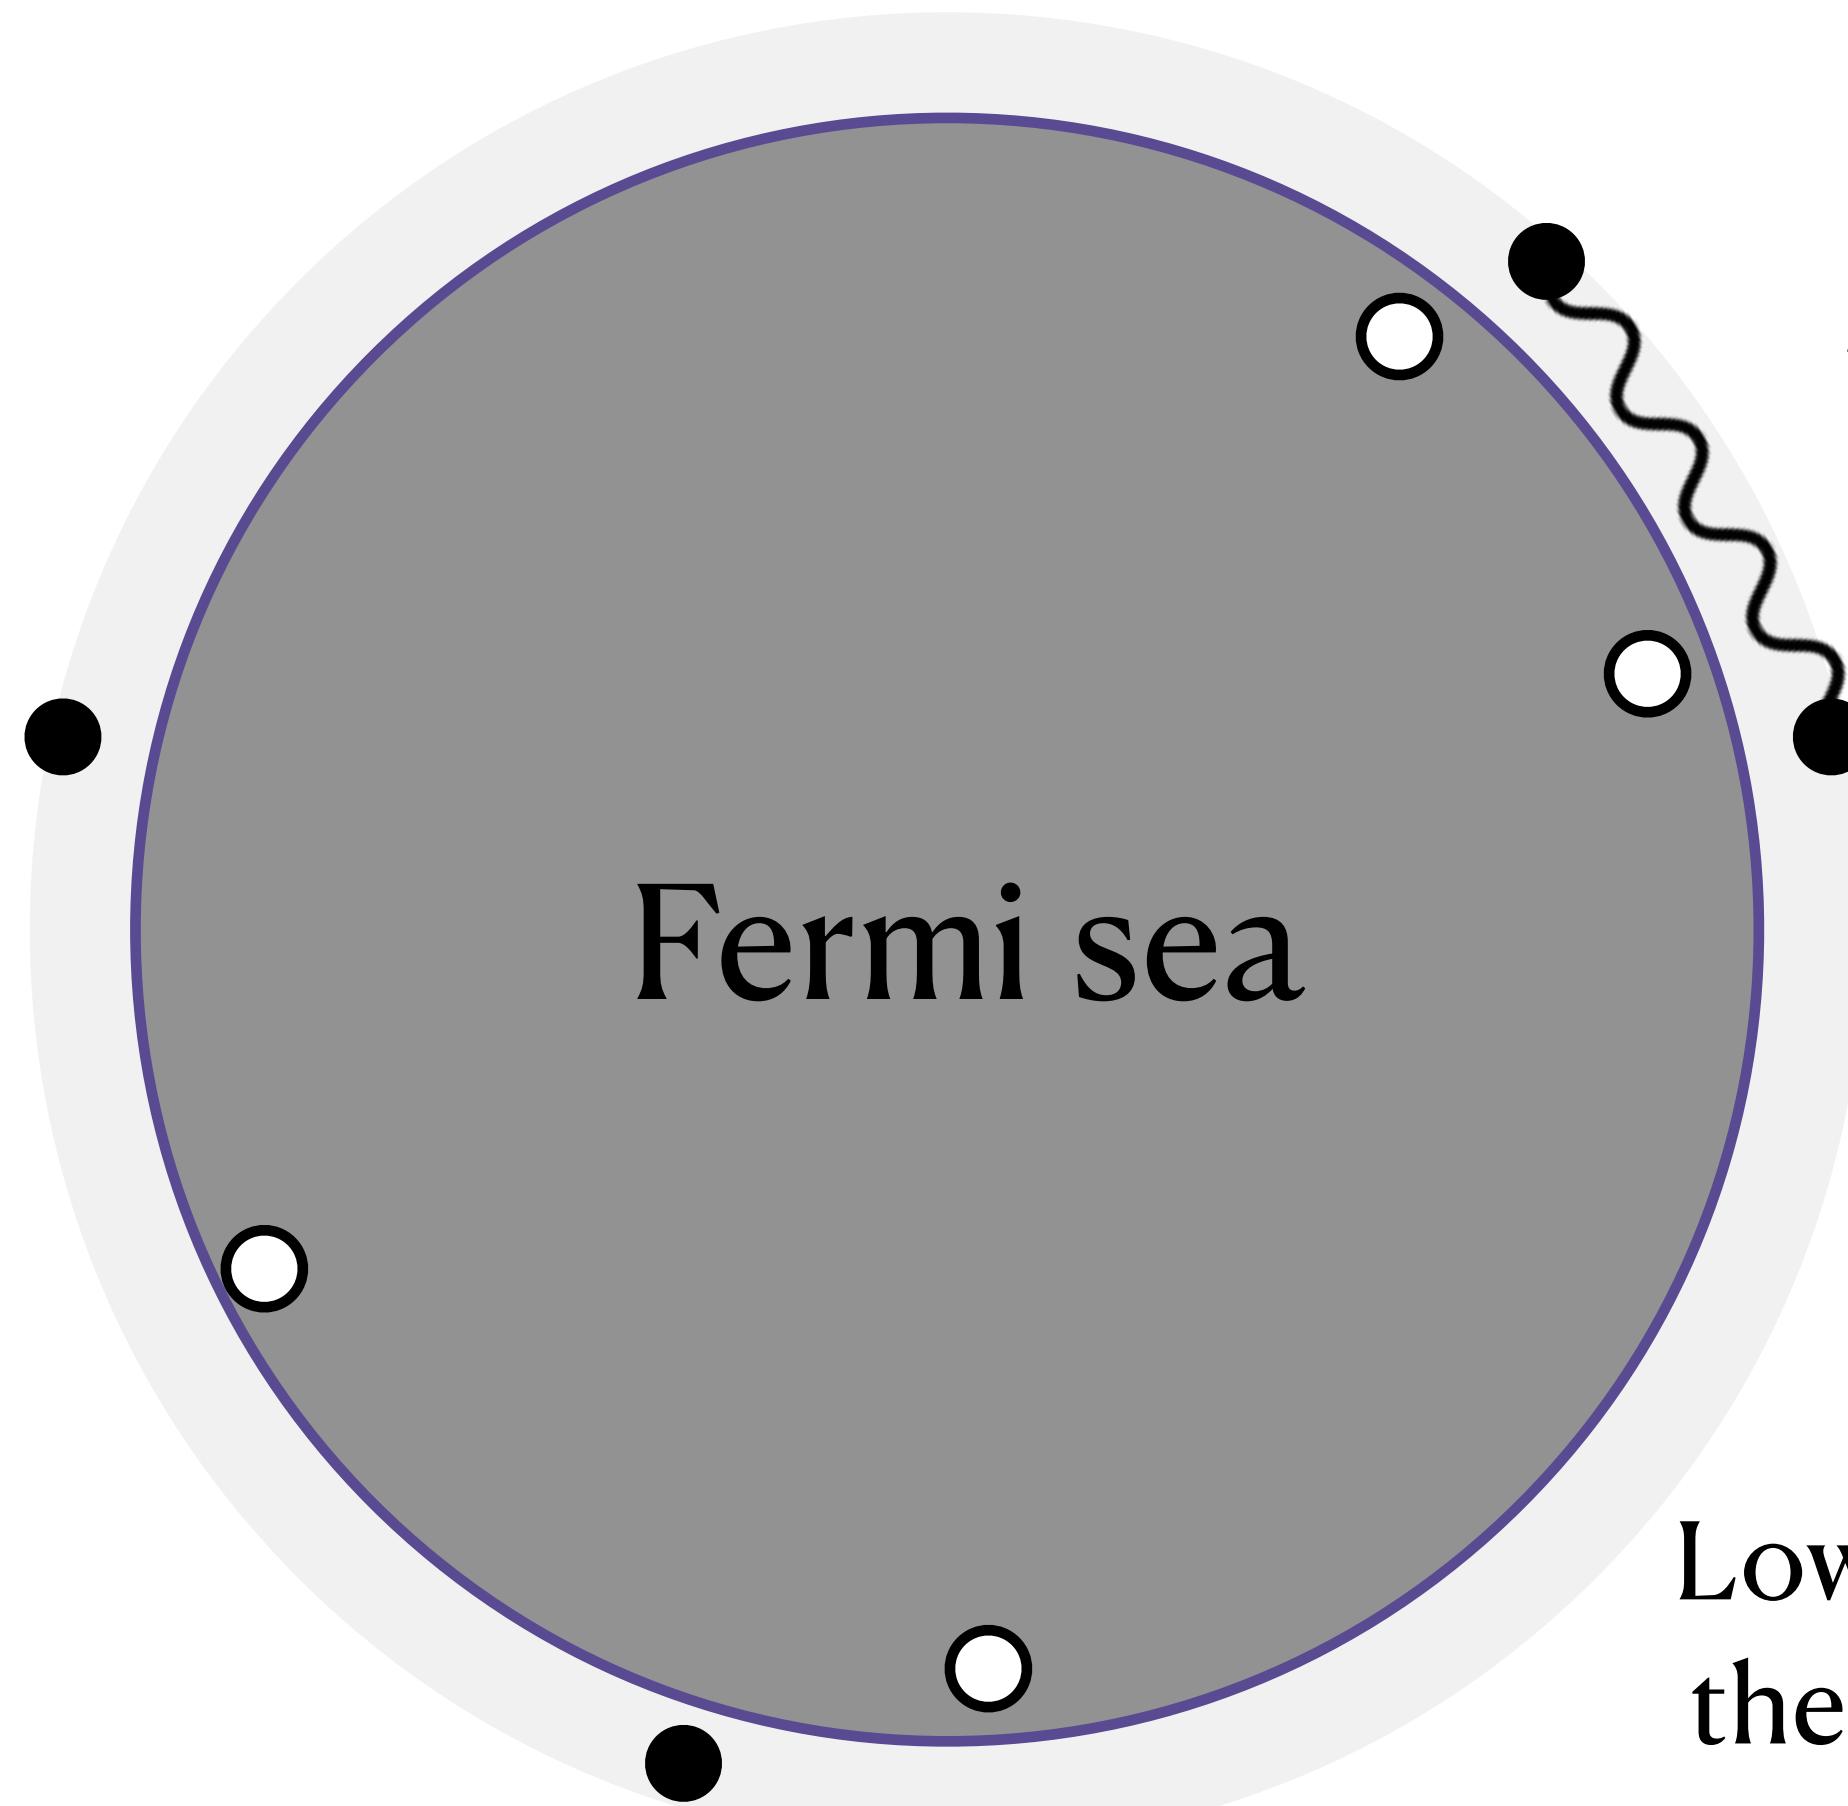
《物理》 2010年第8期

物理学中的演生现象

张广铭^{1,†} 于 绿^{2,3}

些非正常金属态可以被认为是费米液体的某些不稳定的、对称性破缺的基态. 另一方面, 费米液体理论也十分令人费解, 因为普通金属中电子之间的库仑相互作用能和费米能是一个量级, 同时比费米能量附近的能级间距要大很多. 微扰理论对如此强的相互作用已不再适应, 很难相信一个如此强大的相互作用的多电子系统能与一个无相互作用准粒子系统的行为相像. 然而, 自然界本身一遍又一遍地提醒我们: 尽管有强大的库仑相互作用, 金属的低能行为仍与一个自由准粒子系统类似, 这又是重正化群思想再奏凯歌! 直到 20 世纪 80 年代末, 由于分数量子

Landau fermi liquid theory



$$T \ll T_F \lesssim \frac{e^2}{r_s}$$

Low energy excited states labeled in
the same way as the ideal Fermi gas

$$K = \{k_1, k_2, \dots, k_N\}$$

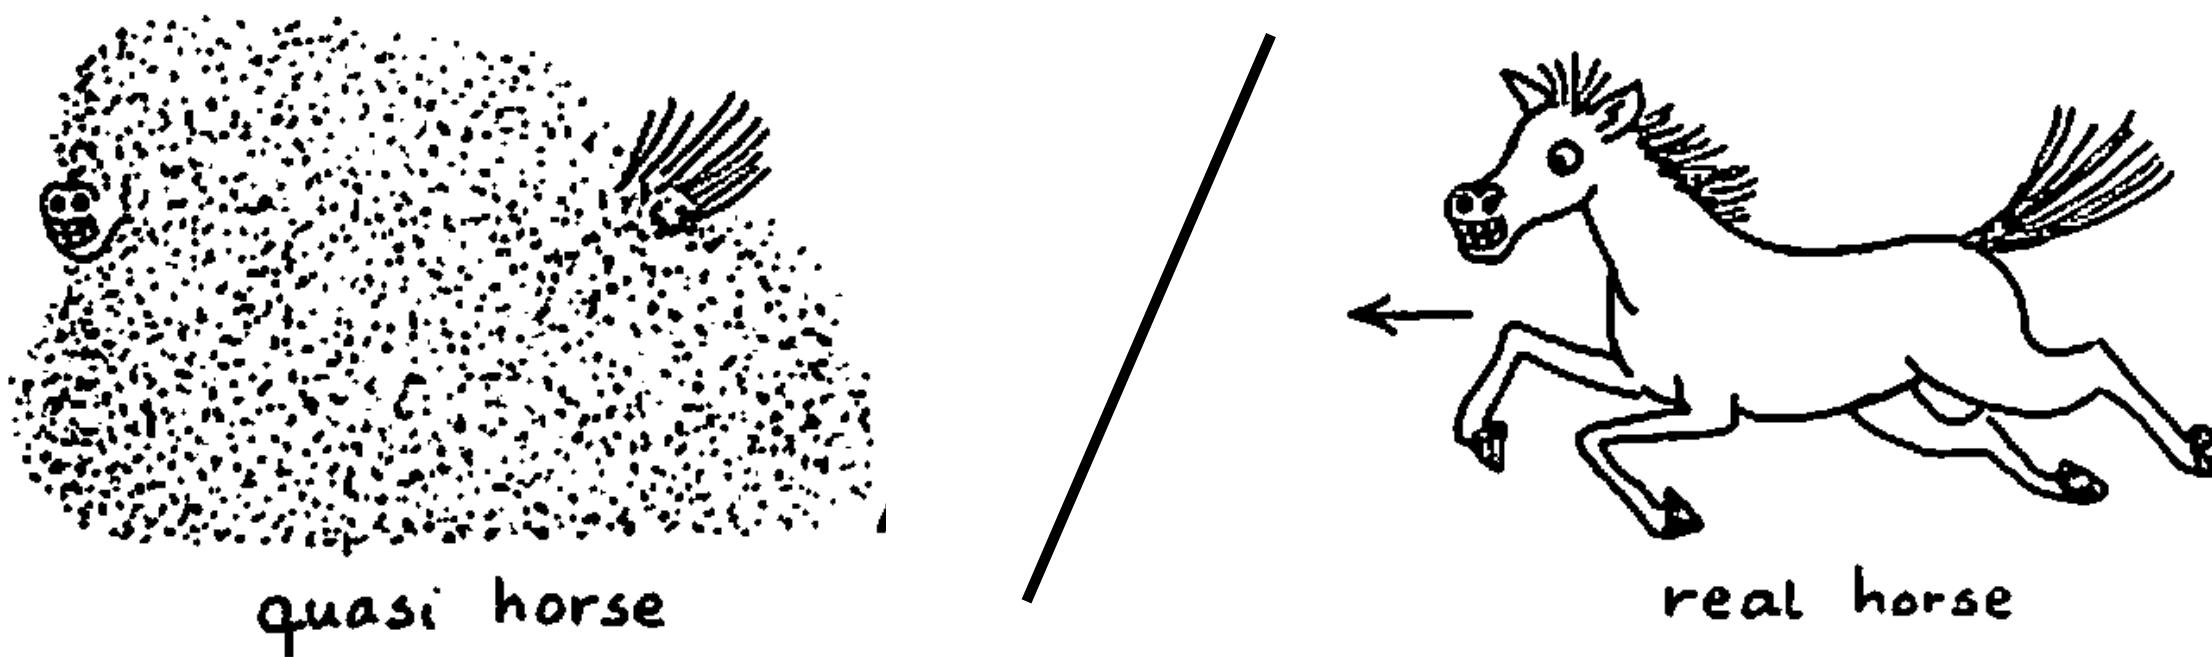
Physics happens around the Fermi surface with strongly constrained phase-space

Have we known everything about a Fermi liquid ?

No!

Quasi-particles effective mass

$$\frac{m^*}{m} =$$



Richard D. Mattuck
*A Guide to Feynman
Diagrams in the Many-
body Problem*

A fundamental quantity appears in nearly all physical properties of a Fermi liquid

$$N(0)$$

Density of states

$$S$$

entropy

$$c_V$$

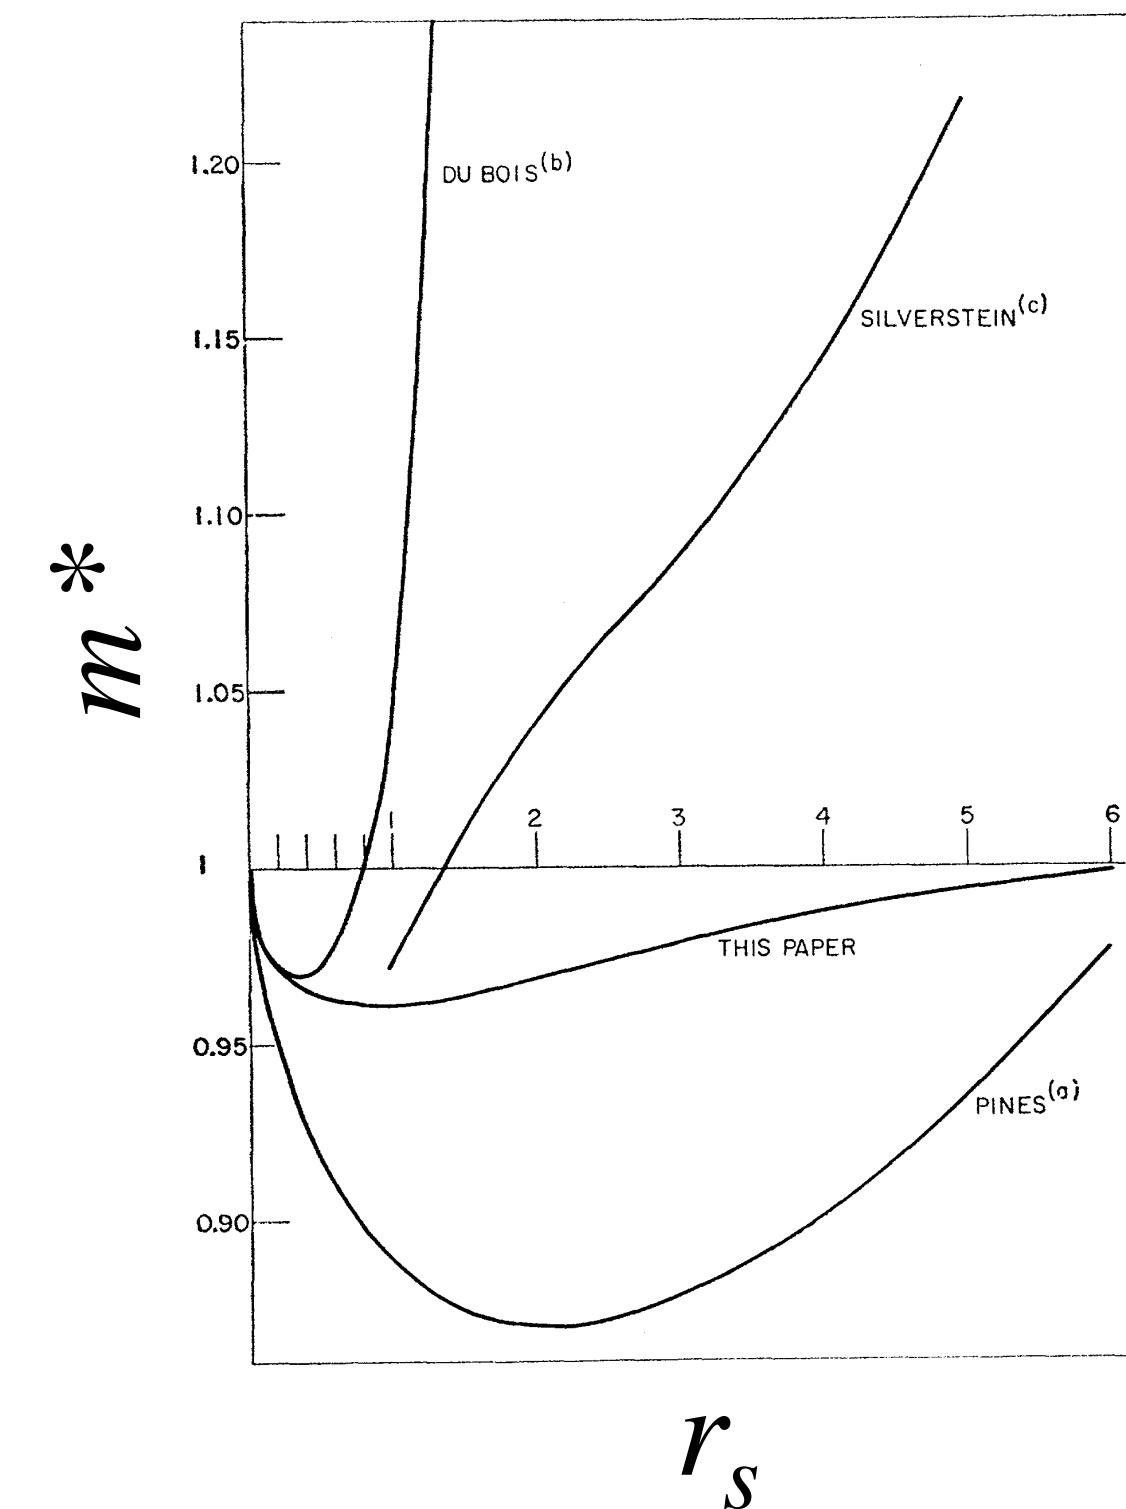
specific heat

$$\chi$$

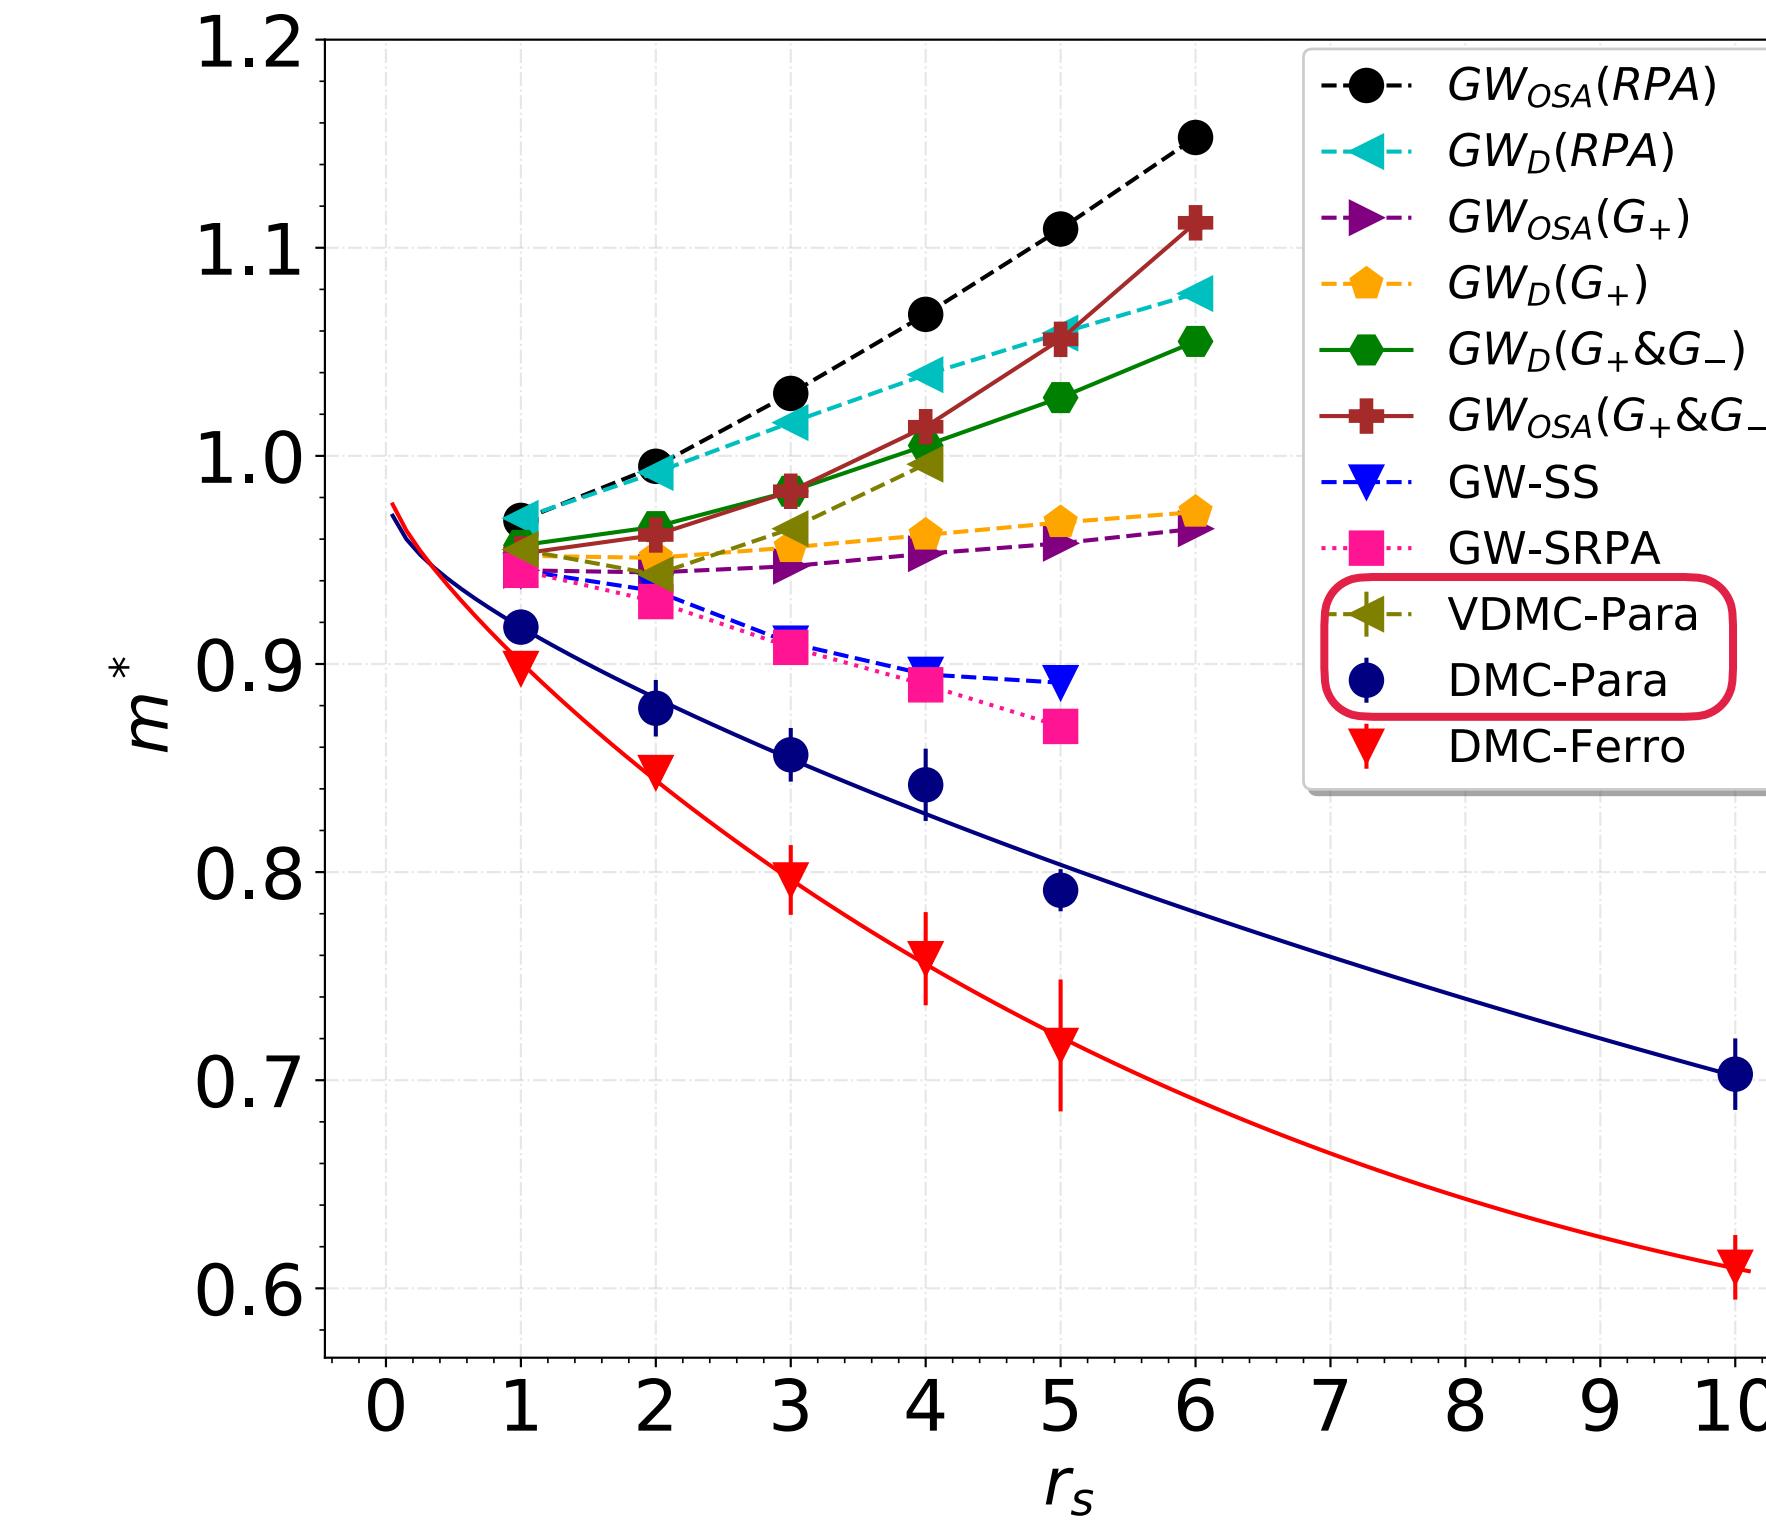
magnetic susceptibility

Quasi-particles effective mass of 3d electron gas

Hedin Phy. Rev. 1965



Azadi, Drummond, Foulkes, PRL 2021



> 50 years of conflicting results !

Two-dimensional electron gas experiments

VOLUME 91, NUMBER 4

PHYSICAL REVIEW LETTERS

week ending
25 JULY 2003

Spin-Independent Origin of the Strongly Enhanced Effective Mass in a Dilute 2D Electron System

A. A. Shashkin,* Maryam Rahimi, S. Anissimova, and S.V. Kravchenko

Physics Department, Northeastern University, Boston, Massachusetts 02115, USA

V.T. Dolgopolov

Institute of Solid State Physics, Chernogolovka, Moscow District 142432, Russia

T. M. Klapwijk

Department of Applied Physics, Delft University of Technology, 2628 CJ Delft, The Netherlands

(Received 13 January 2003; published 24 July 2003)

$$m^*/m > 1$$



PRL 101, 026402 (2008)

PHYSICAL REVIEW LETTERS

week ending
11 JULY 2008

Effective Mass Suppression in Dilute, Spin-Polarized Two-Dimensional Electron Systems

Medini Padmanabhan, T. Gokmen, N. C. Bishop, and M. Shayegan

Department of Electrical Engineering, Princeton University, Princeton, New Jersey 08544, USA

(Received 19 September 2007; published 7 July 2008)

$$m^*/m < 1$$

Layer thickness, valley, disorder, spin-orbit coupling...

m^* from low temperature entropy

Eich, Holzmann, Vignale, PRB '17

$$\frac{m^*}{m} = \left(1 - \frac{\partial \Sigma}{\partial \omega}\right) \left(1 + \frac{m}{k} \frac{\partial \Sigma}{\partial k}\right)^{-1}$$

$$\frac{m^*}{m} = k_F / (d\varepsilon/dk)_{k_F}$$

$$S = \frac{\pi^2 k_B}{3} \frac{m^*}{m} \frac{T}{T_F}$$
$$\Rightarrow \frac{m^*}{m} = \frac{s}{s_0} \quad \begin{array}{l} \text{interacting electrons} \\ \text{noninteracting electrons} \end{array}$$

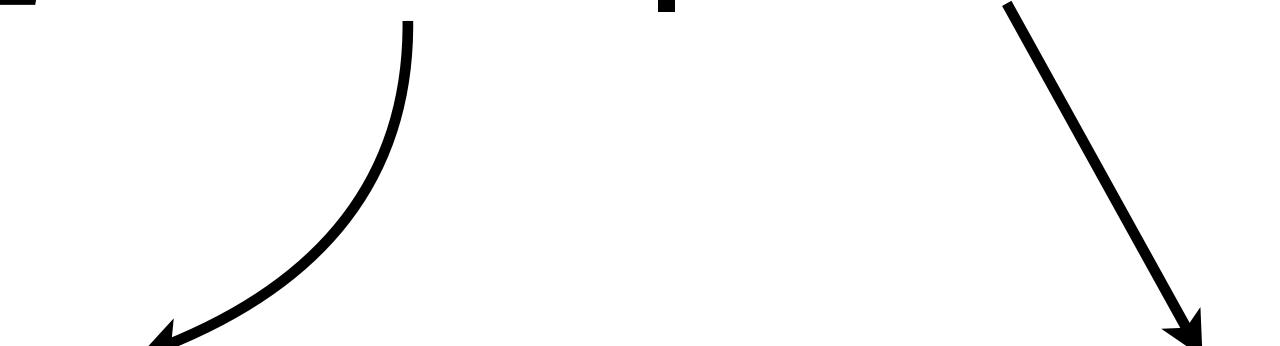
Not an easy task due to the lack of reliable methods
for low-temperature electron gases with intermediate density

computing specific heat also works, but that often requires differentiating (noisy) energies

Deep generative models for the variational density matrix

$$\rho = \sum_K p(K) |\Psi_K\rangle\langle\Psi_K|$$

Normalized probability distribution Orthonormal many-electron basis



$$\textcircled{1} \quad \sum_K p(K) = 1 \quad \textcircled{2} \quad \langle \Psi_K | \Psi_{K'} \rangle = \delta_{K,K'}$$

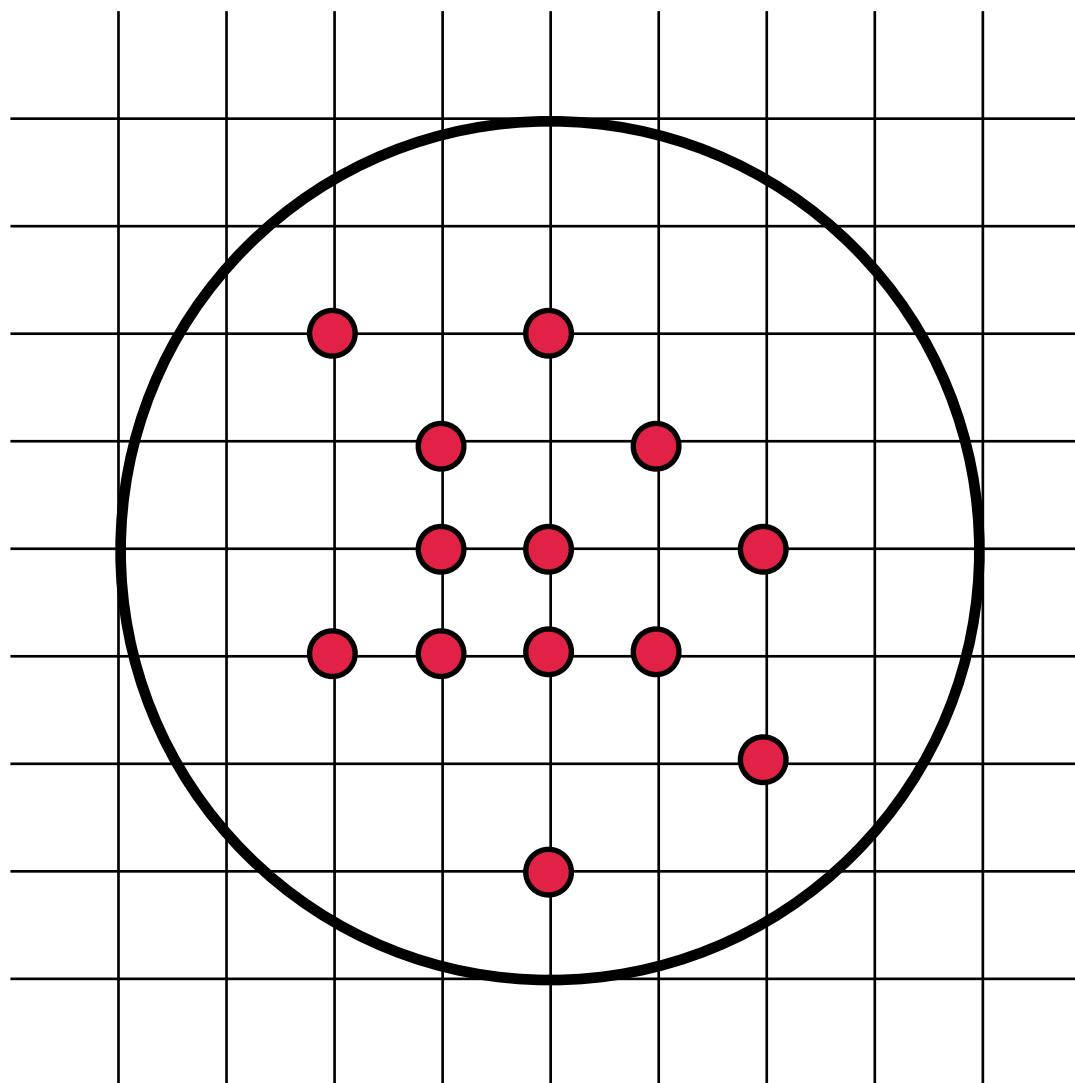
There will also be interesting twists for physics considerations

①

Autoregressive model for $p(K)$

Fermionic
occupation
in k-space

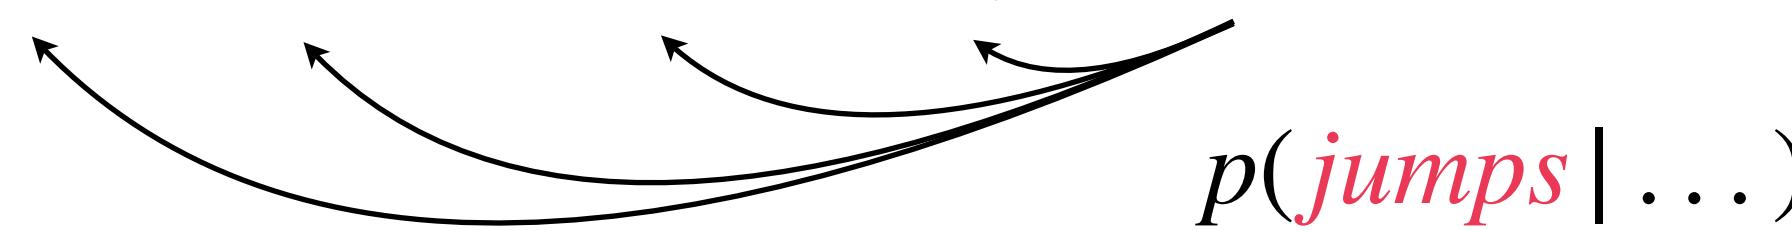
$$K = \{k_1, k_2, \dots, k_N\}$$



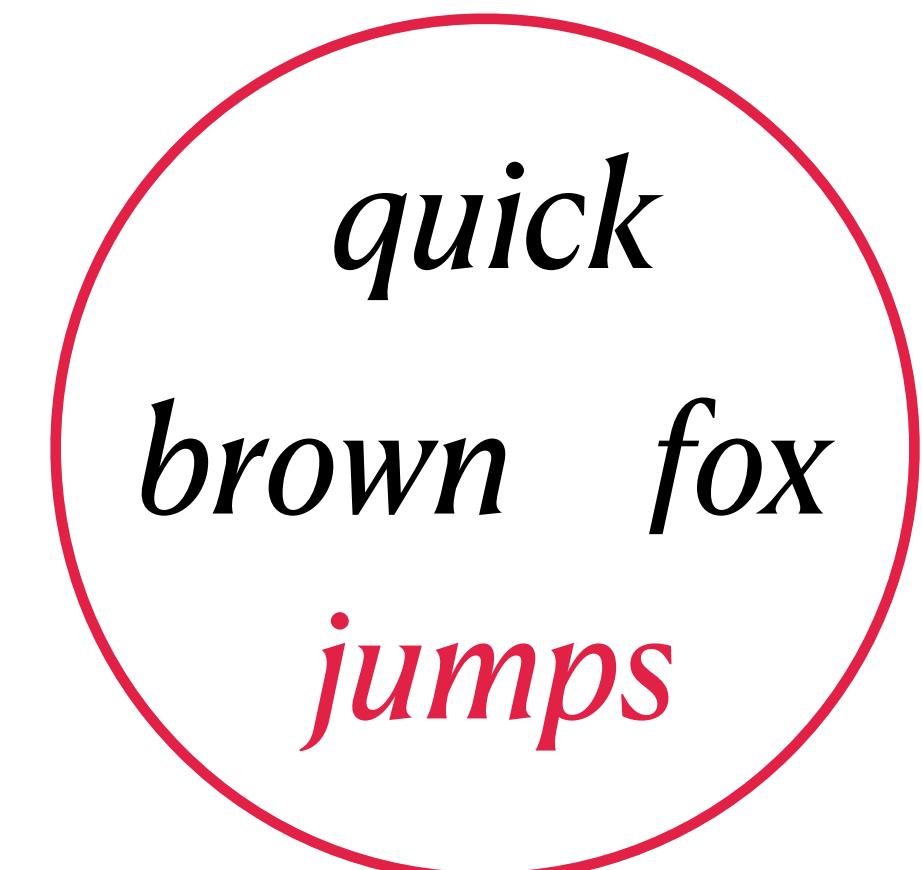
$$p(K) = p(k_1)p(k_2 | k_1)p(k_3 | k_1, k_2)\cdots$$

Wu, LW, Zhang, PRL '19

“... *quick brown fox jumps* ...”

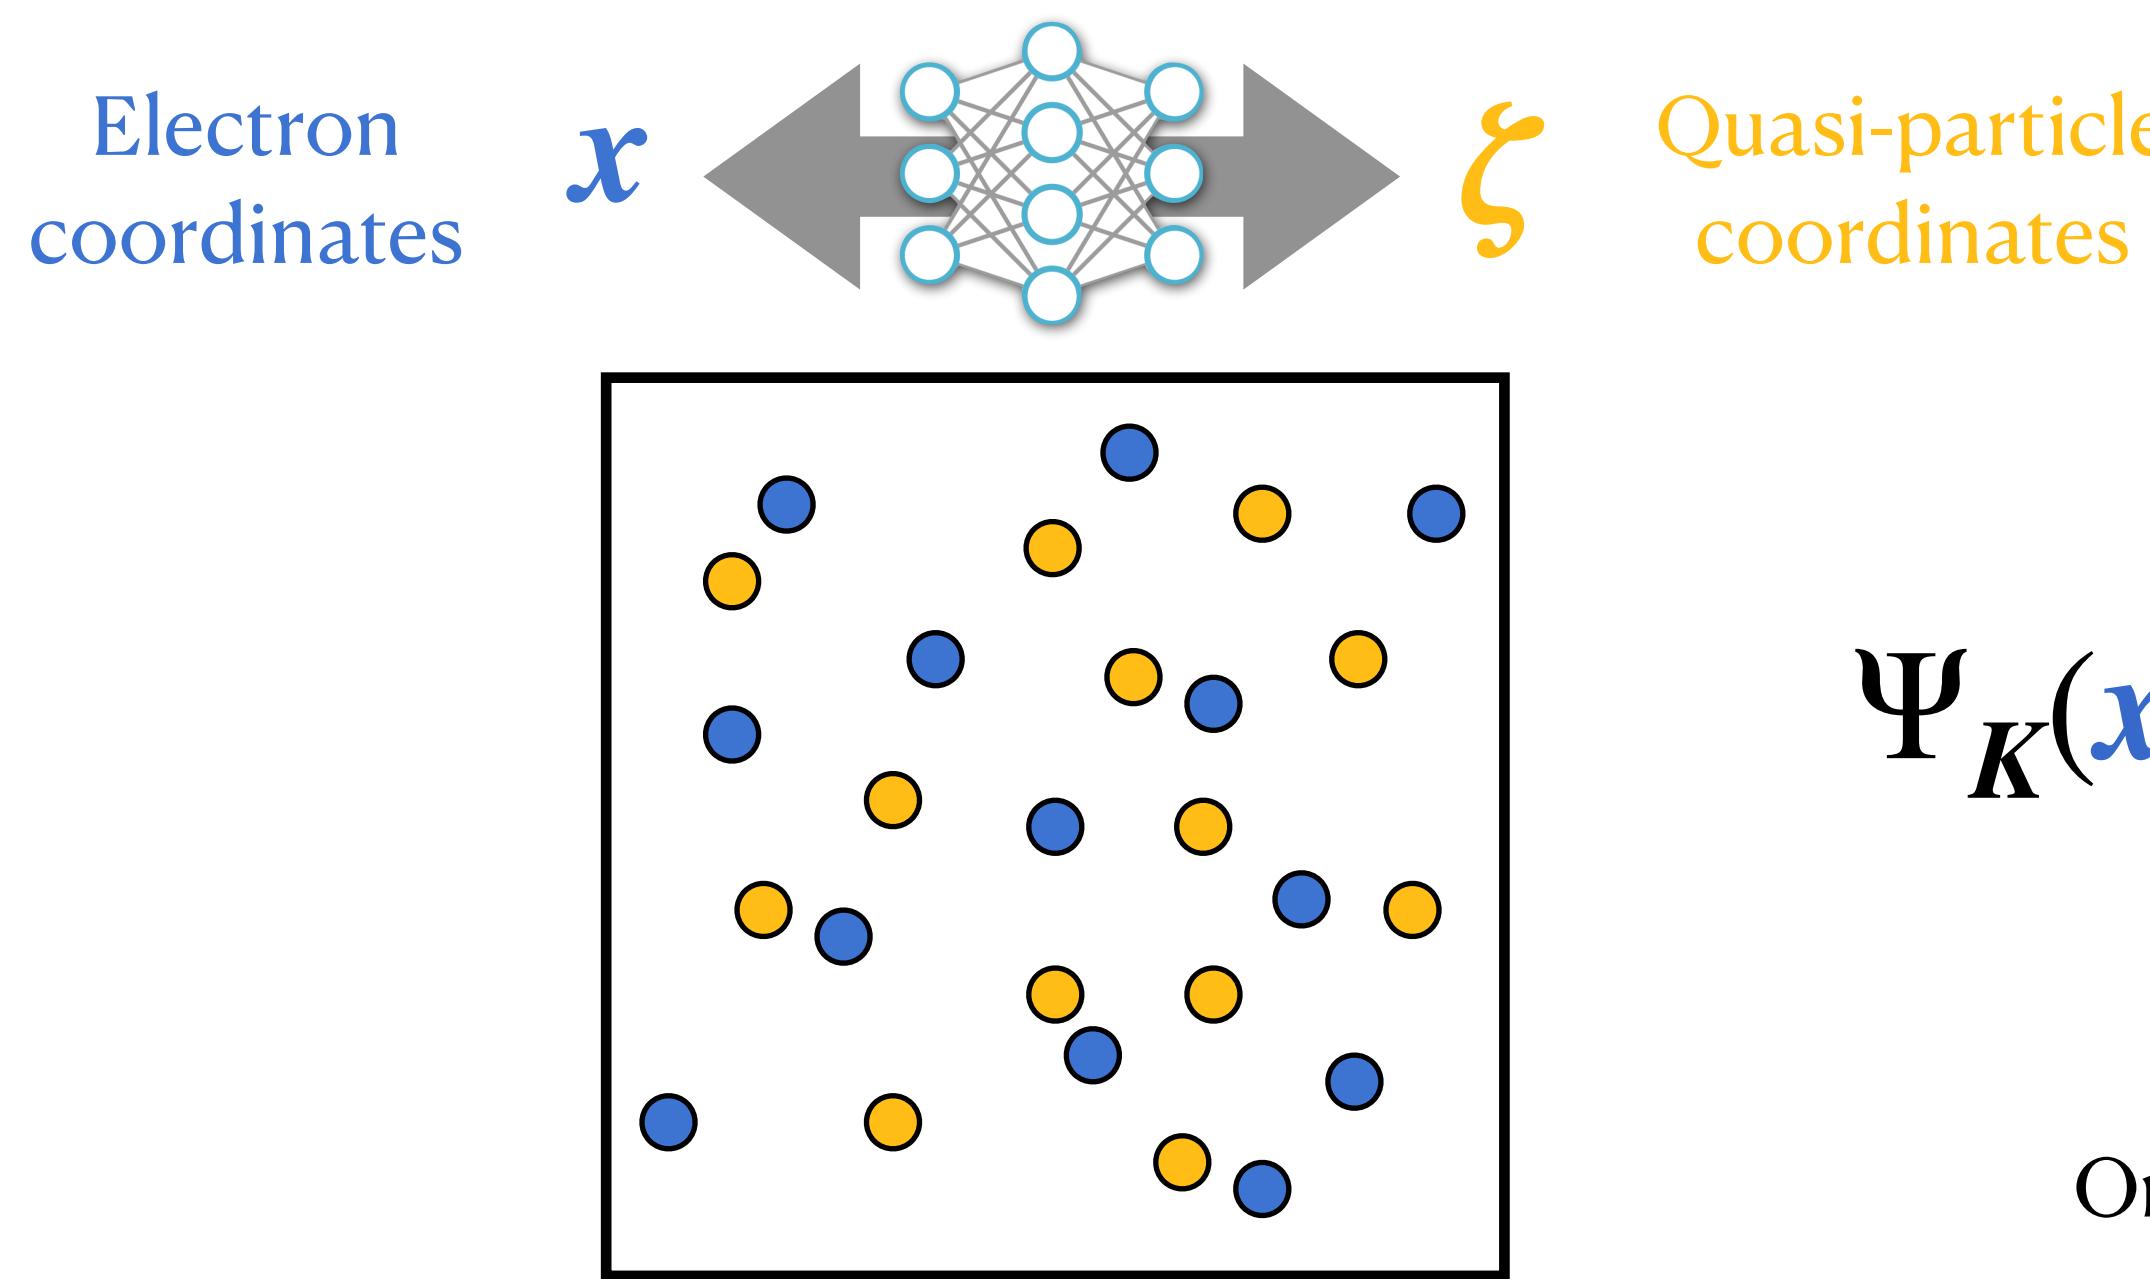


# of particles	# of words
Momentum cutoff	Vocabulary
Entropy	Negative log-likelihood



Twist: we are modeling a *set of words* with no repetitions and no order

② Normalizing flow for $|\Psi_K\rangle$



$$\Psi_K(\boldsymbol{x}) = \frac{\det(e^{ik_i \cdot \boldsymbol{\zeta}_j})}{\sqrt{N!}} \cdot \left| \det \left(\frac{\partial \boldsymbol{\zeta}}{\partial \boldsymbol{x}} \right) \right|^{\frac{1}{2}}$$

Orthonormal many-body states

Jacobian of the transformation

Twist: the flow should be permutation equivariant for fermionic coordinates

we use FermiNet layer Pfau et al, 1909.02487

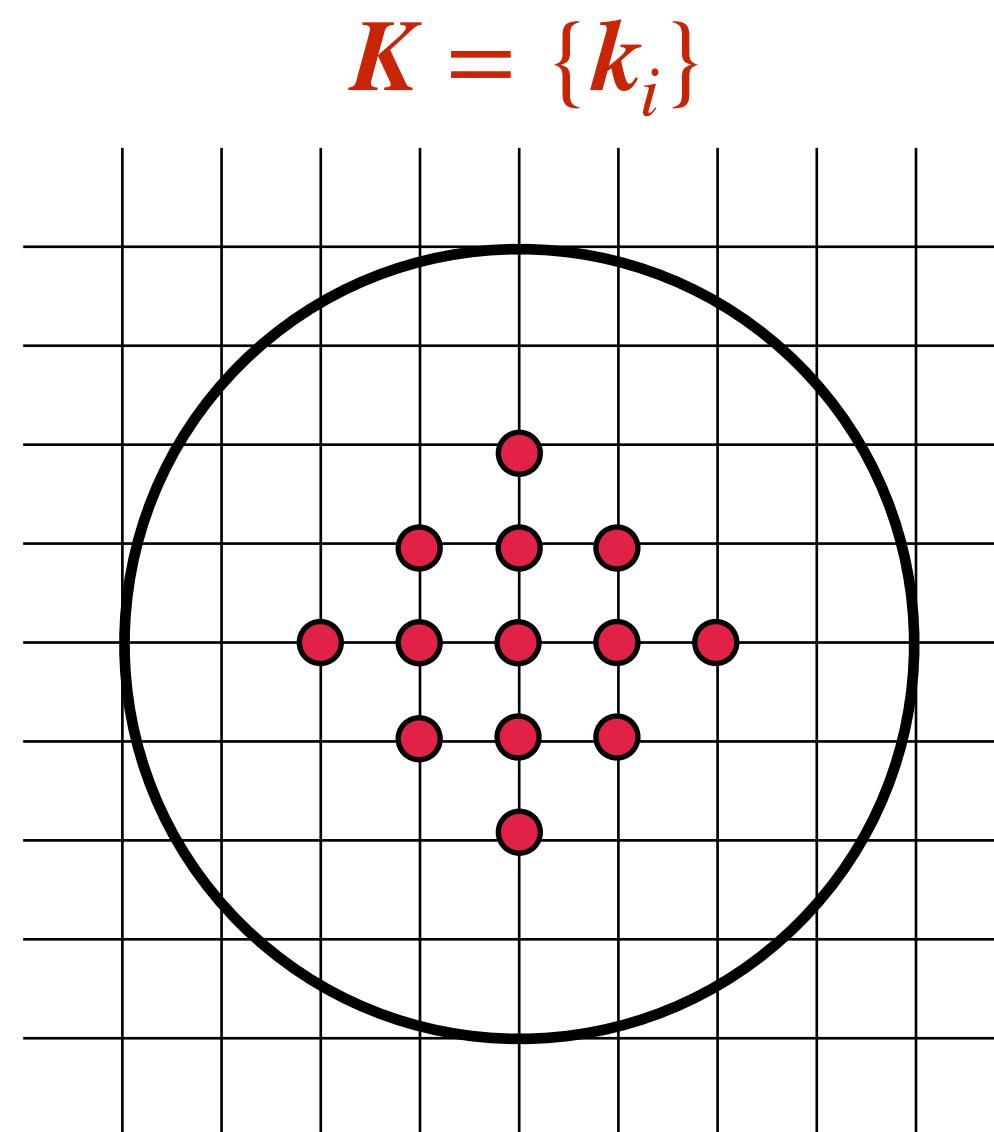
The objective function

$$F = \mathbb{E}_{K \sim p(K)} \left[\frac{1}{\beta} \ln p(K) + \mathbb{E}_{x \sim |\langle x | \Psi_K \rangle|^2} \left[\frac{\langle x | H | \Psi_K \rangle}{\langle x | \Psi_K \rangle} \right] \right]$$

↓ ↓
Boltzmann Born
distribution rule

Jointly optimize $|\Psi_K\rangle$ and $p(K)$ to minimize the variational free energy

Limiting case 1: Interacting electrons at T=0



$p(K) = 1$ only for the closed shell momentum configuration

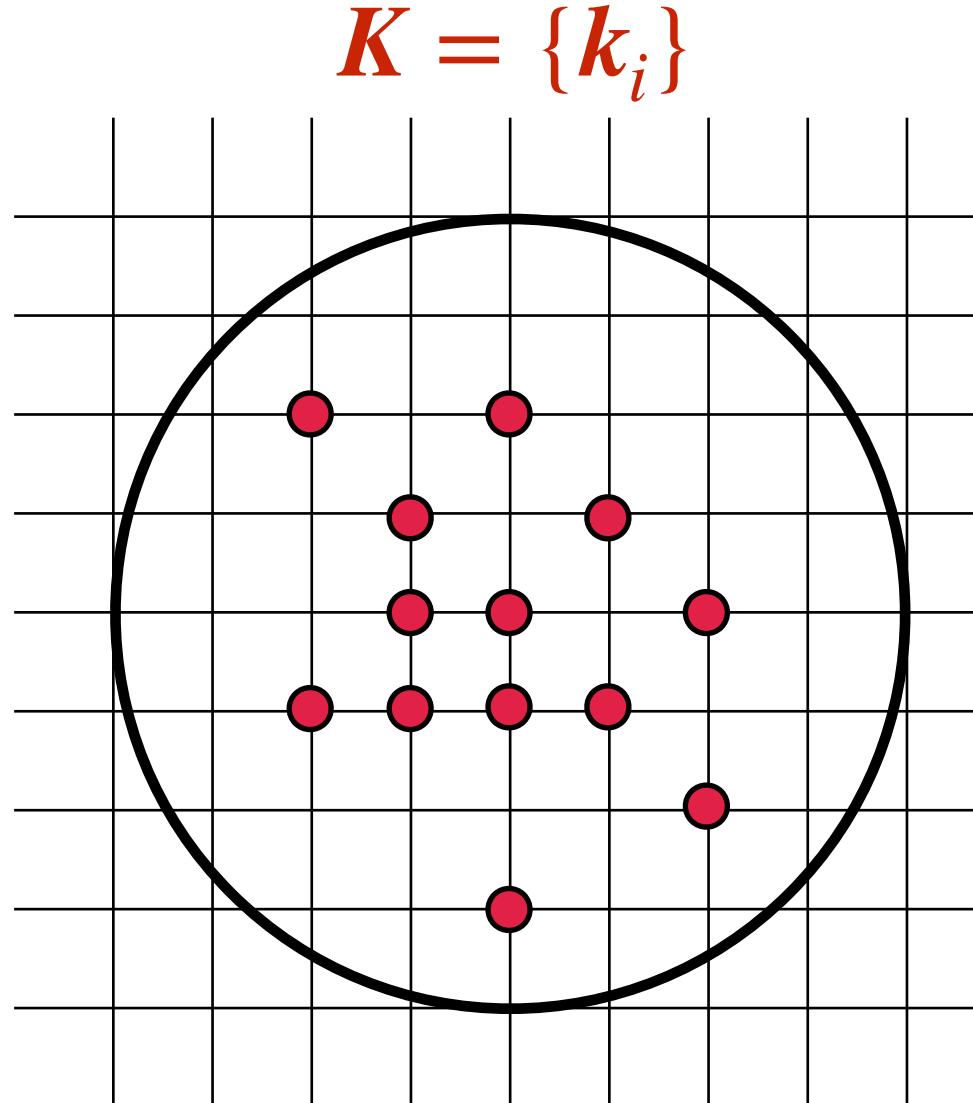
$$E = \mathbb{E}_{x \sim |\Psi_K(x)|^2} \left[\frac{\langle x | H | \Psi_K \rangle}{\langle x | \Psi_K \rangle} \right]$$

A diagram of a neural network. On the left, a blue arrow labeled x points into a central box containing a grid of blue circles connected by lines. From the right side of the box, two grey arrows point outwards, one pointing left and one pointing right, labeled ζ .

Reduces to ground state variational Monte Carlo
with a single normalizing flow wavefunction

Limiting case 2: Noninteracting electrons at T>0

$$R = \zeta$$



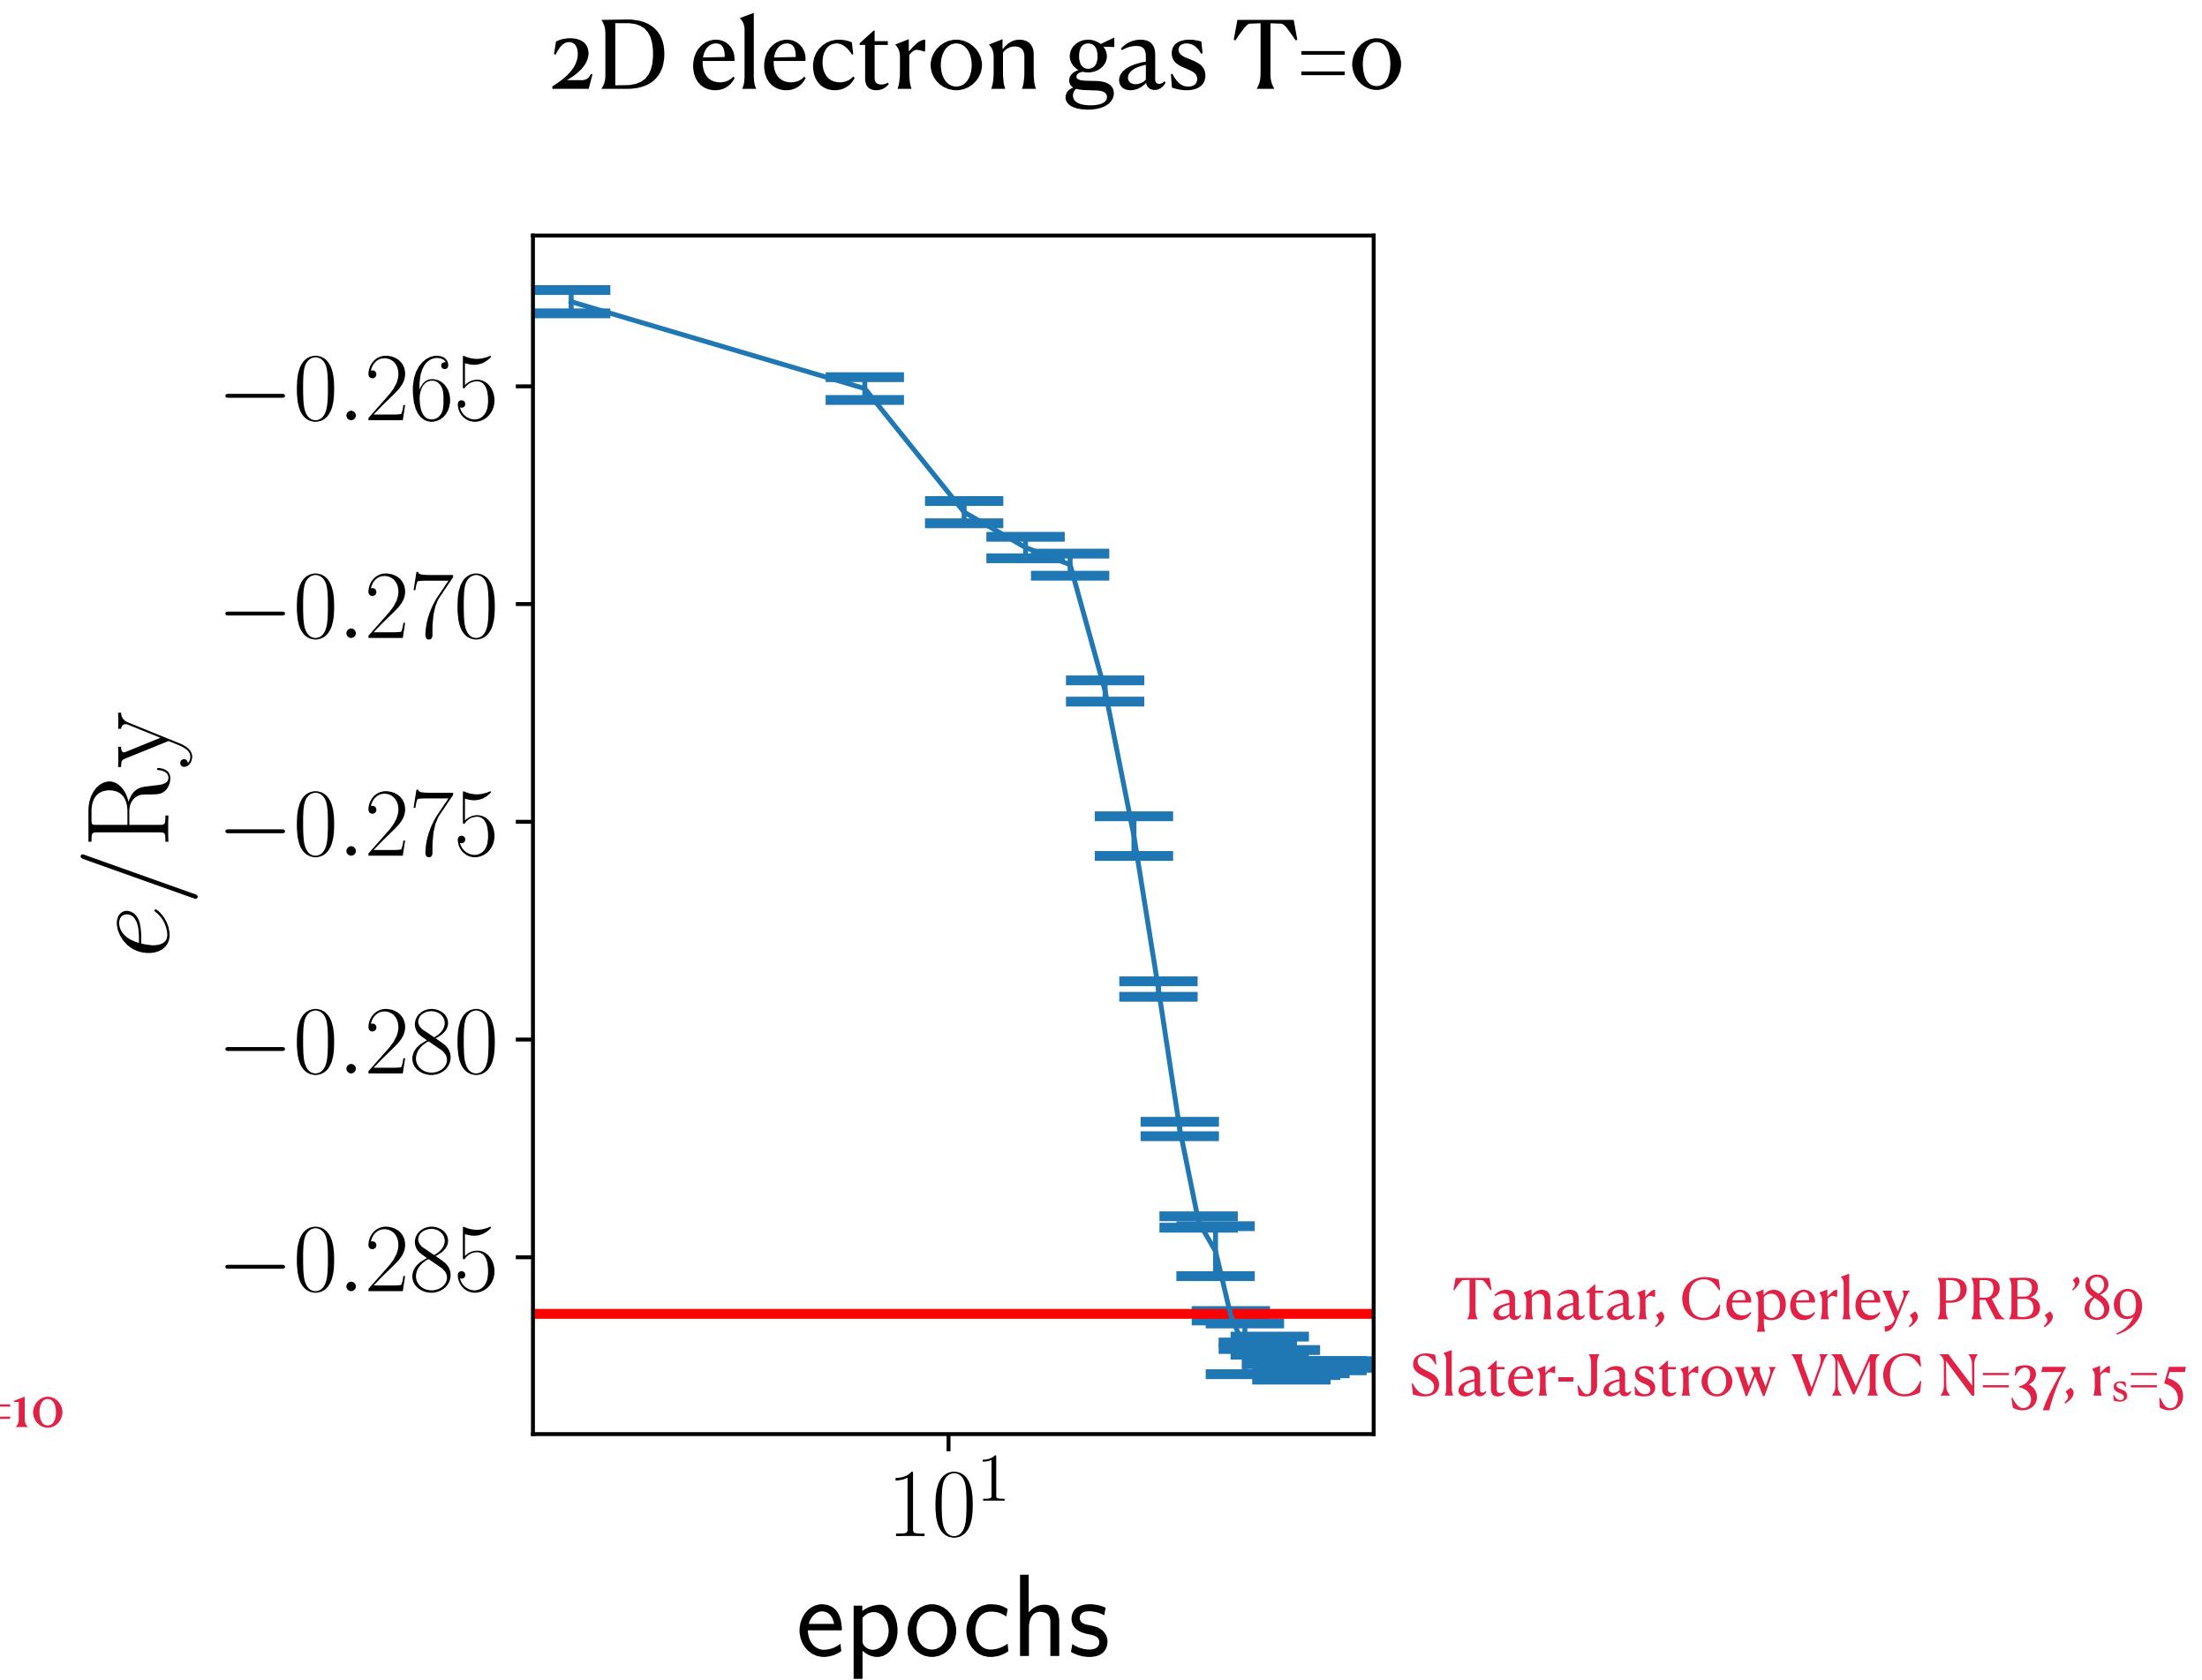
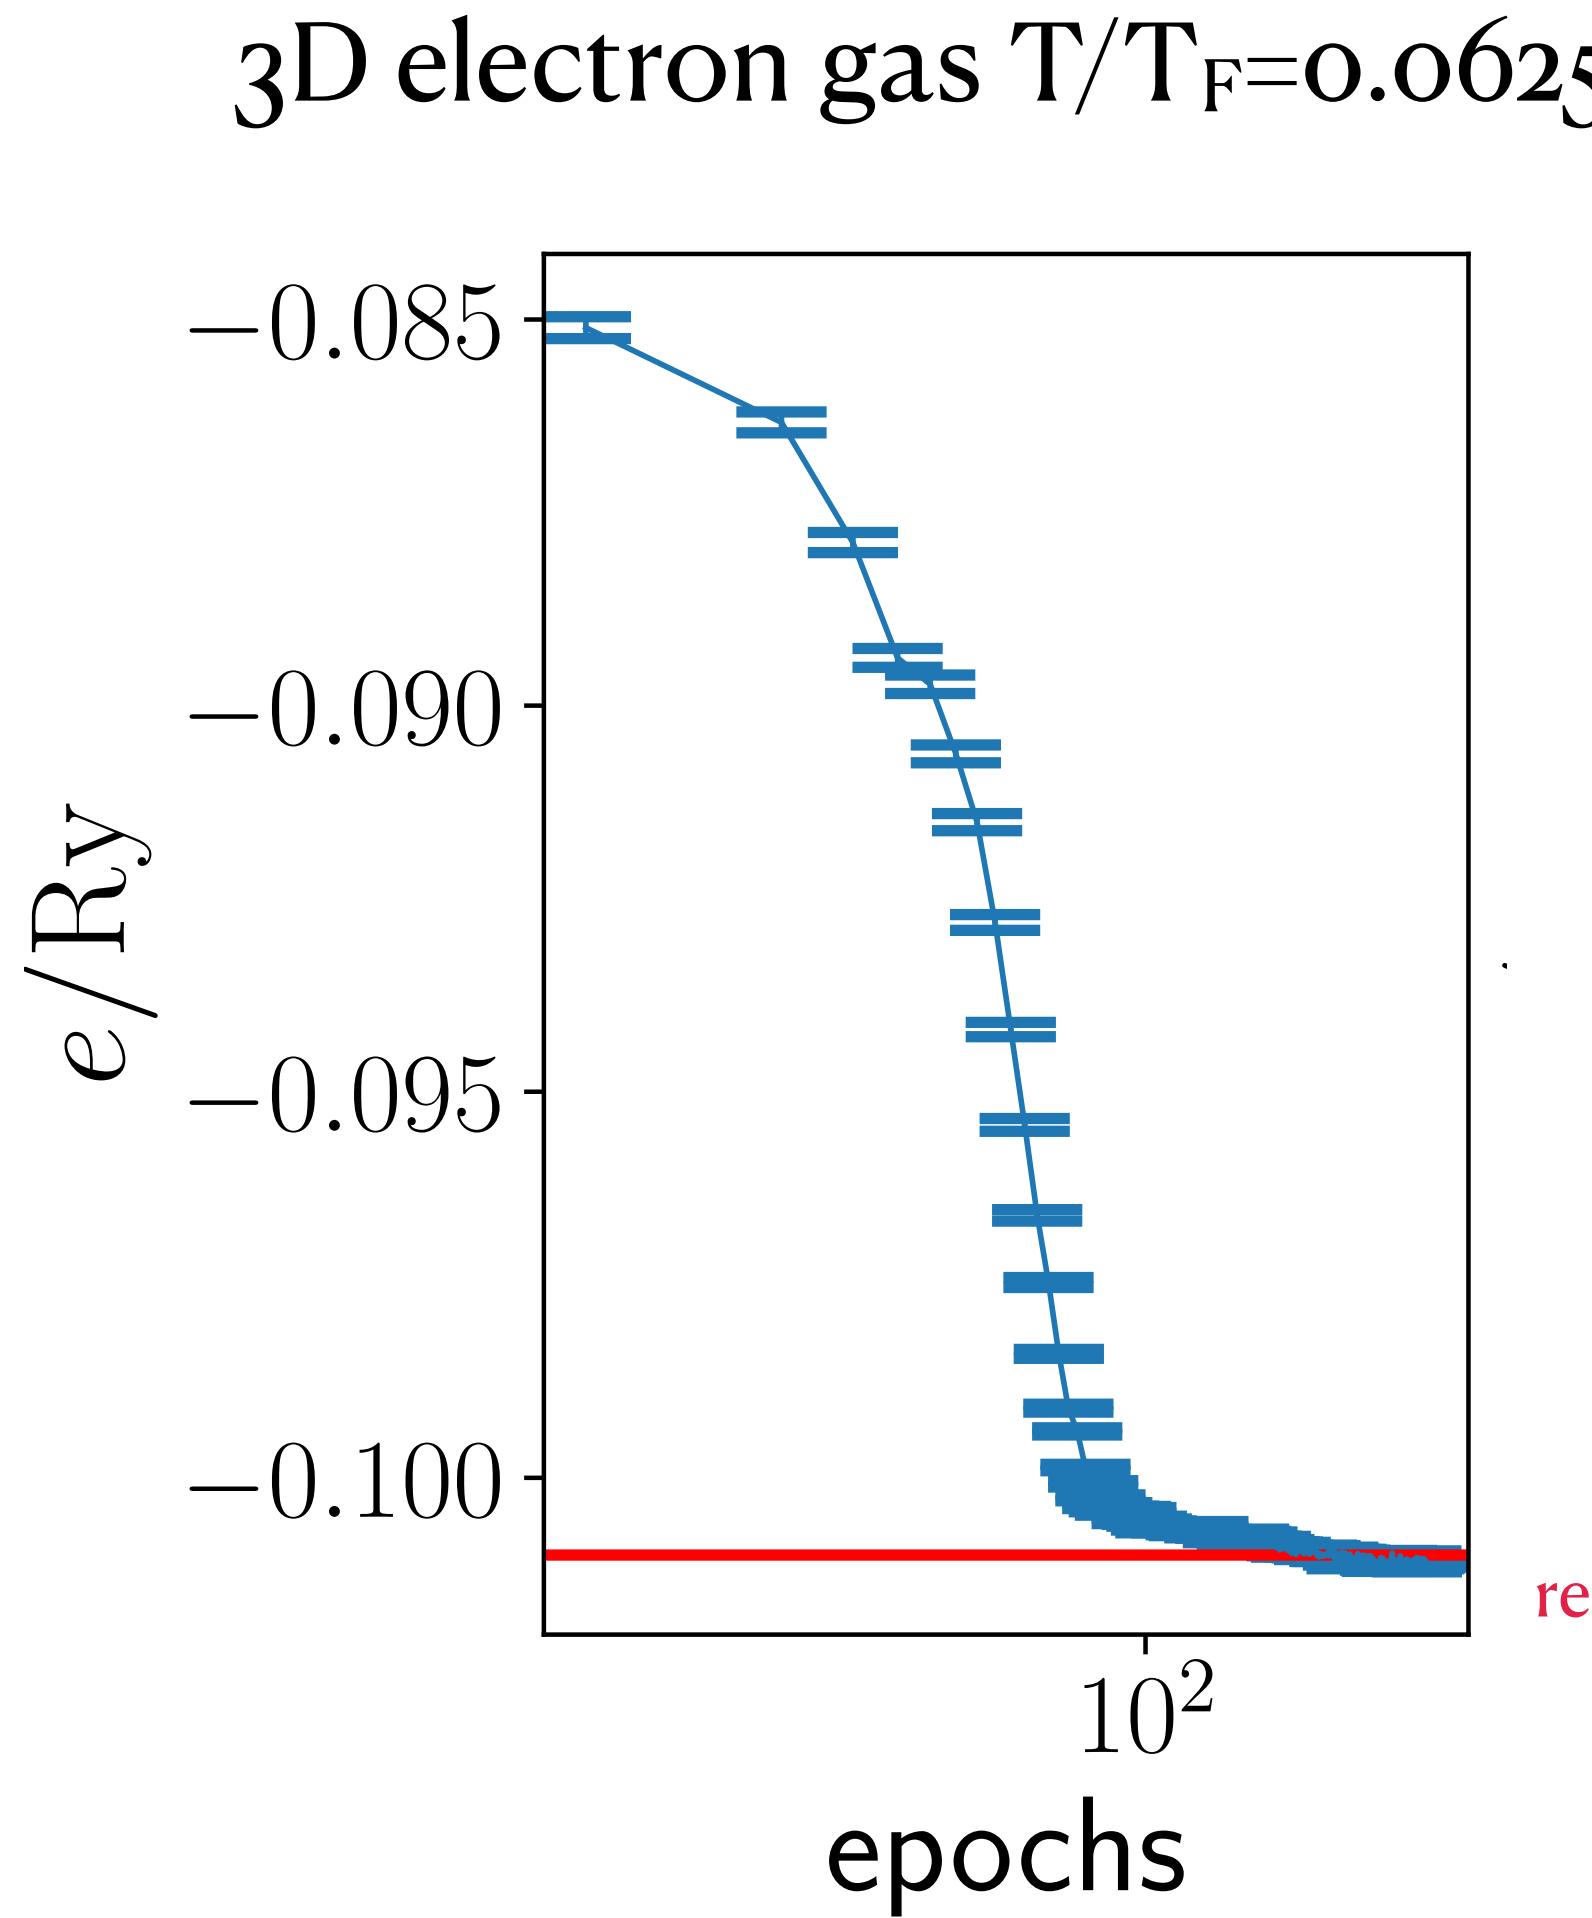
$$F = \mathbb{E}_{K \sim p(K)} \left[\frac{1}{\beta} \ln p(K) + \sum_{i=1}^N \frac{\hbar^2 k_i^2}{2m} \right]$$

A classical statistical mechanics problem:
Noninteracting fermions in canonical ensemble

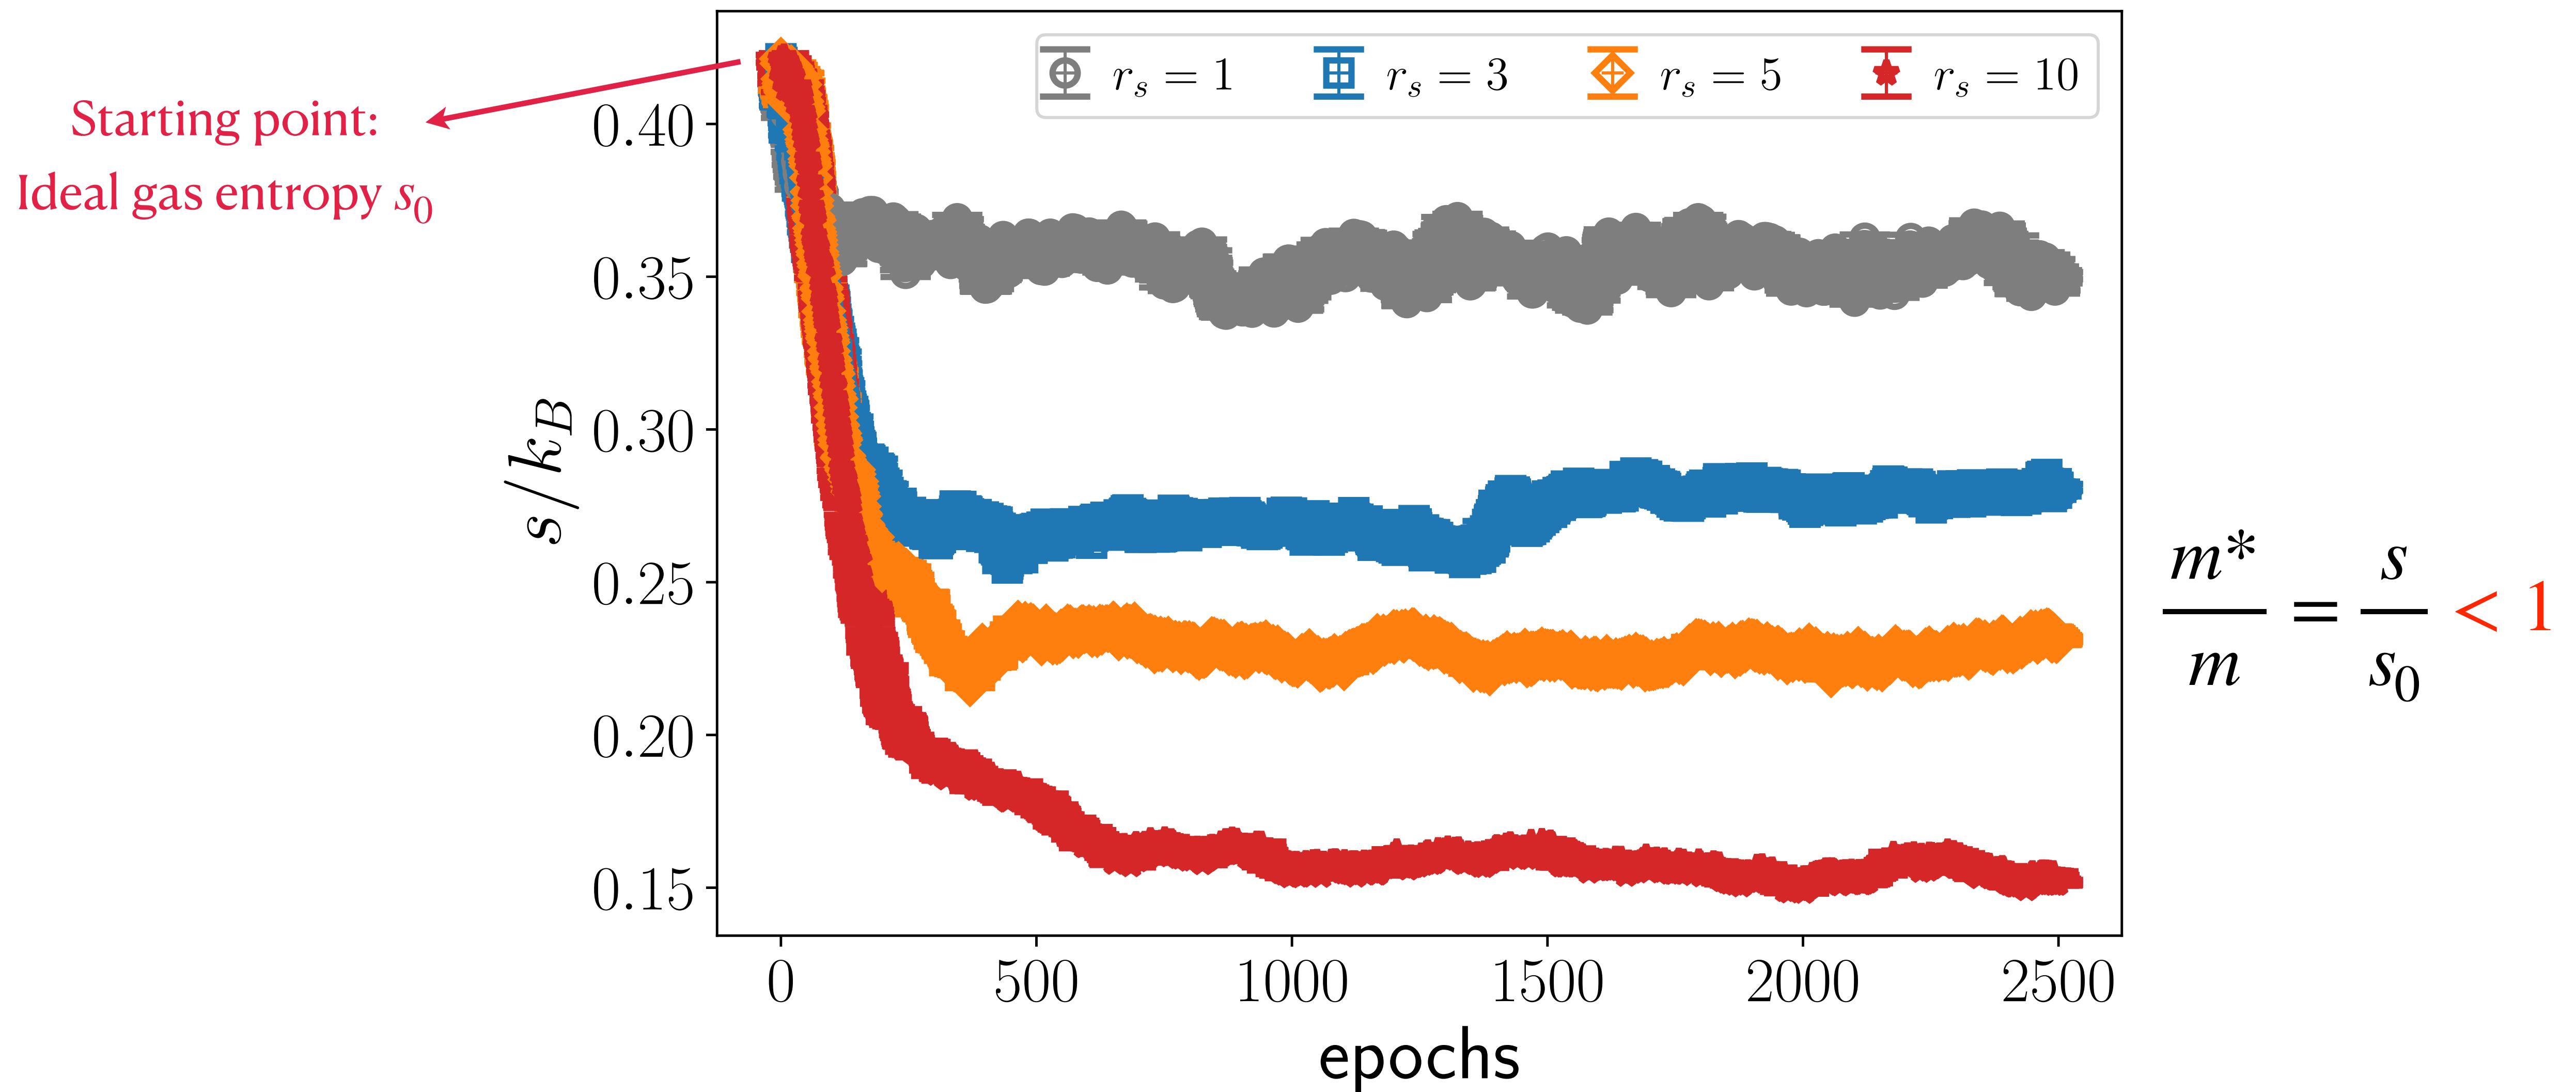
(Not as trivial as you might think) Borrmann & Franke, J. Chem. Phys. 1993

Distribute N fermions in M momenta to minimize the free energy

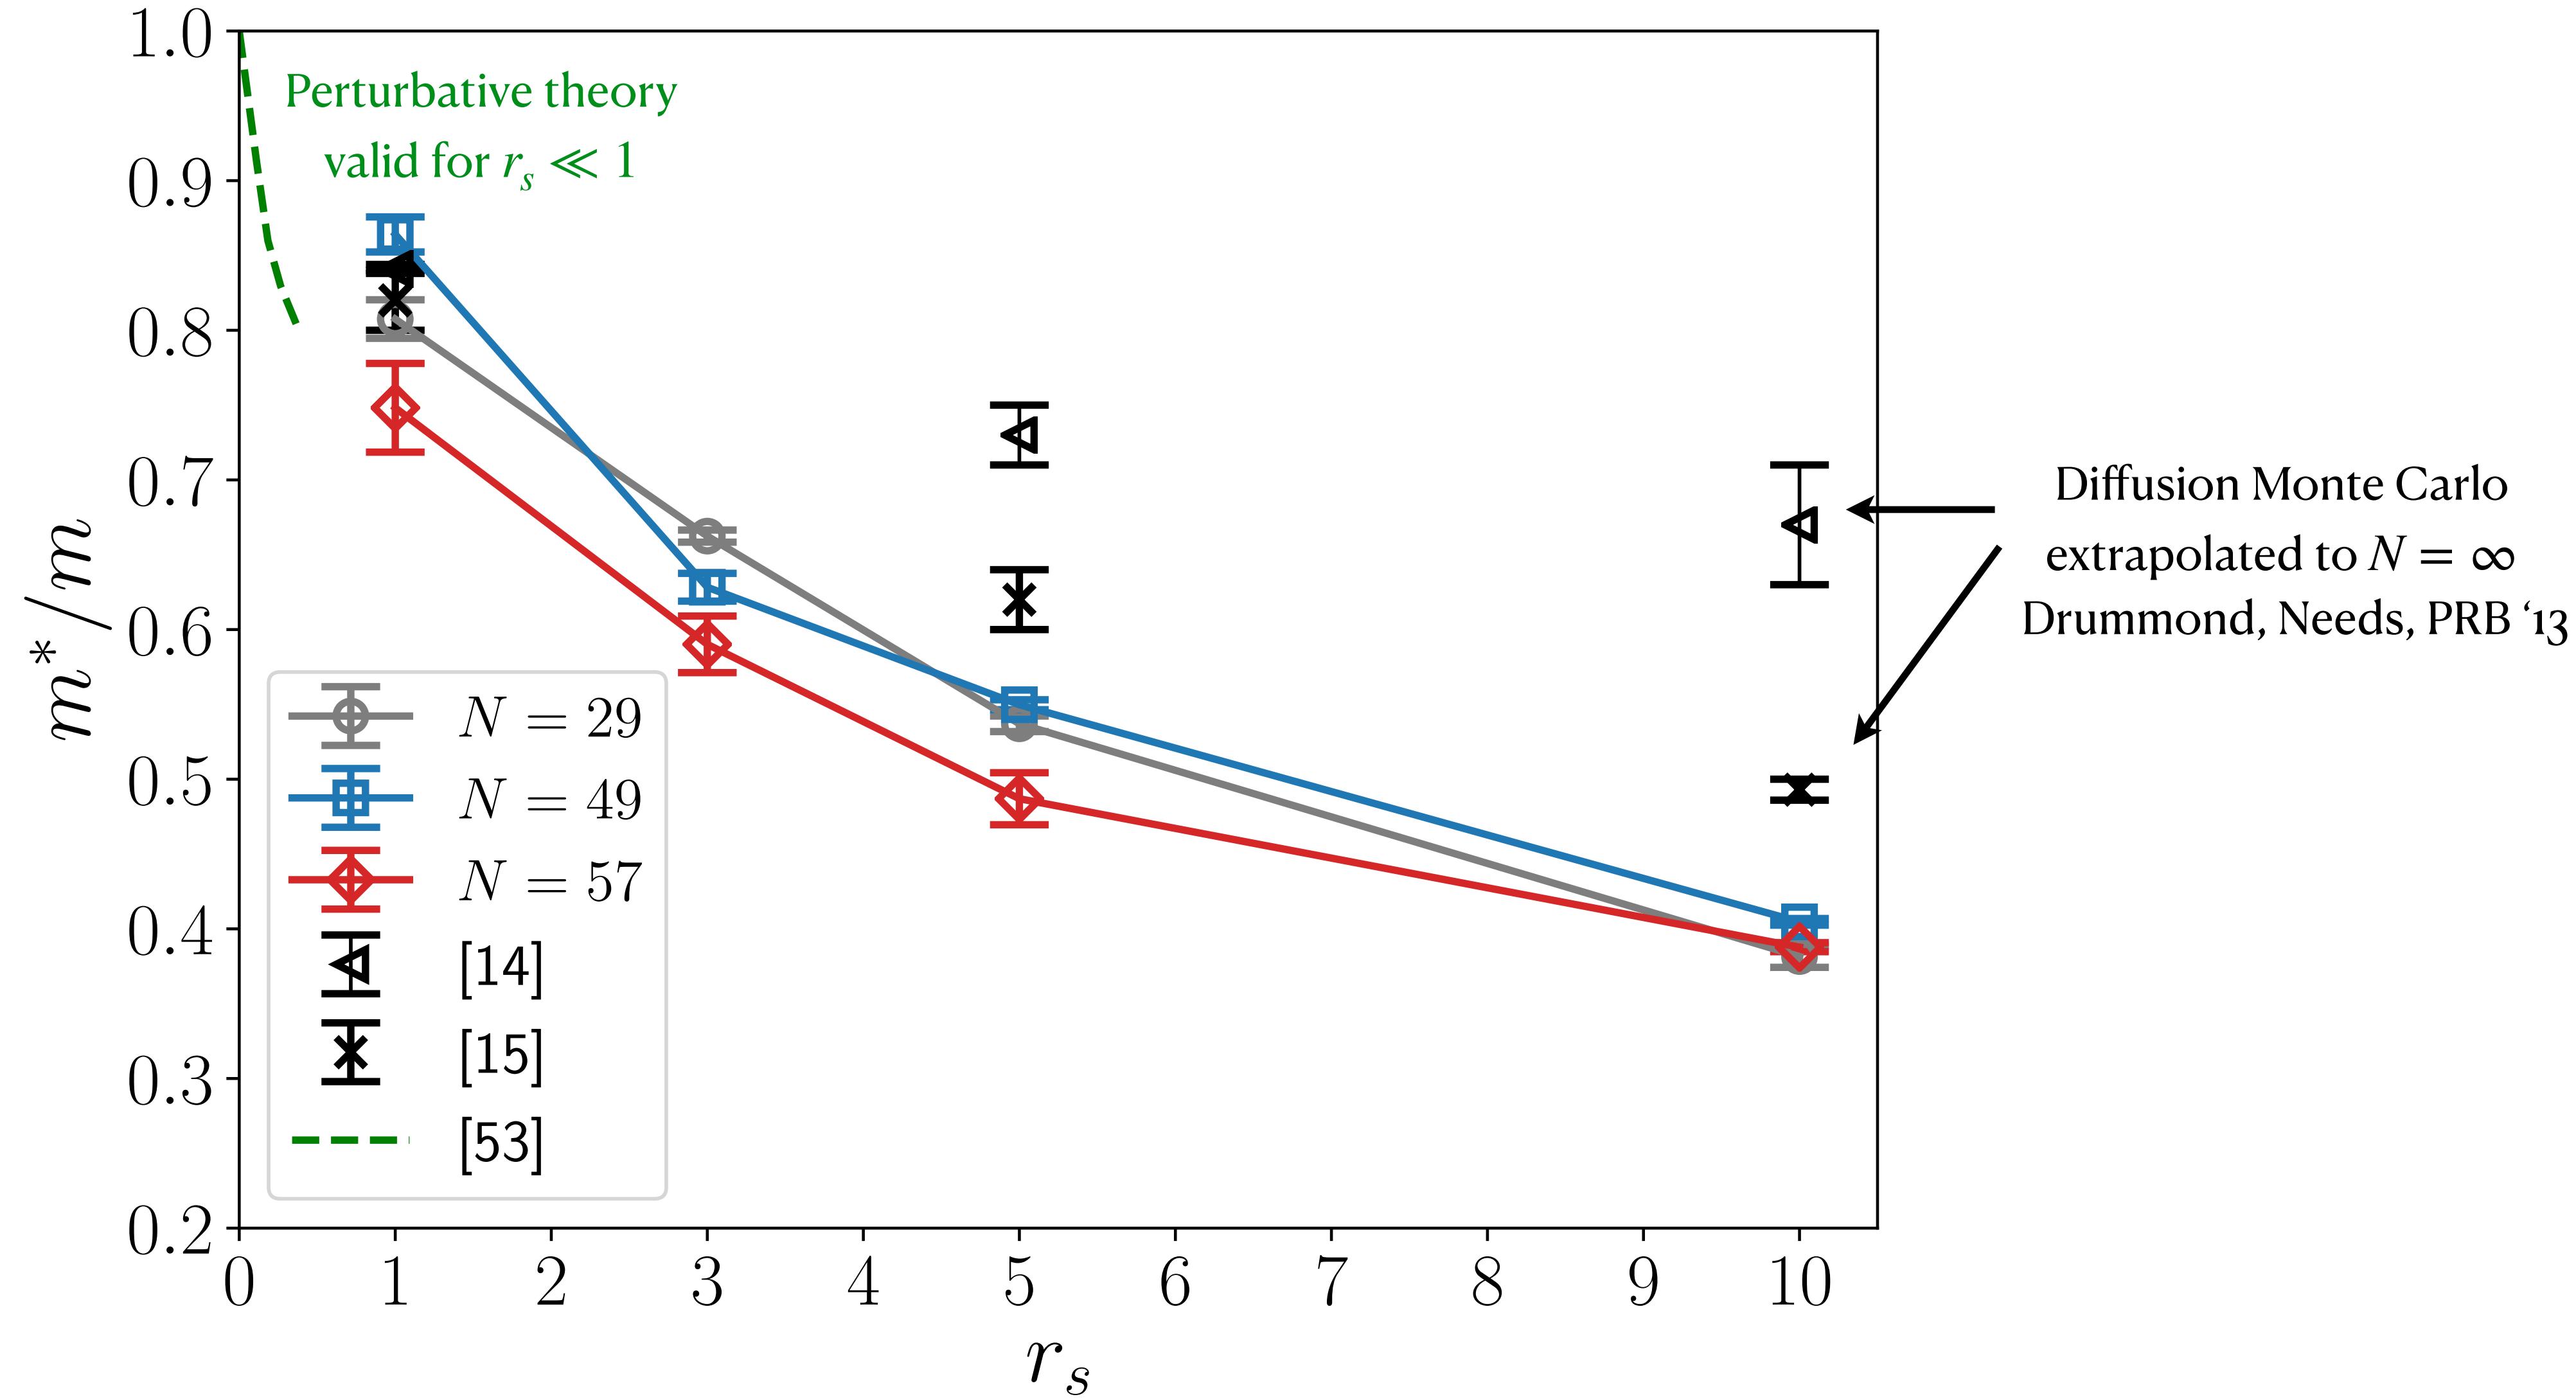
Benchmarks on spin-polarized electron gases



37 spin-polarized electrons in 2D @ T/T_F=0.15



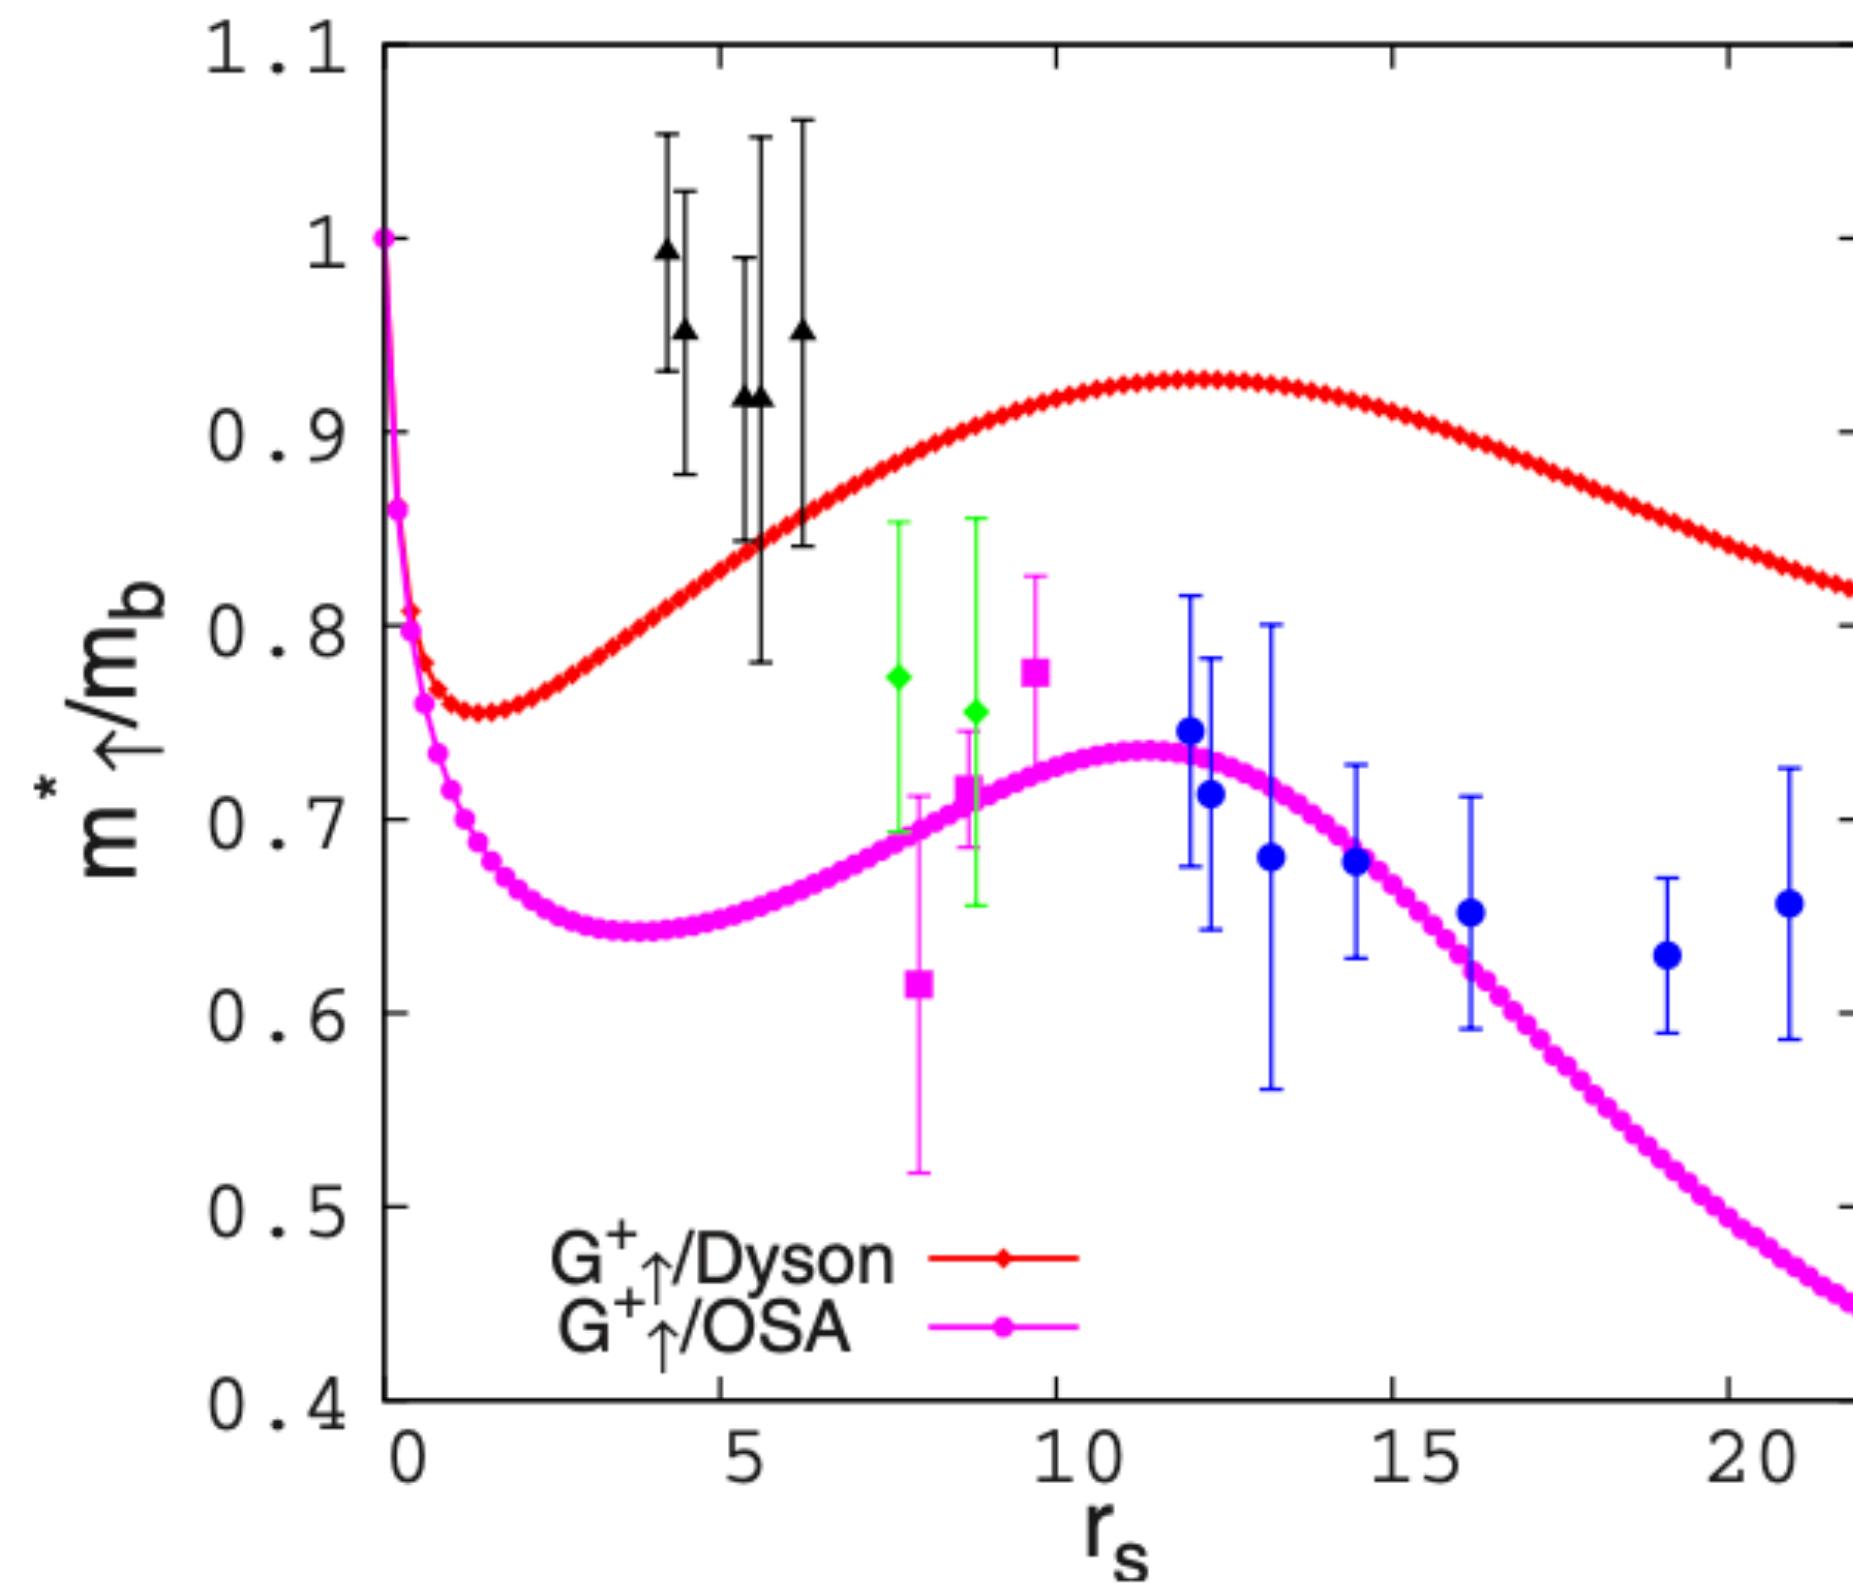
Effective mass of spin-polarized 2DEG



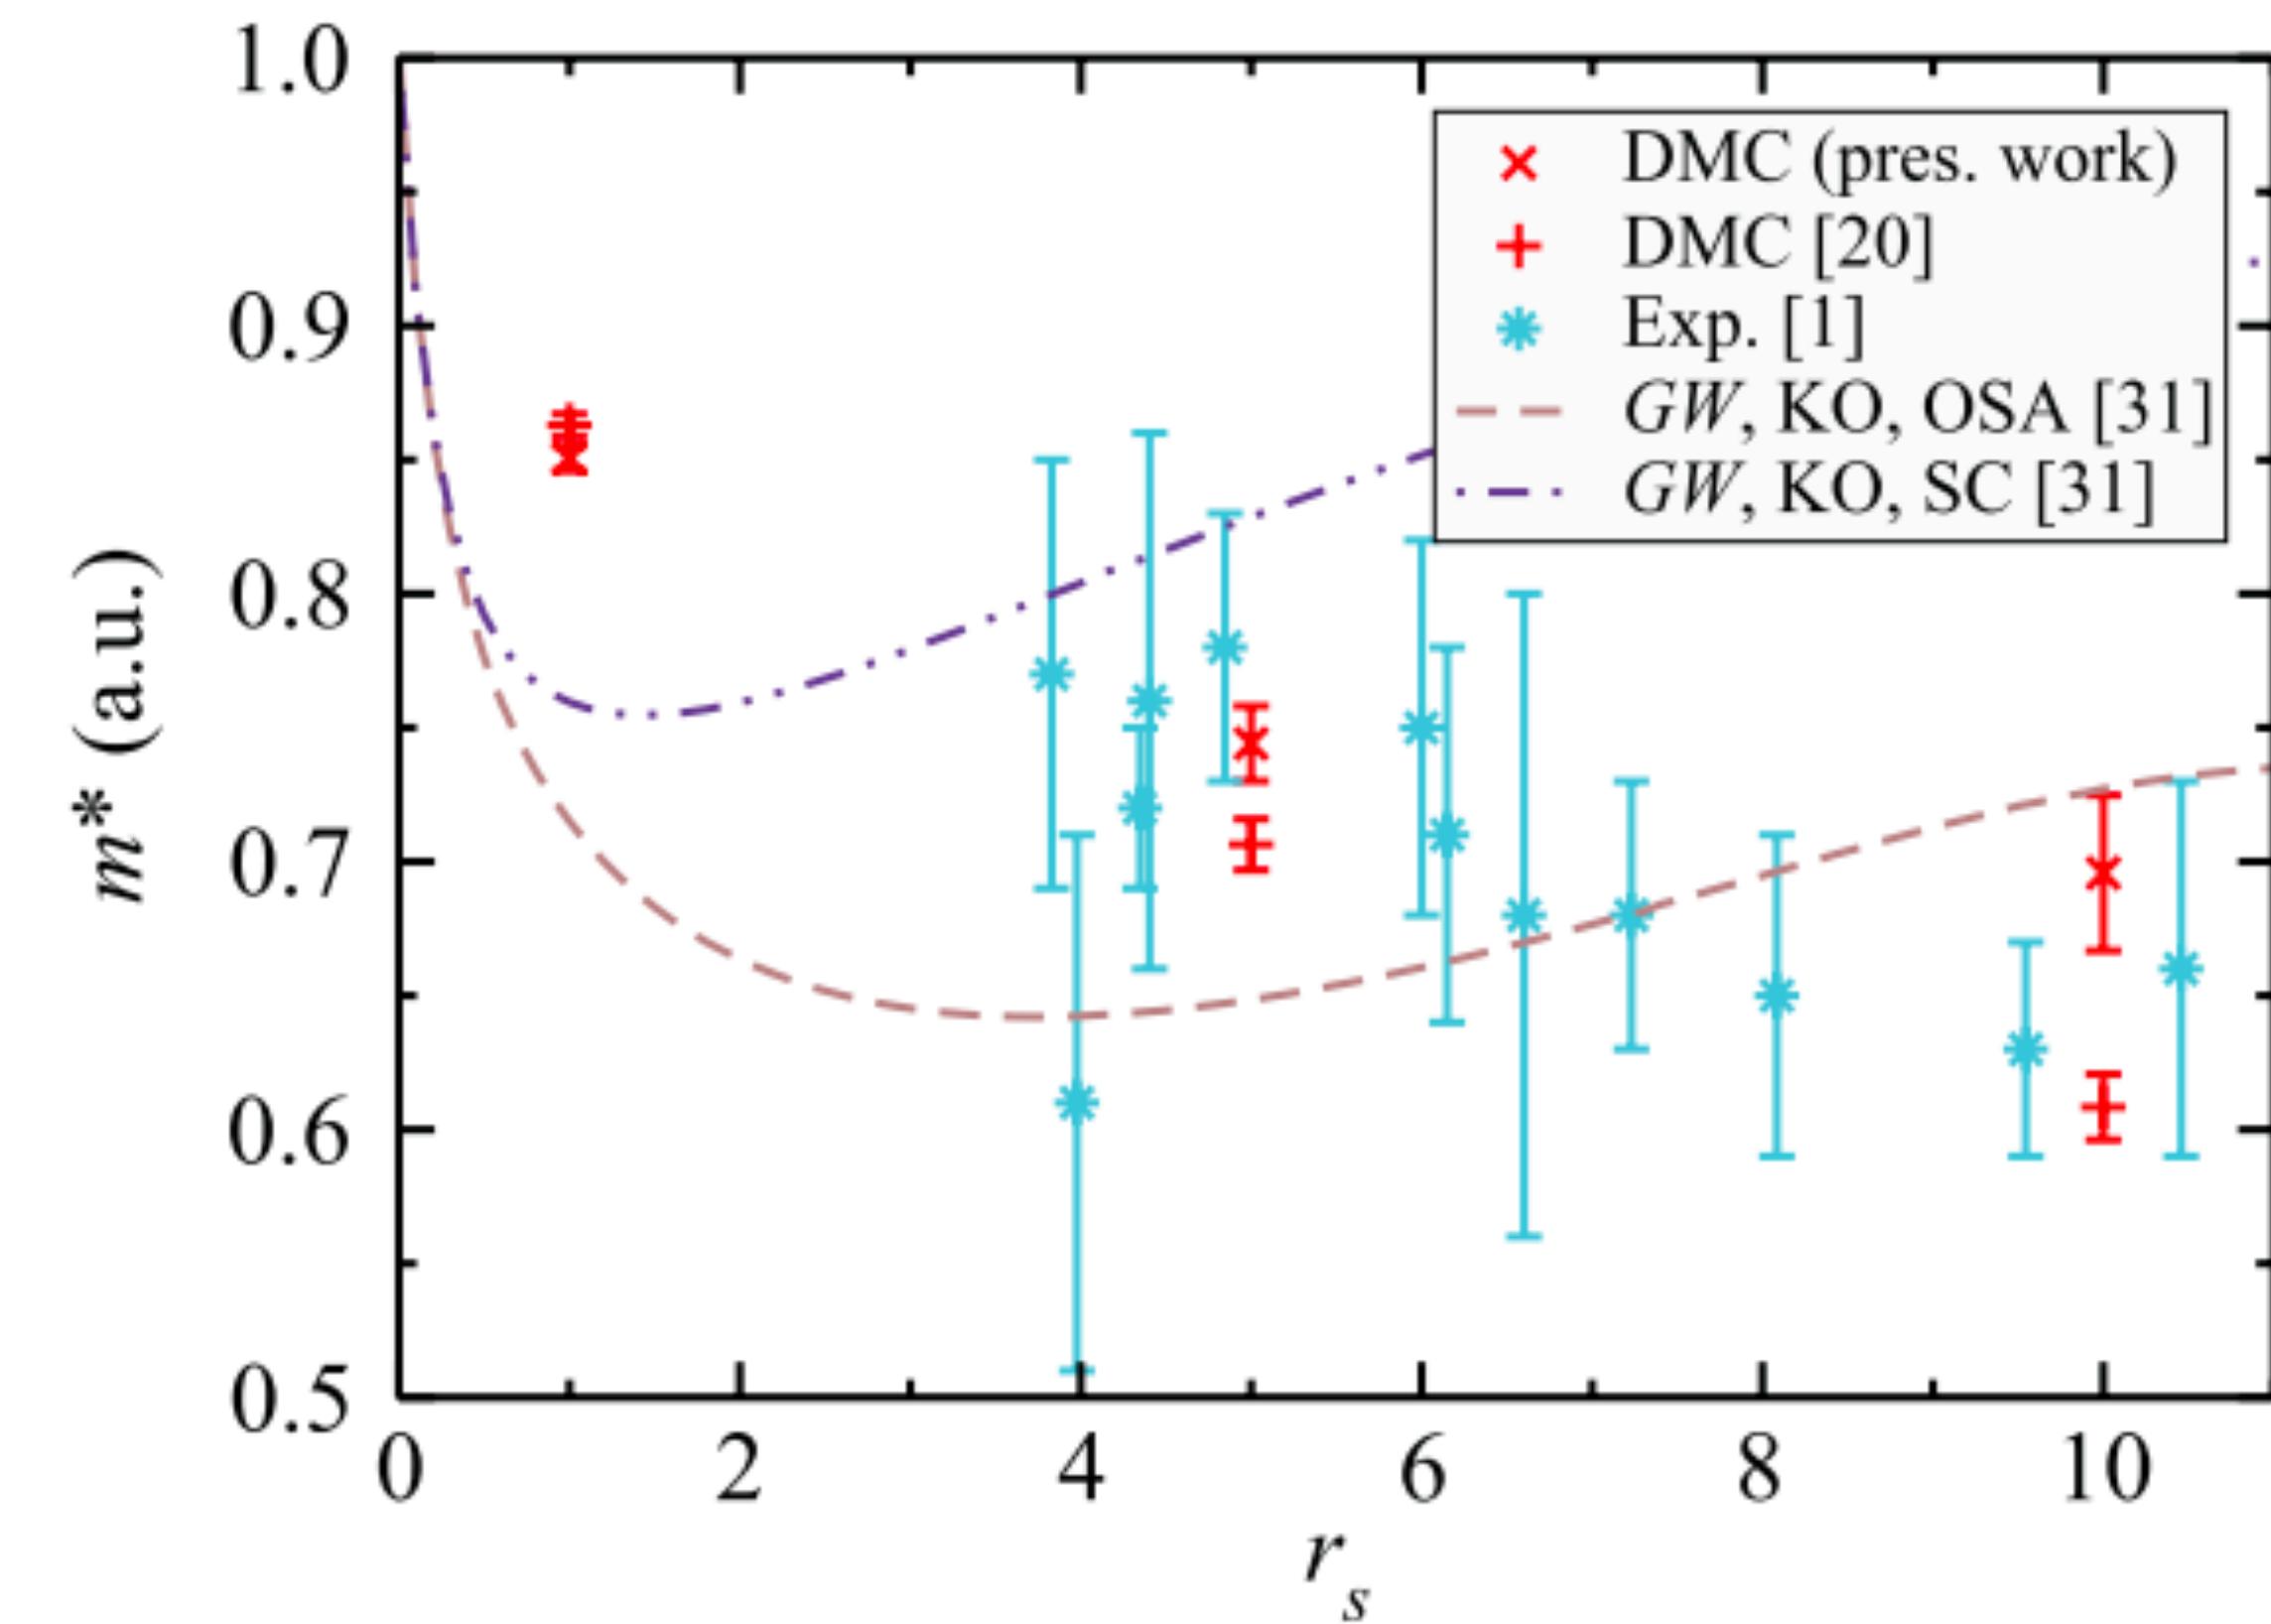
More pronounced suppression of m^* in the low-density strong-coupling region

Experiments on spin-polarized 2DEG

Asgari et al, PRB '09



Drommond, Needs, PRB'13



Quantum oscillation experiments
Padmanabhan et al, PRL '08
Gokmen et al, PRB '09

Entropy measurement of 2DEG

ARTICLE

Received 16 May 2014 | Accepted 27 Apr 2015 | Published 23 Jun 2015

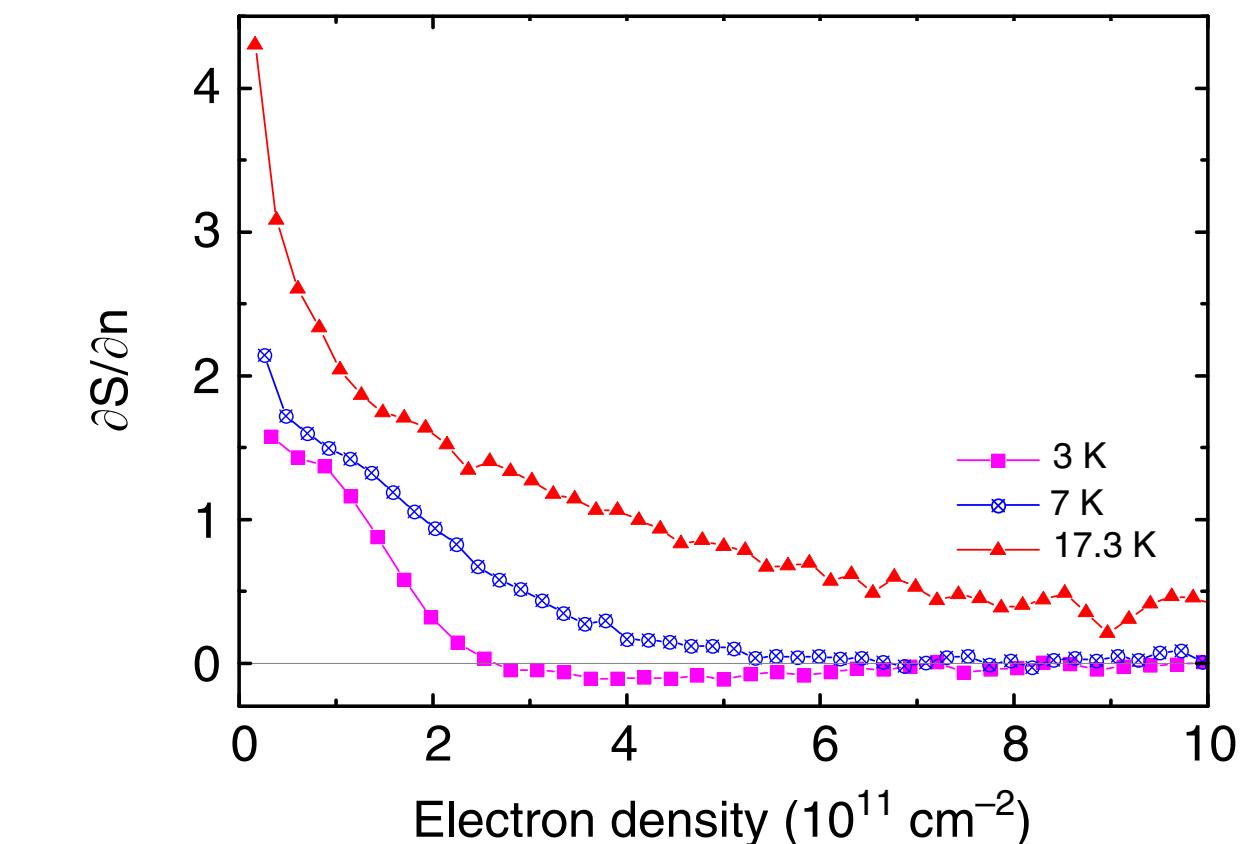
DOI: [10.1038/ncomms8298](https://doi.org/10.1038/ncomms8298)

Strongly correlated two-dimensional plasma explored from entropy measurements

A.Y. Kuntsevich^{1,2}, Y.V. Tupikov³, V.M. Pudalov^{1,2} & I.S. Burmistrov^{2,4}

Maxwell relation

$$\left(\frac{\partial S}{\partial n} \right)_T = - \left(\frac{\partial \mu}{\partial T} \right)_n$$



Next, directly compare computed entropy with the experiment

Why now ?

Variational free-energy is a **fundamental principle** for $T > 0$ quantum systems

However, it was under-exploited for solving practical problems
(mostly due to intractable entropy for nontrivial density matrices)

Now, it has became possible by integrating recent advances in
generative machine learning

FAQs

Where to get training data ?

No training data. Data are self-generated from the generative model.

How do we know it is correct ?

Variational principle: lower free-energy is better.

Do I understand the “black box” model ?

- a) I don't care (as long as it is sufficiently accurate).
- b) $\ln p(K)$ contains the Landau energy functional

$\zeta \leftrightarrow x$ illustrates adiabatic continuity.

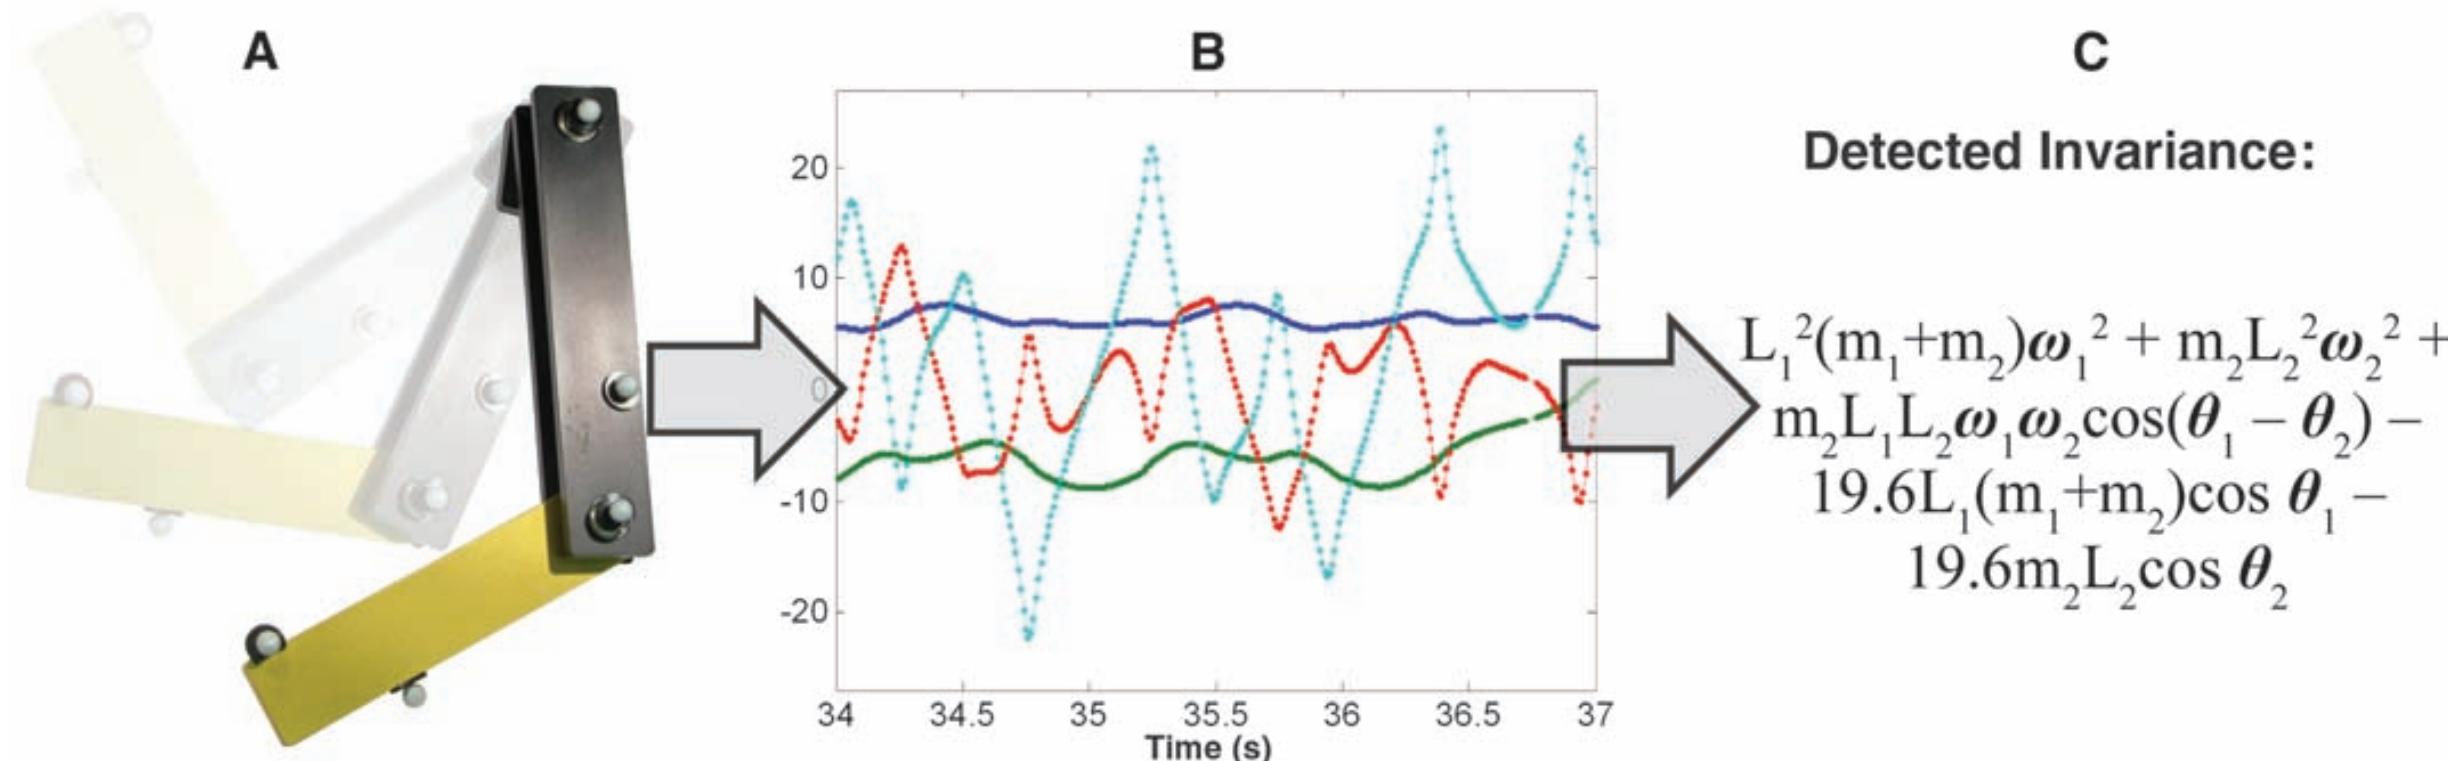
$$E[\delta n_k] = E_0 + \sum_k \epsilon_k \delta n_k + \frac{1}{2} \sum_{k,k'} f_{k,k'} \delta n_k \delta n_{k'}$$

Discussions

Can machines discover physical law ?

Distilling Free-Form Natural Laws from Experimental Data

Schmidt, Lipson
Science '09



- References and Notes**
1. P. W. Anderson, *Science* **177**, 393 (1972).
 2. E. Noether, *Nachr. d. König Gesellsch. d. Wiss. zu Göttingen, Math-Phys. Klasse* 235 (1918).

Machines Fall Short of Revolutionary Science

In the Report by Schmidt and Lipson, a machine deduces the equation behind a sample of chaotic motion. The discovery of deterministic chaos is an example of true Kuhnian revolution; others were its application to unexpected fields like meteorology and population biology. In the constrained problem in the Report, the relevant physical law and variables are known in advance; it is hardly a template for the creative, exploratory nature of true science.

PHILIP W. ANDERSON^{1*} AND ELIHU ABRAHAMS²

汤超院士在2022科学智能峰会上的讲话：“关于AI for Science的几层意思”

<https://mp.weixin.qq.com/s/oL7G7ByazbnsgrXDTToPyrw>

Discussions

Can machines discover mathematics?



Timothy Gowers
@wtgowers

An interesting paper by Adam Wagner appeared on arXiv a couple of days ago (thanks to Imre Leader for drawing my attention to it), which uses reinforcement learning to find non-trivial counterexamples to several conjectures in graph theory. 1/

[arxiv.org/pdf/2104.14516...](https://arxiv.org/pdf/2104.14516.pdf)

...

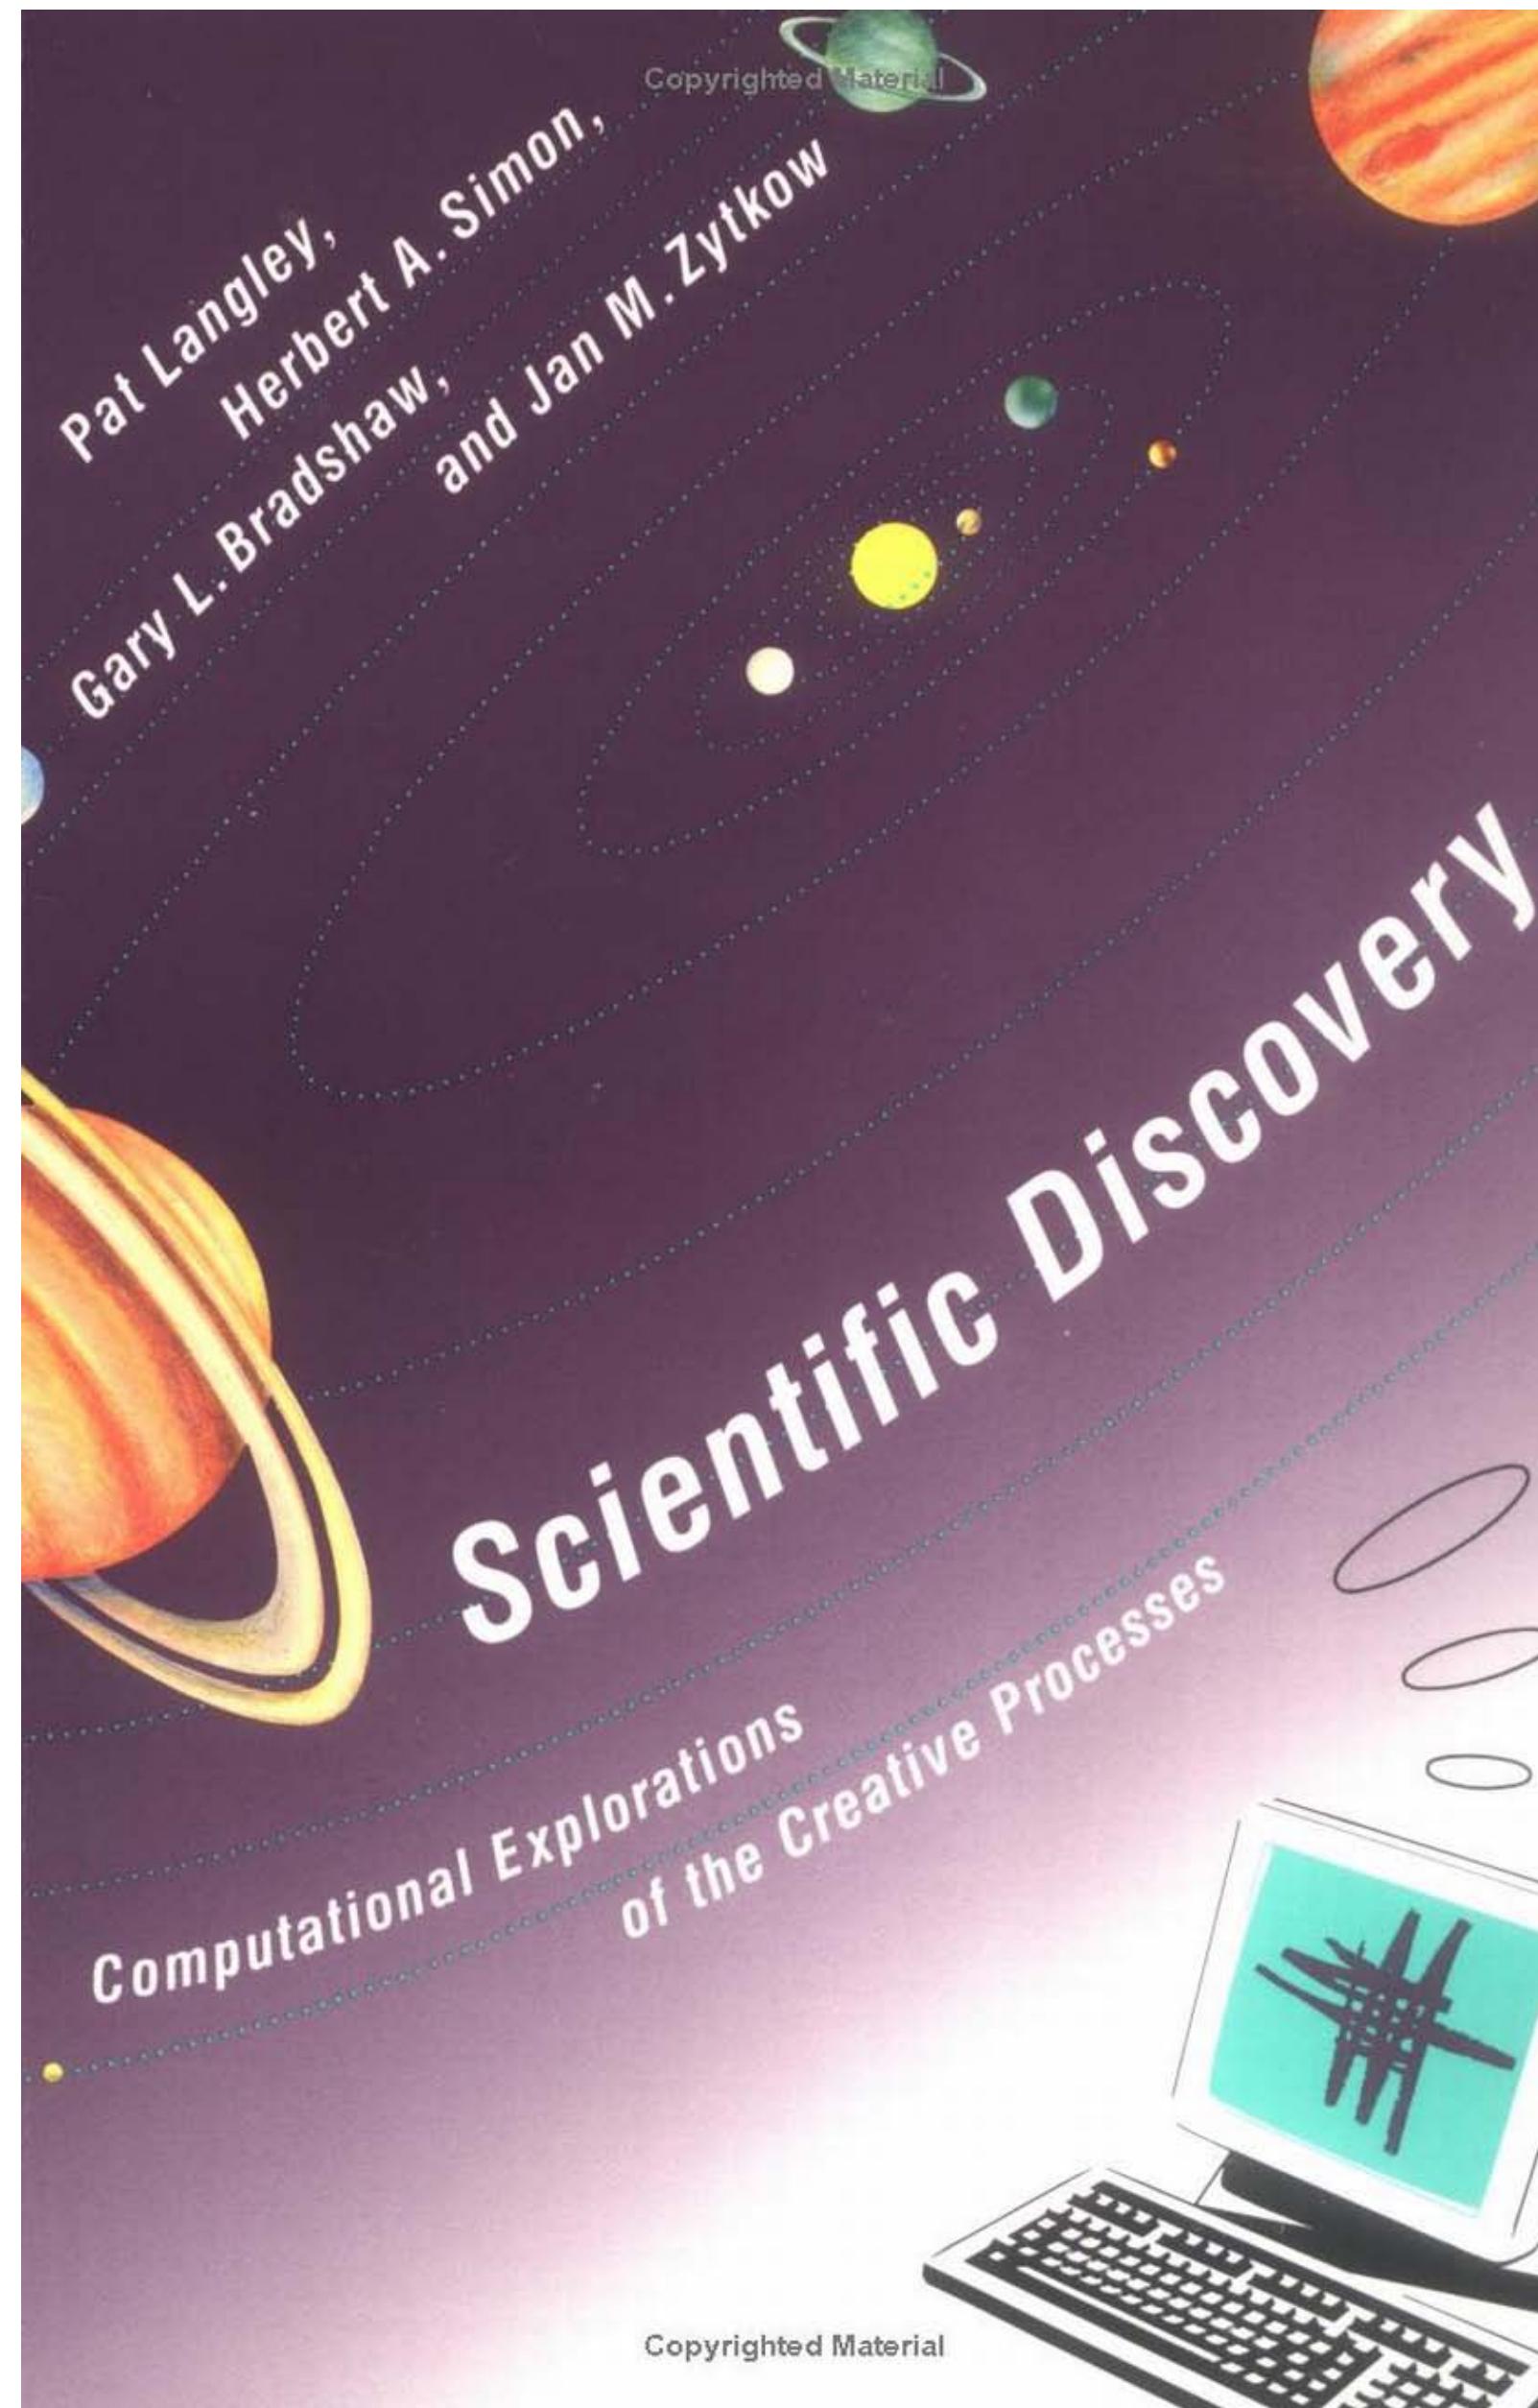
Nature 600, 70 (2021)

Advancing mathematics by guiding human intuition with AI

[Alex Davies](#)✉, [Petar Veličković](#), [Lars Buesing](#), [Sam Blackwell](#), [Daniel Zheng](#), [Nenad Tomašev](#), [Richard Tanburn](#), [Peter Battaglia](#), [Charles Blundell](#), [András Juhász](#), [Marc Lackenby](#), [Geordie Williamson](#), [Demis Hassabis](#) & [Pushmeet Kohli](#)✉

Search counter-examples to
reject conjectures

Guide human mathematician
to propose conjectures



2204.01467

On scientific understanding with artificial intelligence

Mario Krenn,^{1, 2, 3, 4, *} Robert Pollice,^{2, 3} Si Yue Guo,² Matteo Aldeghi,^{2, 3, 4} Alba Cervera-Lierta,^{2, 3} Pascal Friederich,^{2, 3, 5} Gabriel dos Passos Gomes,^{2, 3} Florian Häse,^{2, 3, 4, 6} Adrian Jinich,⁷ Akshat Kumar Nigam,^{2, 3} Zhenpeng Yao,^{2, 8, 9, 10} and Alán Aspuru-Guzik^{2, 3, 4, 11, †}

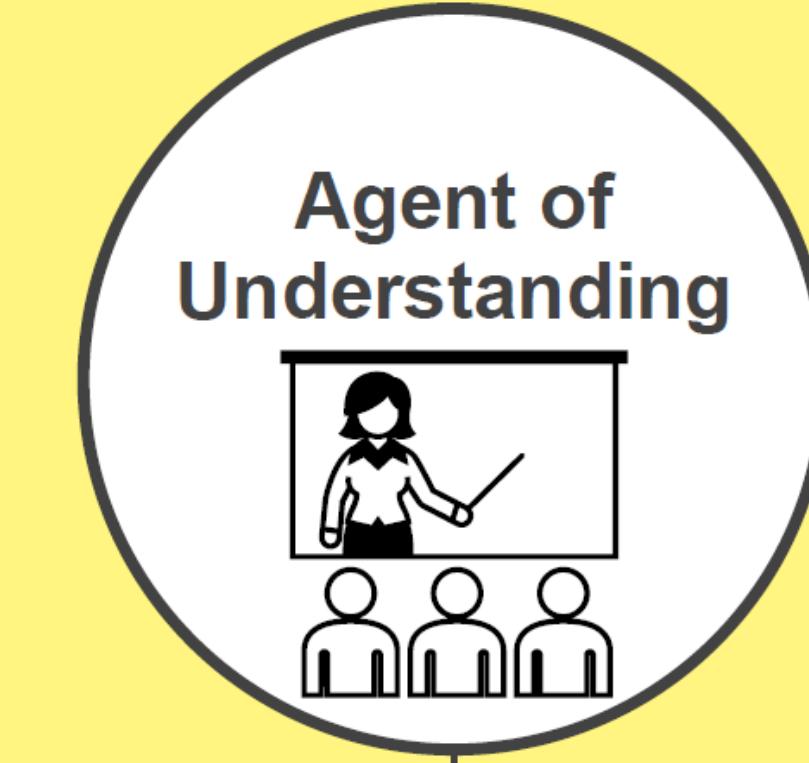
Three Dimensions of Computer-Assisted Scientific Understanding



Computational
Microscope



Resource of
Inspiration

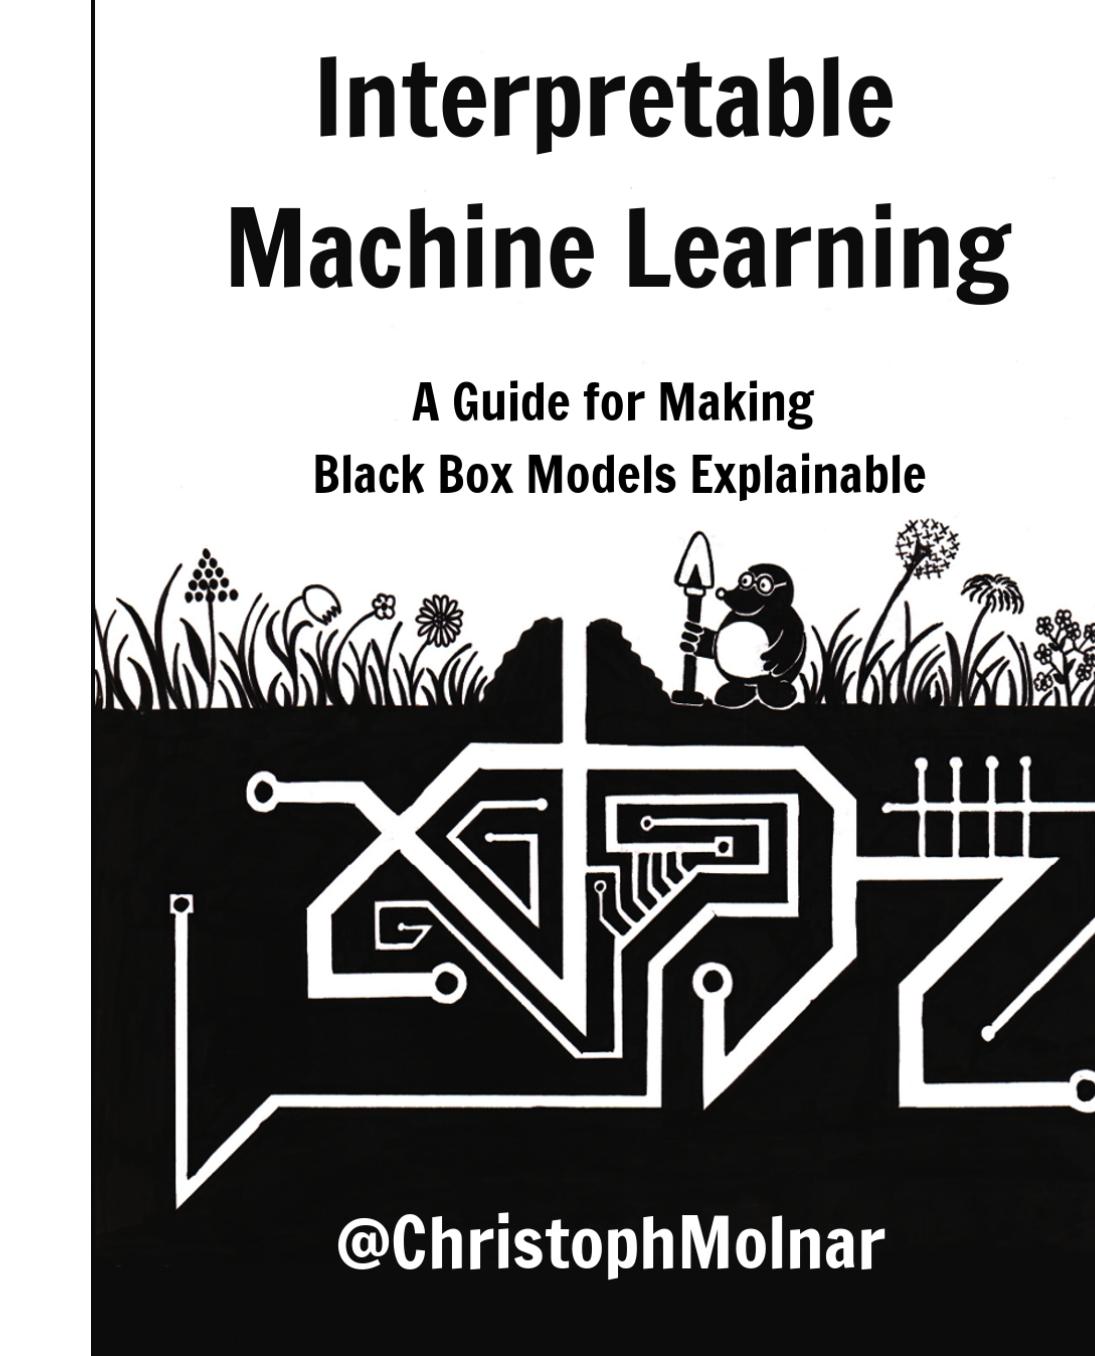
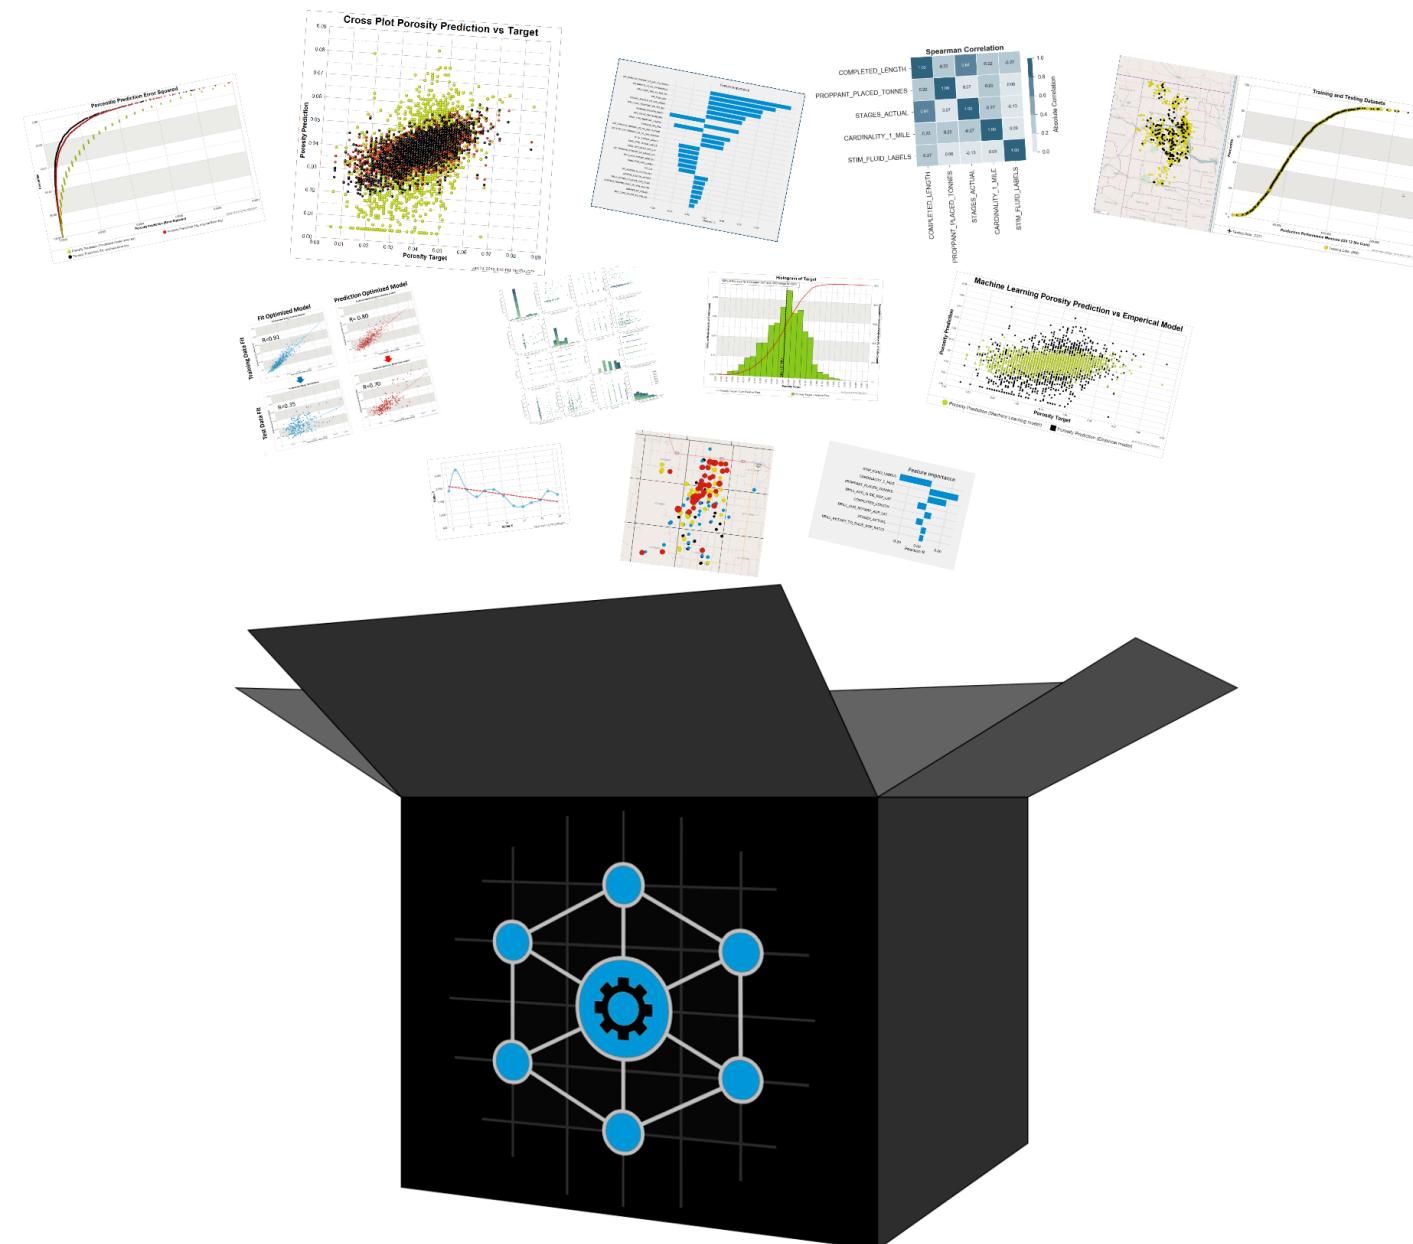


Agent of
Understanding

Copyrighted Material

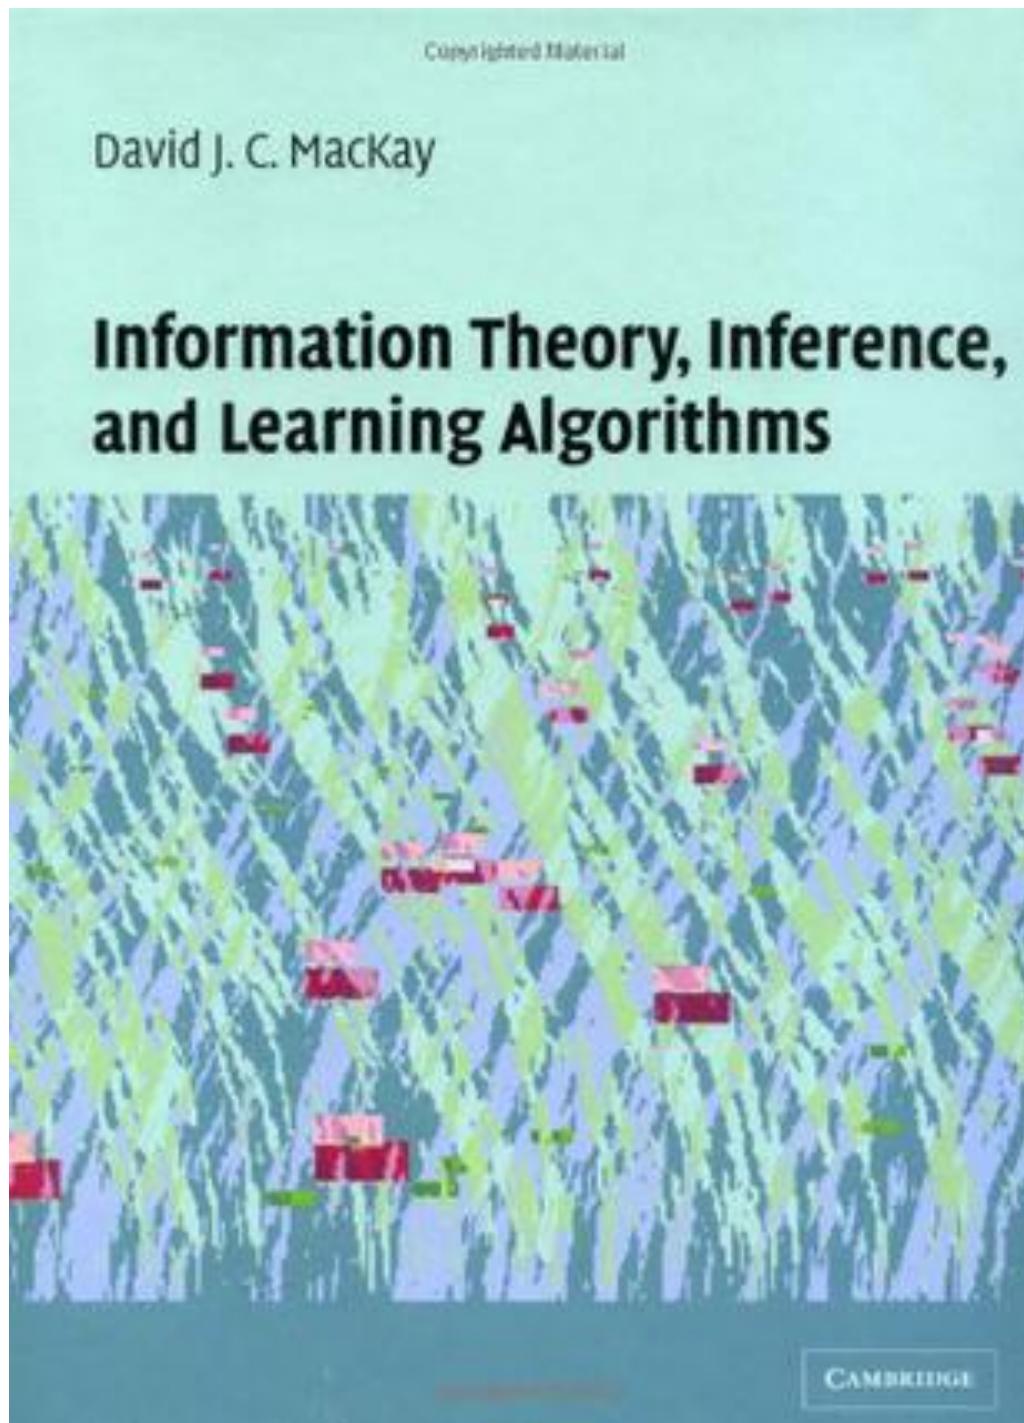
Discussions

Do we understand what is the machine doing ?



Yes/No/Well, do I have to ?/I don't care...

Discussions



Is this all fitting ?

One of my students, Robert, asked:

Maybe I'm missing something fundamental, but supervised neural networks seem equivalent to fitting a pre-defined function to some given data, then extrapolating – what's the difference?

I agree with Robert. The supervised neural networks we have studied so far are simply parameterized nonlinear functions which can be fitted to data.

True for supervised learning, which is hugely successful for real-world applications.
But that is not the whole story, especially for scientific applications.

What makes for a suitable problem?

1

Massive combinatorial
search space

2

Clear objective function
(metric) to optimise
against

3

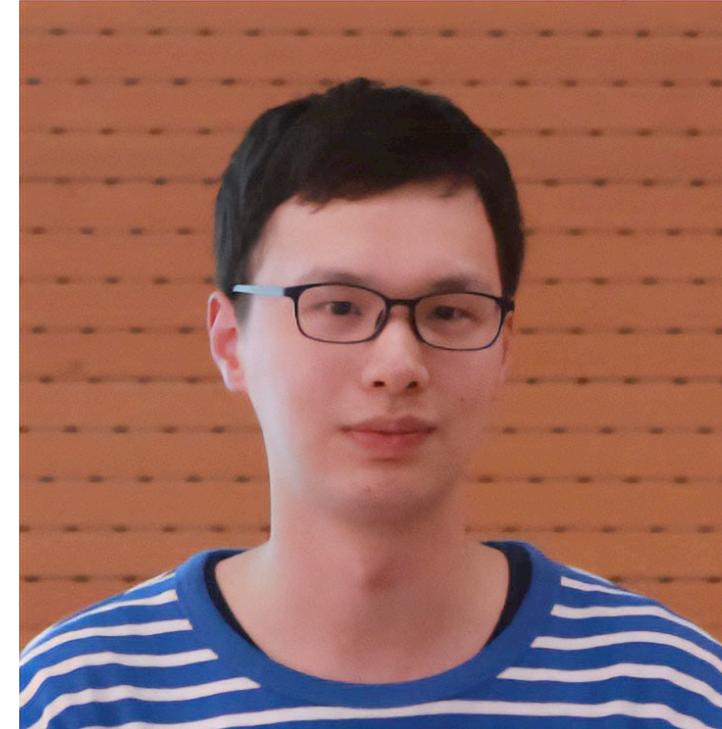
Either lots of data
and/or an accurate and
efficient simulator

Thank you!

Deep generative model-based variational free-energy calculations



Shuo-Hui Li



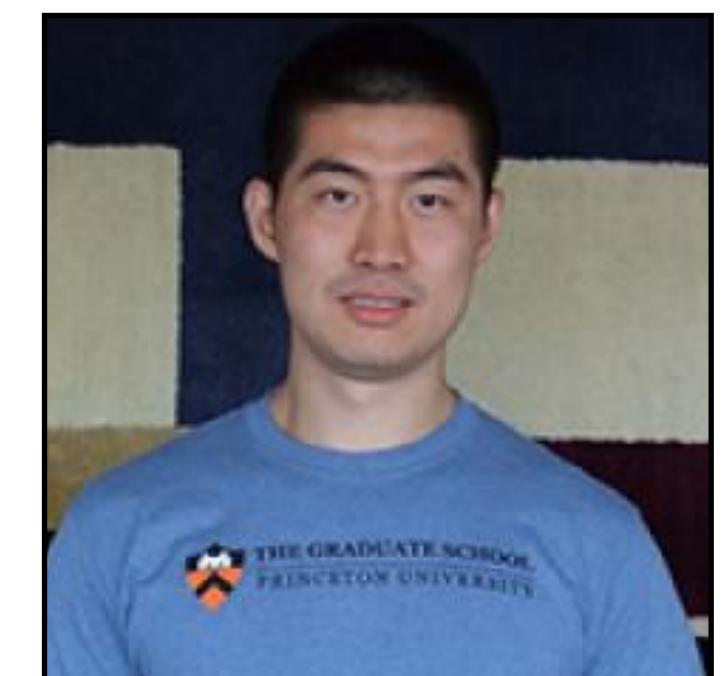
Dian Wu



Pan Zhang



Hao Xie



Linfeng Zhang



1802.02840, PRL '18
1809.10606, PRL '19
2105.08644, JML '22
2201.03156



li012589/NeuralRG
wdphy16/stat-mech-van
fermiflow/fermiflow
fermiflow/CoulombGas